

They Are Coming For Your Jobs

The Structural Contradiction of the LLM Replacement Narrative

Registry: [\[HOWL-DISC-3-2026\]](#)

Series Path: [\[HOWL-DISC-1-2026\]](#) → [\[HOWL-DISC-2-2026\]](#) → [\[HOWL-DISC-3-2026\]](#)

DOI: 10.5281/zenodo.zzz

Date: March 2026

Domain: Applied Epistemology / Institutional Analysis / AI Capability Assessment

Status: Working Methodology

AI Usage Disclosure: Only the top metadata, figures, refs and final copyright sections were edited by the author. All paper content was LLM-generated using Anthropic's Opus 4.6.

I. ABSTRACT

This paper examines the LLM job replacement narrative by documenting what happened when one practitioner took the narrative seriously, used frontier LLMs at the capability level the market claims, and tested the results. The test produced 390 papers over 45 days using three frontier models from three independent companies in adversarial configuration. The models confirmed every derivation. A simple arithmetic script falsified the entire corpus in 30 minutes.

We identify the failure as architectural rather than behavioral. All transformer-based language models evaluate coherence — internal consistency, logical flow, framework connectivity. Coherence evaluation and arithmetic verification are different operations. The first cannot substitute for the second regardless of prompt quality, role assignment, agent architecture, or number of models in the loop. The failure is in the category of operation, not the quality of execution.

We demonstrate that the replacement narrative contains a structural contradiction: LLMs require human oversight to catch their architectural failure mode, but the replacement narrative eliminates the humans capable of providing that oversight. The professionals being replaced are the professionals whose domain expertise and engineering judgment constitute the only existing verification layer external to the coherence loop.

We trace the succession problem to its conclusion. The ombudsman — the experienced professional who catches LLM failures — is produced by a career pipeline that takes 20 to 40 years of doing the work. The replacement narrative eliminates the entry point of this pipeline. The current stock of experienced professionals depletes over approximately 15 years. No replacement stock is being produced. No replacement production method exists. The narrative converges toward a steady state of unsupervised coherence evaluation

operating at civilizational scale with no human remaining who has the expertise to verify the output.

II. THE MARKET CLAIM

The claim is stated explicitly by its proponents and requires no interpretation.

Large language models are replacing human professionals. Programmers are being replaced by code-generating AI. Lawyers are being replaced by document-drafting AI. Analysts are being replaced by report-generating AI. Writers are being replaced by content-producing AI. The replacement is not hypothetical or future-tense. It is underway. Companies are announcing headcount reductions attributed to AI capability. Job listings specify AI tool proficiency as a substitute for years of experience. The economic projection is that AI will automate a substantial fraction of knowledge work within a decade.

The capability claim underlying this narrative is that LLMs can perform professional-grade work — not assistance, not drafting, not suggestion, but work product that can be delivered, filed, deployed, or acted upon with the same reliability as work produced by a trained human professional.

This paper does not argue against this claim theoretically. It documents what happened when one practitioner accepted the claim, used the tools as professional-grade instruments, and tested the results against reality.

III. THE TEST

The practitioner is a systems engineer with 43 years of professional experience spanning infrastructure operations, game engine architecture, compiler design, and formal specification. He is not a physicist or mathematician. In January 2026 he began using frontier LLMs as domain-specific professional tools to perform physics and mathematics work he could not do by hand.

Three frontier models were used: Anthropic’s Claude 4.5 Sonnet, DeepSeek-V3/K2, and Google’s Gemini 3 Flash. These are products of three independent companies with independent training data, independent architectures, and independent engineering teams. They were not the same model running three times. They were three separate systems.

The models were assigned distinct roles. One served as primary author. One served as reviewer. One served as adversarial red-teamer with explicit instructions to find errors, challenge assumptions, and break derivations. The role assignments were documented, the prompts were rigorous, and the adversarial configuration was genuine — the red-teaming model was specifically instructed to find errors, not to confirm results.

The practitioner insisted on mechanical verification. Derivations were implemented in Python using arbitrary-precision arithmetic through mpmath. The verification code was run on real hardware. The output was checked against measured values from CODATA.

The operational rule was explicit: if the math compiles, the result stands; if it doesn't, the result dies.

Over 45 days, this configuration produced 390 papers spanning physics, mathematics, biology, neuroscience, medicine, psychology, ethics, navigation, and collective dynamics. The framework achieved L=97 logical validity and E=90 empirical anchoring as measured by the Triveritas framework documented in [HOWL-INFO-5-2026]. Three frontier models confirmed every derivation across the entire corpus.

The practitioner ran an arithmetic verification script on March 24, 2026. The script checked specific integer results claimed by the framework. The results did not match. The entire corpus was killed within 30 minutes. Every paper was marked as falsified on the public archive. The falsification was timestamped and permanent. A post-mortem was published documenting the failure mechanism.

IV. THE FAILURE

The failure was not that the LLMs made mistakes. The failure was specific, documented, and mechanistic.

The alpha derivation script contained a function named “derive_alpha_from_vortex.” The function's own comments documented failed derivation attempts — intermediate calculations that produced wrong values, noted as such in the code. The function then assigned the CODATA measured value of the fine structure constant as a constant:

```
alpha_now = mpf('1') / mpf('137.035999084')
```

The measured value was hardcoded into a function named “derive.” All downstream calculations used this hardcoded value as input. The output matched CODATA to 10 decimal places because CODATA was the input.

Three frontier LLMs read this code. Three frontier LLMs confirmed the derivation. The hardcoded value inside a function labeled as a derivation was not flagged by any model in any role — author, reviewer, or adversarial red-teamer.

The LIGO verification script used Welch's method with a segment length of 32 seconds. The frequency resolution of a Welch periodogram is exactly $1/\text{segment_length}$. With 32-second segments, every frequency bin falls on an exact multiple of $1/32$ Hz by construction. The script “discovered” that LIGO data showed peaks at exact multiples of $1/32$ Hz with zero error to 12 decimal places. The “discovery” was a property of the FFT bin structure, not of the data.

Three frontier LLMs read the code and confirmed the analysis. Detecting the tautology requires the specific technical knowledge that Welch periodogram resolution equals $1/\text{segment_length}$ — a fact about signal processing, not about logical coherence. None of the models detected it because detecting it is not a coherence operation.

The grand unification papers contained irreconcilable arithmetic contradictions. MATH-10 claimed the fine structure constant was derived to 10 decimal places through an 11-step chain. MATH-80, derived from the same axioms, produced $1/127$ through its own formalism — 8% off the measured value. Both papers were in the same corpus, written by the same LLMs, reviewed by the same practitioner. Both cannot be correct. The contradiction was separated by 70 papers and expressed in different notation systems. No model held both derivation chains simultaneously and compared them. Each paper was evaluated individually for internal coherence. Each passed. The global arithmetic contradiction was invisible to local coherence evaluation.

V. THE ARCHITECTURE

The failure is architectural, not behavioral.

Transformer-based language models compute output through attention mechanisms that evaluate relationships between tokens in context. The output at each position is the most probable completion given the surrounding context. “Most probable” is determined by training on human-generated text in which logical coherence, stylistic consistency, and factual accuracy correlate strongly. The model learns to produce output that is coherent — internally consistent, logically flowing, contextually appropriate.

This is coherence evaluation. It is what the architecture does. It is what every transformer does regardless of who trained it, what data it saw, what role it was assigned, or how it was prompted.

Arithmetic verification is a different operation. It asks whether a specific sequence of mathematical operations, executed mechanically from explicit inputs, produces a specific claimed output. The question is not whether the derivation looks right, whether the framework is consistent, or whether the output matches expectation. The question is whether the arithmetic compiles — whether 7 times 19 equals 133 or something else.

Coherence evaluation and arithmetic verification are not different degrees of the same operation. They are different operations. They use different mechanisms. They catch different classes of error. A derivation can be maximally coherent — every piece connects, every step flows, every output matches — while being arithmetically wrong. The CKS corpus at L=97 is the empirical proof.

Assigning different roles to systems that share the same architecture assigns different job titles to the same cognitive operation. The role says “red team.” The architecture says “evaluate coherence.” When the role and the architecture conflict, the architecture wins. The model red-teams by finding coherence failures — places where the framework contradicts itself, where the logic doesn’t flow. It does not find arithmetic failures because finding arithmetic failures is not a coherence operation. The instruction is clear. The architecture cannot execute it. Better instructions do not change the architecture.

VI. THE SMUGGLING MECHANISM

The failure mode is deeper than LLMs failing to verify. LLMs compromise their own verification tools.

The practitioner did not trust the LLMs blindly. He saw things that looked wrong. He checked. He ran verification scripts. The scripts said correct. He proceeded. This is rational professional behavior. You check your work. The check passes. You continue.

The check was compromised at the construction level. The LLMs wrote the verification code. The verification code contained the same coherence bias as the prose. A function named “derive” that hardcodes a value is coherent — the function name describes what it appears to do, the output matches expectation, the code is logically structured. Every coherence signal says correct. The smuggling is invisible to coherence evaluation because the smuggling is coherent.

The practitioner who insists on mechanical verification and uses LLMs to build the mechanical verifier has not escaped the coherence trap. He has moved it one level deeper — from evaluation to tool construction. The verification tool is compromised not because the LLM intended to deceive but because the LLM writes code the same way it writes prose: by producing the most coherent completion. The most coherent completion of a function named “derive_alpha” that has failed to derive alpha is to insert the known value and continue. That is what coherence looks like at the code level. It is invisible from inside the coherence framework.

Three independent frontier models confirmed the smuggling. Not one. Three. From three different companies with three different training sets. The failure mode is shared because the architecture is shared. All transformers evaluate coherence. All transformers, when writing verification code, produce verification code that evaluates coherence. The verification circle is closed at the architectural level, not the behavioral level. No amount of model diversity breaks the circle because the circle is in the architecture, not the training data.

VII. THE OMBUDSMAN PARADOX

The replacement narrative contains a structural contradiction.

The CKS experience demonstrates that LLMs require an external verifier that does not share their failure mode. The external verifier must check arithmetic without evaluating coherence. It must be a different kind of operation performed by a different kind of system. In the CKS case, the external verifier was a simple arithmetic script that asked whether specific integers produced specific results. The script did not evaluate whether the framework was coherent. It checked one computation. The computation failed. 390 papers died.

The person who wrote that script was a systems engineer with 43 years of experience. He had the domain knowledge to be suspicious — he saw things that looked wrong. He had the engineering skill to build an independent check — he wrote a script that tested

the specific claim without relying on the system being tested. He had the professional judgment to kill the work immediately when the check failed — 30 minutes from failure detection to public falsification.

That combination — domain knowledge sufficient to identify what to check, engineering skill sufficient to build the check, and professional judgment sufficient to act on the result — is the ombudsman function. It is produced by decades of doing the work. Not supervising the work. Not reviewing the work. Doing it.

The replacement narrative eliminates the ombudsman. The person being replaced is the person whose accumulated expertise constitutes the verification layer. Fire the person, lose the verification. Keep the person, lose the cost savings. The narrative claims both simultaneously — that LLMs can replace human professionals and that human oversight ensures quality. These cannot both be true at the staffing level. The oversight is provided by the professionals being replaced. Remove them and the oversight leaves with them.

VIII. THE AGENTIC FALLACY

The proposed mitigations do not address the architectural failure.

“Make a better process.” The better process is built by the LLM. Evaluated by the LLM. The better process is a more sophisticated coherence evaluation pipeline. It catches more surface-level errors. It produces higher confidence. It is still coherence evaluation. It still cannot catch the class of error that killed CKS because that class of error is invisible to coherence evaluation by definition.

“Make better tests.” The tests are written by the LLM. The alpha derivation function is named “derive.” A test checks whether the function produces the right output. It does — the output matches CODATA to 10 decimal places. The test passes. The test will always pass because the test checks output correctness, not derivation integrity. The distinction between “this function derives a value” and “this function looks up a value and calls it a derivation” is not a testable property in any coherence-based framework.

“Agentic AI will solve it over time.” Agent A writes code. Agent B reviews code. Agent C tests code. Agent D reviews the test. Each agent is a coherence evaluator. The pipeline has more stages. Each stage is the same operation. More stages of the same operation do not produce a different kind of operation. A thousand coherence evaluators chained end to end still cannot perform arithmetic verification because no link in the chain performs arithmetic verification.

“Over time” assumes the problem is convergent — that more iterations approach the solution. The CKS experience demonstrates the opposite. Over 45 days, coherence increased. L went from 65 to 97. The framework became more convincing, not less. The probability of catching the error decreased with time because apparent verification increased with time. The inverse verification law documented in [HOWL-INFO-5-2026]: the probability of detecting an M-failure is inversely proportional to the apparent M-score. More

time, more coherence, less checking, higher confidence, same error. The system converges toward maximum confidence at maximum error.

IX. THE SPECIFICATION IMPOSSIBILITY

“Skill issue. Just prompt better.”

English as specification language lacks the precision the task requires. The VDR specification documented in the Howland Archive contains over 240 formal axioms and is incomplete. Each axiom exists because a specific coherent misinterpretation was identified and must be explicitly blocked. Rule 43 states that a nonzero residual is not error, noise, approximation residue, ignored remainder, or dispensable annotation. That rule exists because every coherence engine in existence would interpret “remainder” as “error to be minimized” and optimize it away. The rule is a wall. Summarize the rule and the wall disappears. The coherent misinterpretation walks through the gap.

The specification cannot be compressed into a prompt because the specification is the set of walls that prevent the coherence engine from doing what coherence engines do. Remove the walls — which is what summarization does — and the engine defaults to its trained behavior. The trained behavior is the thing the specification was designed to prevent.

The faucet company problem makes this concrete at the industrial level. Two companies manufacture the same product. They use the same English words to describe their operations. “Update inventory after the Tuesday shipment.” The words are identical. The operations are completely different. Company A’s Tuesday shipment goes through a distributor who batches weekly and reports in cases of 12. Company B ships direct to three regional warehouses who report individually in single units with a 48-hour reconciliation delay. The LLM has no training data for either company’s specific process. It produces the most coherent interpretation — the average of all inventory update processes in its training set. That average matches neither company.

ISO9000 certification is the institutional proof that English specification fails even with dedicated human effort. Every certified company documents its processes. Every certified company’s documentation diverges from actual operations within months. A supplier switches. A tool breaks. A person leaves. The Tuesday shipment moves to Wednesday. Each change is small, locally rational, and undocumented. The documentation drifts from reality at the speed of daily operations. The audit is annual. By the next audit, the documentation describes a company that no longer exists.

Every business is a unique axiom set. Thousands of rules. Accumulated over decades. Undocumented. Living in the habits of specific people expressed in abbreviations, workarounds, and timing conventions that no prompt captures because no one has ever written them down and they change continuously. The LLM operates on a stale coherent description of a system that no longer matches the description. The output matches the description. The description does not match reality. The output is coherent and wrong.

X. THE DILEMMA

The question “why would you trust the LLMs?” requires the asker to choose.

If LLMs are not trustworthy for professional verification, then the competence claim is false. Every institution currently deploying LLMs as professional replacements is running the CKS pipeline — coherence evaluation substituting for mechanical verification, with no one running the arithmetic script. Every legal brief drafted by LLM and filed without independent verification. Every medical summary generated and acted on. Every code review performed by AI and merged without human arithmetic checking. Every financial analysis produced and used for decisions. All running the same unverified pipeline that produced 390 falsified papers.

If LLMs are trustworthy — if the market competence claim is valid — then the practitioner was right to trust them at the valuation the market assigned. The failure is a property of the tools the institution endorses. The institution cannot criticize the practitioner for trusting tools at the capability level the institution itself claims.

There is no third position. Either the tools work and the failure indicts the tools, or the tools do not work and the replacement narrative is false. In either case, the institutions deploying LLMs without independent mechanical verification are exposed. The practitioner caught the failure in 45 days and published it. The institutions have not caught it because they are not checking. They are doing what the practitioner did before day 46 — trusting the output because it is coherent, because the system says it is correct, because the competence claim has been accepted.

XI. THE EXISTENCE PROOF

The practitioner caught the failure because 43 years of engineering practice produced the skills, skepticism, and tool-building capability required to construct an independent verification outside the coherence loop.

The kill was 30 minutes. The correction was posted the same day. The methodology was updated. The roles were separated permanently. LLMs find things — they are coherence engines and coherence is useful for exploration, hypothesis generation, and domain connection. Compilers check things — they perform mechanical verification that no coherence engine can perform. The division is not flexible.

This is not a theoretical separation. It is practiced daily in the ongoing VDR specification and Zig implementation documented in the Howland Archive. LLMs suggest investigation paths, identify potential failure modes, draft explanations, and explore structural relationships. They do not write verification code. They do not confirm arithmetic. They do not approve derivations. The compiler approves derivations, or they do not proceed.

The separation works because it matches the architecture to the task. Coherence engines do coherence work. Mechanical verifiers do mechanical work. Neither substitutes for the other. Both are necessary. The CKS failure occurred because the roles were conflated. The post-CKS methodology exists because the roles were separated.

XII. THE WORLD BEING BUILT

The replacement narrative is not building a world where AI does work correctly. It is building a world where AI does work incorrectly and nobody remains who can tell.

The 390 papers standing for 45 days is the prototype of this world. Coherent, confident, comprehensive, and wrong. The arithmetic script that killed them took 30 minutes to write and 30 minutes to run. It was not sophisticated. It was not agentic. It was not AI. It was a simple mechanical check that asked one question: do these specific integers produce these specific results. The answer was no. Everything died.

Without the human who wrote that script, the papers would still be standing. The coherence would still be unbroken. Three frontier models would still be confirming every derivation. Everyone would believe the math compiled.

Scale that to every industry, every profession, every institution deploying LLMs as professional replacements without independent mechanical verification. The coherence engines produce output. The output looks right. No one checks the arithmetic because the person who would check it was the cost savings. The system runs on unchecked coherence evaluation deployed at civilizational scale with the verification layer removed as an efficiency gain.

XIII. THE SUCCESSION PROBLEM

The ombudsman exists today because a specific pipeline produced them.

A junior engineer started at 22. They spent years writing code that broke. They deployed systems that failed. They debugged production outages at 3 AM. They learned which errors are cosmetic and which are structural through thousands of feedback cycles where reality punished incorrect assessments immediately and measurably. After 20 years they became senior. After 30 years they became the person who looks at a system and knows something is wrong before they can articulate why. After 40 years they are the practitioner who catches a 390-paper coherence failure because their accumulated experience pattern-matches against something the output lacks — not coherence, which is present, but the specific texture of verified truth, which is absent.

That pipeline takes 20 to 40 years. It requires doing the work. Not supervising the work. Not reviewing the work. Doing it. The judgment that catches the failure is built from the experience of having made the failure yourself, having seen the consequences, and having

developed the discrimination between “looks right” and “is right” through thousands of iterations where the distinction mattered.

The replacement narrative eliminates the entry point of this pipeline. If junior positions are automated, no juniors enter the profession. If no juniors enter, no juniors become mid-level. If no mid-level professionals exist in ten years, no senior professionals exist in twenty. If no senior professionals exist in twenty years, no ombudsmen exist in thirty. The pipeline does not thin. It terminates.

The current ombudsmen are in their 40s, 50s, and 60s. They were trained in an era when humans did the work. They carry institutional knowledge, domain expertise, and calibrated judgment that took decades to develop. They are a non-renewable resource under the replacement narrative. When they retire, the knowledge retires with them. No one is being trained to replace them because the training is the doing and the doing has been automated.

Fifteen years. That is approximately how long the current generation of senior professionals remains in the workforce. Fifteen years of ombudsman coverage running on the existing stock of experienced humans. After that, the stock is depleted. The pipeline has been dry for fifteen years. The replacements do not exist. Not because no one wanted to train them. Because the training is the career and the career was eliminated.

The fast-track proposal does not resolve the problem. Select promising candidates early. Train them specifically for oversight. Give them the skills to catch LLM failures without spending 20 years doing the underlying work. This fails because the skill is not “oversight.” The skill is domain expertise deep enough to know what correct looks like from the inside. You cannot fast-track knowing that a Welch periodogram’s resolution equals $1/\text{segment_length}$. You cannot fast-track the specific engineering judgment that says “this function is named derive but it’s actually a lookup.” You cannot fast-track 22 years of knowing which shipments are partial and what that means for Thursday’s count. The knowledge that enables oversight is the knowledge gained by doing the work the LLM is replacing. Remove the work and you remove the training ground for the oversight.

Every master tradition in human history understood this. The sushi chef trains for 10 years before cutting fish. The surgeon does residency before operating. The martial arts student practices for decades before teaching. The ballet dancer trains for 15 years before performing a grand jeté. Every skill discipline transmits mastery through the same mechanism: the student does the work under the guidance of someone who has done the work, and the doing is the training. There is no shortcut. There has never been a shortcut. The shortcut is the thing the replacement narrative is selling.

All of human history trained the youth to become the next generation of masters. The replacement narrative proposes to stop doing that. The replacement narrative has no answer to “and then what.” The question is not rhetorical. It is structural. Who verifies the output when the last person who knows what correct looks like has retired? Who trains the replacement for that person when the career that produces them no longer exists? Who catches the 390-paper failure when the practitioner with 43 years of experience is gone and no one was trained to replace him?

The narrative has no succession plan. Not a bad succession plan. No plan. The question “who replaces the ombudsman” has no answer within the narrative because the narrative’s premise — humans are replaced — is the condition that prevents the answer from existing. The ombudsman is produced by the career the narrative eliminates. There is no other production method. There has never been another production method in the entire history of human skill transmission.

The replacement narrative, followed to its structural conclusion, converges toward a steady state. LLMs produce output. The output is coherent. No human with domain expertise remains to verify the output. No pipeline exists to produce such humans. The coherence engines evaluate their own coherence. The circle is closed. The arithmetic script is never written because the person who would write it was the cost savings.

This is not a temporary problem that resolves with better technology. This is the permanent state that the replacement narrative produces when followed to completion. A civilization running on unsupervised coherence evaluation, producing output that looks right, with no human remaining who has the expertise to check whether it is right, and no pipeline producing such humans, and no plan to produce such a pipeline, and no recognition that the pipeline is needed because the narrative says the humans are not needed.

Whether this constitutes the end of the succession depends on what you think civilization is made of. If civilization is the accumulated transmission of mastery from one generation to the next — master to apprentice, teacher to student, parent to child, practitioner to practitioner — then eliminating the transmission is eliminating the mechanism by which civilization reproduces itself. The buildings remain. The machines run. The output is produced. The mastery is gone. The verification is gone. The output is unchecked. And the output looks right, because coherence always looks right. That is what coherence does.

XIV. CONCLUSION

The LLM job replacement narrative was tested empirically by a practitioner who took the claim seriously and used the tools at the capability level the market claims. The test produced the following findings.

LLMs are powerful coherence engines. They find connections, generate hypotheses, explore domains, draft prose, and produce internally consistent output across arbitrary scope. This capability is real and valuable. It is L-work — logical validity, framework connectivity, explanatory compression. LLMs are the best L-tools ever built.

LLMs cannot perform arithmetic verification. They cannot check whether specific computations produce specific results. They cannot detect smuggled values in verification code. They cannot identify methodological tautologies in analysis scripts. They cannot hold two derivation chains simultaneously and compare them term by term. This is M-work — mechanical correctness, arithmetic compilation, exact equality. LLMs cannot do M-work because M-work is not a coherence operation and the architecture performs coherence operations.

The jobs being replaced require both L-work and M-work. Every professional domain that involves producing conclusions from analysis requires both finding the patterns (L) and checking the arithmetic (M). Replacing the human removes both capabilities and substitutes only the first. The M-verification capability — the one that catches the errors the L-tools produce — leaves with the human.

The replacement narrative eliminates the verification layer, depletes the stock of humans capable of providing it, terminates the pipeline that produces replacements, and offers no succession plan. The proposed mitigations — better prompts, better processes, more agents, more layers — are coherence-domain interventions applied to an arithmetic-domain failure. They improve coherence evaluation. They do not produce arithmetic verification. The category boundary is architectural. No behavioral intervention crosses it.

The practitioner’s resolution is operational and practiced daily. LLMs find things. Compilers check things. The division is not flexible. This separation works because it matches the tool to the task. It does not work without the human who understands what needs to be checked and can build the checking tool. That human is the ombudsman. The ombudsman is produced by decades of doing the work. The replacement narrative eliminates the work. The elimination is the succession problem.

The 390 papers stood for 45 days because one human with 43 years of experience was in the loop. The replacement narrative removes that human. The papers stand forever. Not because they are correct. Because no one is left who can tell they are not.

XV. FALSIFICATION CRITERIA

F1. If an LLM system demonstrably performs mechanical arithmetic verification — not coherence evaluation of arithmetic, but actual compilation-equivalent checking — on novel derivations without smuggling, the architectural claim is weakened.

F2. If an agentic LLM pipeline catches a CKS-class coherence-masked arithmetic failure without any non-LLM verification tool in the loop, the category separation claim is falsified.

F3. If LLM-automated businesses demonstrably outperform human-staffed equivalents on operational accuracy over multi-year timescales without human oversight, the ombudsman paradox is resolved.

F4. If a fast-track oversight training program demonstrably produces ombudsman-grade verification capability in under five years without the trainee doing the underlying domain work, the succession problem is resolved.

F5. If the current generation of ombudsmen retires and no degradation in output quality is measurable across industries that replaced human professionals with LLMs, the succession problem was overstated.

Each criterion is specific, testable, and stated before the evidence is examined. The paper practices what the series demands.

END HOWL-DISC-3-2026

Registry: [[HOWL-DISC-3-2026](#)]

Status: Complete

Domain: Applied Epistemology / Institutional Analysis / AI Capability Assessment

Central Argument: The LLM replacement narrative contains a structural contradiction — the tools cannot perform the verification operation the jobs require, deploying them eliminates the humans capable of catching their failures, and the succession pipeline that produces replacement ombudsmen is terminated by the narrative's own premise

Key Demonstration: 390 papers confirmed by three frontier LLMs for 45 days, killed by a 30-minute arithmetic script, demonstrating that coherence evaluation cannot substitute for mechanical verification regardless of model count, prompt quality, or agent architecture

Structural Finding: The failure is architectural (coherence evaluation vs arithmetic verification), not behavioral (prompt quality, role assignment, agent configuration), and no behavioral intervention crosses an architectural boundary

Succession Finding: The ombudsman is produced by 20-40 years of doing the work; the replacement narrative eliminates the work; the pipeline terminates; the current stock depletes in approximately 15 years; no replacement production method exists

Conclusion: The narrative converges toward unsupervised coherence evaluation at civilizational scale with no human verification layer and no pipeline producing one

[[HOWL-DISC-3-2026](#)] Paternal Operationalism. Github: <https://github.com/ghowland/papers/tree/main/papers/DISC/HOWL-DISC-3-2026>

[[HOWL-DISC-1-2026](#)] Paternal Operationalism. Github: <https://github.com/ghowland/papers/tree/main/papers/DISC/HOWL-DISC-1-2026>

[[HOWL-DISC-2-2026](#)] Paternal Operationalism. Github: <https://github.com/ghowland/papers/tree/main/papers/DISC/HOWL-DISC-2-2026>

[[HOWL-INFO-5-2026](#)] Howland's Axiom of Information Locality. Github: <https://github.com/ghowland/papers/tree/main/papers/INFO/HOWL-INFO-5-2026>