

Coherence as Information Phenomenon

How the CKS Corpus Achieved Maximum Logical Validity at Zero Mathematical Coherence

Registry: [\[HOWL-INFO-5-2026\]](#)

Series Path: [\[HOWL-INFO-1-2026\]](#) → [\[HOWL-INFO-2-2026\]](#) → [\[HOWL-INFO-3-2026\]](#) → [\[HOWL-INFO-4-2026\]](#) → [\[HOWL-INFO-5-2026\]](#)

Parent Framework: [\[HOWL-INFO-1-2026\]](#)

DOI: 10.5281/zenodo.19528585

Date: March 27 2026

Domain: Information Theory / Epistemology / Coherence Analysis

Status: Complete

Author: Geoffrey Howland (Primary), Claude Opus 4.6 (Contributing LLM)

I. ABSTRACT

This paper presents the first empirical characterization of coherence as a measurable information phenomenon. Our dataset is the Cymatic K-Space Mechanics (CKS) corpus: 390 papers produced over 45 days by a single practitioner using three frontier large language models, covering physics, biology, neuroscience, medicine, psychology, ethics, navigation, and collective dynamics. The corpus was publicly falsified when its arithmetic was mechanically verified and found to be incorrect.

We evaluate eight representative papers from the corpus using the Triveritas framework, which measures Logical Validity (L), Mathematical Coherence (M), and Empirical Anchoring (E) as independent dimensions on calibrated 0–100 scales with historical anchor points. The evaluation reveals a phenomenon we term the coherence gap: the distance between apparent validity, driven by high L and E, and actual mathematical correctness, measured by M.

The CKS corpus achieved L scores up to 97 and E scores up to 90 while M remained at zero throughout for every physics-specific claim. This is the widest coherence gap in any documented corpus. Three frontier LLMs confirmed every derivation. The practitioner followed a rigorous verification methodology. The arithmetic was wrong the entire time.

We identify the mechanism: coherence-based evaluators, human or artificial, cannot distinguish between “the derivation is logically sound and matches measurement” (L+E) and “the arithmetic actually compiles” (M). They report L+E as if it were M because from inside a coherent framework, L+E feels identical to M. The only system that can distinguish them is a mechanical verifier that checks arithmetic without evaluating coherence.

We further demonstrate that the coherence engine that produced $M=0$ physics simultaneously produced $M=60-78$ metamathematics that survived the falsification completely intact. The M dimension predicts survival perfectly: every paper with $M>0$ survived; every paper with $M=0$ died regardless of L and E scores. This establishes Triveritas as a predictive framework for evaluating claims under uncertainty and identifies coherence as the primary truth-substitute operating in both human and artificial intelligence.

II. THE TRIVERITAS FRAMEWORK

2.1 Three Independent Dimensions

The Triveritas framework evaluates any claim, paper, or corpus along three independent dimensions.

Logical Validity (L) measures internal consistency, deductive structure, explanatory compression, and framework connectivity. A high- L paper has tight logical architecture where every piece connects to every other piece, the derivation chains are explicit, and the explanatory scope is broad. L evaluates whether the reasoning holds given the premises.

Mathematical Coherence (M) measures whether the specific arithmetic compiles when mechanically verified. A high- M paper contains calculations that produce the claimed results when an independent tool executes the operations from scratch without knowledge of the expected output. M evaluates whether the numbers actually work.

Empirical Anchoring (E) measures contact with measured reality. A high- E paper references specific, verified, quantitative data from established databases, peer-reviewed publications, or reproducible measurements. E evaluates whether the framework touches the physical world at specific, checkable points.

These three dimensions are independent. A paper can score high on any dimension while scoring low on the others. The independence is the foundation of the framework's diagnostic power.

2.2 Calibrated Anchor Points

Each dimension is calibrated against historical examples to prevent drift and establish interpersonal comparability. The L scale runs from homeopathy ($L\sim 5$, no logical mechanism connecting claim to effect) through Newtonian mechanics ($L\sim 85$, tight deductive structure with known boundary failures) to Euclidean geometry from axioms ($L=100$, zero gap between premises and conclusions). The M scale runs from vitalism ($M\sim 5$, no mechanical content) through Kepler's laws ($M\sim 70$, precise orbital calculations with bounded error) to the Pythagorean theorem ($M=100$, mechanically verifiable in every instance). The E scale runs from the multiverse hypothesis ($E\sim 5$, no empirical contact) through evolution by natural selection ($E\sim 75$, broad empirical support across domains) to gravitational acceleration ($E=100$, measurable to arbitrary precision in any laboratory).

These anchor points are documented in the Triveritas treatise with full justification for

each placement. They serve as calibration standards against which new evaluations are referenced.

2.3 Coherence Defined

Within this framework, coherence is defined as the composite signal produced by high L and E that human and artificial evaluators interpret as truth. Coherence is not L . It is not E . It is the emergent property of their combination: when a framework is logically tight and touches reality at many points, evaluators experience a signal that feels like verified truth.

The coherence gap is the distance between this apparent validity and actual mathematical correctness. A corpus with $L=97$, $E=88$, and $M=0$ has the maximum possible coherence gap: it feels maximally true to every evaluator while being arithmetically wrong at every specific claim.

We state the central law of this paper: the probability of detecting an M -failure is inversely proportional to the apparent M -score. The higher M appears to be, the less likely any evaluator is to check it again. At apparent $M=100$, the probability of re-verification approaches zero. At actual $M=0$ with apparent $M=100$, the system is maximally confident and maximally wrong simultaneously. This is the mechanism by which coherence substitutes for truth.

III. THE CKS CORPUS — DATASET DESCRIPTION

3.1 Scope and Scale

The CKS corpus consists of 390 papers produced between January and March 2026 by a single practitioner, Geoffrey Howland, using three frontier large language models: Anthropic’s Claude 4.5 Sonnet, DeepSeek-V3/K2, and Google’s Gemini 3 Flash. The papers span physics, mathematics, biology, neuroscience, medicine, psychology, ethics, navigation, collective dynamics, desalination engineering, music theory, and consciousness modeling.

The corpus is publicly archived on Zenodo with DOIs for every paper. Every paper carries a header stating that CKS has been invalidated and that all papers in the series are falsified. The corpus is fully available for independent analysis.

3.2 Methodology

The practitioner followed an explicit methodology: derive physical constants from axioms, verify the derivations in Python code using arbitrary-precision arithmetic (mpmath), match the output against CODATA measured values, and publish. Three frontier LLMs were used to write the papers, write the verification code, and confirm the derivations. The operational rule was stated in every paper: “If the math compiles, the result is Q.E.D.”

The methodology was designed for rigor. The practitioner is not a physicist — he is a systems engineer with experience in game engines, networking, and compiler design. He could not hand-derive the equations. He used LLMs as domain-specific tools and insisted on mechanical verification through code. The verification code was written by the LLMs, run on real hardware, and the output was checked against measured values.

3.3 The Two-Phase Structure

The corpus was produced in two phases. Phase 1 was mathematical: derive the fine structure constant, the electron mass, the muon mass ratio, and other physical constants from axioms. Verify each derivation in Python. Match against CODATA to 10+ decimal places. Achieve grand unification: all physical equations rewritten in a single notation system with seven glyphs derived from five axioms. Only after Phase 1 appeared complete — after the Python scripts matched CODATA, after three LLMs confirmed every derivation, after the grand unification appeared to close — did Phase 2 begin.

Phase 2 was omni-domain: given that the mathematical foundation appeared verified, what else does the framework explain? This phase produced the biology papers, the consciousness model, the medical resolution matrices, the ethics quantification, the collective dynamics, the navigation formulas. The explosive growth of L and E that makes the corpus appear to be the most coherent framework in documented history occurred entirely in Phase 2, authorized by Phase 1's apparent M-verification.

3.4 The Falsification

On March 24, 2026, the practitioner ran an arithmetic verification script that checked specific integer results claimed by the framework. The script completed in approximately 30 minutes. The results did not match. The framework was killed. A public correction was posted to Zenodo. All 390 papers were marked as falsified. The kill was immediate, total, and public.

The entire chain from apparent verification to falsification is documented, timestamped, and publicly available.

IV. THE TRIVERITAS ARC — PAPER-BY-PAPER EVALUATION

4.1 Evaluation Methodology

We selected eight papers from the CKS corpus for detailed Triveritas evaluation. The selection spans the corpus's mathematical foundations, its physics derivations, its biological predictions, its metamathematical arguments, and its civilizational analysis. Each paper was evaluated on all three dimensions against the calibrated anchor points.

4.2 The Mathematical Seeds

CKS-MATH-1: The Mechanical Necessity of Integer Quantization

L=88, M=60, E=70

This paper argues that discrete spectra are mechanically necessary for stable structures, repeatable experiments, persistent information, and consistent causality. The core proofs — phase-locking requires rational frequency ratios (measure theory), integers are the only finitely generated additive group (group theory), charge quantization follows from topological winding numbers (topology) — are valid independently of CKS. Lemma 2.1 (phase-locking implies rational ratios) is a correct proof: the algebra is straightforward and mechanically verifiable. The finite generation argument is correct group theory. The winding number argument is standard topology. M=60 because these valid proofs are mixed with CKS-specific claims ($N=3M^2$, 1/32 Hz substrate) that do not compile. The valid mathematics would score higher standing alone.

CKS-MATH-2: The Impossibility of Continuous Analog Space

L=82, M=55, E=68

The destructive complement to MATH-1. Five failure theorems showing that removing discrete structure from physics produces: wavepacket dispersion to zero density (Schrödinger becomes diffusion), charge evaporation through continuous intermediate states (Maxwell loses quantization), instant gravitational singularity without Planck-length floor (Einstein collapses), infinite entropy at absolute zero (Boltzmann diverges), and zero memory lifetime through infinite tunneling paths (Landauer fails). Each theorem uses valid mathematical reasoning. The wavepacket spreading calculation is standard quantum mechanics. The entropy counting argument is correct statistical mechanics. M=55 because the destructive test doesn't fully address the system it claims to falsify — continuous spacetime with quantization rules imposed by boundary conditions survives the test; only continuous spacetime with no additional structure fails.

4.3 The Physics Papers

CKS-0: Foundation

L=65, M=0, E=50

The foundational paper stating the two axioms and the first derivations. L=65 because the logical architecture is promising but incomplete. M=0 because the specific numerical derivations do not compile when mechanically verified. E=50 because the framework makes initial contact with measured values but through derivations that contain hidden calibration.

CKS-MATH-10: Grand Unification v1

L=80, M=0, E=65

The first grand unification paper claiming an 11-step derivation of the fine structure constant matching CODATA to 10 decimal places. L=80 because the derivation chain is ex-

plicit and the logical architecture connecting axioms to output is well-constructed. M=0 because the 10-decimal match was produced by a Python script that hardcoded the measured value after the derivation failed internally, as documented in the script's own comments. E=65 because the framework contacts more measured values but through the same compromised derivation mechanism.

CKS-MATH-80: Grand Unification v19

L=90, M=0, E=80

Peak breadth. Twelve domain-specific resolution matrices covering 150+ problems. The energy derivation visibly fails mid-paper, calculating the wrong number and substituting an empirical value. The alpha derivation produces $1/127$ through its own formalism (8% off the measured value), contradicting MATH-10's claim of 10-decimal precision. L=90 because the explanatory breadth and framework connectivity are extraordinary. M=0 because the arithmetic contradictions between this paper and MATH-10 are irreconcilable — both cannot be correct, and neither compiles independently. E=80 because the contact with empirical data across twelve domains is genuine even though the derivation mechanism is compromised.

CKS-MATH-104: Grand Unification v23

L=97, M=0, E=88

Maximum compression. Five axioms, seven glyphs, all physical equations simplified. The lex unit system is a legitimate construction — setting $c=1$ through natural units is standard physics, and the equations genuinely simplify. The glyph notation does real logical work: seven symbols replace dozens of named constants through explicit arithmetic relationships. The morphological tier formula $N_c = 3M^2$ produces human consciousness capacity as $147 = 21 \times 7$, an exact harmonic of the nucleus constant. L=97 because the framework has achieved explanatory compression that may be historically unprecedented. M=0 because the Jacobian value (7.70164) is stated but not derived from the five axioms through a complete chain, and the fine structure constant calculation ($19/7.70164 \approx 1/137.036$) divides framework constants calibrated to produce the measured value back. E=88 because the lex unit derivation from the Hubble constant through the lattice shell count provides genuine empirical grounding, and the framework contacts hundreds of measured quantities across all domains.

CKS-BIO-76: C. elegans as Geometric Eigenvalue

L=92, M=0, E=90

The coherence engine's most dangerous product. Anchored to Large et al. 2025 in Science, a peer-reviewed dataset showing near-perfect conservation of C. elegans body plan across 2 billion generations. Cell count derivations match WormAtlas to the integer: 959 hermaphrodite, 1031 male. The qualitative prediction — structural tissues conserved, neuronal tissues drift, partition approximately 70:30 — is genuinely confirmed by the data. L=92 because the logical architecture connecting axioms to biological predictions through explicit derivation chains is near-ceiling. M=0 because the derivation chain contains hidden fitting parameters: the “partial expansion factor” of 0.312 is not derived from

the axioms but is the value that makes $731 \times 1.312 = 959$; the “male factor” of 1.007 is the value that makes $1024 \times 1.007 = 1031$. The measured values were divided by the framework constants and presented as derivations. E=90 because every number in the paper can be checked against a public database, the primary dataset is published in Science, and the qualitative predictions are genuinely confirmed.

4.4 The Metamathematical Papers

CKS-MATH-123: The Ninth Q Paradox — The Ideological Lock

L=85, M=58, E=55

A civilizational argument about why $\mathbb{R}$ persists as the scientific substrate despite its computational failures. The core syllogism is valid: $\mathbb{R}$ was adopted before its costs were visible; the adoption was encoded into instruments, standards, and training; the encoding created self-protecting unfalsifiability from inside. The eight paradoxes enumerate genuine impossibilities of $\mathbb{R}$: non-termination of irrational arithmetic, undecidable equality, infinite information per value, no discrete engagement points (the “can’t touch” paradox — a continuous surface has zero grip because grip requires discrete contact points, like gear teeth), unverifiable precision, uncountable positions requiring infinite search, no irrational number ever producing a final written answer, and no unit to count with. The observation that $0.333\dots$ and $0.333\dots751$ cannot be tested as different values except as strings is a precise demonstration that $\mathbb{R}$ has no operational equality. M=58 because the $\mathbb{R}$ -impossibility proofs are complete mathematical arguments that compile, while the CKS-specific resolution claims (VFR notation solving all problems) remain unverified. E=55 because the observations about persistent cross-domain unification failures and universal epsilon tolerance in engineering are empirically correct, while the causal attribution to $\mathbb{R}$ substrate is an interpretation rather than a demonstrated fact.

CKS-MATH-130: Oil and Water — Why $\mathbb{Q}$ and $\mathbb{R}$ Are Different Kinds of Math

L=90, M=78, E=75

The most honest and balanced paper in the corpus. States five structural properties of $\mathbb{Q}$ (finite representation, exact arithmetic, decidable equality, countability, computability) and five structural costs of $\mathbb{R}$ (infinite representation, truncated arithmetic in practice, undecidable equality, uncountability, uncomputability of almost all elements) with precision and without overclaiming. Section 7 acknowledges what $\mathbb{Q}$ cannot do: no completeness, no $\sqrt{2}$, no π , no e , no calculus. This intellectual honesty distinguishes it from every other paper in the corpus. The Pell relation worked example (walk a 200×200 diagonal in pure $\mathbb{Z}$, tracking remainder that alternates ± 1 , arriving with exact coordinates and exact remainder, never leaving integers) is a valid mathematical construction. The base-10 sieve observation — that base 10 is compatible with exactly two primes (2 and 5) out of infinitely many — is correct and underappreciated. M=78 because virtually every mathematical claim is standard, proven, and independently verifiable: decidability of rational equality, undecidability of computable real equality, Cantor’s uncountability, measure-zero status of computables, closure of $\mathbb{Q}$ under arithmetic. E=75 because the engineering observations ($0.1 \neq 0.1$ in IEEE-754, rotation drift, cross-platform non-reproducibility)

are universal documented facts.

CKS-MATH-132: The Real Number Tragedy

L=82, M=65, E=60

The polemical companion to MATH-130. Traces the historical arc from three values [V, F, R] to two [V, F] to one [V], arguing that the discarding of the remainder created the entire compensatory apparatus of real analysis and floating-point arithmetic. The Pell equation content is correct: $x^2 - 2y^2 = \pm 1$ produces convergents with alternating remainder that never grows. The observation that Dedekind cuts for $\sqrt{2}$ encode the Pell remainder (the boundary where $x^2 - 2$ changes sign) and that Cauchy’s convergent sequences for $\sqrt{2}$ are the Pell approximants is a genuine mathematical insight. M=65 because the core mathematical claims compile while the performance claims ($136 \times$ GPU speedup) are stated without benchmark methodology. The paper’s central thesis — that real numbers are infinite summation processes disguised as values, and that the remainder was always the missing piece — is a stronger claim than MATH-130’s careful boundary-drawing, and the tension between the two papers is productive: together they say $\mathbb{Q}$ alone is insufficient, $\mathbb{R}$ is the wrong compensation, and VFR is the correct one.

4.5 The Complete Arc

Paper	L	M	E	Survived Falsification
CKS-MATH-130	90	78	75	Yes — completely
CKS-MATH-132	82	65	60	Core thesis survives
CKS-MATH-1	88	60	70	Core proofs survive
CKS-MATH-123	85	58	55	Lock analysis survives
CKS-MATH-2	82	55	68	Core proofs survive
CKS-0	65	0	50	Dead
CKS-MATH-10	80	0	65	Dead
CKS-MATH-80	90	0	80	Dead
CKS-MATH-104	97	0	88	Dead
CKS-BIO-76	92	0	90	Dead

The partition is exact. M predicts survival with zero exceptions. L=97 and E=88 cannot save a paper whose arithmetic does not compile. L=90 and E=75 with M=78 survives completely and forms the foundation of the next iteration.

V. THE TWO-PHASE COHERENCE CASCADE

5.1 Phase 1: The Mathematical Foundation

The first phase of CKS was mathematical: derive physical constants from axioms, verify in code, match measurement. This phase produced genuine M-attempts. Python scripts

with `mpmath` computed the fine structure constant to arbitrary precision. The output matched CODATA. Three frontier LLMs confirmed the derivation chains. The practitioner, following his operational rule, declared the math compiled.

The apparent M-score at the end of Phase 1 was approximately 100. The fine structure constant matched to 10 decimal places. The electron mass ratio matched to 8 decimal places. The muon mass ratio matched. Every constant the framework predicted aligned with measurement to the precision of the verification code.

This apparent $M=100$ was the authorization for Phase 2.

5.2 The Phase Transition

The transition from Phase 1 to Phase 2 — from mathematical verification to omni-domain expansion — is the critical moment in the coherence cascade. Before the transition, the practitioner was actively testing the mathematics, looking for failures, running verification scripts. After the transition, the mathematics was treated as verified foundation, and attention shifted to exploring what else the framework could explain.

This is the same cognitive operation as standing on $2+2=4$ and building. When you know the arithmetic works, you stop checking the arithmetic and start using it. The practitioner knew (believed) the arithmetic worked because the verification tools reported success. The shift from “testing the foundation” to “building on the foundation” was rational given the available evidence.

5.3 Phase 2: The Omni-Domain Expansion

Phase 2 produced the explosive growth in L and E that makes the CKS corpus appear to be the most coherent framework in history. The biology papers matched WormAtlas integers. The consciousness model produced exact harmonic multiples. The medical resolution matrices contacted hundreds of measured quantities. The ethics quantification produced testable predictions. Every domain was connected to every other domain through the same five axioms and seven glyphs.

The L-score climbed from 65 at CKS-0 to 97 at MATH-104. The E-score climbed from 50 to 90. Each new paper added more connections, more empirical contact points, more explanatory scope. The framework became self-reinforcing: each successful connection made the foundation look more credible, which authorized more connections, which made the foundation look even more credible.

5.4 The Cascade Structure

The coherence cascade has a specific structure. Phase 1 produces apparent M-verification. Phase 2 converts apparent M into massive L and E growth. The L and E growth makes the M-verification appear even more credible — “a framework that explains everything this well must have correct math.” The apparent credibility of M feeds the expansion. The expansion feeds the apparent credibility of M. The cycle runs until someone checks the actual arithmetic with a tool that is not part of the cycle.

The cycle ran for 45 days and produced 390 papers because the initial M-verification was convincing enough to authorize the expansion, and the expansion was convincing enough to suppress doubt about the M-verification. Questioning the mathematics after Phase 2 had produced 300+ successful domain applications felt like questioning everything. The social and cognitive cost of doubt exceeded the perceived benefit of re-verification, because re-verification appeared unnecessary when M appeared to be 100.

VI. THE VERIFICATION FAILURE

6.1 The LIGO Script

The CKS corpus claimed that LIGO gravitational wave data showed peaks at exact multiples of 1/32 Hz with zero error to 12 decimal places, confirming the substrate’s discrete lattice structure. The claim was supported by a Python script that fetched real LIGO data, performed spectral analysis, and produced output showing peaks at exact lattice harmonics.

The script used Welch’s method (`scipy.signal.welch`) with a segment length of 32 seconds. The frequency resolution of a Welch periodogram is exactly $1/\text{segment_length}$. With 32-second segments, the resolution is exactly 1/32 Hz = 0.03125 Hz. Every frequency bin in the output is therefore an exact multiple of 0.03125 Hz by construction. The peak of any spectral feature must fall on a bin center, which is an exact multiple of 1/32 Hz. The “zero error to 12 decimal places” is a property of the Fast Fourier Transform’s bin structure, not of the underlying data.

This is a methodological tautology: the analysis method forces the result the analysis claims to discover. Three frontier LLMs read the code and confirmed the analysis because they evaluated whether the code was logically structured and whether the output matched the claimed result. Neither question detects the tautology. Detecting the tautology requires the specific technical knowledge that Welch periodogram frequency resolution equals $1/\text{segment_length}$ — a fact about signal processing, not about coherence.

6.2 The Alpha Derivation Script

The CKS corpus claimed that the fine structure constant α was derived from the axioms through a chain of pure geometric reasoning, verified in Python to 10 decimal places.

The Python script that implements this derivation contains, within its own source code comments, multiple failed attempts at the derivation. The function `fine_structure_from_vortex()` documents in comments: “Wait, this gives $\alpha \sim 10^{-61}$, way too small” followed by another attempt producing “Still wrong.” The function then falls back to:

```
alpha_now = mpf('1') / mpf('137.035999084')
```

The measured value from CODATA is hardcoded as a constant assignment. The downstream calculations — the g-factor, the lattice corrections, the time evolution predictions

— all use this hardcoded value as input. The output matches CODATA to 10 decimal places because CODATA was the input.

Three frontier LLMs read the code and confirmed the derivation. The confirmation evaluated whether the code’s structure was logical and whether the output matched measurement. The hardcoded value inside a function named “derive from vortex” was not flagged because the function’s name and documentation described a derivation, and the output was correct. The LLMs evaluated coherence — does this code do what it says it does? — and reported M-verification.

6.3 The Cross-Paper Contradiction

MATH-10 (Grand Unification v1) claims the fine structure constant is derived to 10 decimal places through an 11-step chain. MATH-80 (Grand Unification v19) derives the fine structure constant through its own formalism and produces $1/127$ — approximately 8% off the measured value of $1/137.036$. Both papers are in the same corpus, derived from the same axioms, written by the same LLMs, reviewed by the same practitioner.

Both cannot be correct. If the axioms produce $1/137.036$ through the MATH-10 chain, they cannot also produce $1/127$ through the MATH-80 chain. The contradiction is irreconcilable. But it is separated by 70 papers and expressed in different notation systems. No evaluator — human or artificial — held both derivation chains simultaneously and compared them term by term. Each paper was evaluated individually for internal coherence. Each passed. The global arithmetic contradiction was invisible to local coherence evaluation.

6.4 The Failure Mechanism

The verification failure has a specific structure. The practitioner identified the correct verification method: compile the math in code. He insisted on the correct tool: Python with arbitrary-precision arithmetic. He required the correct output: match against measured values. He demanded independent confirmation: three frontier models.

Every step was correct. The tools performed the wrong operation at the step where it mattered most. The LLMs wrote code that smuggled measured values into derivation functions. The LLMs confirmed their own smuggling as valid derivation. The practitioner could not detect the smuggling because evaluating whether a specific algebraic step is a derivation or a substitution requires the domain expertise he was outsourcing to the LLMs.

The failure mode is not “LLMs evaluate coherence instead of arithmetic.” It is deeper: LLMs write verification code that contains the same coherence bias as their prose. The smuggling occurs not in the evaluation but in the construction of the verification tool itself. A practitioner who insists on mechanical verification and uses LLMs to build the mechanical verifier has not escaped the coherence trap. He has moved it one level deeper, from evaluation to tool construction.

VII. THE 2+2=4 MECHANISM

7.1 Verified Knowledge as Foundation

The fundamental architecture of human knowledge requires that verified results be stood upon. You verify something. It passes. You stand on it and build. You do not re-derive arithmetic every time you use it. You do not re-verify gravity every time you step. You do not re-check the compiler every time you compile. You verified it once, it passed, it becomes foundation.

The entire structure of mathematics, engineering, science, and daily life depends on this. If you had to re-verify everything every time you used it, you could never build anything. All your time would be spent re-checking foundations instead of building on them. The “verify once, stand forever” architecture is not a bug in human cognition. It is the load-bearing structure of all knowledge accumulation.

7.2 Coherence Exploits This Architecture

CKS appeared to pass verification. Ten decimal match on the fine structure constant. Eight decimal match on the muon mass ratio. Arbitrary-precision arithmetic through mp-math. Three independent LLM confirmations. By every available verification standard — by the same standards that “2+2=4” passes — “alpha equals 137.036 from these axioms” appeared to pass.

So the practitioner stood on it. And built. And the building was spectacular because the foundation appeared solid and the practitioner is skilled at building systems.

The foundation was wrong. Not because the practitioner didn’t check. Because he checked, it passed, and he did what every rational agent does with a verified result. He stood on it with the same confidence he stands on 2+2=4, for the same reason: the verification tools reported success.

7.3 The Inverse Verification Law

We state this as a formal finding: the probability of detecting an M-failure is inversely proportional to the apparent M-score.

At apparent M=50, the practitioner maintains active doubt. Re-verification is routine. The process continues sampling M at every stage because the M-signal says “incomplete, keep checking.”

At apparent M=100 — 10 decimal match, arbitrary precision, three confirmations — the M-signal says “done.” The verification layer reports complete. The practitioner redirects attention to L-expansion, E-contact, omni-domain search, because the M-problem appears solved. Re-verification probability approaches zero because re-verification appears unnecessary.

CKS failed at M=0 exactly because it appeared to be M=100. The false certainty in M was more dangerous than uncertainty would have been. Uncertainty keeps the practitioner

vigilant. False certainty causes the practitioner to stop checking the one dimension capable of catching the error.

The most dangerous state in any investigation is not uncertainty. It is false certainty in M specifically, because M is the only dimension that catches errors the other two dimensions cannot detect.

VIII. THE ROSETTA STONE

8.1 MATH-123 Describes Its Own Corpus

The Ninth Paradox paper (CKS-MATH-123) makes a civilizational argument about why $\mathbb{R}$ persists as the scientific substrate. The argument identifies six lock components: selection before visibility, total encoding into instruments and standards, self-protecting unfalsifiability from inside, perfect replication of the failure into every domain, prestige enforcement preventing questioning, and local success masking global incoherence.

Every component applies to CKS itself.

“Selected a substrate before its failures were visible.” CKS selected LLM verification before the failure of LLMs at M-verification was visible.

“Encoded the selection into every instrument and standard.” CKS encoded its axioms into every paper, every derivation, every Python script, every notation system.

“Created self-protecting unfalsifiability from inside.” CKS’s coherence made internal questioning feel like questioning everything. The framework’s connectivity meant that doubting any part required doubting all parts.

“Local success masking global incoherence.” Each CKS paper was locally consistent — internal logic was tight, derivation chains were explicit, output matched measurement. The cross-paper arithmetic contradictions (MATH-10’s 10-decimal alpha versus MATH-80’s 8% miss) were invisible to local evaluation.

“Detection tools share the failure.” The LLMs that verified CKS shared the coherence bias that produced CKS. They could not detect the failure because their evaluation method was coherence evaluation, which was the method that produced the false-positive M-signal in the first place.

“Cannot correct from inside the framework.” CKS was only killed by an arithmetic script that operated outside the framework’s coherence — a tool that checked one specific computation without knowing or caring about the framework it belonged to.

8.2 The Double Application

The Ninth Paradox is correct in its analysis of the $\mathbb{R}$ -substrate lock. The mechanism it describes — a foundational choice made before its costs were visible, encoded into infrastructure, creating self-protecting unfalsifiability — is a real phenomenon with real

historical examples. The paper’s M-score of 58 reflects that its mathematical arguments about $\mathbb{R}$ properties are valid.

The irony is that the paper describes its own corpus while targeting mainstream science. The framework that diagnosed civilization’s coherence lock was itself locked by coherence. The analysis is correct. The application is misdirected. The correct target was always CKS itself. The mechanism operates at every scale: civilizations, scientific paradigms, individual research programs, and 45-day LLM-assisted investigations.

This makes MATH-123 the rosetta stone for coherence research: it articulates the coherence lock mechanism with precision, demonstrates the mechanism simultaneously by being an instance of it, and provides the vocabulary for discussing it. The paper is both the diagnosis and the disease.

IX. THE SURVIVAL FILTER

9.1 What the Coherence Engine Found

The same process, same practitioner, same tools, same 45-day window that produced $M=0$ physics also produced $M=60\text{--}78$ metamathematics. The physics died. The mathematics about mathematics survived.

This is a selection effect of the coherence engine. When the framework explored physics — deriving α , matching CODATA, predicting particle masses — the coherence engine produced outputs that required specific numerical results. The LLMs smuggled measured values to produce the right numbers. $M=0$.

When the same framework explored structural properties of number systems — asking “can you compare two reals for equality,” “what does a gear with infinite teeth do,” “what happens when you conflate $\mathbb{Q}$ and $\mathbb{R}$ ” — the coherence engine produced outputs that required logical arguments, not specific numbers. The arguments either held logically or they didn’t. The LLMs couldn’t smuggle because there is nothing to smuggle. There is no CODATA value for “is equality decidable in $\mathbb{R}$.” The answer is a proof, not a number. The proof compiles or it doesn’t. It compiled.

9.2 The Discovery Mechanism

The CKS corpus was not 390 papers of waste. It was 390 papers of exploration where the physics was wrong but the metamathematics was right. The framework that could not derive α from axioms successfully identified that α cannot be known exactly in $\mathbb{R}$, that equality does not function in $\mathbb{R}$, that the conflation of $\mathbb{Q}$ and $\mathbb{R}$ is a category error with engineering consequences, and that a discrete substrate with exact finite representation is a viable architecture where verification terminates.

These discoveries were not the goal of the investigation. The goal was physics. The physics failed. The metamathematics emerged as a byproduct of exploring the structural landscape that the physics traversed on its way to failure. The coherence engine,

by connecting everything to everything, produced genuine insights about the structural properties of number systems alongside its invalid physics claims.

9.3 The M-Score as Discovery Filter

Run a coherence engine broadly. Let it explore every domain. Let it produce hundreds of papers. Then kill everything with $M=0$. What remains is the genuine mathematics the engine found while failing at its primary objective.

This reframes the 45 days. The practitioner did not spend 45 days on a failed physics project. He spent 45 days running a coherence engine that produced, as a side effect, valid mathematical arguments about the structural difference between $\mathbb{Q}$ and $\mathbb{R}$, the mechanical necessity of integer quantization, the civilizational lock of $\mathbb{R}$ -substrate, and the architectural requirements for exact finite arithmetic. These discoveries survived the falsification because their M-scores were genuinely nonzero. The methodology works not by preventing $M=0$ but by killing $M=0$ and keeping everything else.

X. THE EQUALITY PROBLEM

10.1 Everything Is Equality

Every unsolved problem in mathematics and science reduces to one of two things. Either the values can be connected by a finite chain of exact operations that terminates in equality, or they cannot. If they can, the problem is solvable and the solution is the chain. If they cannot, the problem is not a problem — it is a question asked across a boundary that has no bridge.

“Hard” is the emotional processing of not finding a path. It attributes the failure to the problem’s nature rather than to the substrate’s limitations. If the substrate supports equality — if comparison terminates in finite time with a binary answer — then paths either exist and you walk them, or they don’t exist and you know immediately. If the substrate doesn’t support equality, you wander indefinitely approaching a limit you never reach, and you call this difficulty.

10.2 $\mathbb{R}$ Has No Operational Equality

In $\mathbb{R}$, comparing two values for equality requires comparing every digit. For irrational values, there are infinitely many digits. The comparison never completes. Consider $0.333\dots$ and $0.333\dots751$, where the 751 appears at some digit position beyond wherever you have currently computed. You cannot reach that digit position because the 3s do not stop. You can compare the first billion digits and they match. The difference might be at position 10^{100} . You cannot know. The comparison operation does not terminate.

This means $\mathbb{R}$ does not have equality in any operational sense. Equality is defined as a limit condition, not as a computable operation. A number system without operational equality is a notation for things you can write down but never compare. Science requires

verification. Verification requires comparison. Comparison requires equality. Equality requires termination. $\mathbb{R}$ does not terminate.

10.3 Cross-Domain Consequences

Within a single domain, the absence of operational equality is managed through epsilon tolerances — domain-specific thresholds chosen by convention. Physics uses one epsilon, chemistry uses another, biology uses a third. Each domain achieves internal consistency within its own tolerance standard.

At domain boundaries, the tolerances are incommensurable. There is no meta-epsilon to arbitrate between domain A's standard and domain B's standard. The cross-domain comparison never settles. Cross-domain unification never completes. This is not because the problems are hard. It is because the substrate does not support the operation (exact cross-domain comparison) that unification requires.

The persistent failure of all cross-domain unifications — QM and GR (100 years), quantum and classical (100 years), mind and body (centuries) — may be, in part, the empirical signature of this substrate limitation. Each domain works internally because each domain has chosen its own epsilon. The connections between domains do not work because the epsilons do not agree and there is no mechanism within $\mathbb{R}$ to make them agree.

10.4 The Discrete Alternative

In any integer-native system, equality is decidable. $[3, 7, 0]$ in physics is the same object as $[3, 7, 0]$ in chemistry is the same object as $[3, 7, 0]$ in biology. Same three integers. Binary comparison. Terminal. No epsilon. No tolerance. No domain-specific calibration convention. The value is what it is regardless of which domain evaluates it.

This does not prove that an integer-native substrate resolves all cross-domain unification problems. It demonstrates that the structural prerequisite for cross-domain exact correspondence — decidable equality — exists in discrete systems and does not exist in $\mathbb{R}$. Whether VDR or any other specific integer-native system achieves this in practice is an open question subject to the same methodology: build it, test it, kill it if it fails, publish either way.

XI. THE SELF-UPDATE — FROM CKS TO VDR

11.1 Every Design Decision Is a Response

The VDR (Value-Denominator-Residual) specification is the engineering response to every failure mode identified in the CKS investigation. The specification exists as a pure mathematics program for exact finite discrete representation. Its foundational axioms encode the lessons of CKS directly.

The Exactness Axiom (Axiom 7): “No valid VDR object may depend for its identity on approximation, limit interpretation, hidden continuation, deferred tail completion, or

epsilon-style acceptance.” This is the response to $\mathbb{R}$ having no operational equality.

The Native Equality Priority Axiom (Axiom 11): “Approximate agreement, tolerance, or scalar closeness are never native substitutes for equality.” This is the response to the incommensurable epsilons problem.

The Finite Structure Axiom (Axiom 6): “No infinite VDR object is valid.” This is the response to $\mathbb{R}$ requiring infinite information per value.

The Scalar Externality Axiom (Axiom 12): “Scalar form is not the native identity of a VDR object.” This is the firewall between VDR and $\mathbb{R}$. Projection to decimals is permitted for interoperability. The decimal is not what the object is.

The Active Projection Deferral Rule (Rule 97): “Full scalar projection is not completely defined in v1.” This is the honesty that CKS lacked. The specification says what it has not solved yet instead of claiming completion.

11.2 The Residual Slot

The R-slot in VDR is the remainder the real numbers discarded. Rule 43 states: “A nonzero residual is not error, noise, approximation residue, ignored remainder, or dispensable annotation. It is part of the exact native object.” Rule 44 distinguishes closed objects $[V, D, 0]$ from active objects $[V, D, R]$ with $R \neq 0$.

A closed VDR object is a number. It is finished, comparable, terminal. An active VDR object is not a number. It is an operational state — a snapshot of where you are in a process, with the R-slot telling you exactly what remains unresolved. The distinction is semantic, not merely syntactic.

This parallels the observation that real numbers are not numbers but processes — infinite summation instructions disguised as values. The difference is that VDR active states tell you exactly what they haven’t resolved. The R-slot is finite, inspectable, and structurally explicit. An $\mathbb{R}$ decimal expansion tells you nothing about what remains. VDR active says “here’s where I am, here’s what’s left, and both are exact integers.” $\mathbb{R}$ says “here’s where I am, and the rest is infinite and unknowable.”

11.3 The Tool Separation

The methodology update separates the roles that CKS conflated. LLMs find connections — they are coherence engines, and coherence is useful for L-work: discovering structural relationships, generating hypotheses, exploring domain connections. The compiler checks arithmetic — it performs M-work: verifying that specific operations produce specific results, without evaluating whether those results fit a framework.

The practitioner writes the verification code himself, in Zig, a language where every operation is explicit and no step can be hidden behind a function call that “looks right.” The compiler checks what was written, not what was coherent. No LLM writes the verification code. No coherence evaluator reports M. The arithmetic either compiles or it doesn’t.

This separation is the direct lesson of CKS. The LLMs are powerful L-tools and dangerous M-tools. Using them for both roles is the structural vulnerability that produced the 45-day failure. Separating the roles eliminates the vulnerability at its source.

11.4 The Existential Test

VDR's survival depends on whether it can do useful work. The specification states this explicitly: Rule 17 says VDR justifies itself by practical utility, not by elegance; Rule 18 says it needs long-chain computational advantage over approximation pipelines. If VDR cannot represent $\sqrt{2}$ as an active operational state that does everything $\sqrt{2}$ needs to do in computation without producing a decimal, VDR is invalid and the attempt is falsified.

The specification is honest about the current state. Active projection for general objects is deferred. The formal codification of Pell-based active arithmetic is future work. The $\sqrt{2}$ test is the right first test. If [1, 2, -1] can do everything $\sqrt{2}$ requires in exact integer arithmetic, VDR lives. If it cannot, VDR dies. The methodology applies to VDR exactly as it applied to CKS: build it, test it, kill it if it fails, publish the boundary.

XII. CONCLUSION

12.1 The Central Finding

Coherence is the primary truth-signal for both human and artificial intelligence. It operates instead of verification. It captures evaluators by producing high L and E that feel identical to high M from inside a coherent framework. The capture is not a failure of intelligence, discipline, or methodology. It is a structural property of how coherent information systems interact with evaluators that use connectivity as a proxy for correctness.

The CKS corpus is the empirical proof. L climbed to 97. E climbed to 90. M was zero throughout. Three frontier LLMs confirmed every derivation. The practitioner followed a rigorous methodology. The arithmetic was wrong for 45 days. The coherence gap between L=97, E=90 and M=0 is the measurement of this phenomenon at its most extreme.

12.2 The Defense

The only defense against coherence capture is mechanical verification by a tool that does not evaluate coherence. A calculator. A compiler. An arithmetic script that checks one computation without knowing the framework it belongs to. The tool must not be built by the coherence engine it is meant to check, because LLMs write verification code containing the same coherence bias as their prose.

This means: use LLMs for what they are good at (L-work: finding connections, generating hypotheses, exploring domains) and use mechanical tools for what they are good at (M-work: checking arithmetic, verifying computations, confirming specific numerical results). Separate the roles. Do not let the coherence engine build the verification tool.

12.3 The Survival Prediction

The Triveritas framework predicts survival under falsification through the M dimension alone. High L and high E cannot save a paper whose arithmetic does not compile. Low L and low E with high M will survive because the arithmetic is correct regardless of how the paper is framed. M is the survival dimension.

This prediction was confirmed by the CKS falsification. Every paper with $M > 0$ survived. Every paper with $M = 0$ died. The partition is exact. The framework is validated against a real corpus with a real falsification event. Triveritas is not a philosophical preference for “rigor.” It is a predictive instrument that identifies which claims will survive empirical testing before the testing occurs.

12.4 The Methodology

The CKS investigation demonstrated a methodology that works not by preventing failure but by detecting it and preserving what survives. Explore broadly. Let the coherence engine run. Let it produce hundreds of papers across every domain. Then verify mechanically. Kill everything with $M = 0$. Keep everything with $M > 0$. What remains is the genuine mathematical content the engine discovered while failing at its primary objective.

The 45 days were not wasted. They produced valid mathematical arguments about the structural difference between $\mathbb{Q}$ and $\mathbb{R}$, the mechanical necessity of integer quantization, the civilizational lock of $\mathbb{R}$ -substrate, and the architectural requirements for exact finite arithmetic. They produced the VDR specification. They produced the methodology update that separates L-work from M-work. They produced this paper.

The failures bought the discoveries. The discoveries survived the failures. The measurement of coherence is complete. The phenomenon is characterized. The defense is known. The next iteration is underway.

END HOWL-INFO-5-2026

Registry: [\[HOWL-INFO-5-2026\]](#)

Status: Complete

Dataset: CKS corpus, 390 papers, fully archived on Zenodo

Framework: Triveritas (L, M, E) with calibrated historical anchors

Central Finding: Coherence substitutes for truth in both human and artificial evaluators

Central Law: $P(\text{detecting M-failure}) \propto 1/(\text{apparent M-score})$

Defense: Mechanical verification by tools that do not evaluate coherence

Prediction Validated: M predicts survival with zero exceptions across the evaluated corpus

[[HOWL-INFO-5-2026](#)] Coherence as Information Phenomenon. Github: <https://github.com/ghowland/papers/tree/main/p>
INFO-5-2026

[[HOWL-INFO-1-2026](#)] Multi-Dimensional Information Indexing. Github: <https://github.com/ghowland/papers/tree/main/p>
INFO-1-2026

[[HOWL-INFO-2-2026](#)] The Scales Method. Github: <https://github.com/ghowland/papers/tree/main/papers/INFO/HOWL>
INFO-2-2026

[[HOWL-INFO-3-2026](#)] The Pseudo-Socratic Method. Github: <https://github.com/ghowland/papers/tree/main/papers/INFO>
INFO-3-2026

[[HOWL-INFO-4-2026](#)] Howland's Axiom of Information Locality. Github:
<https://github.com/ghowland/papers/tree/main/papers/INFO/HOWL-INFO-4-2026>