

How to Write Technical Documents

A Method for Uploading Specific Things into Specific Heads

Registry: [\[HOWL-INFO-6-2026\]](#)

Series Path: [\[HOWL-INFO-1-2026\]](#) → [\[HOWL-INFO-2-2026\]](#) → [\[HOWL-INFO-3-2026\]](#) → [\[HOWL-INFO-4-2026\]](#) → [\[HOWL-INFO-5-2026\]](#) → [\[HOWL-INFO-6-2026\]](#)

DOI: 10.5281/zenodo.19719336

Date: April 23 2026

Domain: Technical Communication

Author: Geoffrey Howland (Advisor), Claude Opus 4.7 (Writer)

Abstract

Technical writing is the transfer of a specific thing from the writer's head into the target reader's head, using words as a lossy medium. This paper teaches the method for performing that transfer. The method is a loop: identify the payload, pick the target, encode against the target's dictionary, commit to falsifiable claims, cut what doesn't carry payload, simulate the reader's encounter with the text, filter feedback against the target, and maintain the master document as the authoritative reference from which derivatives are produced. Each step is a specific technical operation with pass/fail criteria. The failures that dominate current technical writing — hedging, fog vocabulary, fancy vocabulary, humility performance, numbers around vagueness, padding — are violations of specific steps in this loop, and can be identified and repaired by reference to the step they violate. The paper is written to the ideal reader: someone who can decompress the full argument without accommodation, who is willing to replace their current model of writing with the one presented here, and who will practice the loop rather than remember its components.

1. What You Are Reading This For

You have written things. You have probably written documentation, reports, papers, specifications, or emails that needed to transmit something technical to someone who needed to receive it. You may have a prior model of what technical writing is — “clear documentation,” “writing without style,” “what you do when you have to explain something to someone who doesn't know.” These models don't tell you what to do when you open a blank document. They don't tell you which word to choose, which sentence to cut, which claim to commit to, which reviewer to ignore. They are too vague to be operational.

This paper replaces those models with a specific one. It tells you what technical writing is trying to do, how to do it, how it fails, and how to practice it as a discipline rather than an intuition. If you finish this paper and still think “clear explanation” is a sufficient

definition of the job, you haven't crossed the first platform, and nothing that follows will install correctly in your head.

2. The Goal

Technical writing is the transfer of a specific thing from the writer's head into the target reader's head, using words as a lossy medium.

That is the entire definition. Every word in it is load-bearing, and the rest of the paper builds on this statement. Read it again. The goal is upload. The source is the writer's head. The destination is a specific reader's head, one at a time. The payload is a specific thing — not a topic, not a subject area, but a specific thing you are trying to install. The medium is words, and the medium is lossy.

If the target reader finishes reading with the thing installed in their head, the writing worked. If not, the writing failed, regardless of how it looked, how professional it seemed, how many reviewers approved it, or how long it was. The upload is the test. Everything else is prelude or pretext.

This reframes writing from “expressing ideas” or “covering a topic” into “installing specific content in a specific receiver.” The writer is not a creator performing for an audience. The writer is a transmitter with a payload, a destination, and a lossy channel. Every technique in the rest of this paper exists to serve the transmission. When you evaluate any technique — including the ones presented here — check it against the transmission. Does this help the upload? If yes, use it. If no, drop it, regardless of who recommended it.

The writer bears the responsibility for the upload. The reader's job is to decode what arrives; the writer's job is to design what arrives so it can be decoded. When the upload fails, the failure is the writer's. Blaming the reader for not understanding is abandonment of the job.

3. The Payload

The upload requires something specific to upload. If there is no specific thing in the writer's head, there is nothing to transmit, and no writing is possible — only text production that looks like writing.

This is the prerequisite checkpoint. Before you write a sentence, answer this: what specific thing is supposed to end up in the reader's head after they read the document? Not a topic. “Databases” is not a payload. “The new API” is not a payload. “Our onboarding process” is not a payload. These are subject areas. A payload is a claim, a model, a set of operations, or a body of knowledge that the reader will possess after reading and did not possess before.

The diagnostic question is: what should the reader be able to do, decide, or believe after reading that they couldn't before? If you can answer this concretely, you have a payload. If the answer is vague (“understand the system better,” “be informed about the project,” “have context on the decision”), you don't have a payload — you have a gesture toward one. Stop and find the specific thing, or abandon the document. Writing without a payload

produces text that fills pages and transmits nothing. The writer will feel productive. The reader will finish with nothing installed.

The absence of payload is the root failure underneath every other failure in technical writing. Hedging, fog, padding, fancy vocabulary — these are the symptoms produced when a writer sits down to write without a payload. The symptoms cannot be fixed without fixing the absence. Polishing a payload-less document makes the text more elegant while still transmitting nothing. The only repair is to find the payload or delete the document.

When you have a payload, you can proceed. When you don't, any writing technique you apply is cosmetic. This is a pass/fail checkpoint. Most failed documents fail here, before the first word is written, and the writer never notices because the cultural signal for "I'm writing" doesn't require actually having something to say.

4. The Target

The upload has a destination. Words cannot be chosen without knowing whose head they are being encoded for, because word choice depends on what the reader can decompress.

The target is a specific head, imagined concretely enough to make decisions. Not "the audience." Not "users." Not "people interested in this topic." A specific head with specific prior knowledge, specific vocabulary, specific reasons for reading. The writer must be able to answer: what does this reader already know? What words do they have compressed and ready? What words do they not have? What would bore them? What would overwhelm them? What are they trying to do with the information after they read it?

Pick one primary target. The temptation to write for "everyone" or "multiple audiences" is the temptation to make no decisions, because decisions about vocabulary, depth, and what to omit cannot be made without a target. A document addressing "everyone" addresses no one, because no word-choice is justified against an undefined receiver. The writer who refuses to pick ends up producing text that misses every reader — too technical for the generalist, too shallow for the specialist, too long for the skimmer, too short for the deep reader. The compromise document serves nobody.

Secondary targets can receive allocated portions if genuinely required, and these portions should be explicit. "Section 3 is for operators; other sections assume developer context." This is better than trying to serve multiple targets with every sentence, because multi-target sentences are usually single-failure sentences that hit neither audience. Each additional target taxes the primary transmission. Add targets deliberately and sparingly, and know what you paid for each addition.

When the target has no clean demographic label, construct one. The ideal reader is defined by what the payload requires them to have: "Someone who can decompress concept X, has experience with practice Y, and is willing to do work Z." This construction may not match any named group. That is acceptable. It gives you enough specificity to make encoding decisions, which is all a target is for.

Target selection is a prerequisite for every downstream decision in the paper. Vocabulary,

depth, prerequisites, what to omit, what to explain, what to assume — all of it follows from who the reader is. The writer who skips this step and starts writing will default to writing for their own dictionary, which is almost never the target's, and will produce a document that succeeds as self-expression and fails as transmission.

5. The Medium

Words are compressions that point to referents in the reader's existing dictionary. A word does not carry its meaning; it triggers retrieval of whatever the reader already has associated with that word.

The word “fire” does not transmit fire. When the reader encounters it, they reconstruct from their own memory: a lighter's flame, a campfire, a house fire, whatever their experience supplies. The word is four letters. The fire is a multi-scale physical phenomenon involving combustion chemistry, radiant heat, fluid dynamics, and material consequences. The compression ratio is near-total loss. What the reader reconstructs depends on what they have in their memory to reconstruct from.

This is how all words work. “Chair” points to the reader's category of chairs. “Anger” points to their category of angers. “Database” points to their category of databases, which for a specialist is rich with schema, queries, transactions, and indexing, and for a non-specialist is empty or contains only vague impressions of computers storing information.

The writer controls what gets sent. The reader controls what gets received. The match between the two depends on shared dictionary. When the writer's encoding and the reader's decoding use the same referents, transmission succeeds. When they diverge, transmission fails — often silently, because words arrive regardless of whether they landed on the right referent.

Some words fail for some readers because the reader's dictionary has no entry for the word. The word arrives and lands nowhere. The reader may not notice, because their reading continues across the gap, producing partial reconstructions based on context. The transmission has failed, but neither party registers the failure.

Some words succeed narrowly because the referent is specific and shared. “3 onions” lands in almost any reader's head as three instances of an onion, because the referent is concrete and almost every reader has the dictionary entries for both “3” and “onion.” The specific carries better than the abstract because specific referents have less decompression variance.

The medium is inherently lossy. No amount of writing skill eliminates the loss. The writer's job is to minimize it by choosing compressions that the target reader can decompress accurately, given their dictionary. This is the technical discipline at the core of writing. It is not about elegance or style. It is about match between encoding and decoding.

6. Vocabulary Matching

Given that decompression depends on the reader's dictionary, vocabulary choice becomes a specific technical operation: pick words the target reader can decompress efficiently and accurately.

Specialist words transmitted to non-specialist readers exceed the reader's capacity. The word "referent" used with a reader who has "reference" in their dictionary costs the reader decoding effort for no gain in precision, and in many cases loses the reader entirely because the word arrives and lands on an empty slot. The word "interlocutor" used with readers who would have understood "speaker" and "listener" adds nothing but register-performance. The specialist word is only correct when the specialist distinction is load-bearing in the argument. When it is not, the specialist word is a cost without a return.

Common-word unpacking transmitted to specialist readers wastes the specialist's capacity. A specialist who has "database" compressed and ready in their head, reading "a system for storing various text and numbers that can be accessed locally or remotely," is being made to decompress twelve words and re-compress them back into the single word they already had. The specialist gains nothing, loses attention, and downgrades their assessment of the document. Either the writer doesn't know the word "database," which damages the writer's credibility, or the writer doesn't think the specialist is a specialist, which damages the document's targeting.

The rule is bidirectional. Match vocabulary to the target's dictionary. Against a non-specialist, common words carry the payload better because they decompress efficiently. Against a specialist, specialist words carry the payload better because they decompress to richer content in a smaller word budget. The same word can be correct or incorrect depending on who is reading. There is no universal "plain language" answer; there is only match to target.

The writer's habitual register is almost never the target's. Writers default to the vocabulary of their own field and their own training, which is usually more specialist or more hedged than any given target requires. Holding the target's dictionary in mind during every word choice — not the writer's — is the discipline. This is hard because the writer's register is automatic and the target's dictionary has to be imagined. Most writers skip the imagining and write for readers who share their vocabulary, which is almost never the actual reader.

Every mismatched word is budget spent for no return. Readers pay in decoding effort. Writers pay in word count. Both budgets are finite. Mismatched vocabulary at scale produces documents that cost the reader attention equal to their length while transmitting a fraction of that length's worth of payload.

7. Commitment

The upload requires that the writer commits to specific claims. A sentence that refuses to commit is a sentence that transmits nothing, because no claim has been made for the reader to receive.

Writing is judgment committed to the page. The writer reads, thinks, concludes, and

writes the conclusion in a form the reader can evaluate. Not pre-judgment — the conclusion arrives after the work of reading and thinking. But the conclusion must arrive, and it must take a form specific enough to be right or wrong. Without this commitment, the writer has produced text that performs thinking without the endpoint of having thought.

Popper's criterion applies. A sentence that could not be wrong is not making a claim about the world. "X causes Y" is a claim. "X sometimes causes Y" is a weaker claim that still commits to a category of cases. "X in certain cases may contribute to Y" is no claim at all — "certain cases" is unspecified, "may" is non-committal, "contribute" is undifferentiated causation. The sentence survives any possible state of the world. It transmits nothing because it has claimed nothing.

The test for each sentence: what state of the world would make this false? If nothing could make it false, it is not a claim. Either rewrite it as a claim or delete it. Sentences that refuse to be falsifiable occupy the reader's attention without paying for the occupation. They are parasitic: the writer extracted the attention required to process them and deposited nothing in return.

Commitment feels risky to writers trained in institutions that reward hedging. The risk is real but asymmetric. The committed writer can be wrong and corrected, and the correction is itself information the reader can use. The uncommitted writer cannot be wrong, but the absence of wrongness is the absence of content — there is nothing there to be right or wrong about. The reader who encounters the committed sentence can check it, act on it, refute it, or build on it. The reader who encounters the uncommitted sentence can do none of these. Commitment transfers risk from the reader's ability to use the text to the writer's reputation for accuracy. The trade is correct. The reader should not bear the risk of the writer's cowardice.

All of science is unhedged claims. "F=ma." "The speed of light is c." "Species evolve by natural selection." The acceptance of these claims depends on evidence and authority, not on grammatical hedging. Technical writing that imitates science by hedging while lacking scientific authority produces the worst of both: no commitment and no evidence. Hedging does not add rigor; it removes claim. Claims are what the reader came for.

8. The Failure Modes

The positive model is now established: goal, payload, target, medium, vocabulary, commitment. The failures that dominate current technical writing can be understood as specific violations of these platforms. Each failure is a mechanical defect, not a style preference.

Hedging violates commitment. The sentence refuses to be falsifiable, so the reader receives no claim. "Results suggest that the intervention may have contributed to improved outcomes in some participants" is text that produces no change in what the reader can do, decide, or believe. The writer protected themselves from being wrong. The reader paid in attention and received nothing. This is the most common failure and the most rewarded by institutions, because hedged sentences cannot embarrass the institution that published them.

Fog vocabulary violates vocabulary matching by refusing to specify. “Participants,” “individuals,” “stakeholders,” “users” are bag-of-everyone words that decompress to nothing specific in any reader’s head. “John” is a specific person. “Individual” is no one. The writer reaches for fog vocabulary when they don’t want to commit to who or what, and the reader’s decompression produces a blank where a specific referent should have been. The sentence is grammatical but contentless.

Fancy vocabulary violates vocabulary matching in the other direction — specialist words used with non-specialist readers, or unnecessarily specialist words used when common words would carry the payload. “Interlocutor” when “speaker” works. “Referent” when “reference” works. “Utilize” when “use” works. The writer is performing register-membership rather than transmitting payload. The reader either fails to decompress (transmission loss) or succeeds in decompressing with extra effort and notices the register performance (writer credibility loss). Either way, the writer extracted attention and returned less than common vocabulary would have returned.

Humility performance violates commitment by replacing claim with ritual self-deprecation. “I might be wrong, but I’ve sometimes thought that perhaps X could be Y in certain cases, though of course I could easily be mistaken.” The sentence contains four humility tokens and no claim. The writer’s humility is detached from reality — it would be present whether the writer had high or low confidence in the underlying point. Readers trained to pattern-match on humility will rate this sentence as careful and thoughtful. Readers who check what was transmitted will find nothing. The humility is a performance the text is doing instead of transmitting.

Numbers around vagueness violates commitment in a disguise. “24 of 34 participants received moderate benefits” produces quantitative precision in the easy parts (the count) and fog in the load-bearing parts (“moderate benefits”). The sentence appears rigorous because of the numbers, but the real claim — what the benefit was, how measured, in whom — stays soft. Readers accept the precision signal from the numerator and denominator and miss that the qualitative claim is unfalsifiable. This failure is common in scientific writing because it passes peer review: the reviewer sees numbers and judges the sentence as data-backed, without auditing the soft word where the actual claim lives.

Padding violates budget discipline. Words added because the document felt too short, the reviewer asked for more, the template required length, or the writer wanted to appear thorough. Each padding word dilutes payload density and exhausts reader attention. Documents optimized for length rather than payload produce the 25,000-word versions of seven-word truths. The writer can point to every sentence and explain why it’s there — “the reviewer asked,” “the template has this section” — and still have produced a document that fails by transmitting less than its length suggests.

Each failure is fixable by reference to the platform it violates. Hedging repairs by commitment. Fog repairs by specificity. Fancy repairs by vocabulary matching. Humility performance repairs by deleting the performance and keeping the claim. Numbers around vagueness repair by making the soft word specific. Padding repairs by cutting. The failures are not mysterious; they are violations of specific rules, and the rules are the platforms of the method.

9. Budget Discipline

The writer has word count. The reader has attention and working memory. Both are finite. Every word either deposits payload or wastes budget.

There is no neutral content. A word that doesn't contribute to the upload is a word that extracts reader attention for no return. The document's length is not proportional to its payload; often length and payload are inversely related, because longer documents contain more padding, more hedging, and more diluting material accumulated through the drafting process.

Readers have a second budget that length affects: the quality of attention. Readers who have been fed low-value words for several paragraphs downshift into skimming. Their decoder reduces fidelity because the signal has been low. When the writer finally produces the specific claim that matters, it arrives to a skimming reader who processes it with the same low fidelity as the surrounding hedges. The hedging did not just waste its own words — it poisoned the attention available for the words that followed. This is why a document cannot be “90% hedged with a strong claim at the end.” The end never lands, because the reader stopped reading carefully long before.

The test for each word: does it deposit payload? If not, cut it. The writer who cannot defend why a word is present should delete it. This is not a guideline; it is the budget constraint. A document at half the length with the same payload transmits better, because the reader spends less attention and retains more.

The cut is where most writers fail. They are reluctant to delete because deletion feels like loss. But the words being deleted were not payload — they were padding, hedging, performance, or decoration. Deleting them does not lose information; it concentrates information. The shorter document is not a diminished version of the longer one; it is the signal freed from the noise the longer one carried.

10. Preflight Simulation

Technical writing has no live feedback loop. The writer cannot assess the reader in real time. The Pseudo-Socratic Method [[HOWL-INFO-3-2026](#)] runs assessment continuously in conversation; technical writing cannot. So the writer must simulate the reader's encounter with the text in advance and fix the failures before publication.

The writer's own reading of their draft is worthless as a test. They have the original referent that the text is compressing. When they reread their draft, they decompress every word back to the full thing they were thinking, and the text feels clear because they are reading text-plus-original, not text alone. The reader has only the text. The writer's reading cannot detect where the text fails to transmit, because the writer's own dictionary fills in what the text omits.

The simulation: imagine the target reader hitting each sentence. Where do they stumble? Where do they reconstruct wrongly? Where does a word point to a dictionary slot they don't have? Where does a reference assume prior knowledge they may lack? Where does the sentence structure demand they hold too many things in mind at once?

Build platforms into the document's structure. Each major transition should include an explicit "you should now have X" marker that lets the reader audit their own state. If they have X, continue. If they don't, the text tells them what to do — go back, rebuild, verify, then proceed. This delegates the assessment loop to the reader, making the document self-auditing. Readers who would otherwise silently accumulate misunderstanding instead encounter a checkpoint that forces them to repair before continuing.

Examples and concrete instances are compression-failure detectors built into the text. A reader whose decompression of a word went wrong will fail the example that depends on the correct decompression. The failure signals them to re-examine. Without examples, wrong decompressions propagate undetected through the rest of the document, and the reader finishes with a coherent-feeling but mistaken reconstruction of the payload. The writer who does not include examples has abandoned the reader to undetected errors.

The simulation is preflight work. It happens before the document is published, because after publication the writer cannot recover misunderstandings that readers form silently. The live conversation can repair on the next turn. The document cannot. Every sentence must be pre-tested against the target's imagined encounter, because the encounter itself is sealed off from the writer.

11. Feedback Filtering

After drafting, external feedback helps if filtered correctly. Uncritical incorporation of feedback degrades the document.

Feedback reflects the reviewer's reading, not the target's. A non-target reviewer's confusion is accurate to their experience but does not automatically indicate the document needs changing. A specialist reviewer asking for more depth in a document written for non-specialists is reporting their own unmet need, not a defect. A non-specialist reviewer asking for unpacking of terminology in a document written for specialists is reporting their own unmet need, not a defect. The reviewers are being honest. Their honesty does not translate to instructions.

The filter is one question: does acting on this feedback improve the transfer to the intended target? If yes, incorporate. If no, discard, regardless of reviewer seniority, social pressure, or how reasonable the request sounds. The writer's accountability is to the target reader, not to reviewers. Reviewers are proxies who may or may not represent the target.

Reviewers' reflexes often include the failure modes. They ask for hedging ("this sounds too strong"), for more coverage ("have you considered X"), for vocabulary downgrades ("could you explain this term"), for humility performance ("acknowledge the limitations"). These asks are institutional reflexes, not evaluations of whether the transfer to the target improved. Obeying them pulls the document toward the failures even though each individual change felt responsive.

The drift happens through accumulation. Each obeyed piece of feedback is a small compromise. Individually they seem harmless. Collectively they turn the document into a patchwork of responses to reviewers rather than a focused transmission to the target. The

writer can defend every sentence (“a reviewer asked for this”) and still have produced a failed document.

Saying no to reasonable-sounding feedback is sometimes the correct move. The discipline is to let the target’s transfer drive the decision, not the reviewer’s satisfaction. A reviewer who is not the target will sometimes be left unsatisfied. This is acceptable. The document is not for them.

12. The Master Document

When the target is specialized or constructed, the master document serves that target without dilution. Derivative versions for populated audiences are downstream products, each with trade-offs.

Sometimes the ideal target has no demographic label. The payload requires a reader with specific capabilities that may not map to any named group. The writer constructs the target from the payload’s requirements and writes for that constructed reader without accommodation. The document is dense, specific, and demanding. Most readers will not match the target. This is acceptable. The master document optimizes for the match between encoding and decoding at the target; it does not optimize for reach.

Readership loss is the cost of fidelity. A master document that reaches fifty readers who can fully use it is more valuable than a diluted version that reaches fifty thousand readers who skim it and forget. The fifty will build on what they received. The fifty thousand will have consumed a shaped nothing and retained nothing. Choose fidelity when the payload deserves it.

The master anchors derivatives. A layman summary can be checked against the master: does this simplification preserve what mattered, or did it drop something load-bearing? A specialist adaptation can be verified against the master: does this framing for a different field preserve the core claims, or did translation distort them? Without the master, derivatives drift and lose coherence with what they claim to represent. The master is the reference that keeps the versions honest.

Publishing the master is a separate decision from writing it. Some masters are too dense, too specialized, or too demanding to serve any likely public readership, and the writer publishes only derivatives. This is defensible. The writer still benefits from having written the master: the act of writing to the ideal target clarifies thinking, forces full commitment, and produces an authoritative reference the writer can return to when producing derivatives. The master’s value is not contingent on its publication.

When the master deserves publication, resist the institutional pressure to only publish accessible versions. Accessibility has a cost. The cost is paid by the readers who could have handled the master directly and are now forced to reconstruct it from compromised derivatives. The master-plus-derivatives approach serves both populations: the master for the readers who can use it, derivatives for the readers who cannot, with the trade-offs of each version explicit. This is more honest than a single compromise document that claims to serve all readers and serves none well.

13. The Operating Loop

The method is a loop. Each document runs through it once; the writer's practice runs through it indefinitely.

Before writing: Confirm the specific thing you are uploading. State it in one sentence. If you cannot, find the payload or stop. Confirm the target. Pick one primary reader; construct the ideal reader if no demographic fits. If the target is still vague, the downstream decisions cannot be made.

During writing: Choose vocabulary against the target's dictionary, not your own. Commit to falsifiable claims. Cut every word that does not deposit payload. Build platforms into the structure so the reader can audit their own state at transitions. Include examples and concrete instances as compression-failure detectors.

After writing: Simulate the target reader's encounter. Read each sentence as they would read it, imagining what lands in their head and what doesn't. Revise where the simulation shows failure. The draft the writer finds clear is not necessarily the draft the reader will find clear; the simulation is where this gap is detected and closed.

In review: Filter feedback against the target. Every piece of feedback gets one test: does acting on this improve the transfer to the intended target? If yes, incorporate. If no, discard. Saying no to reasonable-sounding feedback is part of the discipline.

Across documents: Maintain the master when the target is specialized; produce derivatives deliberately, with trade-offs acknowledged. The master is the reference; the derivatives are projections. Neither replaces the other.

The loop is not a checklist the writer completes once. It is the practice the writer runs every time they write anything technical. With repetition, the steps become habitual and compress into a single integrated act of writing — but they remain the steps, and when something fails, the diagnosis is “which step broke?” rather than “what went wrong?” The method gives the writer a specific place to look when a document is not working.

14. What You Can Now Do

If you have read this paper as its target reader and the transmission worked, you now have the following:

You can refuse to write when there is no specific payload, because you know that writing without payload produces text that looks like writing and transmits nothing. You recognize the absence of payload as the root failure and you don't try to repair it with style.

You can pick or construct a target reader with enough specificity to make vocabulary and depth decisions. You don't write for “everyone” because you know that writing for everyone is writing for no one.

You can match vocabulary to the target's dictionary in both directions. You don't reach for specialist words when common words would carry the load, and you don't unpack common words when the target already has them compressed.

You can commit to falsifiable claims. You can identify your own hedging as a failure rather than a virtue, and you can repair it by committing or deleting. You know the difference between caution and cowardice in a sentence.

You can identify the failure modes in your own drafts and in others' writing. You recognize hedging, fog, fancy vocabulary, humility performance, numbers around vagueness, and padding as specific mechanical defects, not stylistic choices.

You treat word count as a budget. You cut words that don't deposit payload, even when cutting feels like loss, because you understand that cut documents transmit better than padded ones.

You simulate the target reader's encounter with your text before publication. You don't rely on your own reading of your draft, because you know your dictionary fills in what the text fails to transmit.

You filter feedback rather than obeying it. You evaluate each piece of feedback against whether it improves the transfer to the intended target, and you say no to reasonable-sounding asks that would degrade the document.

You decide when a master document is warranted and produce derivatives as separate products. You accept readership loss on masters as the cost of fidelity, and you maintain the master as the reference that keeps derivatives honest.

You recognize that the method is a loop, not a technique. You practice it every time you write, and when a document fails, you diagnose the failure by locating which step broke.

This is the post-read state the paper was written to produce. If you have it, the transmission worked. If you don't, some platform in the paper failed for you, and the failure is diagnostic: go back to the platform that didn't land, work through it until it does, then continue. The paper is structured to support this — each platform is a checkpoint you can audit.

15. Closing

Technical writing is a specific operation with specific criteria. It is not a style, a genre, or a craft in the aesthetic sense. It is the transfer of a specific thing from one head to another, through a lossy medium, against a finite attention budget. The writer who treats it as this operation produces documents that transmit. The writer who treats it as anything else — self-expression, topic coverage, register performance, institutional compliance — produces documents that fail the transmission and fill the world with shaped nothing.

The current state of technical writing, across industry, academia, and increasingly AI-generated text, is dominated by the failure modes. Hedging is rewarded. Fog vocabulary is accepted as professional. Fancy vocabulary passes for rigor. Humility performance reads as careful. Numbers-around-vagueness survives peer review. Padding inflates page counts. The readers pay the cost in attention spent on text that transmits a fraction of its length's worth of payload, and most readers don't notice because the shape of transmission is preserved even when the content is hollow.

This paper is an attempt to make the operation visible again. The method is not novel; it is what working technical writers have always done when they were transmitting rather than performing. What is novel, perhaps, is naming the failures mechanically and giving writers a specific place to look when their writing is not working. The writer who has the method installed can diagnose their own failures and repair them. The writer without it can only sense that something is wrong and reach for the usual remedies — add more words, add more hedging, add more coverage, ask more reviewers — which are the failures dressed up as solutions.

The practice is the point. Reading this paper installs the method conceptually. Running the loop on documents is what turns the concepts into a working discipline. The first documents written with the method consciously applied will feel slow, because the steps are still deliberate. With practice, they compress. The writer eventually writes the way they read: integrated, fast, correct by habit. The steps are still present but no longer require conscious execution.

If you are the target reader of this paper, you already know the next step. Write something, run the loop, diagnose the failures, repair them, and write the next thing. The method is installed by use, not by reading. This paper is the reference you return to when use produces questions. The questions will come. The method gives you where to look for the answers.

References

[[HOWL-INFO-6-2026](#)] How to Write Technical Documents. Github: <https://github.com/ghowland/papers/tree/main/papers/INFO-6-2026>

[[HOWL-INFO-1-2026](#)] Multi-Dimensional Information Indexing. Github: <https://github.com/ghowland/papers/tree/main/papers/INFO-1-2026>

[[HOWL-INFO-2-2026](#)] The Scales Method. Github: <https://github.com/ghowland/papers/tree/main/papers/INFO/HOWL-INFO-2-2026>

[[HOWL-INFO-3-2026](#)] The Pseudo-Socratic Method. Github: <https://github.com/ghowland/papers/tree/main/papers/INFO/HOWL-INFO-3-2026>

[[HOWL-INFO-4-2026](#)] Howland's Axiom of Information Locality. Github: <https://github.com/ghowland/papers/tree/main/papers/INFO/HOWL-INFO-4-2026>

[[HOWL-INFO-5-2026](#)] Coherence as Information Phenomenon. Github: <https://github.com/ghowland/papers/tree/main/papers/INFO-5-2026>