Achieving High Output and High Quality

Requires Both Experience and Falsifiability

Registry: [HOWL-INFO-9-2026]

Series Path: [HOWL-INFO-1-2026] → [HOWL-INFO-2-2026] → [HOWL-INFO-3-2026] → [HOWL-INFO-4-2026] → [HOWL-INFO-5-2026] → [HOWL-INFO-6-2026] → [HOWL-INFO-7-2026] → [HOWL-INFO-8-2026] → [HOWL-INFO-9-2026]

Date: May 2026

DOI: 10.5281/zenodo.20176605

Domain: Information Theory / Methodology of Production

Status: Complete

Abstract

High output and high quality are compatible if and only if two conditions are jointly satisfied: the writer has experience behind what they write, and the writing is delivered in usable, falsifiable form. Neither condition alone is sufficient. Experience without form produces unusable insight — the writer has something to say but cannot transmit it. Form without experience produces usable-shaped fabrication — the writing commits and admits falsification but describes nothing real. Both factors together produce information the receiver can act on, test, and build from. The rate of production is bounded by the writer's experience availability and their form-discipline consistency, not by volume-per-se. A writer with substantial cross-domain experience who reliably applies form-discipline to every output can produce quality at volume. A writer missing either factor produces failure regardless of how much time they spend per piece. This paper names the two factors, demonstrates the failure modes when either is absent, shows how the two combine to produce quality at rates that look impossible under the standard quality-volume assumption, and gives the writer an operating loop they can run on every piece they produce.

1. What You Are Reading This For

You produce information. You write papers, analyses, documentation, research reports, explanations, arguments. You have a volume question.

Maybe you produce too little. You want to write more but the quality feels like it would drop if you did. Each piece takes what it takes, and accelerating it seems to mean accepting worse output.

Maybe you produce a lot and you're uneasy. The volume is there but you suspect the quality is uneven, or you're not sure what would happen if someone checked the pieces you produced fastest against the ones you labored over longest.

Maybe you watch others produce at volumes that seem impossible and can't tell if the work is real or if you're looking at shaped nothing. The pieces have the right surface features. You don't know how to tell whether the substance matches the shape.

All three versions of the situation share an assumption: quality and volume trade off against each other. You can have one or the other. High output implies low quality; high quality implies low output. The assumption is so standard it usually goes unexamined.

This paper rejects the assumption. High output and high quality are compatible — but only under specific conditions. When those conditions are met, the trade-off disappears. When they aren't met, the trade-off operates, and slowing down doesn't rescue quality that was never there.

The conditions are two factors operating jointly on every piece you produce. The rest of this paper names them, shows why each is necessary, demonstrates why neither alone is sufficient, explains how they combine to allow volume without quality loss, and gives you an operating loop you can run on every output. If you finish this paper and still believe quality and volume are in fundamental opposition, you haven't crossed the first platform, and the disciplines that follow won't install correctly.

2. The Two Standard Corners

Writers mostly occupy one of two corners.

High output, low quality. Content farms, LLM-generated text produced without human input, institutional output produced on publication schedules, commentary that synthesizes other commentary. The volume exists. The substance doesn't. A reader finishes the output and cannot state what they learned, what the writer claimed, or what they should do differently. The shape of writing is preserved. What was supposed to travel through the shape didn't.

Low output, high quality. The serious author producing a book every few years. The careful researcher with one major paper per decade. The craftsman whose output is small but dense. The substance is there. The volume isn't. A reader finishes the output changed — they learned something specific, they can state the writer's claim, they can apply what they received. But the receiver has to wait years between pieces.

Both corners are familiar. Both are standard. The cultural assumption is that these two corners define the space — you pick a corner and pay the cost of the axis you abandoned.

What the assumption misses is *why* quality and volume usually trade off. Nobody specifies the mechanism. The trade-off is treated as a law of nature, like conservation of energy. With the mechanism unspecified, you cannot tell whether the trade-off is a law or an artifact of the conditions most writers operate under.

This paper specifies the mechanism. Once you see it, the trade-off resolves. Writers in the two standard corners are each missing something. The low-output-high-quality writer has what they need but cannot sustain volume because they're doing work the LLM era has a solution for. The high-output-low-quality writer has sustained volume but is missing the factors that would make the volume mean something. Neither is stuck where they are by physics. They're stuck by missing factors.

Name the factors and a third corner becomes available: high output, high quality. Not by accident, not by genius, not by luck. By having both factors present on every piece.

3. Factor One: Experience Behind the Information

The first factor is that you must have encountered what you're writing about. Not read about it. Not reasoned about it. Encountered it — attempted the thing, observed it directly, accumulated obstacle-data and lesson-data from contact with the reality your writing purports to describe.

This is strong. It rules out most of what gets written. Most writing is synthesis of other writing. Someone reads the literature on a topic, synthesizes it, adds their own framing, publishes. The synthesis can be competent and even insightful, but it does not carry experience behind it. The writer has not encountered the thing. They have encountered other people's writing about the thing, which is different.

Three categories you need to recognize in your own work:

Experience-grounded writing. You have lived the thing. You attempted it and it worked or it didn't. You have scars from where it didn't work, corrections you made when you figured out why, specific failures you remember, specific lessons the failures produced. When you write, the residue of contact shows up in the text. You can describe obstacles with specificity nobody could produce from outside. You know where the exceptions are and why. You know what would falsify your position because you know what almost falsified it along the way.

Intellectual writing. You have thought about the thing but not encountered it. Your position came from reasoning, not from observation. You have no obstacle-data because you didn't attempt anything. You have no lesson-data because you had no failures to learn from. Your writing may be internally consistent — reasoning can produce coherent structures — but it lacks the specificity that experience produces. A reader with experience can usually tell the difference, because the intellectual version misses exactly what experience would have taught.

Derivative writing. You have read what others wrote about the thing. Your position is a synthesis of theirs. The experience behind the writing is theirs, not yours, and it has been diluted through your filter. What you produce can be competent but is bounded by what your sources contained. You can't exceed their depth because you didn't generate the depth — they did.

The diagnostic questions that separate these:

- Did you attempt the thing? Not read about others attempting it, not reason about what would happen if you attempted it — did you actually do it?

- What obstacles did you encounter? Specific ones, with enough detail that someone who had not done the thing couldn't have produced the description.
- What lessons did you accumulate? Specific corrections, specific changes in your approach, specific observations that changed what you thought at the start.
- Can you answer follow-up questions about the content in specificity beyond what's on the page? If the page is compression of more you actually know, the answer is yes. If the page is everything you know, the answer is no, and the writing is intellectual or derivative.

Experience-grounded writing transmits with a density the other two cannot match. The reader can feel the difference, even when they can't articulate it. They notice that the piece holds up under their own questions, that it handles edge cases the reader didn't expect it to handle, that it names mechanisms rather than gestures at them. What they're noticing is the residue of contact. The writer encountered the thing. The writing carries what the encounter produced.

Derivative writing is bounded in volume by its sources. You can only synthesize what's already been written. When you run out of synthesizable material, your output stops or starts repeating itself.

Experience-grounded writing's volume is bounded by your experience availability. For someone with broad experience across domains, the available payload is substantial. You can write about what you've done — and if you've done many things across many domains, there's a lot to write about. This is the beginning of how volume becomes compatible with quality: experience-grounded writing has a supply curve that doesn't run out as fast as derivative writing's.

But experience alone doesn't produce quality output. It produces the *payload* that quality output would carry, if you had the second factor.

4. Factor Two: Usable, Falsifiable Form

The second factor is that you must state your information in a form that commits to specific content and provides the substrate for its own verification. The receiver must be able to do something with what you wrote — learn from it, apply it, test it — and they must be able to check it against something.

This is Popper's criterion, extended past the narrow scientific application it usually gets. A scientific claim is falsifiable if there's a state of the world that would make it false. The criterion generalizes to all information acts.

Factual claims check against reality. "The sun rises in the east" is falsifiable because if the sun rose in the west tomorrow, the claim would be dead.

Principles check against repeated application. "You will know the tree by the fruit it bears" is falsifiable because if a tree's fruit consistently failed to reveal its species or health, the principle would be dead.

Fiction checks against internal setup. A revelation in a novel is falsifiable because if it contradicts what the setup established about the character or world, readers detect the contradiction and the revelation fails.

Opinions and preferences check against behavioral record. “I don’t like onions” is falsifiable because if the person eats onions regularly and seeks them out, the claim is wrong regardless of what they say.

Beliefs and values check against costly action. “I value honesty” is falsifiable because if the person lies when lying is easy and truth-telling is costly, the claim is wrong.

Every information type has a verification substrate. None is exempt. The substrate differs by type, but something is always checkable. A writer who stakes a claim in a form that admits no check has not transmitted information — they have transmitted text that has the shape of information.

The self-diagnostic for every sentence you produce: what state of the world would make this false? If nothing would, the sentence has claimed nothing and has to be rewritten or deleted.

Hedging, fog vocabulary, humility performance, and unfalsifiable framing are the standard failures at this factor. They produce sentences that cannot be wrong, which sounds careful and is in fact empty. “Results suggest that in some cases the intervention may contribute to improved outcomes among certain participants” is consistent with any possible state of the world. It has made no claim. The reader cannot check it because there’s nothing to check. They cannot act on it because there’s nothing specific to act on. The writer has taken reader attention and returned nothing for it.

Form is not editorial polish. It is structural commitment. You cannot fix a hedged draft by tightening the prose. The hedging expresses your uncommitment. It’s a symptom, not a style. The fix is to commit to what you actually know, state it specifically, and let it be wrong if it’s wrong.

Commitment feels risky to writers trained in institutional environments that reward hedging. The risk is asymmetric. The committed writer can be wrong and corrected, and the correction is information the reader can use. The hedged writer cannot be wrong, because they claimed nothing, but the absence of wrongness is the absence of content. The reader who reads a committed claim can check it, act on it, refute it, build on it. The reader who reads the hedged version can do none of these. Commitment transfers risk from the reader’s ability to use the writing to the writer’s reputation for accuracy. The trade is correct. The reader should not bear the cost of the writer’s protection of their own record.

Form is the second factor. It is what lets payload travel. Without it, you can have all the experience in the world and produce writing nobody can use.

5. Why Neither Factor Alone Is Sufficient

The claim of the paper is that both factors are necessary. This section demonstrates why each without the other fails.

Experience without form produces unusable insight.

Imagine a person who has genuinely lived something. They spent thirty years doing a specific kind of work. They have obstacle-data, lesson-data, observations that would be valuable to receivers who could use them. But when they write, they produce text like this: “I’ve been thinking about this for thirty years and there’s more going on than people realize. It’s complicated. Different cases come out different ways. The main thing is to pay attention to the specifics.”

The statement may be accurate. The experience is real. The writer has something to say. But what the reader receives is nothing specific — no claim, no mechanism, no test, no application. The writer has not transmitted the thirty years. They have transmitted that they have thirty years, which is different. The payload stays with the writer. The reader walks away with nothing.

This is a common failure of senior practitioners. They know a lot. They cannot state what they know in a form that travels. Their knowledge dies with them, or gets passed to apprentices slowly through observation over years, because they never found the form that would let it be transmitted at the speed of reading. The experience was real. The transmission vehicle didn’t exist.

Form without experience produces usable-shaped fabrication.

Consider this statement: “Eating blueberries allows you limited flight abilities at 6-6:30am on a western facing coastline.”

Run the form checks on it. Is it specific? Yes — specific food, specific time window, specific geography, specific outcome. Is it committed? Yes — it states a factual claim without hedging. Is it falsifiable? Yes — you could follow the protocol exactly and see whether you achieved flight. Is it actionable? Yes — a reader who wanted to fly has a clear set of instructions.

The statement passes every formal check a reader would apply to verify the form. And it is entirely fabricated. Nobody ate blueberries on a western coastline at dawn and took flight. No observation lies behind the specifics. Every detail was constructed — the food, the time, the location, the outcome — and assembled into a statement that has all the virtues of good information form.

This is the pure case of factor two without factor one. The writer has produced something that looks like information at every structural level, and it describes nothing real. A reader who checks only the form cannot distinguish this from a true claim with the same form. They’d have to run the test — eat the blueberries, go to the coastline, wait until dawn — and then discover the statement is false. The form held up until empirical contact. Only then did the absence of experience behind it show.

Most fabrication isn’t this cartoonish. But the mechanism is the same. Someone writes about a method they haven’t tried, a system they haven’t built, a phenomenon they haven’t observed. They produce text with good form — specific, committed, falsifiable, usable-looking — from reasoning about what such text should contain. Readers who evaluate on form alone don't detect the missing factor. Readers who try to apply the text discover the gap. By then, the reader has paid the cost.

The two failure modes look different but produce the same result.

Experience without form: the reader cannot extract anything usable from what the writer knows. The payload exists but doesn't travel.

Form without experience: the reader extracts something usable-looking from what the writer doesn't know. The transmission happens, but what gets transmitted isn't real.

Both fail. Neither can be repaired by adding more of the factor that's already present. Spending more time editing a intellectual piece doesn't add experience. Having more experience doesn't by itself teach you to write in committed form. The factors are orthogonal. You need both, and deficit in either cannot be compensated by excess of the other.

Writers commonly try the compensation and it doesn't work. The experienced writer who struggles with form keeps refining the prose and the piece keeps not landing. The committed formalist who writes confidently about things they haven't done keeps producing plausible-sounding work that fails when tested. Each is missing a factor. Neither can reach quality by doubling down on the factor they have.

6. The Joint Condition

What happens when both factors are present together.

Experience supplies the payload. You have encountered the thing. You have specific observations — not generic ones, specific ones that a person who hadn't encountered the thing couldn't have produced. You have obstacles you've met, lessons you've accumulated, mechanisms you've watched operate. The payload exists before the writing starts, and the writing's job is to transmit it, not generate it.

Form supplies the transmission. You commit to specific claims. You state them so a state of the world could make them false. You provide the verification substrate — the reader knows what would check your claim, and often has enough of the substrate available to check it themselves. You cut hedging that protects you at the reader's expense. You cut fog vocabulary that would obscure what you mean. You maintain the commitment the experience earned you the right to make.

Together they produce writing the reader can use. The experience guarantees what you're saying is real. The form guarantees the reader can receive it, check it, act on it. The reader walks away with something they didn't have — not just information, but information they can verify against reality or against their own experience, and information they can apply.

This is quality output. Not "quality" as a vague aesthetic judgment but as a specific structural property: the piece satisfies both factors on every claim it makes. A piece that satisfies both factors is quality regardless of genre, domain, length, or style. A piece that fails either factor is not quality regardless of how carefully it was written, how long it took, how many revision passes it got, how many reviewers approved it.

The two factors are the quality condition. Satisfying them is what quality *is*, structurally. Not satisfying them is what failure is. There is no third category. There is no “mostly met both factors” that produces partial quality — you either commit or you don’t, you either have experience or you don’t, and the failure modes don’t compose gracefully.

7. The Volume Compatibility

Here is where the paper’s claim about volume becomes stateable.

Volume in writing is usually bounded by time-per-piece. A writer who takes a month per piece produces twelve pieces a year. A writer who takes a year per piece produces one. The assumption is that shortening the time degrades the quality, so the only way to increase volume is to accept quality loss.

Under the two-factor framework, volume is bounded by something different. It is bounded by:

- **Experience availability**: how much you have actually lived across how many domains, available to be drawn on when you write.
- **Form-discipline consistency**: how reliably you apply the commitment-and-falsifiability requirements to every output you produce.

Neither bound is set by the time you spend per piece. A writer with substantial cross-domain experience has a supply curve that doesn’t run out quickly. A writer with reliable form-discipline applies the discipline at the speed of writing — it doesn’t take longer to commit to a claim than to hedge it. The actual time per piece is the time to produce the text, not the time to supply the factors.

Time-per-piece matters if the writer is themselves doing the text production. A human typing at 60 words per minute is bounded by their typing speed. A human editing iteratively through five revision cycles is bounded by the time to re-read and correct.

But text production and factor-supply are separate operations. The human supplies the experience and enforces the form-discipline. Someone or something else can produce the text under those constraints. When the text production is handled separately — by an LLM, by a dictation-to-editor pipeline, by a co-writer — the human’s rate-limit becomes factor-supply rather than text-production.

Factor-supply is not bounded by words per minute. It’s bounded by availability of experience (how much you have) and consistency of discipline (how reliably you apply it). A human with ten years of broad experience across five domains has a payload supply that, if drawn on efficiently, can produce tens or hundreds of pieces per year. A human with reliable form-discipline can enforce the discipline on each piece at the speed of reading, not the speed of writing.

The LLM is the text producer under this arrangement. It can generate words at rates humans cannot. What it cannot do is supply experience it doesn’t have or enforce discipline on its own output — LLMs drift toward hedging, toward synthesis of training data, toward plausible-sounding text without commitment. Without a human supplying both factors, the LLM's output falls into the high-output-low-quality corner. The LLM's speed is real but it's not producing quality alone.

When a human supplies both factors and the LLM produces text under those constraints, the combination operates in the previously-unreachable corner: high output, high quality. The human's factor-supply sets the ceiling on rate. The LLM's text-production sets the floor on time per piece. The two together produce pieces at rates humans alone cannot achieve, at quality LLMs alone cannot produce.

This is not a trick. It is a natural division of labor under the two-factor framework. Each participant does what they can do. The human cannot produce text at LLM speed. The LLM cannot supply experience or enforce discipline. Together they produce both factors jointly on every output, at a rate neither can achieve alone.

The quality-volume trade-off operates under the standard assumption because the standard assumption collapses factor-supply and text-production into a single operation performed by one human. Separate them, and the trade-off resolves into two separate rate-limits that don't compete with each other.

A writer who objects that they have no LLM, or refuses to use one, is not excluded from the framework. They can still achieve quality at their own rate — bounded by the speed at which they type, revise, and produce text. What they cannot achieve is the specific volume that LLM-assisted production unlocks. The two factors still govern quality; the volume ceiling is lower because the text-production rate-limit is stricter.

8. Failure Mode Taxonomy

You need to recognize each failure mode in your own output. The taxonomy:

Missing experience (factor one absent).

- Derivative synthesis masquerading as original work. You read ten papers on a topic, synthesized them, added your framing, published. The experience behind the writing is the ten authors', not yours. Your contribution is the filter. When readers apply your synthesis to reality, they discover the gaps between what the literature says and what happens — gaps that would have been visible to you if you had encountered the thing directly.

- Form-only fabrication. You produce specific, committed, falsifiable claims from reasoning rather than observation. The blueberry-flight case. The form survives until empirical contact.

- Institutional consensus piled up by citation. You cite many sources that cite other sources that cite other sources. At no point does experience enter the chain. The consensus is real as a social fact; it's not grounded in anyone's encounter with the thing being consensed about.

Recognition: what did you actually encounter here, versus what did you read about or reason about? If the answer is "read about" or "reasoned about," you're missing factor one. Acquire the experience or stop writing about the topic.

Missing form (factor two absent).

- Hedging cascade. Every claim is qualified with "may," "could," "sometimes," "in some cases," "under certain conditions," "to some extent." Each hedge individually looks like caution. Collectively the piece has claimed nothing and admits every possible state of the world.
- Fog vocabulary. "Stakeholders," "participants," "individuals," "users," "outcomes," "interventions." Generic terms that decompress to nothing specific in the reader's head. The piece appears to be about something but has no specific referent.
- Humility performance. "I might be wrong, but I've sometimes thought that perhaps X could be Y in certain cases, though of course I could easily be mistaken." The humility is detached from reality — it would be present whether you had high or low confidence. It's ritual, not calibration.
- Uncommitted claims. Sentences that survive any possible state of the world. "Complex phenomena require nuanced approaches." No state of reality makes this false. It has claimed nothing.

Recognition: run the Popper test on every sentence. What state of the world would make this false? If nothing would, the sentence has claimed nothing. Rewrite or delete.

Missing both (the most common mode).

- Intellectual content hedged into unfalsifiability. You reasoned about something you hadn't encountered, produced plausible-sounding text, and hedged every claim so that nothing commits. The piece can't be wrong because it hasn't said anything. It also can't be right.
- Confident fabrication qualified just enough to be unkillable. You produced form-only fabrication but added enough hedges that even the empirical test becomes ambiguous. "In many cases, blueberries may contribute to enhanced states of awareness at dawn in coastal environments." Now failing to fly doesn't kill the claim because the claim didn't quite promise flight.
- Shape-without-substance at institutional volume. Papers, reports, analyses produced on publication schedules by writers with no experience and no discipline. The shape of information production is preserved. What's supposed to travel through the shape isn't there.

Recognition: the characteristic signature is text that sounds authoritative, takes a position, and says nothing specific. When you finish reading your own draft, ask what a reader could do differently based on it. If the answer is nothing specific, you're missing both factors.

The taxonomy is diagnostic. For each piece you produce, you can identify which factor is weak or absent, and the remedy follows from the identification. Missing experience cannot be fixed by revision; you have to go encounter the thing or abandon the piece. Missing form cannot be fixed by adding more content; you have to commit to specific claims and cut the hedges. Missing both cannot be fixed at all as it stands; the piece needs to be rebuilt from factors that were never present.

9. The Dilution Cascade

One failure mode deserves its own section because it's the most common institutional pattern and the one most writers fall into without noticing. The dilution cascade is how committed claims become uncommitted claims across revision, often under external pressure, often with the writer believing they're refining.

The cascade operates in stages.

Stage one: a claim is stated committedly. "X causes Y." The writer has observed or reasoned something specific and written it specifically.

Stage two: the claim encounters counter-evidence or uncomfortable cases. A reviewer objects. An edge case appears. A counterexample surfaces. The claim as stated is wrong under some condition the writer hadn't fully specified.

Stage three: the writer rewrites. The new version is "Sometimes X causes Y." The word "sometimes" accommodates the counterexample without specifying which cases are the exceptions. The claim is now weaker.

Stage four: more counterexamples arrive, or the writer anticipates more. The claim becomes "In some cases X can influence Y." The word "influence" is weaker than "causes." "Some cases" is unspecified. "Can" is modal rather than actual. The claim now survives almost any state of the world.

Stage five: final pressure produces the endpoint. "The relationship between X and Y is complex and context-dependent." This is the floor. The claim has retreated to the statement that X and Y are somehow related, which was never in dispute. Nothing has been claimed. The paper still exists. The writer still has credit for it. The readership still has access to it. But the claim is dead.

Here is the test for recognizing the cascade in your own work: ask whether you would have written the piece at its current statement. A paper whose thesis is "In some cases X can influence Y" would usually not have been written. Nobody sits down to write a paper that says "X and Y are somehow related in unspecified ways under unspecified conditions." That claim is trivially true about almost any X and Y, and it provides no reason for the paper to exist. If the current statement is one you would not have bothered to write, the current statement is the endpoint of a retreat. The earlier version is what you meant. The question is whether the earlier version was defensible.

This is where the distinction between dilution and refinement becomes critical, and this is where the two factors come back into the analysis.

Refinement is what experience-grounded writers do under pressure. The original claim is experience-grounded — you observed X producing Y across many cases. A counterexample arrives. Because you have the observational record, you can specify what distinguished the counterexample from the cases where X did produce Y. The claim becomes: “X causes Y under conditions A, B, C. When condition D is present, mechanism M interferes and X does not produce Y.” The claim is now more specific, not less. It accommodates the counterexample by naming where the counterexample fits in the landscape you already mapped. Specificity increases.

Dilution is what intellectual writers do under pressure. The original claim was not experience-grounded — you reasoned that X should cause Y but never ran the observation. A counterexample arrives. You have no observational ground to stand on, so you cannot specify what distinguishes the counterexample. You can only retreat. “X causes Y” becomes “sometimes X causes Y” because you can’t say when. “Sometimes X causes Y” becomes “in some cases X can influence Y” because you can’t defend “causes.” Specificity decreases at each step.

The direction of the change under pressure tells you which process happened. Refinement increases specificity. Dilution decreases it. A claim that has moved from stronger to weaker while becoming less precise has diluted, not refined. A claim that has moved from simpler to more specific while incorporating exceptions has refined.

The correction for dilution is not to polish the diluted version. It is to go back to the original committed claim and ask whether it was experience-grounded. If it was, refine it — specify the conditions, name the exceptions, describe the mechanisms. If it wasn’t, the original claim was fabrication, and no amount of revision will rescue the piece. Delete it. It was a failure from the start. The dilution was just the failure becoming visible under pressure.

This is why experience matters so much to volume. A writer producing many pieces without experience will encounter pressure on all of them eventually. Every piece will start diluting. After a few years the writer’s entire output is diluted into unfalsifiable hedge, and they cannot write anything committed because they have nothing to commit on. Volume without experience collapses in one generation of pressure. Volume with experience survives pressure by refining — each piece becomes more specific as counterexamples accumulate, not less.

10. The Experience Test

You need to run factor-one diagnostics on your own work. The test is a set of questions you apply to every piece before it goes out.

Did you attempt the thing you’re writing about?

Not read about others attempting it, not reason about what would happen if you attempted it. Did you actually do it? A paper on technical writing should be written by someone who has written technical documents and observed what worked. A paper on a physical rehabilitation method should be written by someone who has rehabilitated physically. A paper on an operating system should be written by someone who has built one. If you cannot answer yes to this question for the central thesis of your piece, you are missing factor one.

Ambiguity around the question — "I've sort of done it, in a way, partially" — usually means no. If you had done it, the question would not feel ambiguous. Partial attempts and adjacent activities are not the thing itself. They produce partial obstacle-data and partial lesson-data, which shows up in the writing as gesturing at specifics that don't quite land.

What specific obstacles did you encounter?

Not generic obstacles. Not obstacles you imagine someone would encounter. Specific obstacles from your attempt, with enough detail that someone who had not done the thing could not have produced the description from outside.

The test: try to write down three obstacles you encountered when attempting the thing your piece is about. Each obstacle should be specific enough that it couldn't be generated from reasoning alone. If you can produce three specific obstacles, your piece is likely experience-grounded at least at the level of attempts. If you cannot, you probably haven't encountered the thing closely enough to write about it.

What specific lessons did you accumulate?

Lessons are corrections you made in response to what the obstacles taught you. Not abstract lessons — "I learned to be patient." Specific corrections — "I learned that when the bread dough resists folding, it's usually because the hydration is under 65 percent and I need to add water before continuing, not knead harder." The second is experience. The first is platitude.

The test: for each obstacle you listed, what correction did it produce? If each obstacle produced a specific correction you can state, you have experience. If the lessons are abstract or absent, you encountered the activity at a level shallower than writing about it requires.

Can you answer follow-up questions beyond what's on the page?

Your piece is a compression. If the compression is of experience, you know more than the page. A reader who asks follow-up questions finds out there's more there — you can explain conditions, give examples, handle edge cases, answer from specificity you didn't fit into the piece.

If the compression is of reasoning only, you know exactly what's on the page and nothing more. Follow-up questions reveal the limit immediately. You gesture, rephrase what you already wrote, or retreat to "it's complex" language.

The test: have someone ask you three questions about the piece that aren't directly answered in it. If your answers come from knowledge you had but didn't fit in, the piece is experience-grounded. If your answers come from on-the-spot reasoning, you were at the edge of your knowledge when you wrote, and the piece is probably intellectual.

Would your position survive a critic who has done the thing?

Not a critic who has read about it. A critic who has done it and has their own obstacle-and-lesson record. When such a critic reads your piece, do they recognize what you're describing? Do they push back on specifics that don't match their observations, in ways that indicate they're engaging with real content rather than objecting to your framing?

If you haven't tested the piece against a domain practitioner, you don't know. When you can, test it. The quickest way to discover that a piece is intellectual is to show it to someone with real experience and watch their reaction. A intellectual piece generates polite confusion — they don't know what to say because the piece isn't engaging with what they'd recognize. An experience-grounded piece generates substantive engagement — they agree, disagree, refine, extend, but they're in the territory.

If the answers to these questions reveal your writing is derivative or intellectual, your options are two. Acquire the experience — attempt the thing, observe it, accumulate the obstacle-and-lesson record — or stop writing about it. There is no third option that produces quality output. Time and revision will not add experience that was never there. Hedging the claims doesn't compensate for missing factor one; it just converts the failure from form-failure to substance-failure while keeping the structural deficit.

11. The Form Test

The parallel diagnostics for factor two.

For each sentence, what state of the world would make this false?

Apply this to every sentence in your draft. If there is no state of the world that would make a sentence false, the sentence has claimed nothing. It's either rhetoric, filler, or hedge. Rewrite it into a form where a specific thing could be the case that would disprove it, or delete it.

This test is uncomfortable because most writers produce a lot of sentences that don't commit. Academic training actively rewards such sentences. Running the test aggressively reveals how much of a given draft is transmitting nothing, and the revelation is unpleasant. But the revelation is also the repair — the sentences that don't commit are the ones to cut or rewrite.

For each claim, what would a receiver need to do to check it?

The verification substrate must be available to the reader. If your claim is "X causes Y under conditions A," the reader checks by looking at cases where conditions A hold and seeing whether X and Y correlate. If your claim is "this method works," the reader checks by applying the method and observing the result. If your claim is "this principle holds," the reader checks by applying the principle across cases they know.

For each claim, identify the substrate and either provide enough of it in the piece or point the reader at where they can find it. If there's no substrate — no way for the reader to check — the claim is unfalsifiable in practice. Either restate it into a form that admits checking or cut it.

For each hedge, is it specifying real conditions or protecting you from being wrong?

Hedges are often legitimate. "X causes Y in mammals but not in reptiles" is a hedge that specifies a real condition. "X causes Y when the temperature is above 15 degrees" is a hedge that specifies a real condition. These hedges make the claim more accurate.

But many hedges are protective. “X may cause Y” where the writer actually believes X causes Y but hedges to avoid being wrong. “X causes Y in some cases” where the writer can’t specify the cases. “X might be related to Y” where the writer doesn’t know the relationship but wants to make a claim.

The test: for each hedge, can you replace it with a specific condition? If yes, do so — the claim becomes more accurate. If no, the hedge is protective, and the repair is either to delete the hedge and state the claim confidently or to delete the claim because you don’t actually have grounds for it.

For each paragraph, if the receiver acted on it, what would they do differently?

Writing exists to produce change in the reader. The change might be knowledge (they know something new), skill (they can do something new), belief (they think differently), or action (they do differently).

For each paragraph in your draft, ask what the reader could do differently after reading it that they couldn’t before. If the answer is nothing specific, the paragraph has not transmitted. It’s been furniture. Either rewrite it to transmit something or cut it.

For the whole piece, could someone summarize your position in one sentence that you’d recognize?

Your piece should compress to a single-sentence thesis. If it doesn’t — if someone asked what your piece was about and you couldn’t answer in one sentence — your piece is diffuse. If the single-sentence summaries people produce don’t match what you meant, your piece is ambiguous. Either way, the piece hasn’t committed to a specific position the reader can extract.

The test: write down what you think your piece’s thesis is in one sentence. Have someone who has read the piece do the same. Compare. If the two match, the piece has transmitted. If they don’t, the piece has transmitted something different from what you meant, and the fix is to rewrite until the thesis is clear enough that someone else can produce it from reading the piece.

If the answers to these questions reveal your writing is hedged, uncommitted, or unfalsifiable, the repair is commitment. Go through the draft and, for each sentence, commit to what you actually mean. Cut the protective hedges. Specify the real conditions. State the thesis in one sentence and make sure every paragraph serves it.

Form-discipline is not editorial polish applied at the end. It is structural commitment built in from the start. A draft produced under hedging pressure cannot be rescued by tightening — the hedging expresses uncommitment that runs through every sentence. The repair is to decide what you actually know and state it. If you don’t know, don’t write.

12. The Operating Loop

The method is a loop you run on every piece.

Before writing: confirm factor one.

Before you start drafting, confirm the piece is experience-grounded. Name the thing you encountered. State the obstacles and the lessons. If you cannot do this, don't write the piece. Acquire the experience first or pick a topic you have experience with. Writing without experience and then trying to add it through research during the draft doesn't work — you end up with a derivative synthesis that has the texture of experience-grounded writing but misses exactly what experience would teach.

Confirmation is fast. Five minutes of writing down what you encountered and what you learned, before any drafting starts. If the five minutes produces specifics, factor one is likely present. If the five minutes produces abstractions, go get the experience.

During writing: enforce factor two.

As you write, enforce commitment. Every sentence should state something specific enough to be false. Every claim should have a verification substrate the reader can access. No fog vocabulary. No ritual humility. No protective hedging where specific conditions would work.

The enforcement is live, not post-hoc. It's harder to edit uncommitted prose into committed prose than to write committed prose from the start. The hedging patterns are sticky — if you write them in, they tend to stay in. Write committed, let the claims be specific from the first draft, and you'll have less to fix at the end.

After writing: run both tests.

Before the piece goes out, run the experience test and the form test on the draft. Use the diagnostic questions in sections 10 and 11.

Experience test: did you attempt the thing? Can you name the obstacles? Can you name the lessons? Can you answer follow-up questions beyond the page? Would a practitioner recognize your description?

Form test: does every sentence commit? Is every claim checkable? Are all hedges specifying conditions rather than protecting you? Does every paragraph transmit something actionable? Can the thesis be stated in one sentence?

If either test fails on your draft, the repair is not revision in the general sense. The repair is to address the specific factor that failed. Missing experience: acquire it or cut the piece. Missing form: commit to specific claims and cut hedges.

Under LLM assistance: maintain both factor roles.

If you're writing with an LLM, you remain responsible for both factors. The LLM does not supply experience — it has none. The LLM does not enforce form-discipline on its own output — it drifts toward hedging, synthesis, and plausible-sounding uncommitment.

Your role under LLM assistance is to supply the payload (experience-grounded observations and claims) and to enforce the discipline (commitment, specificity, falsifiability) on every draft the LLM produces. The LLM handles text production under your constraints. The division of labor is stable only if you do both your parts. If you treat the LLM as a writer and accept its output without factor enforcement, you produce high-output-low-quality — which is the LLM-alone failure mode, not a limitation of LLM-assisted writing.

A concrete enforcement pattern: prompt the LLM with the payload you've identified. Read every claim in the output. For each claim, ask whether it matches what you know from experience and whether it's in committed form. Correct the ones that don't. Re-prompt for sections where the LLM drifted. Keep the claims that match your experience and commit to them explicitly. Discard what doesn't.

This is not automated. It is explicit human enforcement of both factors on every piece. The LLM does the typing. You do the checking. The rate at which you can check bounds the rate at which the LLM-assisted pipeline produces quality output. Most writers can check faster than they can type, so the pipeline produces more than a human alone, while maintaining quality a human supplies through the checks.

Across pieces: maintain both factors on every output.

The framework governs each piece independently. There is no accumulation of factor credit across pieces. A writer who produced ten good pieces does not get to produce one bad piece and have it be carried by the good ones. Each piece either has both factors or it doesn't.

This is why quality at volume is stable under the framework. The volume doesn't degrade the quality because the factors are supplied per-piece, not per-career. A writer with reliable factor-supply can produce many pieces with both factors present; a writer without reliable factor-supply produces few pieces with both factors present or many pieces without. The volume ceiling is set by factor availability, which for broadly experienced writers with reliable discipline is high.

The loop closes by returning to the start on the next piece. Before drafting, confirm factor one. During drafting, enforce factor two. After drafting, run both tests. Under LLM assistance, maintain both roles. Across pieces, keep the standard the same. The loop runs indefinitely. Each piece passes through it or doesn't. The ones that pass through it are quality. The ones that don't are failures.

13. Closing

Quality and volume do not trade off against each other. They trade off against missing factors.

A writer who has experience and maintains form-discipline produces quality at whatever rate they can sustain the combination. A writer missing either factor produces failure at whatever rate they produce, fast or slow. Time spent per piece is not the quality variable. Presence of both factors is the quality variable.

The framework specified here is not novel. Writers with deep craft have always done this implicitly. What is novel is naming the two factors explicitly and showing that they are jointly necessary and jointly sufficient. A writer with the framework installed can diagnose their own failures mechanically and repair them by returning to the missing factor. A writer without it can only sense that something is wrong and reach for the usual remedies — more revision, more hedging, more research, more polish — which do not address either factor and do not produce quality.

The current moment makes this framework particularly usable. LLMs handle text production at rates humans cannot. A writer with broad experience across domains, who maintains reliable form-discipline, can produce pieces at volumes previously available only to institutions and content farms, without the quality collapse those operations exhibit. The division of labor — human supplies factors, LLM produces text — is not a trick or a shortcut. It is the natural allocation of work under the two-factor framework. Each participant does what they can do. Together they produce what neither can produce alone.

What the framework refuses is every version of the assumption that volume itself degrades quality. Volume doesn't. Missing factors do. A writer producing 150 pieces in eight months with both factors present on each piece produces quality. A writer producing one piece in a year with either factor missing produces failure. The volumes are different; the quality depends on the factors, not on the volume.

The practice is the point. Reading this paper installs the framework conceptually. Running the operating loop on real pieces is what turns the framework into a working discipline. The first pieces produced under the framework feel slow because the checks are deliberate. With practice, they compress. The writer eventually runs both factors by habit, and the volume they can sustain at quality is set by their experience availability and their discipline consistency, not by any external limit.

If you are the target reader of this paper, you already know the next step. Confirm the experience behind your next piece before you start. Enforce commitment and specificity on every sentence you write. Run both tests before the piece goes out. Repair failures by addressing the specific factor that failed, not by polishing the symptom. Do it again on the next piece. The volume you can produce at quality is whatever your factor-supply supports. For writers with broad experience and reliable discipline, that volume is higher than the standard quality-volume assumption predicts, because the standard assumption was wrong about the mechanism. The mechanism is two factors. The factors are supplied by you. The volume follows.

References

[**HOWL-INFO-9-2026**] Achieving High Output and High Quality. Github: https://github.com/ghowland/papers/tree/main/papers/INFO/HOWL-INFO-9-2026

[**HOWL-INFO-1-2026**] Multi-Dimensional Information Indexing. Github: https://github.com/ghowland/papers/tree/main/ INFO-1-2026

[**HOWL-INFO-2-2026**] The Scales Method. Github: https://github.com/ghowland/papers/tree/main/papers/INFO/HOWL INFO-2-2026

[**HOWL-INFO-3-2026**] The Pseudo-Socratic Method. Github: https://github.com/ghowland/papers/tree/main/papers/INF
INFO-3-2026

[**HOWL-INFO-4-2026**] Howland's Axiom of Information Locality. Github: https://github.com/ghowland/papers/tree/main/papers/INFO/HOWL-INFO-4-2026

[**HOWL-INFO-5-2026**] Coherence as Information Phenomenon. Github: https://github.com/ghowland/papers/tree/main/p
INFO-5-2026

[**HOWL-INFO-6-2026**] How to Write Technical Documents. Github: https://github.com/ghowland/papers/tree/main/pape
INFO-6-2026

[**HOWL-INFO-7-2026**] How to Review Technical Writing. Github: https://github.com/ghowland/papers/tree/main/papers/
INFO-7-2026

[**HOWL-INFO-8-2026**] LLMs are not Tools, LLMs are Maybe-Tools. Github: https://github.com/ghowland/papers/tree/main/papers/INFO/HOWL-INFO-8-2026