

Exact Arithmetic Meets Logical Provenance

A Specification for Constraint-Grounded Language Models With Full Data Lineage

Registry: [\[HOWL-VDR-5-2026\]](#)

Series Path: [\[HOWL-VDR-1-2026\]](#) → [\[HOWL-VDR-2-2026\]](#) → [\[HOWL-MATH-3-2026\]](#) → [\[HOWL-MATH-4-2026\]](#) → [\[HOWL-VDR-3-2026\]](#) → [\[HOWL-VDR-4-2026\]](#) → [\[HOWL-LLM-1-2026\]](#) → [\[HOWL-VDR-5-2026\]](#)

DOI: 10.5281/zenodo.zzz

Date: May 2026

Domain: Applied Philosophy / Exact Machine Learning / Knowledge Systems

AI Usage Disclosure: Only the top metadata, figures, refs and final copyright sections were edited by the author. All paper content was LLM-generated using Anthropic’s Opus 4.6.

Abstract

This paper specifies VDR-LLM-Prolog, a language model architecture where every value is an exact fraction, every derivation is recorded in a logic programming knowledge base, every constraint is a first-class queryable object, and every piece of knowledge is directly surfaceable to the user without passing through the language model’s token generation. The specification integrates four prior results: VDR exact arithmetic (VDR-1 through VDR-3), the VDR machine learning stack (VDR-4), transcendental constant representation (MATH-3/MATH-4), and a custom Prolog-style knowledge engine designed for LLM provenance.

The system has three layers. The arithmetic layer (VDR) ensures every number is an exact fraction with zero drift and zero silent truncation. The logic layer (Prolog) records how every value was derived, what it depends on, and what constraints it satisfies. The conversation layer manages scoped knowledge bases, working data sets, topic tracking, and constraint activation — giving the language model structured persistent memory that survives topic switches, supports inheritance and shadowing, and is directly queryable by the user.

The central claim is that data provenance, constraint enforcement, and conversational state tracking are not features to be bolted onto a language model after the fact. They are architectural requirements that should be present from the foundation. This paper specifies what that foundation looks like.

1. The Problem

Modern language models have three structural deficiencies that no amount of scale, training data, or fine-tuning can fix.

Values without provenance. When a language model produces a number — a probability, a calculation result, a cited statistic — there is no record of how that number was derived. The model cannot show its work because it has no work to show. The computation that produced the number is a sequence of matrix multiplications through opaque float tensors. The number might be correct. It might be hallucinated. There is no systematic way to tell.

Approximate arithmetic. Every number inside a standard language model is a 16-bit or 32-bit float. Every operation silently truncates. After a few hundred operations, the accumulated rounding error is unmeasured and unmeasurable. Two runs of the same model on the same input can produce different outputs because float rounding is platform-dependent. The model’s internal arithmetic is fundamentally unreliable, and the unreliability is invisible.

Stateless conversation. A language model processes each turn by reading the entire conversation history as a token sequence. It has no structured memory, no scoped variables, no persistent working data. If you discuss two topics and switch between them, the model must re-derive everything from the token context. Facts stated thirty turns ago may have scrolled out of the context window. Facts from different topics are mixed in a flat sequence with no scoping, leading to confusion and cross-contamination.

These three deficiencies are independent. Each could be addressed separately. This paper addresses all three simultaneously because the solutions reinforce each other: exact arithmetic makes provenance meaningful (the recorded derivation chain is exact, not approximate), provenance makes constraint checking possible (you can verify that a value satisfies a constraint by tracing its derivation), and scoped knowledge bases make conversation tracking reliable (variables are stored in the right scope, not lost in a token stream).




Figure 1: Fig. 3: VDR-LLM-Prolog — Three layers, 705 tests, 27 modules, exact fractions throughout.

2. The Foundation: VDR Exact Arithmetic

VDR is an exact arithmetic system where every value is a finite tree of integer triples $[V, D, R]$ — value, denominator, remainder. The system was introduced in VDR-1 [HOWL-VDR-1-2026] and tested across 23 mathematical domains in VDR-2 and VDR-3 with 507 tests and zero computation errors.

2.1 What “Exact” Means

In VDR, the number one-half is $[1, 2, 0]$. The number one-third is $[1, 3, 0]$. Adding them gives $[5, 6, 0]$, which is exactly $5/6$. Not $0.833333\dots$ with trailing truncation error. The fraction $5/6$, with numerator 5 and denominator 6, stored as exact integers.

If you add $1/7$ and $1/13$ a hundred times, then subtract $1/13$ and $1/7$ a hundred times, you get back exactly $1/7$. Two hundred operations. Zero drift. The value at the end is structurally identical to the value at the start.

No floating-point system can do this. IEEE 754 double-precision floats accumulate rounding error at every operation. After 200 operations on $1/7$, the float result differs from $1/7$ by approximately 10^{-16} . This error is small but real, unmeasured, and cumulative.

2.2 The Remainder Slot

The third slot R — the remainder — is what makes VDR more than ordinary fraction arithmetic. When a value cannot be expressed cleanly in a given denominator frame, the remainder carries the exact leftover structure. The object `[1, 3, [1, 2, 0]]` means “one-third, with exact completion one-half.” The remainder is not error. It is not rounding residue. It is exact structure that a scalar system would discard.

This is what enables exact discrete calculus: the discretization artifact at every step size is an exact, inspectable, algebraically manipulable object. The discrete derivative of x^2 at $x=3$ with step $h=1/1000$ is exactly $6001/1000$ — not a float near 6, but the exact rational showing the discretization term as an algebraic fact.

2.3 Transcendental Reach

VDR handles transcendental constants through two mechanisms. Functional remainders wrap convergent rational series — each depth produces an exact rational that approaches the transcendental to arbitrary precision. The Q335 basis from MATH-4 [HOWL-MATH-4-2026] represents 22 fundamental constants (π , e , $\ln(2)$, $\sqrt{2}$, $\zeta(3)$, and 17 others) as single integers over the shared denominator 2^{335} , verified at 100 digits against mpmath. Adding $\pi + e$ under Q335 is one integer addition. The rounding error is 10^{66} times smaller than the Planck length.

2.4 What VDR Provides to This System

Every number in the VDR-LLM-Prolog system is an exact fraction. Every model weight. Every gradient. Every attention score. Every output probability. Every intermediate activation. Every training update. None of these are floats. None are approximations. None accumulate silent rounding error. When the system records that a weight has value $31/140$, the weight is exactly $31/140$, and the derivation chain that produced it (initialization at $1/4$, gradient of $-3/7$, SGD update with learning rate $1/10$) can be verified by exact arithmetic: $1/4 - (1/10)(-3/7) = 1/4 + 3/70 = 35/140 + 6/140 = \dots$ wait, let me recompute: $1/4 + (1/10)(3/7) = 1/4 + 3/70 = 17/70 + 3/70 \dots$ no. The SGD rule is $w_{\text{new}} = w_{\text{old}} - \text{lr} * \text{gradient}$. So $1/4 - (1/10)(-3/7) = 1/4 + 3/70 = 35/140 + 6/140 = 41/140$. If the recorded value is $31/140$, the derivation has an inconsistency, and the system can detect it exactly. With floats, this verification is impossible because the expected and actual values differ by accumulated rounding.

3. The Machine Learning Stack: VDR-4

VDR-4 [HOWL-VDR-4-2026] extended the arithmetic library into a complete machine learning system: 24 Python modules implementing exact-fraction softmax, reverse-mode autodiff, trainable layers, optimizers, attention, a transformer architecture, token sampling, and checkpointing. 198 tests, 196 passed, zero VDR computation errors.

3.1 Exact Softmax

The softmax function requires exponentials, which are transcendental. VDR computes $\exp$ as a truncated exact-fraction Taylor series: $\exp_N(x) = \sum x^k/k!$ for $k=0..N$. Every partial sum is an exact fraction. Softmax is then the ratio of exact exponentials over an exact sum. The probabilities sum to exactly 1 — not approximately, but as the fraction $1/1$.

For logits $[1, 2, 3]$, the softmax outputs are $64826368/720042809$, $176214841/720042809$, and $479001600/720042809$. Their sum is $720042809/720042809 = 1$. A rational surrogate softmax using a square-shift kernel (no exponentials) is also available: for the same logits with shift $c=4$, outputs are $4/29$, $9/29$, $16/29$. Sum is $29/29 = 1$.




Figure 2: Fig. 6: Softmax Exact Bars — VDR fractions sum to exactly 1. Float64 fractions do not.

3.2 Exact Autodiff

Reverse-mode automatic differentiation over VDR computation graphs. Every gradient is an exact fraction. $d(x^2)/dx$ at $x=3$ is exactly 6, not $5.999999\dots$. The chain rule, product rule, and quotient rule all produce exact fractions. MSE loss gradients are exact. The autodiff layer supports general computation graphs, not just fixed formulas.




Figure 3: Fig. 2: Gradient Chain — Forward exact fractions and backward exact gradients flowing through a 2-layer network.

3.3 Exact Transformer

A working tiny transformer language model: embedding lookup, self-attention with exact softmax, feedforward with ReLU, residual connections, logits head. Every attention weight sums to exactly 1 per row. Every intermediate value is an inspectable exact fraction. The model runs forward passes, computes exact logits, and produces exact output probability distributions.

3.4 What VDR-4 Provides to This System

The complete forward and backward pass of a language model, with every value exact. This means provenance recording is meaningful — the recorded values are the actual values, not approximations of the actual values. It means constraint checking is exact — “attention weights sum to 1” is verified by exact fraction addition. And it means checkpoints are bit-identical across platforms — the same model produces the same outputs everywhere, because there is no platform-dependent float rounding.

4. The Logic Layer: VDR-Prolog

The logic layer is a custom Prolog-style knowledge engine designed specifically for LLM provenance and constraint management. It is not a general-purpose Prolog implementation. It is a targeted system that stores exact VDR fractions as native term types, records derivation chains for every computed value, manages hierarchical scoped knowledge bases, and enforces declared constraints exactly.

4.1 Terms

A Term is the fundamental data unit. Every fact, rule, and query is built from Terms. The VDR-LLM-Prolog term types are:

Core Prolog types. Atoms (named constants like “softmax” or “layer_2”), variables (unification targets like ?X or ?Weight), and lists (ordered collections of terms).

VDR arithmetic types. Fractions (exact p/q with arbitrary-precision integers), fraction vectors (exact rational vectors for embeddings and hidden states), and fraction matrices (exact rational matrices for weight matrices and attention scores).

Q-basis types. Single integers over a shared power-of-two denominator, for compact representation of transcendental constants and compressed model parameters.

Structural references. Parameter paths (“layer.1.weight[0][0]”), layer references, token ids, token sequences — typed references into the model and data structures.

Provenance types. Derivation records (operation + inputs + output), constraints (type + bound + scope + status), checkpoints (step number + parameter state), gradients, losses, and training step numbers.

4.2 Facts

A Fact is a predicate with arguments. It asserts something is true.

```
parameter_value("layer.1.weight[0][0]", step(0), fraction(1, 4)).
```

This says: at training step 0, the weight at position [0][0] in layer 1 has the exact value 1/4.

```
attention_weight(doc(42), position(3), position(7), fraction(43545600, 59565131)).
```

This says: in document 42, the attention weight from position 3 to position 7 is exactly 43545600/59565131.

```
derived_from(attention_weight(doc(42), pos(3), pos(7)), softmax,
```

```
[score(doc(42), pos(3), pos(0)), score(doc(42), pos(3), pos(1)), ...]).
```

This says: the attention weight was derived by applying softmax to the list of attention scores.

4.3 Rules

A Rule is a logical implication: head :- body. If all the body conditions are satisfied, the head is true.

```
depends_on(X, Y) :- derived_from(X, _, Sources), member(Y, Sources).
```

```
depends_on(X, Y) :- derived_from(X, _, Sources), member(Z, Sources), depends_
```

This says: X depends on Y if Y is a direct input to X, or if Y is an input to something that X depends on. This is transitive dependency — the system can trace any value back to its ultimate sources through any number of intermediate steps.

```
weight_consistent(Param, Step) :-
```

```
parameter_value(Param, Step, V1),
```

```
PrevStep is Step - 1,
```

```
parameter_value(Param, PrevStep, V0),
```

```
gradient_at(Param, PrevStep, G),
```

```
updated_by(Param, PrevStep, sgd, lr(LR)),
```

```
V1 ::= V0 - LR * G.
```

This says: a weight is consistent at step N if its value equals the previous step's value minus the learning rate times the gradient. Because every value is an exact VDR fraction, this check is exact. It either holds or it doesn't. No tolerance.

4.4 Knowledge Bases

A Knowledge Base is a named collection of facts and rules. Knowledge bases are organized in a tree that mirrors the topic structure of the conversation. Each topic has its own KB. Child topics inherit from parent topics. The active topic determines which KBs are in the search scope.

4.5 Unification

When a query is evaluated, terms are unified by exact comparison. Two fraction terms unify if and only if they represent the same rational number (checked by cross-multiplication of exact integers). Two atom terms unify if and only if they have the same name. A variable term unifies with anything, binding the variable to the matched value. This is standard Prolog unification, extended with exact rational comparison instead of float comparison.

5. Scoped Knowledge Bases

5.1 The Scoping Principle

Every piece of knowledge has a home — a specific KB in a specific place in the topic tree. When the system searches for a fact, it searches only the KBs in the current scope: the active topic's KB, its parent's KB, up to the root, plus the global KB. Out-of-scope KBs are not searched at all.

This is lexical scoping applied to knowledge. If you are discussing story B, the facts from story A are not in scope. Not deprioritized. Not ranked lower. Not searched. The system cannot confuse the two because it never sees both simultaneously.




Figure 4: Fig. 5: KB Tree Scoping — Active topic vdr_llm illuminates its ancestry path. Story KBs are invisible to queries.

5.2 The KB Tree

kb_global (always in scope)

├─ axioms, system constraints, operational rules

|

├─ kb_project_vdr

```

|   |— language=python, version=3.8
|   |— kb_vdr_core (24 modules, 705 tests)
|   |— kb_vdr_llm (softmax, autodiff, transformer)
|   |— kb_math_series (Q335, 22 constants)
|
|— kb_story_a
|   |— setting=Pacific Northwest, year=2024
|   |— kb_characters_a (bob_age=32, alice_age=28)
|   |— kb_plot_a (chapter=3, conflict=missing artifact)
|
|— kb_story_b
|   |— setting=London, year=2026
|   |— kb_characters_b (bob_age=59, margaret_age=45)
|   |— kb_plot_b (chapter=1, conflict=double agent)

```

When story_b is active, queries resolve against kb_story_b, kb_characters_b, kb_plot_b, and kb_global. The entire story_a subtree is invisible. “How old is Bob?” returns 59, from kb_characters_b. No disambiguation needed. No risk of returning 32 from the wrong story.

5.3 Automatic Disambiguation

Ambiguous terms resolve by scope. “Bank” in a finance KB means financial institution. “Bank” in a geography KB means river edge. “Field” in a physics KB means force field. “Field” in an agriculture KB means farmland. The active scope selects the meaning. No heuristics. No “did you mean...” The scope is the disambiguator.

5.4 Explicit Cross-Scope Queries

When the user genuinely needs to reach across scopes, they can:

```
?- query_in(kb_characters_a, "bob_age", Age) .
```

```
% Explicitly query story A's Bob while story B is active

?- query_across("bob_age", Results).

% Search all KBs, return tagged results:

% [kb_characters_a: 32, kb_characters_b: 59]
```

Cross-scope queries are explicit and tagged. The results identify which KB each value came from. No silent mixing.

5.5 KB Activation

When topics change, KBs activate and deactivate. Switching from `story_a` to `story_b` deactivates `kb_story_a` and its children, activates `kb_story_b` and its children. The parent chain up to global stays active. Deactivation does not delete anything — the facts remain, they are just out of scope. Switch back and they reappear instantly.

6. Working Data Sets

6.1 Scoped Variable Storage

A working data set is a named, scoped collection of variable bindings attached to a topic. It is a nested dictionary where keys are names, values are exact VDR terms, and scope determines visibility.

When the user says “Bob is 32 years old,” the system does not just add tokens to the context window. It asserts a binding:

```
binding(kb_characters_a, "bob_age", number(32)).
```

This binding is stored in the correct scope, retrievable by exact lookup, overridable by a child scope, diffable against other scopes, and persistent until explicitly changed.

6.2 Inheritance and Shadowing

Working data sets form a tree that mirrors the topic tree. Variable lookup walks from the current dataset up to the root, like lexical scoping in a programming language.

If the project-level dataset says `language = python` and the `linalg` sub-dataset does not override it, then querying “language” from the `linalg` context returns “python” — inherited from the parent. If the `linalg` dataset sets `max_matrix_size = 50`, that value is local to `linalg` and does not appear in the parent scope.

6.3 History

Every binding set is logged with its turn number:

```
binding_history(kb_characters_a, "bob_age", number(32), turn(5)).
```

The system knows when every value was set. “When did we decide Bob was 32?” has an exact answer: turn 5.

6.4 Snapshots and Diffs

Working data sets can be snapshotted (frozen copy of all bindings at a point in time) and diffed (structured comparison between two datasets or two snapshots):

```
diff(snapshot_step_0, snapshot_step_1) →
```

```
  Changed: layer1_weight_0_0: 1/4 → 31/140
```

```
  Changed: layer1_bias_0: 0 → 3/70
```

```
  Unchanged: 4291 bindings
```

6.5 Type Discipline

Bindings carry exact VDR types. The system can enforce type constraints: “bob_age is always a number” prevents accidentally setting it to an atom. Schema enforcement on the working data set, checked by the constraint system.

7. The Constraint System

7.1 Constraints as First-Class Objects

A constraint is not a boolean flag or a comment in a system prompt. It is a structured object stored in the knowledge base, with a name, scope, status, condition, violation policy, activation time, and source.

```
constraint("sum_to_one", scope("axiom"), active,  
condition(forall(distribution(D), sum(D, 1))),  
on_violation("error"),  
source("mathematics")).
```

7.2 Four Constraint Domains

Operational rules. System-level constraints that are always on: exact arithmetic, Python 3.8 compatibility, runtime limits. These are verified before generating output.

Axioms. Mathematical invariants: probabilities sum to 1, attention weights are non-negative, gradient derivations are consistent. Axioms cannot be suspended. A violation is a bug.

Legal and policy constraints. External requirements that can be activated or deactivated by context: no medical diagnosis, cite sources, GDPR compliance. Transitions between active and suspended states are logged.

Project constraints. Specific to the current work: Zig 0.14 syntax, prefer i32, no changes beyond requested, use “remainder” not “residual.” These are the user preferences and project rules that currently exist only in system prompts and are easily forgotten.

7.3 Constraint Sets and Set Operations

Constraints form named sets. Sets can be enabled, disabled, intersected, unioned, and differenced as groups:

```
enable_set("vdr_project").  
  
disable_set("story_constraints").  
  
active_in_both("safety", "project", Result).
```

This gives the system the ability to reason about groups of constraints. “Enable all the VDR project constraints but disable the 30-second runtime limit for this specific task” is a set operation on constraint groups.

7.4 Exact Verification

Because every value is an exact VDR fraction, constraint checking is exact. “All attention weights in row 3 sum to 1” is checked by adding the exact fractions. The sum is either 1/1 or it is not. No tolerance. No epsilon. No “close enough.”

8. Topic and Conversation Tracking

8.1 Topics as Structured Objects

A topic has a name, a status (open, closed, parked, branched), timestamps for lifecycle events, a parent topic, child topics, and a list of pending items.

```
topic("vdr_llm", open, opened_at(turn(20)), parent("vdr_core")).
```

```
pending("vdr_llm", "implement_gaussian_elimination").  
pending("vdr_llm", "implement_cross_entropy_loss").
```

8.2 Lifecycle Management

Topics follow a lifecycle: open $\rightarrow$ work $\rightarrow$ close, or open $\rightarrow$ interrupted $\rightarrow$ park $\rightarrow$ resume $\rightarrow$ work $\rightarrow$ close. The system tracks this explicitly with rules:

A topic should close when it has no pending items and no open children. A topic should be parked when it has been open for many turns with no recent activity. Dangling topics — open but stale — can be identified by query.

8.3 Wrap and Unwrap

Parking a topic wraps it: the current state is summarized, the constraint set is deactivated, the working data set is preserved. Resuming a topic unwraps it: the state is restored, constraints are reactivated, pending items are listed.

The user says “let’s go back to the gym testing.” The system finds that topic parked, unwraps it, reactivates its constraint set, and lists pending items: “fix maxflow BFS, fix gym21 decay threshold.” The conversation continues from where it was parked, with full context, because the working data set was preserved exactly.

8.4 Triggered Functions

Topic state changes can trigger functions: loading relevant constraints when a topic opens, verifying all pending items when a topic closes, snapshotting state when a topic is parked. These triggers are Prolog rules, declarative and inspectable.

9. First-Class Knowledge Surfacing

9.1 The Principle

Nothing in the system is hidden from the system itself or from an authorized user. Every KB, every fact, every binding, every constraint, every derivation is queryable and directly surfaceable.

9.2 Direct Data Output

The system can return structured data without passing it through the language model’s token generation. When the user asks for data, the KB produces it directly. The LLM provides framing text. The data and the framing are visually and semantically distinct.

```
LLM: "Here's what we have for Bob in the London story:"
```

KB [kb_characters_b]:

bob_age: 59

bob_town: London

bob_occupation: retired professor

LLM: "Margaret is also in this story – want to see her details?"

The data block is a direct read from the KB. The LLM did not generate the number 59. The KB produced it. The data cannot be hallucinated because it is not generated — it is retrieved.

9.3 Addressable References

Every fact has an address: the KB path plus the fact identifier. The LLM can embed live references in its output:

LLM: "The learning rate is "

Reference: [kb_vdr_training/lr → 1/10]

LLM: ", set during training configuration."

The reference is a live link. If the value changes, the reference resolves to the new value. The LLM text is static commentary. The reference is dynamic truth.

9.4 Surfacing Modes

Narrative mode. LLM text with embedded KB references. Default conversational mode.

Table mode. Direct structured dump of a KB or query result. No LLM framing.

Tree mode. The KB hierarchy showing active and inactive scopes.

Provenance mode. The complete derivation chain for a specific value, from current state back to initialization or raw input.

Constraint mode. Active constraints with verification status.

Diff mode. Structured comparison between two points in time or two scopes.

9.5 The Permission Model

The owner of the system sees everything. End users see only what the permission rules allow. The LLM reasons over all KBs (it needs full access to produce correct outputs). The output filter determines what the user sees. For the owner, the filter is identity — everything is visible.

9.6 Self-Reference in Reasoning

The LLM queries its own KBs during reasoning. It does not need to remember constraints from a system prompt — it queries them from the constraint KB. It does not need to guess whether code is Python 3.8 compatible — it checks the constraint condition against the actual code. The reasoning is grounded in structured facts, not in pattern-matched context tokens.

10. The Integration Architecture

10.1 Three Layers

Layer 1: Arithmetic (VDR). Every number is an exact fraction. Zero drift. Zero silent truncation. Provides the computational foundation.

Layer 2: Logic (Prolog). Every value has a derivation. Every derivation is a chain of facts and rules. Every constraint is declared and verified. Provides the reasoning and provenance foundation.

Layer 3: Conversation (Scoped KBs + Working Data + Topics + Constraints). Knowledge is scoped to topics. Variables persist in working data sets. Topics have life-cycle state. Constraints are activated and deactivated with scope changes. Provides the structured memory foundation.




Figure 5: Fig. 4: Architecture Stack — VDR arithmetic, Prolog logic, and conversation management with data flowing between layers.

10.2 The Data Flow

User input (text)

↓

Tokenization → token ids (exact integers)

↓

Embedding lookup → exact fraction vectors

↓ (fact asserted: embedding derived from token + table)

Attention scores → exact fraction matrix

↓ (fact asserted: scores derived from $Q \cdot K^T$)

Softmax → exact fraction weights, sum exactly to 1

↓ (fact asserted: weights derived from softmax(scores))

↓ (constraint checked: sum_to_one verified)

Value mixing → exact fraction vectors

↓ (fact asserted: mixed derived from weights · values)

Feedforward → exact fraction vectors

↓ (fact asserted: ff_out derived from linear + relu)

Logits → exact fraction scores over vocabulary

↓ (fact asserted: logits derived from logits_head · hidden)

Output probabilities → exact fractions summing to 1

↓ (constraint checked: valid_distribution verified)

Token sampling → selected token

↓ (fact asserted: sampled token, seed, probability)

Output to user

↓ (KB data surfaced directly where applicable)

↓ (LLM generates framing text around KB data)

Every arrow is an exact VDR operation. Every parenthetical is a Prolog fact assertion.
Every constraint check is exact.




Figure 6: Fig. 1: Provenance Chain — Raw text through tokenization, attention, softmax to output, with exact fractions at every stage.

10.3 The Training Flow

Forward pass (as above, all exact)

↓

Loss computation → exact fraction

↓ (fact asserted: loss value, derivation from predictions and targets)

Backward pass → exact fraction gradients for every parameter

↓ (fact asserted: gradient for each parameter at this step)

Optimizer step → exact fraction parameter updates

↓ (fact asserted: new parameter value, update derivation)

↓ (constraint checked: weight_consistent verified)

↓ (constraint checked: denominator_bound verified)

Checkpoint (if policy says so)

↓ (all parameter values serialized as exact fractions)

↓ (provenance chain complete from initialization to current step)

10.4 The Conversation Flow

User speaks

↓

Topic identification → active topic determined

↓ (scoped KBs activated, out-of-scope KBs invisible)

Working data set → current bindings available

↓

Constraint set → active constraints loaded

↓

Intent recognition → is this a KB query or a generation request?

↓

If KB query:

Run Prolog query against in-scope KBs

Surface structured results directly

LLM generates framing text

↓

If generation request:

LLM generates response

References KB facts where relevant

Constraint system checks output

New bindings asserted if user declares facts

Pending items updated

Topic state updated

11. Implementation Specification

11.1 Term Types

```
TermType = enum {  
    // Core Prolog  
    atom,  
    variable,  
    list,  
    // VDR Arithmetic  
    fraction,  
    fraction_vec,  
    fraction_mat,  
    // Q-Basis Compressed  
    qbasis,  
    qbasis_vec,  
    qbasis_mat,  
    // References  
    parameter,  
    layer,  
    token,  
    token_seq,
```

```

// Provenance
derivation,
constraint,
checkpoint,
gradient,
loss,
step,
// Conversation
topic,
binding,
scope,
constraint_set,
};

```

11.2 Fact Structure

```

Fact = struct {
predicate: Text,
args: []Term,
kb_source: Text,          // which KB this fact belongs to
asserted_at: i32,         // turn or step when asserted
derivation: ?Derivation, // how this fact was produced (optional)
};

```

11.3 Rule Structure

```
Rule = struct {  
    head: Fact,          // conclusion (full Fact, not just name)  
    body: []Fact,        // conditions (all must be satisfied)  
    kb_source: Text,      // which KB this rule belongs to  
};
```

11.4 Knowledge Base Structure

```
KnowledgeBase = struct {  
    name: Text,  
    facts: FactSet,  
    rules: RuleSet,  
    parent: ?Text,        // parent KB name, or null for root  
    children: []Text,     // child KB names  
    topic: ?Text,         // associated topic, if any  
    visibility: enum { public, internal, owner_only },  
    frozen: bool,         // if true, read-only (snapshot)  
};
```

11.5 Working Data Set Structure

```
WorkingDataSet = struct {  
    name: Text,  
    topic: Text,  
    parent: ?Text,
```

```

bindings: HashMap(Text, Term),

history: []BindingEvent, // timestamped log of all changes

frozen: bool,

};

BindingEvent = struct {

key: Text,

old_value: ?Term,

new_value: Term,

turn: i32,

};

```

11.6 Constraint Structure

```

Constraint = struct {

name: Text,

scope: enum { operational, axiom, legal, project, conversation },

status: enum { active, suspended, violated, satisfied, parked },

condition: Fact,           // the logical condition to check

on_violation: enum { warn, block, log, escalate, error },

activated_at: i32,

last_checked: i32,

source: Text,              // who or what declared this

};

```

11.7 Topic Structure


```

Topic = struct {
    name: Text,
    status: enum { open, closed, parked, branched },
    opened_at: i32,
    closed_at: ?i32,
    parked_at: ?i32,
    parent: ?Text,
    children: []Text,
    pending: []Text,
    constraint_set: ?Text,    // associated constraint set
    working_data: ?Text,     // associated working data set
};

```

11.8 The Query Engine

The query engine is a standard Prolog-style depth-first search with backtracking, modified for scoped KB search:

1. Determine in-scope KBs from the active topic's ancestry chain.
2. For each KB in scope order (current $\rightarrow$ parent $\rightarrow$... $\rightarrow$ global), search for facts matching the query predicate and attempt unification.
3. First successful unification in scope order wins (cut semantics for scoped search).
4. For cross-scope queries, search all KBs and tag results with their source.
5. For rule evaluation, recursively evaluate body conditions using the same scoped search.

11.9 The Bridge

The Python-Prolog bridge passes exact rationals between VDR (Python) and the Prolog knowledge engine. VDR Fractions map directly to Prolog fraction terms. The bridge is the integration point where VDR computations assert their results as Prolog facts and Prolog constraints trigger VDR recomputations.

12. Memory Management




Figure 7: Fig. 8: Memory Tiers — Five levels from persistent (KB) to transient (pruned), with Prolog retention rules at each tier.

12.1 Tiered Retention

Tier 1: Persistent. Model architecture, rules, constraints, hyperparameters. Never freed.

Tier 2: Checkpoint. Parameter values at declared checkpoint steps. Grows linearly with checkpoint count.

Tier 3: Step-transient. Gradients, updates, activations for the current training step. Pruned after the step is committed, unless the retention policy says otherwise.

Tier 4: Batch-transient. Per-batch data provenance. Pruned after gradient propagation, unless flagged for audit.

12.2 Retention Policies as Rules

```
retain_step(Step) :- Step mod checkpoint_interval == 0.
retain_step(Step) :- loss_anomaly_at(Step).
retain_step(Step) :- constraint_violated_at(_, Step).
retain_batch(Batch) :- contains_unknown_token(Batch).
```

```
retain_batch(Batch) :- output_anomaly(Batch) .
```

```
prune_step(Step) :- not(retain_step(Step)) .
```

```
prune_batch(Batch) :- not(retain_batch(Batch)) .
```

The pruning policy is itself a set of Prolog rules. It is declarative, inspectable, and modifiable without changing code.

13. What This System Provides

13.1 For Arithmetic

Zero drift. Zero silent truncation. Every number is what it is. The number you see is the number that was computed is the number that should have been computed.

13.2 For Provenance

Every value has a derivation chain. The chain is complete, exact, and queryable. Any value can be traced back to its ultimate sources through any number of intermediate steps.

13.3 For Constraints

Every constraint is a declared, verified, first-class object. Constraints can be grouped, activated, suspended, and queried. Verification is exact. Violations are detectable and traceable.

13.4 For Conversation

Topics are tracked objects with lifecycle state. Working data persists in scoped datasets. Nothing is forgotten unless explicitly pruned. Topic switches preserve context exactly. Variable resolution is scoped and unambiguous.

13.5 For Trust

The user can verify any output by querying its provenance. The data surfaced from KBs is not generated by the LLM — it is retrieved from structured storage. The LLM provides framing. The KB provides truth. The two are clearly separated in the output.

14. What This System Does Not Provide

14.1 Scale

Exact fractions grow in denominator size through operations. A production-scale language model with billions of parameters is not practical in exact fractions without aggressive Q-basis compression and periodic reprojection. This system is designed for research-scale models, auditability demonstrations, and the establishment of architectural patterns that could be adapted to larger scale with controlled precision trade-offs.

14.2 Speed

Exact fraction arithmetic is slower than float arithmetic on hardware with native float units. The system trades speed for exactness, provenance, and trust. For applications where trust and auditability matter more than throughput — regulatory compliance, safety-critical systems, scientific computation — this trade-off is correct.

14.3 Standard Transformer Fidelity

The architecture uses rational approximations of transcendental functions (truncated Taylor exp, rational surrogate softmax, ReLU instead of GELU). These are not identical to standard transformer operations. A VDR-native architecture optimized for exact fractions may look different from a standard transformer, and that may be the better path.

15. Falsification Criteria

F1. If any VDR computation produces an incorrect exact rational from correct inputs, the arithmetic layer is wrong. 705 tests across 23 domains have not produced one such error.

F2. If a provenance chain recorded in the Prolog KB is inconsistent with the actual computation (the recorded derivation does not match the recorded values under exact arithmetic), the provenance layer has a bug.

F3. If a constraint marked as “satisfied” is in fact violated (the exact values do not satisfy the exact condition), the constraint checking layer has a bug.

F4. If a scoped KB query returns results from an out-of-scope KB without explicit cross-scope request, the scoping layer has a bug.

F5. If a working data set binding is lost or corrupted during a topic park/resume cycle, the persistence layer has a bug.

Each criterion is testable by exact comparison of exact values. No tolerances. No “close enough.”

16. Implementation Roadmap

Phase 1: Core Prolog Engine

Port the Term, Fact, Rule, KnowledgeBase structures. Implement unification with exact VDR fraction comparison. Implement depth-first search with backtracking. Implement scoped KB search. Test with simple fact assertion and query.

Phase 2: Constraint System

Implement Constraint structure. Implement constraint checking against exact VDR values. Implement constraint sets with enable/disable. Test constraint verification on softmax outputs and weight update consistency.

Phase 3: VDR-ML Integration

Instrument the VDR transformer forward pass to assert provenance facts. Instrument the backward pass and optimizer. Test that the recorded provenance chain is consistent with the exact computed values.

Phase 4: Working Data Sets and Topics

Implement WorkingDataSet with scoped bindings and inheritance. Implement Topic with lifecycle management. Implement wrap/unwrap. Test that topic switches preserve working data exactly.

Phase 5: Surfacing Layer

Implement direct KB data output alongside LLM text. Implement addressable references. Implement surfacing modes (table, tree, provenance, constraint, diff). Implement permission model.

Phase 6: Integration Testing

End-to-end test: raw text input → tokenization → forward pass → provenance recording → constraint verification → output with KB surfacing → topic tracking → working data update. Every step exact. Every fact queryable. Every constraint verified.

17. Conclusion

This paper specifies a language model architecture where knowledge has provenance, arithmetic has exactness, constraints have enforcement, and conversation has structure. The specification integrates four years of work: exact arithmetic that has passed 705 tests with zero computation errors, a machine learning stack that computes exact softmax and exact gradients, a transcendental basis that represents 22 constants as integers, and a logic engine that records derivation chains and verifies constraints.

The result is not a faster language model. It is not a more capable language model. It is a more trustworthy language model — one where every value can be traced to its source, every constraint can be verified exactly, every piece of knowledge has a home and a scope, and the user can inspect any part of the system’s state by querying the knowledge base directly.

The specification is complete. The arithmetic foundation exists and is tested. The ML stack exists and is tested. The logic engine is specified and ready for implementation. The five phases of the implementation roadmap are concrete engineering work packages with clear deliverables and testable outcomes.

Data with provenance. Arithmetic with exactness. Constraints with enforcement. Conversation with structure. That is what this system provides.

Addendum to VDR-5: Constraints as KB-Local and Account Hierarchy

A1. Constraints Belong to Knowledge Bases

The VDR-5 specification described constraints as a separate system that references KBs. This addendum corrects that design: constraints belong inside the KB they govern, not in a separate registry.

A1.1 The Principle

A knowledge base is a self-contained unit of knowledge. It has facts (what is true), rules (what follows from what), and constraints (what must hold). Separating constraints from their KB creates the same problem that separating documentation from code creates — they drift apart. The constraints for a model’s attention weights should live in the same KB as the attention weight facts. The constraints for a user’s spending limits should live in the same KB as the user’s financial data.

A1.2 Revised KB Structure

```
KnowledgeBase = struct {  
  
  name: Text,  
  
  facts: FactSet,  
  
  rules: RuleSet,
```

```

constraints: []Constraint, // constraints local to this KB

parent: ?Text,

children: []Text,

topic: ?Text,

working_data: ?WorkingDataSet,

visibility: enum { public, internal, owner_only },

frozen: bool,

};

```

Constraints are now a field of the KB, not a separate global structure. When a KB is activated, its constraints are activated. When a KB is deactivated (topic switch), its constraints are deactivated. When a KB is snapshotted, its constraints are included. When a KB is pruned, its constraints go with it.

A1.3 Constraint Inheritance

Constraints inherit through the KB tree, the same way facts and bindings inherit. A child KB's constraints add to (and can override) its parent's constraints. The effective constraint set at any point is the union of constraints from the current KB up to the root.

```

Rule: effective_constraints(KB, All) :-
    kb_constraints(KB, Local),
    kb_parent(KB, Parent),
    (Parent = none ->
        All = Local ;
        effective_constraints(Parent, ParentConstraints),
        merge_constraints(Local, ParentConstraints, All)).

Rule: merge_constraints(Local, Parent, Merged) :-
    % Local constraints with same name override parent

```

```

findall(C, (member(C, Local)), LocalList),

findall(C, (member(C, Parent),

            constraint_name(C, N),

            not((member(L, Local), constraint_name(L, N)))),

        InheritedList),

append(LocalList, InheritedList, Merged).

```

A global KB has the constraint “all fractions must be exact.” A project KB inherits that and adds “Python 3.8 compatible.” A sub-project KB inherits both and adds “runtime under 30 seconds.” At the sub-project level, all three constraints are active. Disabling the 30-second constraint at the sub-project level does not affect the parent’s constraints.

A1.4 Constraint Scoping Follows KB Scoping

When the active topic changes, the active KBs change, and therefore the active constraints change. This is automatic. There is no separate constraint activation step. The constraint set is always exactly the effective constraints of the in-scope KBs.

Switching from the VDR project to story B deactivates the VDR project’s constraints (Python 3.8 compat, exact arithmetic, naming conventions) and activates story B’s constraints (character consistency, plot continuity, setting accuracy). The global constraints (system operational rules) remain active because the global KB is always in scope.

A1.5 Portable KBs

Because a KB now contains its own constraints, it is fully self-describing. You can export a KB (facts + rules + constraints + working data) and import it into a different system. The receiving system does not need to know what constraints to apply — they travel with the data. A model checkpoint KB includes not just the parameter values but the invariants those parameters must satisfy. A dataset KB includes not just the data but the validation rules for that data.

A2. User Accounts as Knowledge Bases




Figure 8: Fig. 7: Constraint Inheritance — Global through org, dept, team to user. 9 effective constraints from 5 hierarchy levels.

A2.1 The Insight

A user account is a collection of facts (name, preferences, history), rules (what this user is allowed to do), and constraints (what limits apply). This is exactly what a KB is. Therefore a user account should be a KB.

A2.2 Account KB Structure

```
KB: user_alice

facts:

user_name("Alice Chen")

user_role(end_user)

user_created(2026-03-15)

preference("language", "English")

preference("formatting", "minimal")
```

```

preference("expertise", "expert_zig")

constraints:

constraint("rate_limit", scope("operational"), active,
    condition(queries_per_hour =< 100),
    on_violation("throttle"))

constraint("no_internal_access", scope("policy"), active,
    condition(not(accesses(kb, visibility(internal)))),
    on_violation("block"))

constraint("max_context", scope("operational"), active,
    condition(context_tokens =< 200000),
    on_violation("truncate"))

working_data:

current_project: "vdr_development"

active_topics: ["vdr_llm", "prolog_integration"]

parked_topics: ["story_b"]

children: [kb_alice_vdr_project, kb_alice_story_a, kb_alice_story_b]

```

Everything about Alice — her identity, her permissions, her preferences, her constraints, her active work, her conversation history — is in her KB. The KB is the account.

A2.3 Group Accounts as Parent KBs

Groups work by KB inheritance. A group account is a parent KB. User accounts are child KBs. Group constraints apply to all members through inheritance.

```

KB: group_engineering_team

    facts:

group_name("Engineering Team")

group_created(2026-01-01)


    constraints:

constraint("code_review_required", scope("policy"), active,
    condition(all_code_changes_reviewed),
    on_violation("block"))

constraint("no_production_access", scope("policy"), active,
    condition(not(accesses(production_db))),
    on_violation("block"))

constraint("approved_languages", scope("project"), active,
    condition(language_in([python, zig, rust])),
    on_violation("warn"))


    children: [user_alice, user_bob_engineer, user_carol]

```

Alice inherits the engineering team's constraints. She cannot push code without review. She cannot access the production database. She can only use approved languages. These constraints come from her parent KB (the group), not from a separate access control system. They are facts in the KB tree, queryable and inspectable.

A2.4 Nested Group Hierarchy

Groups can nest. A company has departments. Departments have teams. Teams have members. Each level is a KB with constraints that propagate downward.

```
KB: org_acme_corp
```

```

└─ constraint: "gdpr_compliance" (active, legal)
└─ constraint: "data_retention_7_years" (active, legal)
|
└─ KB: dept_engineering
|   └─ constraint: "code_review_required" (active, policy)
|   └─ constraint: "approved_languages" (active, project)
|       |
|       └─ KB: team_backend
|           └─ constraint: "database_access_logged" (active, operational)
|               └─ KB: user_alice (inherits all above)
|                   └─ KB: user_bob_engineer (inherits all above)
|                       |
|                       └─ KB: team_frontend
|                           └─ constraint: "no_database_direct" (active, operational)
|                               └─ KB: user_carol (inherits all above)
|
└─ KB: dept_legal
    └─ constraint: "client_privilege" (active, legal)
        └─ KB: user_dave_lawyer

```

(inherits org + legal constraints, NOT engineering constraints)

Alice (backend engineer) inherits: GDPR compliance, data retention, code review, approved languages, database access logged. Five constraints from four levels of the hierarchy.

Carol (frontend engineer) inherits: GDPR compliance, data retention, code review, approved languages, no direct database access. The last constraint differs from Alice's because Carol is in a different team KB.

Dave (lawyer) inherits: GDPR compliance, data retention, client privilege. He does not inherit engineering constraints because his KB path goes through dept_legal, not dept_engineering.

The constraint set for any user is determined entirely by their position in the KB tree. No separate ACL system. No role-based access control tables. The KB tree is the access control structure.

A2.5 Constraint Override at Lower Levels

A child KB can override a parent constraint by declaring a constraint with the same name and a different configuration.

```
KB: team_backend

    constraint: "approved_languages" overrides parent
condition(language_in([python, zig, rust, go]))

// backend team also allows Go
```

The merge rule from A1.3 handles this: local constraints with the same name as parent constraints replace the parent version. The backend team's approved language list includes Go. The frontend team's does not, because it inherits the department-level list unchanged.

Override is logged as a fact:

```
constraint_override("team_backend", "approved_languages",
overrides("dept_engineering"),
reason("Go needed for microservices")).
```

The override has provenance. "Why can the backend team use Go?" is answered by querying the override fact.

A2.6 Per-User Constraint Customization

Individual users can have constraints added to their account KB by administrators, by policy rules, or by their own preferences.

```
KB: user_alice

    constraint: "30_sec_runtime" (active, preference, source(self))

// Alice's personal preference
```

```

    constraint: "elevated_access_production" (active, policy, source(admin))
// Admin granted Alice production access, overriding team default
condition(accesses(production_db, read_only))
on_violation("log")

    constraint: "training_required_ml" (active, policy, source(hr_system))
// HR system added this after Alice moved to ML team
condition(completed_training("ml_safety_2026"))
on_violation("block_ml_features")

```

Some constraints come from Alice herself (personal preferences). Some come from administrators (access grants). Some come from automated systems (training requirements). All are facts in Alice's KB. All are queryable. All have provenance showing who added them and when.

A2.7 Querying the Account Hierarchy

```

?- effective_constraints(user_alice, Constraints).

% Returns all constraints Alice is subject to:

% [gdpr_compliance, data_retention, code_review,
% approved_languages(with Go), database_access_logged,
% 30_sec_runtime, elevated_access_production, training_required_ml]

?- constraint_source(user_alice, "elevated_access_production", Source).

% Source = admin, granted(2026-04-15)

?- users_with_constraint("no_production_access", Users).

% Users = [bob_engineer, carol]      (not Alice -
hers was overridden)

?- constraints_from_level(user_alice, "org_acme_corp", OrgConstraints).

```

```
% OrgConstraints = [gdpr_compliance, data_retention]

?- constraint_violations(user_alice, Violations).

% Checks all effective constraints, returns any that are violated
```

A2.8 Account KBs Carry Conversation State

Because the user account is a KB, the user’s conversation state (active topics, parked topics, working data sets) lives inside the KB. When Alice logs in, her KB is loaded, her active topics are restored, her working data sets are available, and her constraints are in effect. When she switches to a different project, her project-level KB activates and the old one deactivates, with all the scoping behavior described in VDR-5 Section 5.

If Alice has two simultaneous sessions (desktop and mobile), both sessions reference the same KB. Changes in one session are visible in the other. Working data is shared. Topic state is shared. Constraints are shared. The KB is the single source of truth for Alice’s account state.

A3. Implications

A3.1 Self-Describing Data

Every KB carries its own constraints. A KB exported to another system includes the rules for validating its contents. The receiving system does not need external documentation to understand what the data must satisfy. The constraints are inside.

A3.2 Auditable Access Control

Every constraint has provenance. “Why can’t Carol access the database?” traces through the KB tree: team_frontend has constraint “no_database_direct,” inherited from dept_engineering’s policy, which was declared by the CTO on a specific date. The access control decision is not a black-box policy engine. It is a queryable fact chain.

A3.3 Portable Accounts

A user account KB can be exported and imported. Alice changes companies. Her personal preference constraints (formatting, language, runtime limits) travel with her. Her company-specific constraints (GDPR compliance, code review) are stripped because they belonged to parent KBs she is no longer a child of. The KB tree handles this naturally — detaching from a parent KB removes the inherited constraints.

A3.4 Testable Constraint Configurations

Because constraint sets are computable from the KB tree, they can be tested before deployment. “If we add this constraint to the engineering department KB, which users are affected?” is a query:

```
?- hypothetical_add(dept_engineering, new_constraint, Affected).
```

```
% Returns all users under dept_engineering who would gain this constraint
```

“If we override this constraint for team_backend, does it create a conflict with the legal department’s requirements?” is another query:

```
?- hypothetical_override(team_backend, "data_access", new_condition, Conflict).
```

```
% Checks whether the new condition violates any ancestor constraints
```

These are not hypothetical features. They are direct consequences of constraints being facts in a logic programming knowledge base. The query engine already exists. The scoping rules already exist. The provenance tracking already exists. Making constraints KB-local and making accounts KBs is a structural rearrangement, not a new capability.

A4. Revised Specification Summary

The VDR-5 specification is amended as follows:

Constraints move from a separate system into the KB structure. Every KB has a constraints field. Constraints inherit through the KB tree. Effective constraints at any point are the union of local and inherited constraints, with same-name overrides.

User accounts are KBs. A user’s identity, preferences, permissions, conversation state, and working data are facts in their account KB. Group accounts are parent KBs. The organizational hierarchy is a KB tree. Access control is constraint inheritance.

Three new fields in the KB structure: `constraints: []Constraint`, `working_data: ?WorkingDataSet`, and `owner: ?Text` (the user or group that owns this KB).

One new invariant: A constraint in a child KB with the same name as a constraint in a parent KB overrides the parent constraint for that child and all its descendants. The override is logged with provenance.

Everything else in VDR-5 remains unchanged. The scoping rules, the query engine, the surfacing modes, the topic management, the provenance tracking — all work exactly as specified, now with constraints embedded in the KBs they govern and accounts embedded in the KB tree they belong to.

END ADDENDUM

Appendix A: Cumulative Project Statistics

Paper	Contribution	Tests	Errors
VDR-1	Core arithmetic, discrete calculus	68	0
VDR-2	15 domain gyms	282	0
VDR-3	8 domain gyms, MATH integration	157	0
VDR-4	24-module ML stack	198	0
VDR-5	Architecture specification	—	—
Total		705	0

Appendix B: Module Inventory

Module	Layer	Purpose
vdr.py	Arithmetic	Core triple, remainder, normalization
active mul.py	Arithmetic	Active multiplication/division
fn.py	Arithmetic	Functional remainders, discrete calculus
linalg.py	Arithmetic	Vec, Mat, parse, serialize
export.py	Arithmetic	Lossy boundary
exp.py	Transcendental	Exact-fraction exponential
logarithm.py	Transcendental	Exact-fraction logarithm
softmax.py	ML	Exact softmax, rational surrogate
autodiff.py	ML	Reverse-mode exact differentiation
nn.py	ML	Linear, ReLU, Sequential
losses.py	ML	MSE, L1, hinge
optim.py	ML	SGD, Momentum
rng.py	Infrastructure	Deterministic PRNG
init.py	Infrastructure	Rational initialization
sampling.py	Infrastructure	Categorical, top-k, nucleus
datasets.py	Infrastructure	Tokenization, batching
metrics.py	Infrastructure	Accuracy, denominator tracking
checkpoint.py	Infrastructure	Exact save/load
basis.py	Infrastructure	Q-basis shared denominator

Module	Layer	Purpose
tensor.py	Infrastructure	Batched operations
attention.py	Architecture	Scores, masking, weighting
transformer.py	Architecture	Embedding, blocks, LM head
trainer.py	Architecture	Training loops
prolog.py	Logic (VDR-5)	KB engine, constraints, provenance
conversation.py	Logic (VDR-5)	Topics, working data, scoping
surfacing.py	Logic (VDR-5)	Direct output, references, modes

24 existing modules. 3 new modules specified by VDR-5. 27 total.

Appendix C: Paper Series

Paper	Registry	Central Result
VDR-1	[HOWL-VDR-1-2026]	Exact arithmetic in irreducible triple form
VDR-2	[HOWL-VDR-2-2026]	15 domains, 282 tests, chaos boundary
VDR-3	[HOWL-VDR-3-2026]	23 domains, transcendental integration
VDR-4	[HOWL-VDR-4-2026]	24-module ML stack, working transformer
VDR-5	[HOWL-VDR-5-2026]	Prolog provenance, constraints, scoped KBs
MATH-3	[HOWL-MATH-3-2026]	Elliptic integrals, Borwein acceleration
MATH-4	[HOWL-MATH-4-2026]	Q335 universal basis, 22 constants

END HOWL-VDR-5-2026

Registry: [\[HOWL-VDR-5-2026\]](#)

Status: Specification complete. Implementation roadmap defined.

Domain: Applied Philosophy / Exact Machine Learning / Knowledge Systems

Central Result: Architecture specification for a language model with exact arithmetic, logical provenance, first-class constraints, scoped knowledge bases, structured conversation tracking, and direct data surfacing.

Foundation: VDR-1 through VDR-4, MATH-3, MATH-4

Key Claim: Data provenance, constraint enforcement, and conversational state tracking are architectural requirements, not afterthought features. This paper specifies what that architecture looks like.

Falsification: Five specific criteria, all testable by exact comparison.

VDR-5 Appendix Tables

Complete Reference for VDR-LLM-Prolog Implementation

Appendix D: Complete Term Type Reference

D.1 Term Types and Storage

TermType	Category	Primary Storage	Secondary Storage	Example
atom	Core Prolog	Text name	—	atom(“softmax”), atom(“layer 2”)
variable	Core Prolog	Text name	—	variable(“?X”), variable(“?Weight”)
list	Core Prolog	[]Term	—	list([atom(“a”), number(3)])
fraction	VDR Arithmetic	BigInt numerator	BigInt denominator	fraction(31, 140)
fraction vec	VDR Arithmetic	[]Fraction	—	fraction vec([1/2, 1/3, 1/7])
fraction mat	VDR Arithmetic	[][]Fraction	—	fraction mat([[1/2, 1/3], [1/7, 1/4]])
qbasis	Q-Basis	BigInt numerator	i32 exponent k	qbasis(219886425873..., 335)
qbasis vec	Q-Basis	[]BigInt	i32 exponent k	qbasis vec([p1, p2, p3], 335)

TermType	Category	Primary Storage	Secondary Storage	Example
qbasis mat	Q-Basis	[]BigInt	i32 exponent k	qbasis mat([[p1,p2],[p3,p4]], 335)
parameter layer	Reference	Text path	—	parameter(“layer.1.weight[0][0]”)
	Reference	Text name	i32 index	layer(“attention”, 2)
token	Reference	i32 id	Text surface	token(42, “the”)
token seq	Reference	[]i32	—	token seq([42, 17, 93, 5])
derivation	Provenance	Text operation	[]Term inputs	derivation(“softmax”, [score1, score2])
constraint	Provenance	Constraint struct	—	constraint(“sum to one”, axiom, active)
checkpoint gradient	Provenance	i32 step	—	checkpoint(100)
	Provenance	Fraction value	Text param path	gradient(- 3/7, “layer.1.weight[0][0]”)
loss	Provenance	Fraction value	Text loss type	loss(17/2, “mse”)
step topic	Provenance	i32 step number	—	step(42)
	Conversation	Topic struct	—	topic(“vdr llm”, open)
binding	Conversation	Text key + Term val	Text dataset	binding(“bob age”, num- ber(32))
scope	Conversation	Text name	—	scope(“legal”), scope(“project”)
constraint set	Conversation	Text name + []Text	—	constraint set(“safety”, [...])

D.2 Unification Rules by Term Type

TermType	Unification Method	Exact?	Notes
atom	String equality	Yes	“softmax” = “softmax”

TermType	Unification Method	Exact?	Notes
variable	Binds to any term	Yes	?X unifies with anything, binds ?X
list	Element-wise recursive	Yes	Same length, each element unifies
fraction	Cross-multiplication: $ad == bc$	Yes	$1/2 = 2/4$ by exact integer comparison
fraction vec	Element-wise fraction comparison	Yes	Same dimension, each element exact
fraction mat	Element-wise fraction comparison	Yes	Same shape, each element exact
qbasis	Integer equality + exponent equality	Yes	Same p and same k
qbasis vec	Element-wise integer + shared exponent	Yes	Same dimension, same k, each p equal
qbasis mat	Element-wise integer + shared exponent	Yes	Same shape, same k
parameter	Path string equality	Yes	"layer.1.weight[0][0]" = same string
layer	Name + index equality	Yes	Both must match
token	Integer id equality	Yes	Token 42 = token 42
token seq	Element-wise integer equality	Yes	Same length, each id matches
derivation	Operation + input list equality	Yes	Recursive on inputs
constraint	Name equality	Yes	For override resolution

TermType	Unification Method	Exact?	Notes
checkpoint	Step number equality	Yes	Step 100 = step 100
gradient	Fraction value + path equality	Yes	Both must match
loss	Fraction value + type equality	Yes	Both must match
step	Integer equality	Yes	Step 42 = step 42
topic	Name equality	Yes	For lookup and state tracking
binding	Key equality (value separate)	Yes	Key match, value via unification
scope	Name equality	Yes	“legal” = “legal”
constraint set	Name equality	Yes	For set operations

Appendix E: Complete Fact Predicate Reference

E.1 Model Architecture Facts

Predicate	Arity	Arg Types	Example	KB Location
model name	1	atom	model name(“tiny transformer v1”)	kb model
model vocab size	1	number	model vocab size(3)	kb model
model embed dim	1	number	model embed dim(2)	kb model
model num layers	1	number	model num layers(1)	kb model
model num heads	1	number	model num heads(1)	kb model
model context length	1	number	model context length(3)	kb model

Predicate	Arity	Arg Types	Example	KB Location
model activation	1	atom	model activation("relu")	kb model
model softmax type	1	atom	model softmax type("surrogate")	kb model
model softmax depth	1	number	model softmax depth(12)	kb model
layer type	2	number, atom	layer type(0, "attention")	kb model
layer input dim	2	number, number	layer input dim(0, 2)	kb model
layer output dim	2	number, number	layer output dim(0, 2)	kb model

E.2 Parameter Facts

Predicate	Arity	Arg Types	Example	KB Location
parameter value	3	parameter, step, fraction	parameter value("layer.1.weight[0][0]", step(0), fraction(1,4))	kb params
parameter shape	2	parameter, list	parameter shape("layer.1.weight", [2, 2])	kb params
parameter count	1	number	parameter count(6)	kb params
initialized from	3	parameter, atom, number	initialized from("layer.1.weight[0][0]", "xavier", seed(42))	kb params
gradient at	3	parameter, step, fraction	gradient at("layer.1.weight[0][0]", step(0), fraction(-3,7))	kb training
updated by	4	parameter, step, atom, fraction	updated by("layer.1.weight[0][0]", step(0), "sgd", lr(fraction(1,10)))	kb training

Predicate	Arity	Arg Types	Example	KB Location
param denominator	3	parameter, step, number	param denominator("layer.1.weight[0][0]", step(1), 140)	kb training

E.3 Data Provenance Facts

Predicate	Arity	Arg Types	Example	KB Location
raw text	2	atom, atom	raw text(doc(1), "the cat sat")	kb data
tokenized	2	atom, token seq	tokenized(doc(1), token seq([42, 17, 93]))	kb data
token id	3	atom, number, number	token id(doc(1), pos(0), id(42))	kb data
embedding	3	atom, number, fraction vec	embedding(doc(1), pos(0), fraction vec([3/7, 1/2]))	kb inference
attention score	4	atom, number, number, fraction	attention score(doc(1), pos(0), pos(1), fraction(5,12))	kb inference
attention weight	4	atom, number, number, fraction	attention weight(doc(1), pos(0), pos(1), fraction(43545600,59565131))	kb inference
hidden state	3	atom, number, fraction vec	hidden state(doc(1), pos(0), fraction vec([...]))	kb inference

Predicate	Arity	Arg Types	Example	KB Location
logits	3	atom, number, fraction vec	logits(doc(1), pos(2), fraction vec([3/7, 1/5, 2/9]))	kb inference
output prob	4	atom, number, token, fraction	output prob(doc(1), pos(2), token(17), fraction(3,7))	kb inference
sampled token	4	atom, number, token, number	sampled token(doc(1), pos(2), token(17), seed(1234))	kb inference

E.4 Derivation Facts

Predicate	Arity	Arg Types	Example	KB Location
derived from	3	term, atom, list	derived from(attention weight(d,p,q,w), "softmax", [scores...])	kb inference
operation type	2	atom, atom	operation type("softmax", "nonlinear")	kb model
operation inputs	2	atom, list	operation inputs(op id(42), [input ref 1, input ref 2])	kb inference
operation output	2	atom, term	operation output(op id(42), fraction(5, 12))	kb inference

Predicate	Arity	Arg Types	Example	KB Location
computation step	3	atom, number, atom	computation step(op id(42), order(7), “attention score”)	kb inference

E.5 Loss and Training Facts

Predicate	Arity	Arg Types	Example	KB Location
loss at	2	step, fraction	loss at(step(0), frac- tion(17,2))	kb training
loss type	1	atom	loss type(“mse”)	kb training
learning rate	2	step, fraction	learning rate(step(0), frac- tion(1,10))	kb training
optimizer type	1	atom	optimizer type(“sgd”)	kb training
momentum coeff	1	fraction	momentum co- eff(fraction(9,10))	kb training
batch size	1	number	batch size(2)	kb training
epoch	2	number, number	epoch(1, steps(3))	kb training
checkpoint at	1	step	checkpoint at(step(100))	kb training

Appendix F: Complete Rule Reference

F.1 Provenance Rules

Rule	Head	Body	Purpose
Direct dependency	depends on(X, Y)	derived from(X, , Sources), member(Y, Sources)	X directly uses Y
Transitive dependency	depends on(X, Y)	derived from(X, , Sources), member(Z, Sources), depends on(Z, Y)	X indirectly uses Y
Affected by parameter	affected by param(P, V)	parameter value(P, ,), depends on(V, P)	V depends on parameter P
Root sources	root source(X, S)	depends on(X, S), not(derived from(S, ,))	S is an underived input to X
Derivation depth	derivation depth(X, 0)	not(derived from(X, ,))	Underived fact has depth 0
Derivation depth	derivation depth(X, D)	derived from(X, , Sources), max depth(Sources, MD), D is MD + 1	Derived depth is max input depth + 1

F.2 Consistency Rules

Rule	Head	Body	Purpose
Weight consistent	weight consistent(P, S)	parameter value(P, S, V1), S1 is S-1, parameter value(P, S1, V0), gradient at(P, S1, G), learning rate(S1, LR), V1 ::= V0 - LR * G	SGD update is exact
Loss decreasing	loss decreasing(S)	loss at(S, L1), S0 is S-1, loss at(S0, L0), L1 < L0	Training is improving
Distribution valid	valid distribution(D)	sum list(D, 1), forall(member(P, D), P >= 0)	Probabilities are valid

Rule	Head	Body	Purpose
Attention valid	attention valid(Doc, Pos)	findall(W, attention weight(Doc, Pos, , W), Ws), valid distribution(Ws)	Attention row is valid
Gradient bounded	gradient bounded(P, S, B)	gradient at(P, S, G), abs(G) =< B	Gradient within bound
Denominator bounded	denom bounded(P, S, B)	parameter value(P, S, fraction(, D)), D =< B	Denominator within bound

F.3 Scoping Rules

Rule	Head	Body	Purpose
In scope	in scope(KB)	active topic(T), kb belongs to(KB, T)	Current topic's KB
In scope parent	in scope(KB)	active topic(T), topic(T, , , parent(P)), P = none, kb belongs to(KB, P)	Parent topic's KB
In scope global	in scope(KB)	kb belongs to(KB, "global")	Global always in scope
In scope secondary	in scope(KB)	secondary scope(T), kb belongs to(KB, T)	Explicitly activated scope
Out of scope	out of scope(KB)	kb exists(KB), not(in scope(KB))	Everything else
Scoped query	query(Pred, Args, Result)	in scope(KB), fact in(KB, Pred, Args, Result), !	First in-scope match
Cross query	query across(Pred, Args, Results)	findall(KB-R, (kb exists(KB), fact in(KB, Pred, Args, R))), Results)	All KBs, tagged
Qualified query	query in(KB, Pred, Args, Result)	fact in(KB, Pred, Args, Result)	Specific KB, bypass scope

F.4 Working Data Rules

Rule	Head	Body	Purpose
Resolve binding	resolve(Var, Val)	active topic(T), active dataset(T, DS), resolve in(Var, DS, Val)	Scoped lookup
Local resolve	resolve in(Var, DS, Val)	binding(DS, Var, Val)	Found locally
Inherited resolve	resolve in(Var, DS, Val)	not(binding(DS, Var,)), dataset(DS, , parent(P)), P = none, resolve in(Var, P, Val)	Walk to parent
Unbound	resolve in(Var, DS, unbound)	not(binding(DS, Var,)), dataset(DS, , parent(none))	Reached root, not found
Set binding	set binding(Var, Val)	active dataset(, DS), retrac- tall(binding(DS, Var,)), as- sert(binding(DS, Var, Val)), assert(binding history(DS, Var, Val, current turn))	Store with history
Visible bindings	visible bindings(All)	active dataset(, DS), collect all(DS, All)	Everything in scope

Rule	Head	Body	Purpose
Diff	diff(DS1, DS2, Add, Rem, Chg)	findall(K-V, (binding(DS2,K,V), not(binding(DS1,K,))), Add), findall(K-V, (binding(DS1,K,V), not(binding(DS2,K,))), Rem), findall(K-V1-V2, (binding(DS1,K,V1), binding(DS2,K,V2), V1=V2), Chg)	Structured comparison

F.5 Topic Management Rules

Rule	Head	Body	Purpose
Should close	should close(T)	topic(T, open, ,), not(pending(T,)), not((topic(C, open, , parent(T))))	No pending, no open children
Should park	should park(T)	topic(T, open, O, current turn(N), N-O > 20, not(recently discussed(T, 5))	Stale topic
Recently discussed	recently discussed(T, W)	current turn(N), last mentioned(T, M), N-M =< W	Active within window
Dangling	dangling topics(Ts)	findall(T, (topic(T, open, ,), should park(T)), Ts)	Open but stale
Topic depth	topic depth(T, 0)	topic(T, , , parent(none))	Root topic
Topic depth	topic depth(T, D)	topic(T, , , parent(P)), topic depth(P, PD), D is PD+1	Nested depth

Rule	Head	Body	Purpose
Wrap	wrap(T)	summarize(T, S), re- tract(topic(T,open, ,)), as- sert(topic(T,parked, ,)), log(wrapped(T,S))	Park with summary
Unwrap	unwrap(T)	retract(topic(T,parked, ,)), as- sert(topic(T,open, ,)), list pending(T), log(unwrapped(T))	Resume from park

F.6 Constraint Rules

Rule	Head	Body	Purpose
Effective constraints	effective constraints(KB, All)	kb constraints(KB, Local), kb parent(KB, P), (P=none -> All=Local ; effective constraints(P, PC), merge constraints(Local, PC, All))	Inherited union
Merge with override	merge constraints(L, P, M)	findall(C, member(C,L), LL), findall(C, (member(C,P), name(C,N), not((member(X,L), name(X,N))))), IL), append(LL, IL, M)	Local overrides parent
All satisfied	all satisfied(KB)	effective constraints(KB, Cs), forall(member(C, Cs), constraint satisfied(C))	Every constraint passes

Rule	Head	Body	Purpose
Violations	violations(KB, Vs)	effective constraints(KB, Cs), findall(C, (member(C,Cs), not(constraint satisfied(C)))), Vs)	Failed constraints
Enable set	enable set(Name)	constraint set(Name, Ms), forall(member(C,Ms), activate(C))	Group activation
Disable set	disable set(Name)	constraint set(Name, Ms), forall(member(C,Ms), suspend(C))	Group suspension
Set union	set union(S1, S2, R)	constraint set(S1,A), constraint set(S2,B), union(A,B,R)	Combine sets
Set intersect	set intersect(S1, S2, R)	constraint set(S1,A), constraint set(S2,B), intersection(A,B,R)	Common constraints
Set diff	set diff(S1, S2, R)	constraint set(S1,A), constraint set(S2,B), subtract(A,B,R)	Remove set from set

Appendix G: Complete Constraint Type Reference

G.1 Constraint Types

Type	Domain	Can Suspend?	Violation Action	Verification Method
sum to one	axiom	No	error	Exact fraction sum == 1/1
non-negative	axiom	No	error	Every fraction >= 0/1

G.2 Constraint Status Transitions

From	To	Trigger	Logged?
active	satisfied	Verification passes	Yes
active	violated	Verification fails	Yes
active	suspended	User or system suspends	Yes
active	parked	Topic parked	Yes
suspended	active	User or system reactivates	Yes
parked	active	Topic resumed	Yes
violated	active	Underlying issue fixed, reverify	Yes
violated	suspended	User acknowledges and suspends	Yes
any	any	—	All transitions logged with turn, source, reason

G.3 Constraint Scope Hierarchy

Scope	Applies To	Inherited?	Overridable?	Example
axiom	All KBs	Yes, always	No	sum to one
operational	System-level	Yes	By admin only	rate limit
legal	Organization-level	Yes	By legal authority only	gdpr compliance
policy	Group/team-level	Yes	By group admin	code review
project	Project-level	Yes	By project lead	approved languages
conversation	Topic-level	No (local only)	By user	character consistent

Appendix H: Complete KB Hierarchy Reference

H.1 KB Node Structure

Field	Type	Required	Default	Purpose
name	Text	Yes	—	Unique identifier
facts	FactSet	Yes	empty	Stored facts
rules	RuleSet	Yes	empty	Logical rules
constraints	[]Constraint	Yes	empty	Local constraints
parent	?Text	No	none	Parent KB for inheritance

Field	Type	Required	Default	Purpose
children	[]Text	No	[]	Child KB names
topic	?Text	No	none	Associated topic
working data	?WorkingDataSet	No	none	Scoped variable bindings
visibility	enum	Yes	public	Access control level
frozen	bool	Yes	false	Read-only flag
owner	?Text	No	none	Owning user/group
created at	i32	Yes	current turn	Creation timestamp
last modified	i32	Yes	current turn	Modification timestamp

H.2 KB Types by Purpose

KB Type	Parent Pattern	Typical Contents	Retention	Example
Global	none (root)	Axioms, system rules	Permanent	kb global
Organization	Global	Legal constraints, org policies	Permanent	kb acme corp
Department	Organization	Team policies, access rules	Permanent	kb engi-neering
Team	Department	Team-specific constraints	Permanent	kb backend
User account	Team or Organization	Identity, preferences, conversation state	Permanent	kb user alice
Project	User or Team	Project facts, project constraints	Long-term	kb project vdr
Sub-project	Project	Module-specific facts	Long-term	kb vdr llm
Model	Project	Architecture, hyperparameters	Long-term	kb model
Training run	Model	Parameter history, gradients, losses	Checkpoint-retained	kb training run 1
Inference batch	Training run	Activations, attention, outputs	Batch-transient	kb batch 42
Story/creative	User	Characters, plot, setting	User-managed	kb story a
Snapshot	Any	Frozen copy of another KB at a point	Read-only	kb snapshot step 100

H.3 KB Visibility Levels

Level	Owner Sees	Operator Sees	End User Sees	LLM Reasons Over
public	Yes	Yes	Yes	Yes
internal	Yes	Yes	No	Yes
owner only	Yes	No	No	Yes

H.4 KB Search Order

Priority	KB Source	Condition	Stop on Match?
1	Active topic's KB	Always searched first	Yes (cut)
2	Parent topic's KB	If not found at level 1	Yes (cut)
3	Grandparent topic's KB	If not found at level 2	Yes (cut)
...	Ancestors up to root	Walk up tree	Yes (cut)
N	Global KB	Always searched last	Yes (cut)
—	Secondary scope KBs	Only if explicitly activated	Yes (cut)
—	Out-of-scope KBs	Only via query in() or query across()	N/A

Appendix I: Complete Working Data Set Reference

I.1 Binding Operations

Operation	Signature	Effect	Logged?
set	set binding(Key, Value)	Assert binding in active dataset	Yes: key, old val, new val, turn
get	get binding(Key) → Value	Resolve with inheritance	No (read-only)
delete	delete binding(Key)	Retract binding from active dataset	Yes: key, old val, turn
list	visible bindings() → All	Collect with inheritance	No (read-only)
list local	local bindings() → Local	Current dataset only, no inheritance	No (read-only)
snapshot	snapshot(Name)	Frozen copy of all visible bindings	Yes: name, turn, binding count

Operation	Signature	Effect	Logged?
restore	restore(Snapshot, Target)	Copy all bindings from snapshot to target	Yes: source, target, turn
diff	diff(DS1, DS2) → Changes	Added, removed, changed bindings	No (read-only)
merge	merge(Source, Target)	Copy source bindings into target, source wins conflicts	Yes: source, target, turn, conflict count
clear	clear(DS)	Remove all bindings from dataset	Yes: dataset, turn, count

I.2 Binding History Event Types

Event	Fields	Example
created	key, value, turn, source	created("bob age", number(32), turn(5), "user")
updated	key, old value, new value, turn, source	updated("bob age", number(32), number(33), turn(12), "user")
deleted	key, old value, turn, source	deleted("bob age", number(33), turn(20), "user")
inherited	key, value, from dataset, turn	inherited("language", atom("python"), "kb project", turn(1))
shadowed	key, local value, parent value, turn	shadowed("max n", number(50), number(6), turn(15))
snapshot	snapshot name, turn, binding count	snapshot("before refactor", turn(30), 42)
restored	snapshot name, turn	restored("before refactor", turn(35))

I.3 Type Constraints for Bindings

Type	Check	Example Valid	Example Invalid
number	Term.type == number	number(32)	atom("thirty-two")
fraction	Term.type == fraction	fraction(1, 4)	number(0.25)
atom	Term.type == atom	atom("Portland")	number(42)
fraction vec	Term.type == fraction vec	fraction vec([1/2, 1/3])	list([1, 2])
fraction mat	Term.type == fraction mat	fraction mat([[1/2]])	fraction vec([1/2])
token	Term.type == token	token(42, "the")	number(42)
any	Always passes	—	—

Appendix J: Complete Topic Lifecycle Reference

J.1 Topic States

State	Meaning	KB Active?	Constraints Active?	Working Data?	Can Query?
open	Currently active or accessible	Yes	Yes	Read/write	Yes
closed	Completed, no pending items	No	No	Read-only	Via query in only
parked	Temporarily suspended	No	No	Preserved (read-only)	Via query in only
branched	Split into child topics	Yes (parent)	Yes (parent)	Read-only (parent)	Yes (parent)

J.2 State Transitions

From	To	Trigger	Precondition	Side Effects
—	open	User opens topic	—	Activate KB tree, load constraints, create working data
open	closed	User closes or auto-close rule	No pending items, no open children	Deactivate KB tree, summarize, freeze working data
open	parked	User parks or auto-park rule	—	Snapshot state, deactivate KB tree, preserve working data
open	branched	User creates subtopic	—	Create child topic with own KB, parent stays open
parked	open	User resumes	—	Restore state, activate KB tree, list pending
closed	open	User reopens	—	Unfreeze working data, reactivate KB tree
branched	open	All children closed	No open children	Resume normal open state
any	any	—	—	Transition logged with turn, source, reason

J.3 Topic Triggered Actions

Trigger	Action	Purpose
on open	activate kb tree(Topic)	Make facts visible
on open	load constraints(Topic)	Apply topic's constraints
on open	restore working data(Topic)	Make bindings available
on close	verify no pending(Topic)	Ensure completeness
on close	summarize(Topic)	Record what was accomplished
on close	freeze working data(Topic)	Prevent accidental modification
on close	deactivate kb tree(Topic)	Remove from scope
on park	snapshot state(Topic)	Preserve for later resume
on park	log park reason(Topic)	Record why parked
on park	deactivate kb tree(Topic)	Remove from scope
on resume	restore state(Topic)	Bring back snapshot
on resume	activate kb tree(Topic)	Make facts visible again

Trigger	Action	Purpose
on resume	list pending(Topic)	Remind user of open items

Appendix K: Complete Surfacing Mode Reference

K.1 Surfacing Modes

Mode	Trigger	Output Type	LLM Involvement	Data Source
Narrative	Default conversation	LLM text + KB references	LLM generates framing	KB provides data
Table	“/show KB” or data request	Structured rows	None (direct)	KB dump
Tree	“/tree”	Hierarchical KB display	None (direct)	KB structure
Provenance	“/provenance ref”	Derivation chain	Optional summary	KB derivation facts
Constraint	“/constraints”	Status table	None (direct)	Constraint facts
Diff	“/diff A B”	Added/removed/changed	Optional summary	KB comparison
Query	“/query ...”	Prolog query results	None (direct)	Query engine
History	“/history key”	Binding change log	None (direct)	Binding history
Pending	“/pending”	Open items list	None (direct)	Topic pending lists
Context	“/context”	Active scope summary	Optional summary	KB tree + topic state

K.2 Output Block Types

Block Type	Source	Can Hallucinate?	Editable?	Live?
LLM text	Language model generation	Yes	No	No (static)
KB data	Direct KB read	No	Via KB operations	Yes (resolves current value)

Block Type	Source	Can Hallucinate?	Editable?	Live?
Reference	KB address + current value	No	Via KB operations	Yes (updates on change)
Query result	Prolog query execution	No	No	Snapshot at query time
Constraint status	Constraint verifica- tion	No	Via constraint operations	Yes (re- verifiable)

K.3 Permission Filtering

User Role	Sees LLM Text	Sees Public KB Data	Sees Internal KB Data	Sees Owner-Only Data	Sees Reasoning Trace
Owner	Yes	Yes	Yes	Yes	Yes
Operator	Yes	Yes	Yes	No	Optional
End User	Yes	Yes	No	No	No
Auditor	Yes	Yes	Yes	Yes (read-only)	Yes (read-only)

Appendix L: Complete Account Hierarchy Reference

L.1 Account KB Structure

Field	User Account	Group Account	Organization Account
name	user alice	team backend	org acme corp
facts	Identity, preferences, history	Group metadata, policies	Org metadata, legal facts
rules	Personal rules	Team rules	Org-wide rules
constraints	Personal + inherited	Team + inherited	Org + global
parent children	Group or org KB Project KBs, story KBs	Org KB User KBs	Global KB Department KBs
working data	Current session state	Shared team state	Shared org state

Field	User Account	Group Account	Organization Account
visibility	owner only	internal	internal
owner	self	group admin	org admin

L.2 Constraint Inheritance Through Account Hierarchy

Level	Example KB	Constraints Defined Here	Inherited From Above	Total Effective
Global	kb global	exact arithmetic, sum to one	—	2
Organization	kb macme corp	gdpr compliance, data retention	2 from global	4
Department	kb engi-neering	code review, approved languages	4 from above	6
Team	kb backend	database logged, go allowed (override)	6 from above	7 (one override)
User	kb user alice	30 sec runtime, elevated access	7 from above	9

L.3 Constraint Override Rules

Scenario	Parent Constraint	Child Override	Effective at Child	Logged As
No override	approved languages([py,zig,rust])	—	approved languages([py,zig,rust])	inherited
Expand	approved languages([py,zig,rust])	approved languages([py,zig,rust,go])	approved languages([py,zig,rust,go])	overridden(expand)
Restrict	rate limit(100/hr)	rate limit(50/hr)	rate limit(50/hr)	overridden(restrict)
Suspend	code review(active)	code review(suspended)	code review(suspended)	overridden(suspend)
Cannot override	sum to one (axiom)	—	sum to one (axiom)	override blocked
Cannot override	gdpr compliance (legal)	—	gdpr compliance (legal)	override blocked

L.4 Account Operations

Operation	Who Can Do It	Effect	Logged?
Create user	Org admin, group admin	New user KB as child of group	Yes
Delete user	Org admin	Remove user KB and all children	Yes
Move user to group	Org admin	Reparent user KB	Yes: old parent, new parent
Add constraint to user	Admin, automated system	Assert constraint in user KB	Yes: source, reason
Remove constraint from user	Admin with sufficient level	Retract constraint	Yes: source, reason
Override group constraint	Group admin	Assert same-name constraint in child KB	Yes: override details
Grant KB access	KB owner, admin	Assert kb granted(KB, User)	Yes
Revoke KB access	KB owner, admin	Retract kb granted(KB, User)	Yes
Export account	User (own), admin (any)	Serialize user KB and children	Yes
Import account	Admin	Deserialize KB, attach to parent	Yes

Appendix M: Memory Tier Reference

M.1 Retention Tiers

Tier	Contents	Lifetime	Pruning Policy	Typical Size
1: Persistent	Architecture, rules, axioms, accounts, org structure	Never pruned	—	Small (KB)

Tier	Contents	Lifetime	Pruning Policy	Typical Size
2: Check-point	Parameter snapshots, model state at checkpoint steps	Until explicitly deleted	Keep every Nth step + anomaly steps	Medium (MB per checkpoint)
3: Step-transient	Gradients, updates, intermediate state for current step	Until step committed	Prune unless retain step rule matches	Small per step (KB)
4: Batch-transient	Per-batch inference: embeddings, attention, logits, outputs	Until batch processed	Prune unless retain batch rule matches	Medium per batch (KB-MB)
5: Conversation transient	Per-turn working state, draft outputs, reasoning traces	Until turn complete	Prune unless flagged	Small per turn (KB)

M.2 Retention Policy Rules

Rule	Condition	Effect	Override
retain step(S)	S mod checkpoint interval == 0	Keep full step state	—
retain step(S)	loss anomaly at(S)	Keep for diagnosis	—
retain step(S)	constraint violated at(, S)	Keep violation context	—
retain step(S)	user flagged(S)	User explicitly requested	—
retain batch(B)	contains unknown token(B)	Keep for analysis	—
retain batch(B)	output anomaly(B)	Keep for diagnosis	—
retain batch(B)	user flagged(B)	User explicitly requested	—
prune step(S)	not(retain step(S))	Free step-transient data	Cannot override

Rule	Condition	Effect	Override
prune batch(B)	not(retain batch(B))	Free batch-transient data	Cannot override

M.3 Storage Estimates

Component	Per-Parameter	Per-Step	Per-Batch	Per-Checkpoint
Parameter value	~20 bytes (fraction)	—	—	~20 bytes
Gradient	—	~20 bytes	—	—
Update record	—	~50 bytes	—	—
Attention weight (per position pair)	—	—	~20 bytes	—
Embedding (per token)	—	—	~dim*20 bytes	—
Logits (per position)	—	—	~vocab*20 bytes	—
Derivation fact	—	—	~100 bytes	—
Constraint check	—	~50 bytes	~50 bytes	—

For a tiny model (6 parameter groups, 2-dim, 3 vocab): ~120 bytes per parameter per checkpoint, ~500 bytes per step, ~2KB per batch inference with full provenance.

Appendix N: Query Pattern Reference

N.1 Weight Archaeology Queries

Query	Purpose	Returns
parameter value(P, step(S), V)	Value at specific step	Exact fraction
parameter value(P, step(0), V), parameter value(P, step(100), V)	Steps where value equals initial	Step numbers or failure
findall(S, (parameter value(P, step(S), V), V > fraction(1,1)), Steps)	Steps where weight exceeds 1	List of step numbers
gradient at(P, step(S), G), abs(G) > fraction(100,1)	Gradient explosion detection	Parameter + step pairs
param denominator(P, step(S), D), D > 2^64	Denominator growth detection	Parameter + step + denominator

N.2 Data Lineage Queries

Query	Purpose	Returns
depends on(output prob(D,P,T,), embedding(D,Pos))	Which input positions influenced output	Position list
root source(attention weight(D,P,Q,), S)	Ultimate inputs to an attention weight	Raw sources
affected by param(P, output prob(D,Pos,T,))	Does this parameter affect this output	Yes/no
derivation depth(V, D)	How many computation steps produced V	Integer depth

N.3 Constraint Queries

Query	Purpose	Returns
effective constraints(KB, Cs)	All constraints in scope	Constraint list
violations(KB, Vs)	Currently violated constraints	Violation list
constraint source(User, C, Source)	Who added this constraint	Source + timestamp
users with constraint(C, Users)	Who is subject to constraint C	User list
hypothetical add(KB, C, Affected)	Impact of adding new constraint	Affected KB list

N.4 Conversation Queries

Query	Purpose	Returns
dangling topics(Ts)	Open but stale topics	Topic list
findall(P, pending(, P), All)	All pending items globally	Item list
topic(T, open, , parent(P))	Open subtopics of P	Topic list
binding history(DS, Key, , Turn)	When was key last changed	Turn number
visible bindings(All)	Everything in current scope	Key-value list
diff(snap1, snap2, A, R, C)	Changes between snapshots	Added, removed, changed

N.5 Account Queries

Query	Purpose	Returns
user can see(User, KB)	Access check	Yes/no
effective constraints(user alice, Cs)	Alice's full constraint set	Constraint list with inheritance chain
constraint override(Team, C, Details)	Override details	Original + override + reason
users under(Group, Users)	All users in group hierarchy	User list
kb owner(KB, Owner)	Who owns this KB	Owner reference

Appendix O: Falsification Test Matrix

O.1 Arithmetic Layer Tests

Test	Input	Expected	Falsifies If
Fraction addition	$1/2 + 1/3$	$5/6$ exactly	Result $\neq 5/6$
200-op roundtrip	+100, -100 on $1/7$	$1/7$ exactly	Result $\neq 1/7$
Hilbert inverse	$H4 * H4^{-1}$	Identity exactly	Any element $\neq 0$ or 1
Softmax sum	$\text{softmax}([1,2,3])$	Sum exactly 1	Sum $\neq 1$
Gradient x^2	$d(x^2)/dx$ at $x=3$	Exactly 6	Result $\neq 6$
705 existing tests	Various	Various	Any VDR computation error

O.2 Logic Layer Tests

Test	Input	Expected	Falsifies If
Fact assert/query	Assert fact, query it	Fact found	Fact not found or different
Unification exact	$\text{fraction}(1,2)$ vs $\text{fraction}(2,4)$	Unify (same rational)	Fails to unify

Test	Input	Expected	Falsifies If
Scope isolation	Query out-of-scope KB	Not found	Found (scope leak)
Inheritance	Query binding from parent KB	Found via parent	Not found
Shadowing	Local binding overrides parent	Local value returned	Parent value returned
Derivation chain	Trace value to source	Complete chain	Missing link
Constraint check	Sum-to-one on exact distribution	Satisfied	Violated

O.3 Conversation Layer Tests

Test	Input	Expected	Falsifies If
Topic park/resume	Park topic, resume later	Exact state restored	Any binding changed or lost
Working data scope	Switch topic, query binding	Correct scope value	Wrong scope value
Constraint inheritance	Child KB query effective constraints	Union of ancestor chain	Missing or extra constraint
Override	Same-name constraint in child	Child version wins	Parent version wins
KB surfacing	/show kb name	Exact KB contents	Missing or modified data
Cross-scope query	query in(other kb, ...)	Tagged result from specified KB	Result from wrong KB

O.4 Account Layer Tests

Test	Input	Expected	Falsifies If
Permission check	End user queries internal KB	Blocked	Data returned
Owner full access	Owner queries any KB	All data returned	Any data blocked

Test	Input	Expected	Falsifies If
Group inheritance	User inherits group constraint	Constraint effective	Constraint missing
Override logged	Team overrides dept constraint	Override fact recorded	No record
Account export/import	Export user KB, import elsewhere	Identical KB	Any difference

END VDR-5 APPENDIX TABLES

[[HOWL-VDR-5-2026](#)] Exact Arithmetic Meets Logical Provenance. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-5-2026>

[[HOWL-VDR-1-2026](#)] VDR Arithmetic: Value, Decimal, Remainder. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-1-2026>

[[HOWL-VDR-2-2026](#)] VDR Gym: Exact Arithmetic Across Fifteen Domains. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-2-2026>

[[HOWL-MATH-3-2026](#)] The Transcendental Hierarchy. Github: <https://github.com/ghowland/papers/tree/main/papers/MATH-3-2026>

[[HOWL-MATH-4-2026](#)] The Universal Power-of-Two Basis. Github: <https://github.com/ghowland/papers/tree/main/papers/MATH-4-2026>

[[HOWL-VDR-3-2026](#)] VDR Gym Extension. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-3-2026>

[[HOWL-VDR-4-2026](#)] Exact-Fraction Language Model Architecture. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-4-2026>

[[HOWL-LLM-1-2026](#)] Integer LLM with Prolog Knowledge Base. Github: <https://github.com/ghowland/papers/tree/main/papers/LLM/HOWL-LLM-1-2026>