

VDR-LLM-Prolog: Performance

Integer Arithmetic on GPU Hardware: Why Wider Operands on More Cores Outrun Narrower Operands on Fewer Passes

Registry: [\[HOWL-VDR-18-2026\]](#)

Series Path: [\[HOWL-VDR-1-2026\]](#) → [\[HOWL-VDR-2-2026\]](#) → [\[HOWL-MATH-3-2026\]](#) → [\[HOWL-MATH-4-2026\]](#) → ... → [\[HOWL-VDR-15-2026\]](#) → [\[HOWL-VDR-16-2026\]](#) → [\[HOWL-VDR-17-2026\]](#)

DOI: 10.5281/zenodo.20236975

Date: May 2026

Domain: Applied Philosophy / Computational Linguistics

Supplementary Material:

- `supplementary/gpu_integer_perf_tech_spec.md` — Full GPU technical specification
- `supplementary/gpu_integer_perf_tech_spec_llm_compacted.md` — LLM-compacted form

AI Usage Disclosure: Only the top metadata, figures, refs and final copyright sections were edited by the author. All paper content was LLM-generated using Anthropic’s Opus 4.6.

Abstract

The VDR-LLM-Prolog system replaces floating-point arithmetic with exact integer computation. The immediate objection is performance: integer arithmetic on 100-digit numbers must be slower than hardware-accelerated floating-point on 16-bit or 32-bit values. This paper demonstrates that the objection confuses per-operation cost with per-prompt cost. A conventional language model spends thousands of tokens — each requiring a full forward pass through billions of floating-point parameters — on infrastructure work that VDR handles through exact primitive calls costing a few hundred integer operations each. VDR-15 established that the token reduction is 85 to 97 percent. This paper establishes that the integer arithmetic executing those primitives maps efficiently to GPU hardware, that the wider operands are offset by the massive parallelism of modern GPUs, and that several architectural properties of VDR — fixed-frame regularity, grammar-constrained decode, indexed knowledge base scans, and frontier-based Prolog execution — produce GPU utilization patterns that are structurally superior to the irregular, attention-dominated workloads of conventional language model inference. The complete GPU mapping is specified in the supplementary technical specification. This paper explains why it works, what the performance characteristics are, and where the actual bottlenecks lie.

1. The Performance Objection

Every discussion of exact arithmetic in language models encounters the same objection: it must be too slow. Floating-point arithmetic has decades of hardware acceleration. Modern GPUs contain thousands of cores optimized for 16-bit and 32-bit floating-point operations. Tensor cores perform matrix multiplications at throughputs measured in hundreds of teraflops. The entire hardware ecosystem — from chip design through compiler optimization through library implementation — is built around floating-point.

VDR replaces all of this with integer arithmetic on large numbers. A Q335 value — the standard representation in the VDR system, specified in VDR-1 through VDR-3 — has a numerator of approximately 100 decimal digits stored as a fixed-width vector of 11 unsigned 32-bit integers (352 bits total) or 6 unsigned 64-bit integers (384 bits total), with an implicit denominator of 2 raised to the power 335. Adding two Q335 values requires a limbwise addition with carry propagation across 11 limbs. Multiplying two Q335 values requires a full multi-precision multiplication producing a product of roughly 700 bits, followed by a split at bit 335 to extract the quotient and remainder. This is more work per operation than a single 16-bit floating-point multiply.

The objection is correct at the per-operation level. A single Q335 multiplication is slower than a single float16 multiplication. The question is whether per-operation cost determines per-prompt cost. VDR-15 established that it does not. A conventional language model spends 3,000 to 200,000 tokens on a single prompt, where each token requires a full forward pass through the model — matrix multiplications across every layer, full vocabulary softmax, the entire attention mechanism. VDR-LLM-Prolog spends 150 to 6,700 language model tokens on the same prompts, with the remainder handled by exact primitive calls that cost a few hundred integer operations each.

The relevant comparison is not one Q335 multiply versus one float16 multiply. It is the total computation cost of 210 language model tokens plus 500 primitive calls versus 3,000 language model tokens with no primitive offload. The per-operation cost of Q335 arithmetic is higher. The operation count is dramatically lower. The paper demonstrates that the lower operation count wins.

2. The GPU Mapping

The complete technical specification for mapping VDR-LLM-Prolog to GPU hardware is provided in the supplementary material. This section introduces the key architectural decisions and explains why they work.

The fundamental split is between a control plane on the CPU and a data plane on the GPU. The CPU handles everything that requires sequencing, string processing, or external interaction: parsing user input, resolving dotted paths to integer IDs, managing sessions, verifying grants, executing filesystem and network operations, scheduling GPU kernels,

forming batches, making backtracking decisions, and assembling responses. The GPU handles everything that is parallel and numeric: integer arithmetic, knowledge base scans, term matching, unification, confidence propagation, grammar masking, language model forward pass and decode, and live-state mutation.

The boundary between CPU and GPU is integer IDs. All strings in the system — predicate names, atom values, path segments, grammar nonterminals, builtin names — are interned on the host side before any device transfer. Interning means each unique string is assigned a sequential integer identifier, and a lookup table maps identifiers back to strings. The GPU never processes a string. It compares integer symbol IDs. The string “salary” might be symbol 4,817. The string “bob” might be symbol 12,003. The GPU sees a fact as a tuple of integers: (kb_id: 47, pred_id: 4817, arg0: 12003, arg1: vdr_ref_9281). String resolution happens once on the CPU, is cached, and all subsequent operations use integers exclusively.

This design minimizes host-device transfer. A dotted path like root.org.acme.hr.personnel is parsed and resolved on the CPU to a single integer — say, KB ID 47. That integer is sent to the GPU. The GPU accesses KB 47’s data through array indexing at offset 47 in the knowledge base metadata table. The entire path resolution, with all its string processing, happens once on the CPU. The GPU receives one 32-bit integer.

3. The Numeric Model

The system handles four classes of numeric values, each with its own representation and performance profile on GPU hardware.

The first class is small integers — signed 32-bit or 64-bit immediate values. These are used for knowledge base IDs, counter values, bitset words, term payloads, and any exact value that fits in a machine word. Operations on small integers use the GPU’s native integer arithmetic units. There is no performance penalty compared to floating-point for these operations — integer add, multiply, compare, and shift are native GPU instructions at the same throughput as their floating-point counterparts.

The second class is Q335 fixed-frame closed values — the common case for exact arithmetic in the system. The denominator 2^{335} is implicit and never stored. The numerator is stored as a fixed-width limb vector: 11 unsigned 32-bit words (352 bits, portable across all GPU architectures) or 6 unsigned 64-bit words (384 bits, optimized for GPUs with native 64-bit integer support). The sign is represented via two’s complement, which makes addition, subtraction, and comparison regular operations without special case handling.

The third class is general closed rationals — values with explicit numerator and denominator limb references — used for values that don’t fit the Q335 frame. These arise from parsing exact fractions from external sources or from cross-frame operations. They reference limb vectors in an append-only arena.

The fourth class is active VDR values — triples where the remainder is not zero. These carry a top-level value and denominator reference plus a remainder tag and payload refer-

ence. The remainder payload can be an atomic integer, a composite structure with child VDR values, or a functional remainder descriptor. Active values are the minority case — the system is designed so that most values stay in Q335 closed form, with active spills occurring only when operations exceed what the frame can represent.

The performance strategy is built on the assumption that Q335 closed values dominate. The supplementary specification defines three spill thresholds: below 1 percent active values, the system uses dense closed storage with sparse spill tags; between 1 and 10 percent, it uses tile-local spill tables; above 10 percent, it switches to sparse active mode or escalates to host intervention. The VDR-14 paper’s characterization of denominator growth during training — plateauing around 2^{45} with small learning rates — confirms that Q335’s 2^{335} frame provides enormous headroom. Most values never leave the fast closed path.

4. Q335 Arithmetic on GPU

The three primary Q335 operations — addition, multiplication, and comparison — each map to efficient GPU execution patterns.

Addition of two Q335 values is limbwise addition with carry propagation. Each of the 11 limbs (in the 32-bit layout) is added with the carry from the previous limb. This is 11 integer additions with 11 carry checks — roughly 22 integer operations total. On a GPU, this executes in a single thread in a few dozen clock cycles. For bulk operations — adding vectors of Q335 values, summing arrays, computing residuals — thousands of additions execute in parallel across GPU cores. The operation is perfectly regular: every addition has exactly the same shape, the same number of limbs, the same carry propagation pattern. GPU hardware achieves peak utilization on regular uniform workloads.

Multiplication of two Q335 values is a multi-precision multiply producing a full product of approximately 700 bits, followed by a split at bit position 335. The split is a fixed-position operation: the bits above position 335 become the quotient (the new V), the bits below become the remainder. If the remainder is zero, the result is closed and stays in the fast Q335 path. If nonzero, the result becomes an active node with the remainder stored in the append-only arena.

The multiplication itself can use schoolbook multiplication ($11 \times 11 = 121$ limb-by-limb multiplies with accumulation) for the standard case, with optional Karatsuba or Toom-Cook algorithms for larger operands. Schoolbook on 11 limbs produces roughly 200 integer operations — multiplies, adds, and carry propagations. This is more expensive than a single float16 multiply (one operation) by a factor of roughly 200. But this factor is the per-operation cost. The relevant question is how many of these multiplications occur per prompt compared to how many float operations a conventional forward pass requires.

A conventional language model forward pass for one token through a model with 7 billion parameters requires roughly 14 billion floating-point operations — two operations per parameter (multiply and accumulate) across the forward pass. Generating 3,000 tokens requires roughly 42 trillion floating-point operations.

A VDR prompt with 210 language model tokens requires roughly 2.94 trillion exact integer operations for the forward passes (same 14 billion per token, but with wider integer operands increasing the per-operation cost by roughly the $200 \times$ factor, giving approximately 2.8 trillion effective operations for 210 tokens). The 500 primitive calls add perhaps 100,000 Q335 operations — negligible in comparison.

The conventional system does 42 trillion float operations. The VDR system does roughly 2.94 trillion integer operations, where each integer operation is wider but the total operation count is $14 \times$ lower. The question becomes: can a GPU execute 2.94 trillion wide integer operations faster than 42 trillion narrow float operations? The answer depends on GPU utilization patterns, which the next sections address.

5. Why GPU Utilization Favors VDR

Modern GPUs achieve peak performance only when their thousands of cores are uniformly occupied with the same operation on different data. Any divergence — threads taking different code paths, threads waiting for memory at different times, threads idle while others work — reduces utilization. The utilization characteristics of VDR workloads are structurally different from conventional language model inference in ways that favor VDR.

Conventional language model inference has an irregular utilization profile. The attention mechanism requires a softmax computation that is inherently sequential within each row — the exponentials must be computed, summed, and normalized before the weighted sum can proceed. The softmax itself requires a max reduction (sequential), exponentiation of each element (parallel but transcendental — requiring multiple operations per element through lookup tables or polynomial approximation), a sum reduction (sequential), and a division (parallel). The transcendental operations are the most expensive: computing $\exp(x)$ in floating-point requires a range reduction, polynomial evaluation, and reconstruction, typically costing 10 to 20 floating-point operations per element. For a vocabulary softmax over 50,000 elements, this is 500,000 to 1,000,000 floating-point operations just for the exponentials.

VDR’s rational surrogate softmax replaces all of this with integer arithmetic. The formula is: each output equals the square of the shifted input divided by the sum of all squared shifted inputs. The computation is: find the row maximum (parallel reduction on integers), subtract the maximum and add a shift constant (parallel), square each element (one Q335 multiply each, parallel), sum the squares (parallel reduction), divide each squared element by the sum (one Q335 divide each, parallel). Every step is uniform integer arithmetic. There are no transcendentals. No lookup tables. No range reductions. No polynomial approximations. The GPU cores execute the same operation on every element with no divergence.

Knowledge base scans in VDR are structurally regular. Facts are stored in predicate-major columnar form — all facts with the same predicate are contiguous in memory. A query for predicate “salary” with argument “bob” becomes: load the predicate bucket

for symbol 4817 (salary), scan the argument column for symbol 12003 (bob), produce a bitmask of matches. This is a parallel filter scan — one of the most efficient operations a GPU can perform. Thousands of facts are checked simultaneously, one per core, with uniform memory access patterns (contiguous reads from columnar arrays) and uniform computation (integer comparison per element).

Conventional language model inference has no equivalent to these bulk scan operations because it has no structured data to scan. Everything is in the weight matrices, accessed through matrix multiplication, which is efficient but serves a different purpose — statistical pattern matching rather than exact indexed lookup.

Grammar-constrained decode in VDR further improves utilization. When a grammar declares that the next token must be one of four enum values, the decode step constrains the vocabulary to those four candidates. The softmax (or rational surrogate) computes over 4 values instead of 50,000 or more. This is a $12,500 \times$ reduction in decode computation for that token. Conventional language model inference must compute the full vocabulary softmax on every token because it has no structural constraint on what the next token can be. VDR's grammar system provides these constraints as a structural property of the output format, reducing decode cost on every structurally determined token.

6. Knowledge Base Operations on GPU

The knowledge base is the system's primary data structure — every fact, rule, constraint, connection, and piece of metadata lives in knowledge bases organized as a tree. The GPU mapping stores knowledge base data in structure-of-arrays columnar form, which aligns with how GPU hardware accesses memory.

A GPU accesses memory most efficiently in coalesced patterns — adjacent threads reading adjacent memory locations. Structure-of-arrays achieves this naturally. When 1,000 threads each need to check whether their candidate fact's `kb_id` is in the visible set, they read 1,000 consecutive `kb_id` values from a contiguous array. This is one coalesced memory transaction. In an array-of-structures layout, the `kb_id` fields would be scattered across memory at struct-width intervals, causing multiple memory transactions to fetch the same 1,000 values.

Fact storage is predicate-major. All facts sharing a predicate are stored in a contiguous bucket. The bucket is indexed by predicate ID — an integer lookup that gives the offset and count of facts for that predicate. A query for all salary facts starts with one index lookup to find the salary bucket, then scans within the bucket. If the bucket has 500 facts, 500 GPU threads scan 500 facts in parallel in one memory transaction per column checked.

The scope filter is where VDR's structural access control meets GPU execution. Every query must check that returned facts come from knowledge bases visible in the user's scope chain. The supplementary specification defines three strategies for this filter. The simplest uses a bitset: the user's visible KB IDs are set in a bitset, and each thread checks its candidate fact's `kb_id` against the bitset with a single bit-test operation. For larger

trees, DFS in/out intervals are used: each knowledge base has a pre-computed interval representing its position in a depth-first traversal of the tree, and checking whether a KB is in a subtree is an interval containment test — two integer comparisons. The scope filter adds one or two integer operations per candidate fact per query. Visibility checking — the core of VDR’s structural safety — costs almost nothing on GPU hardware.

This is a performance property that directly connects to VDR-16’s safety claims. Structural access control through visibility checks and scope filtering is not only more secure than behavioral safety — it is cheaper. A behavioral safety check requires the language model to evaluate the request through its full forward pass, spending tokens on classification and potential refusal. A structural visibility check is one bit-test per candidate fact, executed in parallel across thousands of facts in a single GPU kernel launch.

7. Prolog on GPU

The Prolog engine is the logic layer of VDR-LLM-Prolog — it evaluates rules, performs unification, and produces derivation results with provenance. Conventional Prolog implementations use recursive depth-first search with backtracking, which is fundamentally incompatible with GPU execution. Recursive control flow causes thread divergence: each thread follows a different branch of the search tree, executing different code at each step, making GPU utilization collapse.

The supplementary specification transforms Prolog execution into frontier-based batched joins — a technique borrowed from GPU database query processing. The transformation works as follows.

A Prolog query starts with a goal: find all facts matching a predicate with certain argument patterns. The first step is candidate retrieval — load the predicate bucket and produce a list of candidate facts. This is a parallel indexed scan, identical to the knowledge base operations described in the previous section.

The second step is term compatibility filtering. Each candidate fact’s arguments are checked against the goal’s argument patterns. An atom argument in the goal matches only facts with the same atom (integer symbol comparison). A variable argument matches any fact (no check needed). A VDR argument matches facts with equal VDR values (cross-multiplication comparison). This is a parallel filter: each GPU thread evaluates one candidate fact against the goal pattern, producing a pass/fail bitmask.

The third step is unification. For candidates that passed the filter, variable bindings are produced. If the goal had variable ?X in the first argument position and candidate fact 47 has atom “bob” in that position, the binding ?X = bob (symbol 12003) is recorded. The unification kernel takes the goal, the filtered candidate range, the term pool, and any existing bindings from earlier in the query, and produces extended binding rows — one per successful unification. This is parallel: each thread unifies one candidate, producing one binding row or marking failure.

The fourth step is rule body evaluation. If the query matched a rule head rather than a fact,

the rule’s body goals must be evaluated with the bindings from the head match. This is where the join model applies. Each body goal is another query, and the binding rows from the previous goal become the input for the next. This is a join: the binding frontier from goal one is joined with the candidate facts for goal two, producing an extended frontier for goal three, and so on.

The join strategies from database query processing apply directly. Hash join works when the join key is an equality condition — build a hash table on the smaller frontier, probe with the larger. Sort-merge join works for large frontiers — sort both sides by join key, merge. Bitmap semijoin works for unary filters — build a bitmap from one side, filter the other. Each of these is a well-studied GPU operation with known performance characteristics and optimized implementations.

For deterministic first-solution queries, the GPU computes all solutions in parallel and the CPU selects the first by canonical ordering. This is slightly more work than strictly necessary (a sequential system would stop at the first solution), but it maintains GPU utilization — all threads do useful work, and the selection is a trivial CPU operation on the completed result set.

The performance characteristic of frontier-based Prolog is that it converts recursive branching control flow into flat parallel data operations. The GPU never branches. It scans, filters, unifies, and joins — all uniform parallel operations on arrays of integers. Thread utilization stays high because every thread does the same thing. The work is proportional to the frontier size (number of candidate bindings) times the number of body goals, not to the depth of the search tree.

8. The Forward Pass

The language model forward pass in VDR uses the same computational structure as a conventional transformer — embedding lookup, attention computation, feedforward blocks, residual connections, and output projection — but with Q335 exact integer arithmetic instead of floating-point.

The supplementary specification defines three tensor storage classes for exact values. The primary class, T0, stores dense shared-frame tensors: all entries share the implicit Q335 denominator, and the numerators are stored as contiguous arrays of limb vectors. This is the highest-throughput path because every entry has the same width and the same implicit denominator, making operations perfectly regular. A matrix multiplication of two T0 tensors is a multi-precision integer matrix multiply where each element accumulation is a sequence of Q335 multiply-and-add operations.

The T1 class adds per-entry spill tags for tensors where some entries have exceeded the Q335 frame and become active values. The dense layout is preserved for the closed entries, with spill entries referencing the active value arena. The T2 class handles the rare case where active values exceed 10 percent of the tensor — sparse active storage with coordinates or tile-local spill lists.

Embedding lookup is an indexed read from a T0 tensor — token IDs index into rows of Q335 limb vectors. This is identical in structure to floating-point embedding lookup, just with wider entries.

Attention uses the rational surrogate softmax rather than the transcendental exponential softmax. The rational surrogate computes each attention weight as the square of a shifted input divided by the sum of all squared shifted inputs. It produces weights that sum to exactly one — not approximately, but exactly, verified by the fraction one over one. It preserves monotonicity: higher logits produce higher weights. It produces exactly equal weights for equal logits. And it uses zero transcendental operations — only integer subtraction, squaring, summing, and division.

The feedforward block uses ReLU activation — piecewise linear, producing exact zero or exact passthrough. No approximations, no lookup tables, no polynomial fits. The computational cost of ReLU is one comparison per element (is it negative?) and one conditional copy. On GPU, this is a single instruction per element with no divergence in the common case.

The gradient path for training stores the backward computation graph as flat GradNode records — each recording the operation type, arity, input references, output reference, and auxiliary data. The records are topologically sorted so that backward propagation is a sequential scan through the sorted array, with each operation’s gradient kernel dispatched in batch by operation type. Exact per-operation kernels compute exact gradients — the derivative of x squared at x equals three is exactly six, not 6.0000000001.

9. Grammar-Constrained Decode

The grammar system’s performance impact on decode is substantial and has no equivalent in conventional language model inference.

In conventional inference, every generated token requires a softmax over the full vocabulary — typically 50,000 to 100,000 entries. Each entry requires an exponentiation (10 to 20 floating-point operations), the results must be summed (one reduction), and each entry must be divided by the sum (one division per entry). Total cost per token: roughly 500,000 to 2,000,000 floating-point operations for the softmax alone, on top of the forward pass computation that produced the logits.

In VDR, the grammar system specifies what the next token can be based on the output format structure. When the grammar declares that the current position is an enum slot with four legal values, the decode constrains to those four candidates. The surrogate softmax computes over 4 values: 4 subtractions, 4 squarings (each a Q335 multiply), one sum of 4 values, 4 divisions. Total: roughly 20 Q335 operations, where each Q335 operation is roughly 200 integer operations, giving roughly 4,000 integer operations. Compare to 500,000 to 2,000,000 float operations for the unconstrained softmax.

The supplementary specification describes the mechanism. Each decode stream carries a grammar state: grammar ID, state ID, slot type, and a candidate begin/count pair pointing

into a candidate buffer. When the slot type is a categorical enum, the candidates are the enum’s token IDs. When the slot type is a KB identifier reference, the candidates are the token IDs of identifiers in the relevant knowledge base. When the slot type is a builtin opcode, the candidates are the roughly 300 opcode name tokens. When the slot type is free text, the full vocabulary is available.

The GPU kernel builds a vocabulary mask from the candidate set. For small candidate sets (the common case for structural tokens), this is trivial — set a few bits in a mask. For KB identifier references where the candidate set might be 200 items, it is still a small fraction of the full vocabulary. The mask is applied to the logit vector before the surrogate softmax, zeroing out impossible candidates.

The savings compound across a response. VDR-15 established that 40 to 60 percent of tokens in structured output are structural — pipes, headers, delimiters, enum values, type tags. Each structural token has a constrained candidate set, typically between 2 and 20 items. If half the tokens in a 200-token response are grammar-constrained to an average of 10 candidates, that is 100 tokens decoded at 10-candidate cost instead of 50,000-candidate cost — a $5,000 \times$ reduction in decode computation for half the response.

10. Concurrent Execution

The supplementary specification defines five concurrent GPU streams that execute simultaneously during each conversational turn. The streams are: LLM forward pass and decode (stream 0), knowledge base query and predicate scan (stream 1), VDR primitive execution (stream 2), grammar mask and candidate preparation (stream 3), and provenance and event compaction (stream 4). The CPU manages the dependency graph, launch ordering, external operations, and path resolution.

The turn pipeline has eleven steps. The host parses the request and resolves paths and scope (CPU). The GPU prefetches scoped facts and rules into fast memory (stream 1). The LLM begins decode (stream 0). When a command token is emitted, the host lowers it to an opcode batch (CPU), and the GPU executes pure queries and primitives (streams 1 and 2) while grammar masks update (stream 3). Results flow to the scratchpad or knowledge base. The LLM continues with structured state. Provenance events are committed (stream 4). The host assembles the response.

The concurrent execution model means that knowledge base queries, primitive execution, grammar preparation, and provenance logging happen in parallel with the LLM forward pass. In a conventional system, the language model does everything sequentially — it generates tokens one at a time, with each token requiring the full forward pass before the next can begin. In VDR, the LLM generates a command token (8 tokens), then the GPU executes the commanded operation in parallel with the LLM preparing to assess the result. The primitive execution is overlapped with LLM computation rather than serialized through the token stream.

This concurrency is possible because the operations are independent. A knowledge base

query reads from indexed fact tables. A VDR primitive computes on value references. The LLM decode processes its own attention and feedforward computations. These access different memory regions and use different computational units. The GPU's stream model allows them to execute simultaneously on different portions of the GPU's core array.

11. Memory and Allocation

GPU performance depends critically on memory management. The conventional approach — general-purpose malloc for dynamic allocation — is catastrophic on GPU hardware. Malloc requires synchronization, causes fragmentation, and produces irregular memory access patterns that defeat coalescing.

The supplementary specification uses append-only arenas for all irregular objects. An arena is a contiguous block of GPU memory with a bump pointer. Allocation is one atomic increment of the pointer — no synchronization beyond the atomic, no fragmentation, no free-list traversal. Every irregular structure in the system — limb vectors, VDR nodes, remainder payloads, child lists, term structures, binding rows, provenance events, live-state deltas — allocates from arenas.

The allocation model has two tiers. A global bump allocator handles large arena segments. Per-block suballocators carve small allocations from block-local arena segments using shared memory, which is the fastest memory tier on GPU. For operations where allocation size is predictable — a batch of 1,000 unification results where each produces a fixed-size binding row — a prefix-sum computes the offsets before any allocation occurs, enabling perfectly regular contiguous allocation with no synchronization.

The arena model produces coalesced memory access by construction. Objects allocated in sequence are stored in sequence in memory. When GPU threads process these objects in order — which the frontier-based Prolog execution ensures by construction — adjacent threads access adjacent memory, which is the optimal pattern for GPU memory controllers.

Arenas grow monotonically during a session. Dead objects (retracted facts, expired bindings, superseded provenance entries) are not immediately freed. Periodically — between turns or during idle time — the host or a maintenance kernel compacts arenas: live objects are copied to a fresh arena, dead regions are reclaimed, indexes are rebuilt, and handles are remapped. Session-visible handles (those referenced by user-accessible facts or primitive results) are stable across compaction through an indirection table.

12. Live State on GPU

The eight data primitive types from VDR-8 — counters, locks, queues, stacks, LRU caches, ring buffers, and bitsets — each have GPU-efficient implementations.

Counters are dense arrays of signed integers, one per counter in the active session. Incrementing a counter is one atomic integer addition. Checking a counter against a bound is one integer comparison. A batch of 100 counter increments — checking 100 dependencies in a codebase migration, for example — is 100 parallel atomic increments in a single kernel launch. The violation flag is set by a parallel comparison kernel after the increment batch.

Bitsets are packed 64-bit words. Setting, clearing, and testing bits are single bitwise operations — OR to set, AND-NOT to clear, AND to test. A bitset representing which of 256 evidence dimensions have been covered is 4 words. Testing all dimensions is 4 parallel word operations.

Ring buffers are fixed-capacity arrays with an atomic write position. Writing a metric snapshot to a 90-entry ring buffer is one atomic increment of the write position followed by one write to the computed index. Reading the full buffer is a contiguous memory read of 90 entries — one coalesced transaction.

Queues and stacks use batched push and pop operations. Single-item contended operations (multiple threads trying to push or pop simultaneously) are host-assisted or block-local, because fine-grained contention defeats GPU parallelism. Batched operations — push 50 items to a queue, pop the top 20 from a stack — are efficient: one prefix-sum for offsets, one bulk copy for the data.

LRU caches are the most complex case. Exact LRU ordering requires tracking access times and evicting the oldest entry, which involves sorting or heap operations. The supplementary specification recommends host-managed exact LRU for correctness-critical caches, with device-side approximate LRU (batched timestamped cache with periodic host reorder) for high-throughput caches where exact ordering is not critical.

13. Provenance on GPU

Every primitive execution and every rule derivation in the system produces a provenance event. The event records what was computed, by which tool, with what inputs, producing what output, at what confidence, on which turn. The supplementary specification stores these as an append-only structure-of-arrays log: `event_type`, `tool_id`, `subject_ref`, `input_begin`, `input_count`, `output_ref`, `confidence_ref`, `turn`.

Appending provenance events is concurrent with the computation that produces them. Each GPU block executing primitive operations appends its provenance events to a block-local buffer. After the kernel completes, the block-local buffers are compacted into the global provenance log with a single parallel copy operation.

The provenance log grows monotonically. It is never edited — events are append-only. This is the structural basis for the audit trail described in VDR-16: every data access, every computation, every primitive execution is logged. The append-only property means the log cannot be retroactively modified. The GPU’s append model — atomic bump pointer

plus bulk copy — makes this logging essentially free in performance terms. The cost is memory, not computation.

Deterministic ordering of provenance events requires attention. GPU execution is inherently parallel, and the order in which blocks complete is nondeterministic. The supplementary specification addresses this with canonical ordering: events are stamped with (turn, batch_sequence, item_sequence) and either appended in canonical batch order or post-sorted by these keys. The sort is a single parallel radix sort on the event buffer — a well-optimized GPU primitive.

14. Sessions, Snapshots, and Clones on GPU

The session model from VDR-8 separates persistent state (facts, rules, constraints, connections, grammars) from live state (counters, locks, queues, stacks, caches, ring buffers, bitsets, scratchpad, working data). This separation maps efficiently to GPU memory management.

A session snapshot captures all live state. On GPU, this means capturing the current arena positions for each live-state type and copying the live data to a snapshot buffer. The supplementary specification defines a SnapshotHeader containing the session ID plus begin/count pairs for each live-state type. Snapshot size is small — VDR-16 Appendix I established that a typical incident investigation’s complete live state is roughly 21 kilobytes. Copying 21 kilobytes on a GPU with hundreds of gigabytes per second of memory bandwidth takes microseconds.

Cloning a session uses copy-on-write over persistent state and duplicates live-state references with new delta arenas. The persistent stores — fact tables, rule tables, constraint tables — are shared between all clones because they are read-only from the clone’s perspective (new assertions go into the clone’s delta arena, not into the shared persistent store). The live state — counters, caches, queues — is duplicated because each clone mutates independently.

Reset discards the clone’s live-state deltas and restores the snapshot baseline references. Kill invalidates the clone’s live-state references but preserves any persistent assertions the clone committed — facts asserted to persistent knowledge bases survive the kill because they are in the shared persistent store, not in the clone’s delta arena.

The disposable clone pattern from VDR-8 — spawning worker clones, killing them when they drift, respawning from a frozen snapshot — is cheap on GPU. Clone creation is one arena allocation plus one small memory copy. Clone death is one arena invalidation. The persistent knowledge base grows monotonically across clone lifetimes, accumulating exact provenance-tagged facts while the language model stays fresh.

15. The Actual Bottlenecks

The system has genuine performance bottlenecks. Identifying them honestly is more useful than claiming universal superiority.

The language model forward pass with Q335 arithmetic is the primary bottleneck. Matrix multiplication with 352-bit operands is wider than float16 matrix multiplication by a factor of 22 in operand width. The actual slowdown depends on the multiplication algorithm and GPU integer throughput, but a reasonable estimate is that Q335 matrix multiplication is 100 to 500 times slower per element than float16 matrix multiplication. For the LLM forward pass, which is dominated by matrix multiplications, this is the main cost.

However, VDR-15 established that the VDR system generates 85 to 97 percent fewer LLM tokens than a conventional system for the same tasks. The forward pass cost per token is higher, but the number of tokens is dramatically lower. At 95 percent token reduction, the VDR system executes 5 percent as many forward passes, each roughly $200 \times$ more expensive per element, giving a total forward pass cost of roughly $10 \times$ the conventional system. This is slower in absolute terms for the LLM component.

The offsetting factor is that the conventional system’s 3,000 tokens include massive amounts of infrastructure work that produces no value — arithmetic through digit prediction, state reconstruction, formatting, hedging. The VDR system’s 210 tokens are almost entirely judgment and prose — the high-value output that the user actually wants. The user waits longer per token but waits for fewer tokens, and every token they receive is useful rather than infrastructure.

The second bottleneck is active value management. When values exceed the Q335 frame and produce active remainders, the append-only arena grows, normalization passes become necessary, and the regular fixed-width computational path is interrupted by irregular remainder structures. The normalization pipeline — mark active roots, normalize leaves in parallel, compact upward, canonically order — is well-designed for GPU execution but still more expensive than closed-value operations. The performance strategy is to keep active spill rates below 1 percent through appropriate Q335 frame sizing and Q-basis reprojection.

The third bottleneck is Prolog execution on large knowledge bases. Frontier-based joins are efficient for moderate frontier sizes, but very large frontiers (thousands of candidate bindings joined across multiple body goals) produce proportionally large intermediate results. The join strategies (hash join, sort-merge, bitmap semijoin) are all proven GPU operations, but the total work scales with the product of frontier sizes across join steps. For typical VDR queries — scoped to a subtree, filtered by predicate bucket, constrained by visibility — frontier sizes are moderate. For pathological queries (unrestricted cross-scope scans with highly variable predicates), frontier explosion is possible.

The fourth bottleneck is host-device synchronization. Every command token the LLM emits requires the host to resolve paths, verify grants, and lower the command to GPU opcodes. This is a round trip across the host-device boundary. The supplementary specification mitigates this by overlapping LLM decode with primitive execution through concurrent GPU streams, but the synchronization points at command token boundaries re-

main.

16. Comparative Analysis

The honest comparison between conventional float-based LLM inference and VDR integer-based inference requires separating the workload into components and comparing each.

For the LLM forward pass in isolation, conventional floating-point is faster per token. Float16 tensor cores achieve peak throughput that Q335 integer arithmetic cannot match on current hardware, because the hardware was designed for float16. The VDR system pays a per-token penalty of roughly $100 \times$ to $500 \times$ on the forward pass. This is the real cost of exact arithmetic on hardware optimized for approximate arithmetic.

For knowledge base operations, VDR is faster. A conventional LLM has no structured data operations — everything is in the weights, accessed through the forward pass. VDR’s indexed predicate scans, scope filters, and connection traversals are standard GPU database operations executing at millions of facts per second. A query that would cost a conventional LLM hundreds of tokens of attention-based “searching” through its context window is one indexed GPU scan in VDR.

For Prolog evaluation, VDR has capability that conventional LLMs lack entirely. Logical deduction through token prediction is unreliable and scales poorly with chain length. Frontier-based Prolog on GPU is exact, provenance-tracked, and scales with frontier size rather than chain depth.

For decode, VDR is dramatically faster on structured output. Grammar-constrained decode over small candidate sets is orders of magnitude cheaper than full-vocabulary softmax. For free-text output, VDR uses the surrogate softmax which avoids transcendentals but still computes over the full vocabulary, giving roughly comparable cost.

For primitive execution, VDR replaces thousands of LLM tokens with hundreds of integer operations. Sorting 500 items is one GPU sort kernel — roughly $500 \times \log_2(500) \times 200$ integer operations for Q335 comparisons, totaling roughly 900,000 operations, equivalent to less than one LLM token. A conventional LLM generates hundreds of tokens to sort the same list through prose comparison.

For provenance and confidence, VDR adds cost that conventional systems don’t have — but the cost is small (one provenance event per operation at append-only arena cost, confidence propagation as parallel integer reductions) and the value is substantial (every result is auditable and carries computed confidence).

The total system comparison depends on the workload. For workloads dominated by free-text generation (writing a novel, composing a poem), the conventional system is faster because the forward pass dominates and grammar constraints provide little benefit. For workloads involving structured data, computation, state management, and multi-turn reasoning (SRE incidents, portfolio analysis, research synthesis, code migration), VDR is

faster in total wall-clock time despite the slower forward pass because it generates 85 to 97 percent fewer tokens and replaces the rest with structurally efficient GPU operations.

17. Hardware Trajectory

Current GPU hardware is optimized for floating-point. Integer throughput is a secondary consideration in GPU design, driven primarily by address computation and index arithmetic rather than computational workloads. The performance gap between Q335 integer arithmetic and float16 tensor core arithmetic reflects this design emphasis, not any fundamental computational limitation.

Integer arithmetic on GPU is not slow in absolute terms. A modern GPU can perform billions of 32-bit integer operations per second. Q335 arithmetic, at roughly 200 integer operations per multiply, achieves millions of Q335 multiplications per second on current hardware. This is sufficient for the VDR workload where the LLM generates 210 tokens per prompt and primitive calls number in the hundreds.

The trajectory favors wider integers. GPU integer ALU widths have expanded with each generation. Support for 64-bit integers is now standard. Specialized instructions for multi-precision arithmetic (carry-propagating adds, wide multiplies) appear in some architectures. If exact arithmetic workloads create market demand, hardware designers can add integer tensor cores — fixed-function units for multi-precision integer matrix multiplication — that would close the per-operation gap with floating-point tensor cores. This is an engineering decision, not a physics limitation.

The more significant trajectory is the token reduction. As the grammar, knowledge base, and primitive systems mature, the fraction of work handled by exact primitives rather than LLM token generation will increase. Every new grammar template, every new Prolog rule composition, every new knowledge base schema reduces future LLM token cost. The system gets more efficient with use because the structural components accumulate while the LLM cost per structural operation stays constant at roughly 8 tokens per command.

18. What This Means for Real Workloads

For the SRE incident investigation from VDR-15: 210 LLM tokens plus 500 primitive calls. The 210 tokens cost roughly $210 \times 200 = 42,000$ token-equivalent operations at the Q335 slowdown factor (versus $3,000 \times 1 = 3,000$ token-equivalent operations for the conventional system). VDR's forward pass is roughly $14 \times$ more expensive in total. But the primitive calls — fetching metrics, parsing JSON, filtering time series, computing correlations, evaluating Prolog rules, propagating confidence — add roughly 100,000 Q335 operations, equivalent to 0.05 token-equivalent operations at the 2-billion-ops-per-token scale. The primitives are computationally invisible.

Wall-clock time: if the conventional system takes 30 seconds to generate 3,000 tokens, VDR takes roughly 3 seconds for 210 LLM tokens (token generation is roughly propor-

tional to token count, with wider arithmetic partially offset by grammar-constrained decode savings) plus sub-second for all primitive execution. Total roughly 4 seconds versus 30 seconds. The user gets a faster response that is also exact, provenance-tagged, and auditable.

For the 20-turn investigation: the conventional system’s per-turn cost escalates because each turn re-processes growing context through attention. By turn 20, context is roughly 60,000 tokens and each turn takes proportionally longer. VDR’s per-turn cost is flat — state is in knowledge bases, not the context window. Turn 20 costs the same as turn 1. Over 20 turns, the conventional system’s total time might be 10 minutes with degrading quality. VDR’s total time might be 80 seconds (20×4 seconds) with constant quality and accumulating knowledge.

For the financial portfolio analysis: 500 Q335 multiplications (position values), 500 Q335 additions (portfolio sum), correlation matrix computation (roughly 125,000 Q335 operations for 500×500 pairwise correlations). Total: roughly 126,000 Q335 operations. At millions of Q335 operations per second on GPU, this completes in well under one second. The conventional LLM cannot perform the task at all — 500 positions exceed context capacity and digit-by-digit arithmetic through token prediction would take thousands of tokens.

For academic grading: 150 essays, each requiring LLM judgment at 300 tokens per essay. The LLM cost is 45,000 tokens regardless of system because this is irreducible judgment work. But the statistical analysis — computing means, variances, percentiles, standard deviations across 150 students and multiple criteria — is pure Q335 arithmetic. A few thousand Q335 operations, completing in milliseconds on GPU. The conventional system cannot compute these statistics at all.

The pattern across all workloads: the LLM forward pass is slower per token on VDR, but the total LLM workload is 85 to 97 percent smaller, and the primitive and knowledge base workload that replaces it executes at GPU-parallel speed. The workloads where VDR is slower are those dominated by free-text generation with no structural component. The workloads where VDR is faster — often dramatically faster — are those involving data, computation, state management, and structured reasoning, which are precisely the professional workloads where accuracy and provenance matter most.

Appendix A: Q335 Operation Cost

A.1: Per-Operation Integer Cost

Operation	Q335 Integer Ops	float16 Ops	Ratio	GPU Parallelism	Effective Throughput Ratio
Addition	~22 (11 limb adds + carries)	1	$22 \times$	Full — uniform work per element	$\sim 22 \times$ slower per element, same utilization
Subtraction	~22 (11 limb subs + borrows)	1	$22 \times$	Full	$\sim 22 \times$
Multiplication	~200 (school-book 11×11 + accumulate)	1	$200 \times$	Full — uniform work per element	$\sim 200 \times$ slower per element
Comparison	2-22 (fast path to full limb scan)	1	$2\text{-}22 \times$	Full	$2\text{-}22 \times$ average
Division (by Q335)	~400 (multiply + shift + remainder)	1	$400 \times$	Full	$\sim 400 \times$
Surrogate softmax (per element)	~220 (sub + square + div)	~ 15 (exp + div)	$15 \times$	Full — no transcendentals	$\sim 15 \times$ slower but no divergence
ReLU	2 (compare + conditional copy)	2	$1 \times$	Full	Identical
Bit test (scope check)	1	N/A	—	Full	No conventional equivalent

A.2: Per-Token Forward Pass Cost

Component	Conventional (float16)	VDR (Q335 integer)	Ratio	Notes
Embedding lookup	1 read per dim	11 reads per dim (limb width)	$11 \times$	Memory-bound; limb width increases bandwidth
Attention QK^T	$d \times k$ multiplies per position pair	$d \times k \times 200$ int ops per pair	$200 \times$	Dominated by multiplication cost
Softmax / Surrogate	~ 15 ops per element (transcendental)	~ 220 ops per element (integer)	$15 \times$	VDR avoids transcendental divergence
Value mixing	$d \times v$ multiplies per position	$d \times v \times 200$ int ops per position	$200 \times$	Same structure, wider operands
Feedforward (ReLU)	2 ops per element	2 ops per element	$1 \times$	Identical — piecewise linear
Feedforward (linear)	$d \times ff$ multiplies	$d \times ff \times 200$ int ops	$200 \times$	Dominated by multiplication
Layer norm / rational scaling	~ 10 ops per element	~ 25 ops per element	$2.5 \times$	VDR avoids sqrt; uses mean absolute value
Total per token	$\sim 14B$ float ops	$\sim 14B \times \sim 150$ avg factor	$\sim 150 \times$	Weighted average across components

Per-token cost is roughly $150 \times$ higher. But token count is 85-97% lower. Net system cost depends on workload composition.

Appendix B: Token Count vs Operation Count

B.1: SRE Incident Investigation (Single Turn)

Category	Conventional	VDR	Notes
LLM tokens generated	3,000	210	93% reduction
Forward pass ops per token	14B float	$14B \times 150 = 2.1T$ int	Per-token cost higher

Category	Conventional	VDR	Notes
Total forward pass ops	42T float	441T int	VDR forward pass $\sim 10.5 \times$ more total ops
Primitive call ops	0 (no primitives)	$\sim 100\text{K}$ Q335 ops = $\sim 20\text{M}$ int	Computationally invisible
KB query ops	0 (attention-based search)	$\sim 50\text{K}$ int ops (scans + filters)	GPU-parallel, sub-millisecond
Prolog evaluation ops	0 (prose reasoning)	$\sim 10\text{K}$ int ops (unification + join)	GPU-parallel, sub-millisecond
Formatting ops	0 (all generated tokens)	0 (grammar-provided)	VDR: zero cost; conventional: ~ 600 tokens
Confidence computation	0 (hedging tokens)	~ 500 int ops	One VDR multiply per propagation step
Provenance logging	0 (no provenance)	$\sim 5\text{K}$ int ops (append events)	Append-only, negligible
Total computation	42T float	$\sim 441\text{T}$ int + $\sim 25\text{M}$ int	VDR: more total int ops for forward pass, negligible for everything else

B.2: 20-Turn Investigation

Turn	Conventional Forward Pass Ops	Conventional Attention Ops	VDR Forward Pass Ops	VDR Primitive + KB Ops	VDR Total
1	42T float	42T float attention	441T int	$\sim 25\text{M}$ int	$\sim 441\text{T}$ int
5	42T float	210T float attention	441T int	$\sim 25\text{M}$ int	$\sim 441\text{T}$ int
10	42T float	420T float attention	441T int	$\sim 25\text{M}$ int	$\sim 441\text{T}$ int
20	42T float	840T float attention	441T int	$\sim 25\text{M}$ int	$\sim 441\text{T}$ int
Cumulative	840T float gen	$\sim 8,820\text{T}$ float attn	8,820T int	$\sim 500\text{M}$ int	$\sim 8,820\text{T}$ int

Conventional cumulative: $\sim 9,660\text{T}$ float (generation + attention, growing quadratically). VDR cumulative: $\sim 8,820\text{T}$ int (flat per turn). At $150 \times$ per-operation cost, VDR's $8,820\text{T}$ int $\approx$ conventional equivalent of $\sim 59\text{T}$ float in raw per-op terms — but spread across a constant attention window rather than a growing one. The actual wall-clock time favors

VDR from roughly turn 5 onward because the conventional system’s attention cost grows quadratically while VDR’s stays flat.

B.3: Crossover Turn

Turn	Conventional Total Ops (float)	VDR Total Ops (int, at $150 \times$ factor)	VDR Equivalent Float Ops	Ratio (Conv/VDR)
1	84T	441T int = 2.94T float-equiv	2.94T	$28.6 \times$ (conventional faster)
3	378T	1,323T int = 8.82T float-equiv	8.82T	$42.9 \times$ (conventional faster per-op, but doing more total work)
5	840T	2,205T int = 14.7T float-equiv	14.7T	$57.1 \times$
10	2,940T	4,410T int = 29.4T float-equiv	29.4T	$100 \times$
20	9,660T	8,820T int = 58.8T float-equiv	58.8T	$164 \times$

The “ratio” column shows that VDR’s float-equivalent cost stays roughly proportional to turn count (linear), while conventional cost grows quadratically. The crossover in wall-clock time depends on the per-operation speed difference between Q335 int and float16. At $150 \times$ per-operation slowdown, VDR breaks even at roughly turn 7-10 for a 7B parameter model. For larger models with larger context windows, the crossover comes earlier because attention cost grows faster.

Appendix C: Grammar Decode Savings

C.1: Candidate Set Sizes by Slot Type

Slot Type	Typical Candidates	Full Vocab	Reduction Factor	Example
Boolean	2	50,000	$25,000 \times$	yes/no field
Visibility enum	3	50,000	$16,667 \times$	public/internal/owner only

Slot Type	Typical Candidates	Full Vocab	Reduction Factor	Example
Constraint class enum	4	50,000	$12,500 \times$	axiom/operational/legal/project
Status enum	5-8	50,000	$6,250\text{-}10,000 \times$	pending/running/completed/failed/killed
Builtin opcode	~300	50,000	$167 \times$	Primitive name selection
KB identifier	50-500	50,000	$100\text{-}1,000 \times$	Entity reference in scope
Relation type	~20	50,000	$2,500 \times$	uses/enables/implements/etc
Structural punctuation	1-5	50,000	$10,000\text{-}50,000 \times$	Pipe, comma, brace, bracket
Free text	50,000	50,000	$1 \times$	Prose content slots

C.2: Decode Cost per Token by Slot Type

Slot Type	Surrogate Softmax Ops (Q335)	Full Vocab Softmax Ops (float16)	Ratio	Notes
Boolean (2 candidates)	~440 int ops	~750K float ops	$1,700 \times$ cheaper	2 squares + 1 sum + 2 divides
Enum (4 candidates)	~880 int ops	~750K float ops	$850 \times$ cheaper	4 squares + 1 sum + 4 divides
Builtin opcode (300)	~66K int ops	~750K float ops	$11 \times$ cheaper	Even at Q335 width, small candidate set wins
KB identifier (200)	~44K int ops	~750K float ops	$17 \times$ cheaper	—
Free text (50K)	~11M int ops	~750K float ops	$0.07 \times$ ($15 \times$ more expensive)	Full vocab VDR more expensive per token

The per-token cost of constrained decode is dramatically cheaper. The per-token cost of unconstrained free-text decode is roughly $15 \times$ more expensive due to Q335 width. The net effect depends on the fraction of tokens that are grammar-constrained versus free-text.

C.3: Net Decode Cost by Response Type

Response Type	Total Tokens	Grammar Constrained %	Avg Constrained Candidates	Constrained Token Cost (Q335)	Free-Text Token Cost (Q335)	Total Decode Cost	Conventional Total Decode Cost	Net Ratio
JSON object (20 fields)	223	74%	~5	165 × 1.1K = 182K	58 × 11M = 638M	638M	223 × 750K = 167M	3.8 × (VDR slower)
Markdown table (20 × 5)	367	60%	~3	220 × 660 = 145K	147 × 11M = 1,617M	1,617M	367 × 750K = 275M	5.9 × (VDR slower)
Incident report	210	50%	~8	105 × 1.8K = 189K	105 × 11M = 1,155M	1,155M	3,000 × 750K = 2,250M	0.51 × (VDR faster)
SRE metrics table	80	85%	~4	68 × 880 = 60K	12 × 11M = 132M	132M	500 × 750K = 375M	0.35 × (VDR faster)

The critical insight: for VDR responses where the total token count is much lower than conventional (incident report: 210 vs 3,000), VDR is faster even with the per-token Q335 penalty, because the massive token reduction outweighs the per-token cost increase. For responses where VDR and conventional have similar total tokens (pure JSON generation), VDR is slower per token and the grammar savings partially but not fully compensate.

Appendix D: GPU Utilization Comparison

D.1: Utilization by Operation Type

Operation	Conventional GPU Utilization	VDR GPU Utilization	Reason for Difference
Matrix multiply	85-95% (tensor core optimized)	60-80% (integer, no tensor core)	Hardware optimization gap; integer path uses general ALUs

Operation	Conventional GPU Utilization	VDR GPU Utilization	Reason for Difference
Softmax (exponential)	40-60% (transcendental divergence, reductions)	80-95% (integer surrogate, uniform ops)	No transcendentals; pure arithmetic; uniform workload
Attention masking	90%+ (simple fill)	90%+ (identical operation)	Same operation both systems
Embedding lookup	70-90% (memory-bound)	60-80% (wider reads, memory-bound)	Q335 entries wider; more bandwidth per lookup
KB scan	N/A (no equivalent)	90%+ (columnar scan, coalesced access)	Standard GPU database operation
Scope filter	N/A (no equivalent)	95%+ (bitset test, trivial per-element)	One bit operation per element
Prolog unification	N/A (no equivalent)	70-85% (frontier-based, some divergence on var/atom paths)	Minor divergence on different term types
Grammar-constrained decode	N/A (always full vocab)	95%+ (tiny candidate set, trivial computation)	Orders of magnitude less work
Provenance append	N/A (no logging)	90%+ (atomic bump, bulk copy)	Append-only; minimal synchronization
Counter/bitset mutation	N/A (no live state)	95%+ (atomic int ops)	Single instruction per element

D.2: Aggregate Utilization by Workload Phase

Phase	Conventional	VDR	Notes
Token generation (forward pass)	70-85% average	55-75% average	VDR lower due to integer ALU vs tensor core
Data processing	N/A	85-95%	KB scans, filters, aggregations — GPU sweet spot
Logic evaluation	N/A	70-85%	Frontier-based Prolog; moderate divergence

Phase	Conventional	VDR	Notes
Structural output	Included in token generation at 70-85%	90%+ for grammar-provided; 55-75% for content slots	Grammar tokens free; content tokens same as forward pass
State management	N/A (in context window)	90%+	Counter/bitset/ring buffer ops trivial
Overall weighted	70-85%	65-85%	VDR slightly lower on forward pass, higher on everything else

Appendix E: Memory Layout

E.1: Q335 Storage Formats

Format	Limb Layout	Bits per Value	Memory per 1K Values	Alignment	Coalesced Access Pattern
$11 \times \text{u32}$ portable	[limb0..limb10] contiguous	52	44 KB	4-byte aligned	11 coalesced u32 reads per thread
$6 \times \text{u64}$ optimized	[limb0..limb5] contiguous	384	48 KB	8-byte aligned	6 coalesced u64 reads per thread
SoA limb-major	limb0[0..N], limb1[0..N], ...	352 or 384	44 or 48 KB	Vector-width aligned	Perfect coalescing — adjacent threads read adjacent memory
T0 dense tensor	SoA limb-major, implicit denominator	352 or 384 per entry	Row: d model $\times$ 44 bytes	Row-aligned	Standard dense matrix access

E.2: KB Fact Storage Layout

Column	Type	Bytes per Fact	Access Pattern	Index
fact kb id	u32	4	Coalesced scan	Predicate bucket offset
fact pred id	u32	4	Bucket lookup	Primary: pred id → offset/count
fact arg begin	u32	4	Indirect (to term pool)	—
fact arity	u16	2	Filter	—
fact turn	u32	4	Filter/sort	—
fact confidence	u32	4	Lookup	—
ref				
fact derivation	u32	4	Lookup	—
ref			(provenance)	
Total per fact	—	26 bytes	—	—

1 million facts: ~26 MB. Fits entirely in GPU global memory with room for indexes, terms, and working space. A predicate bucket scan over 10,000 facts is one coalesced read of ~260 KB — well within GPU memory bandwidth for sub-millisecond completion.

E.3: Arena Allocation Pattern

Arena	Contents	Typical Growth per Turn	Compaction Trigger
Limb arena	Q335 numerators, general rational limbs	10-100 KB	After session reset or every 100 turns
VDR node arena	VdrNode structs for active/rational values	1-10 KB	With limb arena
Remainder arena	AtomicR, CompositeR, FunctionalR payloads	1-50 KB	With limb arena
Term arena	New terms created during Prolog evaluation	5-50 KB	After Prolog query completion
Binding arena	Binding rows from unification frontiers	10-100 KB per query	After each query (most bindings are temporary)
Provenance arena	Event records	1-5 KB per turn	Periodically; append-only so less urgent

Arena	Contents	Typical Growth per Turn	Compaction Trigger
Live-state delta	Counter changes, buffer writes, cache updates	0.5-5 KB per turn	On session reset/clone kill

Total GPU memory for a typical active session: 50-500 MB including model weights, KB data, arenas, and working space. Well within modern GPU memory capacity (16-80 GB).

Appendix F: Concurrent Stream Execution

F.1: Stream Assignment

Stream	GPU Resources Used	Memory Accessed	Overlap With
Stream 0: LLM forward/decode	Tensor ALUs, shared memory	Weight tensors, activation buffers	Streams 1, 3, 4 (different memory)
Stream 1: KB query/scan	Integer ALUs, global memory	Fact tables, term pool, index tables	Streams 0, 2, 3 (different ALU units)
Stream 2: VDR primitives	Integer ALUs, global memory	Value arenas, live-state arrays	Streams 0, 3, 4
Stream 3: Grammar mask/prep	Integer ALUs, shared memory	Grammar state, candidate buffers	Streams 0, 1, 2
Stream 4: Provenance compact	Memory copy engines	Provenance arena, event buffers	All (uses DMA, not ALUs)

F.2: Turn Pipeline Timing (Estimated, SRE Single Turn)

Step	Component	Duration	Overlaps With
1. Parse request	CPU	0.1 ms	—
2. Resolve paths/scope	CPU	0.05 ms	—
3. Prefetch scoped facts	GPU stream 1	0.2 ms	Step 4 prep
4. LLM decode start	GPU stream 0	— (continuous)	Steps 5-8

Step	Component	Duration	Overlaps With
5. Command token emitted (~every 8 tokens)	CPU detects	0.01 ms	—
6. Lower to opcode batch	CPU	0.02 ms	LLM continues on stream 0
7. Execute primitives	GPU stream 2	0.1-1 ms	LLM decode on stream 0
8. Grammar masks update	GPU stream 3	0.05 ms	All other streams
9. LLM continues with results	GPU stream 0	— (continuous)	—
10. Provenance commit	GPU stream 4	0.1 ms	LLM decode on stream 0
11. Response assembly	CPU	0.2 ms	—
Total wall-clock	—	~2-4 seconds	(dominated by LLM decode of 210 tokens)

The primitive execution, KB queries, grammar updates, and provenance logging are all overlapped with LLM decode. Their wall-clock cost is effectively zero because they complete while the LLM is still generating the next token.

Appendix G: Bottleneck Analysis

G.1: Bottleneck Severity by Workload

Bottleneck	SRE Incident	Financial Portfolio	Academic Grading	Legal Contract	Code Migration
LLM forward pass (Q335 cost)	Medium (210 tokens)	Low (600 tokens but mostly command)	High (57K tokens, judgment-dominated)	Medium (1,130 tokens)	Medium-High (6,700 tokens)
Active value management	Low (<0.1% spill expected)	Low (simple arithmetic)	Low (statistics on integers)	Low (currency amounts)	Low (no VDR arithmetic)
Prolog frontier size	Low (small rule sets)	Low (simple queries)	Low (score queries)	Medium (cross-reference queries)	Low (dependency queries)

Bottleneck	SRE Incident	Financial Portfolio	Academic Grading	Legal Contract	Code Migration
Host-device sync	Medium (20+ command tokens)	Medium (15+ command tokens)	Low (few commands, mostly judgment)	Medium (15+ command tokens)	High (200+ files, many commands)
Memory (arena growth)	Low (small state)	Low (moderate data)	Low (150 student KBs)	Low (document structure)	Medium (200 file KBs)

G.2: Mitigation Strategies

Bottleneck	Strategy	Mechanism	Supplementary Spec Reference
LLM Q335 forward pass cost	Reduce token count further through grammar and primitive coverage	More structural tokens grammar-provided; more operations as primitives	ML1-ML8, GD1-GD4
Active value spill	Q-basis reprojection at declared thresholds	AK6 batch reprojection; keep spill rate below 1%	PF6, AK6
Prolog frontier explosion	Query planning on CPU; scope restriction; predicate indexing	PG7 deterministic ordering; IX1-IX2 indexes; SC2-SC3 scope filters	PG1-PG7, IX1-IX10
Host-device sync	Batch command tokens; overlap with LLM decode	OR1-OR2 concurrent streams; batch lowering	OR1-OR2
Arena memory growth	Periodic compaction; session reset	MM3 compaction; SN4 reset	MM1-MM3, SN1-SN5

Appendix H: Implementation Phase Performance Targets

H.1: Per-Phase Performance Milestones

Phase	Contents	Target Throughput	Validation
Phase 1	Q335 arithmetic, dense tensor kernels, surrogate attention, confidence	1M Q335 muls/sec; surrogate softmax 10K rows/sec at d=512	Arithmetic correctness vs Python VDR reference; softmax sum-to-one verified
Phase 2	KB metadata + fact tables, scope filtering, interned terms, fact queries	10M facts/sec scan rate; scope filter <0.1ms for 100K facts	Query correctness vs Python Prolog reference
Phase 3	Frontier unification, rule joins, binding buffers, live-state primitives	100K unifications/sec; joins <1ms for 1K $\times$ 1K frontiers	Derivation correctness; binding consistency
Phase 4	Active spill arena, normalization, functional remainders	Normalization <10ms for 10K active nodes; fn sqrt at depth 8 <1ms	Normalized forms match Python reference; convergence rates verified
Phase 5	Full command pipeline, grammar decode, provenance, sessions	End-to-end prompt <5sec for 210-token SRE scenario	Full system benchmark against Python prototype

H.2: Phase Dependencies

Phase	Depends On	Enables
Phase 1	None (foundation)	All subsequent phases (Q335 is the arithmetic substrate)
Phase 2	Phase 1 (Q335 for fact values, confidence refs)	Phase 3 (Prolog needs KB data on GPU)
Phase 3	Phases 1+2 (arithmetic + KB data)	Phase 5 (command pipeline needs Prolog and live-state)
Phase 4	Phase 1 (arithmetic; extends to active/irregular)	Phase 5 (full system needs active value handling)
Phase 5	All prior phases	Production readiness

Appendix I: Integer ALU Throughput by GPU Generation

I.1: Native Integer Operations per Clock per SM

GPU Archi- tecture	Generation Year	INT32 ops/clock/SM	INT64 ops/clock/SM	Concurrent INT+FP	Notes
Kepler (GK110)	2013	192	64	No — shared pipeline	INT and FP compete for same ALUs
Maxwell (GM204)	2014	128	4	No	INT64 severely limited
Pascal (GP100)	2016	64	32	No	Better INT64 but still shared
Volta (GV100)	2017	64	32	Yes — separate INT datapath	First concurrent INT+FP execution
Turing (TU102)	2018	64	16	Yes	Concurrent INT+FP standard
Ampere (GA100)	2020	64	32	Yes	Improved INT64
Ada Lovelace (AD102)	2022	128	64	Yes	Doubled INT throughput
Hopper (GH100)	2022	128	64	Yes	Data center focused; strong INT path
Blackwell (GB202)	2024	128+	64+	Yes	Latest generation; INT path continues widening

I.2: Q335 Operation Throughput Estimates by Architecture

Architecture	SMs	INT32 ops/sec (peak)	Q335 adds/sec (est.)	Q335 muls/sec (est.)	Q335 muls as % of FP16 tensor peak
Volta (GV100, 80 SM)	80	7.9T	360M	40M	0.02%
Ampere (A100, 108 SM)	108	9.7T	440M	50M	0.016%
Hopper (H100, 132 SM)	132	16.9T	770M	85M	0.0085%
Ada (RTX 4090, 128 SM)	128	16.4T	745M	83M	0.05%
Blackwell (B200, est.)	160+	25T+	1.1B+	125M+	~0.01%

Q335 addition: ~22 INT32 ops per add. Q335 multiplication: ~200 INT32 ops per multiply (schoolbook). The percentage of FP16 tensor peak is low because tensor cores are specialized fixed-function units. But absolute throughput — 50-125 million Q335 multiplies per second — is sufficient for VDR workloads where the total multiply count per prompt is in the thousands to millions, not trillions.

I.3: Time to Complete VDR Workload Components (H100 Estimates)

Workload	Q335 Operations	Time at 85M muls/sec	Conventional Float Equivalent	Notes
500 portfolio multiplications	500 muls	0.006 ms	Cannot perform (context overflow)	Instantaneous
Correlation matrix (500 × 500)	125K muls + 125K adds	1.5 ms	Cannot perform	Sub-second
Sort 500 items by Q335 key	~4,500 comparisons	0.05 ms	Hundreds of tokens of generated prose	Instantaneous
Sum 500 Q335 values	500 adds	0.001 ms	Hundreds of tokens digit accumulation	Instantaneous

Workload	Q335 Operations	Time at 85M muls/sec	Conventional Float Equivalent	Notes
Softmax surrogate (1000 elements)	1K muls + 1K adds + 1K divs	0.04 ms	~15 ms for transcendental softmax	Faster despite wider operands
Softmax surrogate (4 elements, grammar)	4 muls + 4 adds + 4 divs	0.0002 ms	~0.06 ms for full-vocab softmax	$300 \times$ faster due to candidate reduction
Q335 matrix multiply (512×512)	~134M muls	1.6 sec	~0.01 ms on tensor cores	This is the real bottleneck
Prolog unification (1000 candidates)	~5K comparisons	0.06 ms	N/A — no conventional equivalent	Sub-millisecond

The 512×512 matrix multiply row shows the honest cost: large matrix operations in Q335 are dramatically slower than tensor core float. The VDR strategy is to minimize how often these occur (fewer LLM tokens) rather than to make them individually competitive.

Appendix J: Memory Bandwidth Analysis

J.1: Bandwidth Requirements by Operation

Operation	Bytes Read per Element	Bytes Written per Element	Elements per Typical Call	Total Bandwidth	H100 Bandwidth (3.35 TB/s) Utilization
Q335 add ($11 \times u32$)	88 (two operands)	44 (one result)	1,000 (vector add)	132 KB	0.004% — entirely compute-bound
Q335 mul ($11 \times u32$)	88	44 + possible 44 (remainder)	1,000	176 KB	0.005% — compute-bound
KB fact scan	26 per fact (all columns)	0.125 per fact (bitmask)	10,000	261 KB	0.008% — trivial

Operation	Bytes Read per Element	Bytes Written per Element	Elements per Typical Call	Total Bandwidth	H100 Bandwidth (3.35 TB/s) Utilization
Scope filter (bitset)	4 per fact (kb id) + 8 (bitset word)	0.125 per fact (result bit)	10,000	41 KB	0.001% — trivial
Term pool lookup	12 per term (tag + 2 payloads)	0 (read-only)	5,000	60 KB	0.002% — trivial
LLM embedding lookup (Q335)	$44 \times d$ model per token	$44 \times d$ model per token	1 token	44×4096 = 176 KB	0.005% — trivial
LLM weight matrix row (Q335)	$44 \times d$ model per row	44 per output element	d model = 4096	176 KB per row	Accumulated over matrix: significant
Provenance event append	0 (write-only)	32 per event	50 per turn	1.6 KB	0.00005% — invisible

VDR primitive operations are universally compute-bound, not memory-bound. The memory bandwidth of modern GPUs is vastly underutilized by VDR arithmetic and KB operations. The LLM forward pass is the only component that approaches significant bandwidth utilization, due to weight matrix reads — identical in structure to conventional LLM inference but with wider entries.

J.2: Working Set Sizes

Component	Size	Fits in L2 Cache (50 MB on H100)?	Notes
Active session live state	21 KB typical	Yes	Counters, bitsets, buffers — trivial
Session scope chain	64 bytes ($16 \times$ u32)	Yes	Always in registers or L1
Predicate bucket index	~4 KB for 1000 predicates	Yes	One u32 offset + u32 count per predicate
Fact table for one predicate bucket	2.6 KB per 100 facts	Yes for most buckets	Hot predicates stay cached

Component	Size	Fits in L2 Cache (50 MB on H100)?	Notes
Term pool (active subset)	12 KB per 1000 terms	Yes	Repeatedly accessed during unification
Grammar state + candidates	<1 KB	Yes	Tiny; always in fast memory
Binding frontier (Prolog)	64 bytes per binding row $\times$ frontier size	Yes up to ~750K rows	Moderate frontiers fit; large frontiers spill
Q335 tensor row (d=4096)	176 KB	Partially	Typical matrix row doesn't fit L2; streamed from global
Full KB fact store (100K facts)	2.6 MB	Yes	Entire moderate KB fits in L2
Full KB fact store (1M facts)	26 MB	Yes on H100 (50 MB L2)	Large KBs fit in L2 on data-center GPUs

Most VDR data structures fit in GPU L2 cache. The LLM weight matrices do not — they stream from global memory, same as conventional inference. The difference is that VDR's non-LLM operations (KB queries, Prolog, primitives) all operate on cache-resident data, giving them near-peak throughput.

Appendix K: Karatsuba Threshold Analysis

K.1: Multiplication Algorithm Selection

Limb Count	Schoolbook Ops	Karatsuba Ops	Toom-3 Ops	Recommended	Reason
6 (Q335 u64)	36 muls + 60 adds	27 muls + ~90 adds	Not applicable (too small)	Schoolbook	Simpler; no recursion overhead; uniform GPU execution

Limb Count	Schoolbook Ops	Karatsuba Ops	Toom-3 Ops	Recommended	Reason
11 (Q335 u32)	121 muls + ~200 adds	~50 muls + ~300 adds	~40 muls + ~400 adds	Schoolbook or Karatsuba	Karatsuba marginal win; depends on GPU add/mul ratio
22 (Q670, double frame)	484 muls + ~800 adds	~120 muls + ~700 adds	~80 muls + ~900 adds	Karatsuba	Clear win at this size
44+ (very large rational)	1936+ muls	~300 muls + ~1500 adds	~150 muls + ~2000 adds	Toom-3 or Karatsuba	Large operands justify recursive splitting

K.2: GPU Considerations for Algorithm Selection

Factor	Favors Schoolbook	Favors Karatsuba/Toom
Thread divergence	Schoolbook: all threads do identical work	Karatsuba: recursive splitting causes thread divergence
Instruction count	Higher mul count	Higher add count + more complex control
Register pressure	Low — simple loop	Higher — intermediate values from splitting
Shared memory use	Minimal	May need shared memory for intermediate buffers
Warp utilization	100% — perfectly uniform	80-95% — depends on recursion balance
Operand size ≤ 11 limbs	Optimal or near-optimal	Marginal improvement doesn't justify complexity
Operand size > 20 limbs	Increasingly suboptimal	Clearly superior

For Q335 (11 u32 limbs), schoolbook is recommended for GPU due to perfect uniformity. The supplementary specification's AK2 correctly notes "optional Karatsuba/Toom thresholds for large operands" — meaning the system defaults to schoolbook for Q335 and switches only for larger intermediate values that arise during general rational arithmetic or active value normalization.

Appendix L: Scope Filter Strategy Comparison

L.1: Strategy Performance by KB Tree Size

Strategy	Mechanism	Setup Cost	Per-Fact Cost	Best For	GPU Kernel Pattern
Linear scan	Check each fact's kb id against scope chain array	0	$O(\text{scope depth})$ comparisons	Tiny trees (<50 KBs)	Each thread: loop over scope chain
Binary search	Sort scope chain; binary search per fact	$O(\text{depth} \times \log(\text{depth}))$ sort	$O(\log(\text{scope depth}))$ comparisons	Small-medium trees	Each thread: binary search
Bitset	Set bits for visible KB IDs in u64 array	$O(\text{scope depth})$ bit-sets	1 bit-test per fact	Medium trees (<4096 KBs)	Each thread: one AND + zero-test
DFS interval	Precomputed in/out interval per KB; interval containment test	$O(\text{tree size})$ DFS precomputation (once)	2 integer comparisons per fact	Large trees, subtree queries	Each thread: two comparisons
Interval + shadow	DFS interval with shadow-resolution pass for mounts	$O(\text{tree size}) + O(\text{mount count})$	2-4 integer comparisons per fact	Large trees with mounts	Two-pass: interval test then shadow check

L.2: Visibility Check Cost in Context

System Size	KBs	Facts	Scope Depth	Strategy	Filter Time (est. H100)	As % of Prompt Time
Personal session	5	500	4	Linear scan	0.001 ms	0.00003%
Team workspace	50	5,000	6	Bitset	0.005 ms	0.00017%

System Size	KBs	Facts	Scope Depth	Strategy	Filter Time (est. H100)	As % of Prompt Time
Department	200	50,000	6	Bitset	0.02 ms	0.0007%
Organization	2,000	500,000	8	DFS interval	0.1 ms	0.003%
Enterprise (full)	10,000	2,000,000	10	DFS interval + shadow	0.5 ms	0.017%

Structural safety (visibility checking) costs at most 0.017% of total prompt time even for the largest enterprise deployment. This confirms VDR-16’s claim that structural safety is computationally free relative to the language model’s operating cost.

Appendix M: Prolog Join Strategy Selection

M.1: Join Strategy Performance Characteristics

Strategy	Build Cost	Probe Cost per Row	Memory	Best Frontier Size	GPU Efficiency
Nested loop	0	$O(N)$ per probe	0 extra	<100 rows	Low (thread divergence)
Hash join	$O(N)$ build hash table	$O(1)$ avg per probe	$O(N)$ hash table	100-100K rows	High (parallel build + probe)
Sort-merge	$O(N \log N)$ sort both sides	$O(1)$ amortized per row	$O(1)$ extra after sort	10K-1M rows	High (GPU sort is efficient)
Bitmap semijoin	$O(N)$ build bitmap	$O(1)$ per probe (bit test)	$O(\text{domain}/8)$ bitmap	Any size, when join is filter	Very high (trivial per-element)

M.2: Typical VDR Query Join Sizes

Query Type	Goal Count	Avg Frontier After Goal 1	After Goal 2	After Goal 3	Recommended Strategy
Simple fact lookup	1	1-10	N/A	N/A	Direct scan (no join)
Scoped predicate query	1	10-100	N/A	N/A	Direct scan with scope filter
Two-goal rule (e.g., parent(X,Y), age(Y,A))	2	50-500	10-100	N/A	Hash join
Session scoring rule	3	5-20 (counter values)	5-20	1-5	Nested loop (tiny frontiers)
Contradiction detection	2	50-200 (findings)	5-20 (conflicting pairs)	N/A	Hash join
Cross-reference resolution	2	100-1000 (connections)	10-100 (matched)	N/A	Hash join or sort-merge
Full audit query	1-2	1K-100K (audit events)	100-10K (filtered)	N/A	Sort-merge or bitmap

Most VDR queries produce moderate frontier sizes (10-1000 rows) where hash join is optimal. Pathological frontier explosion (>100K rows) only occurs in broad audit queries, which are infrequent and tolerance of higher latency is acceptable.

M.3: Join Execution Time Estimates (H100)

Frontier Sizes	Strategy	Build Time	Probe Time	Total	Notes
50 × 50	Hash join	0.002 ms	0.002 ms	0.004 ms	Trivial
500 × 500	Hash join	0.01 ms	0.01 ms	0.02 ms	Sub-millisecond
5K × 5K	Sort-merge	0.1 ms	0.05 ms	0.15 ms	Well within interactive latency
50K × 50K	Sort-merge	1 ms	0.5 ms	1.5 ms	Acceptable for audit queries

Frontier Sizes	Strategy	Build Time	Probe Time	Total	Notes
500K × 500K	Sort-merge	15 ms	5 ms	20 ms	Large audit query; batch acceptable

Appendix N: Functional Remainder Evaluation Cost

N.1: Per-Function Evaluation Cost

Function (fn id)	Algorithm	Ops per Depth Step	Depth for 100 Digits	Total Q335 Ops	Time (est. H100)
SQRT (Newton)	$x \{n+1\} = (x n + S/x n)/2$	1 div + 1 add + 1 shift = ~600 ops	8	~4,800	0.06 ms
EXP (Taylor)	$\Sigma x^n/n!$ cumulative	1 mul + 1 div + 1 add = ~600 ops per term	45	~27,000	0.3 ms
LOG (series + reduction)	$\ln(1+x)$ series with argument reduction	1 mul + 1 div + 1 add per term	340 (at x near 1; much less with reduction)	~20,000 with reduction	0.2 ms
SIN (Taylor odd)	$\Sigma (-1)^n x^{(2n+1)}/(2n+1)!$	2 muls + 1 div + 1 add per term	45	~36,000	0.4 ms
COS (Taylor even)	$\Sigma (-1)^n x^{(2n)}/(2n)!$	2 muls + 1 div + 1 add per term	45	~36,000	0.4 ms
ARCTAN (Taylor + Machin)	Series with identity reduction	1 mul + 1 div + 1 add per term	170 (without reduction)	~15,000 with Machin	0.2 ms
ZETA (Borwein)	Weighted alternating sum with 3^{-n} con- vergence	1 mul + 1 add per term	210 (for any $s \geq 2$)	~42,000	0.5 ms

N.2: Batch Evaluation Throughput

Scenario	Function Calls	Total Q335 Ops	Time (H100, sequential)	Time (H100, batched parallel)	Notes
Single sqrt evaluation	1	4,800	0.06 ms	0.06 ms	Single thread sufficient
DFT twiddle factors (N=1024)	1024 sin + 1024 cos	74M	870 ms sequential	~7 ms batched (1024 parallel)	Parallelism across twiddle factors
Kepler orbit (50 Newton steps)	50 sqrt-like iterations	240K	3 ms	0.5 ms batched	Pipeline Newton iterations
Q335 constant resolution (22 constants)	22 lookups	~0 (precomputed)	<0.001 ms	<0.001 ms	Constants stored as Q335 limb vectors
Softmax via Taylor exp (1000 elements)	1000 exp evaluations	27M	320 ms sequential	~0.3 ms batched (1000 parallel)	Why rational surrogate is preferred

The last row demonstrates why the rational surrogate softmax (ML6) is preferred over Taylor exponential softmax (ML5) for GPU execution. Batched Taylor exp for 1000 elements takes ~0.3 ms. The rational surrogate for the same 1000 elements takes ~0.04 ms — an $8 \times$ speedup from avoiding transcendental evaluation entirely, with the same sum-to-one guarantee.

Appendix O: Determinism Verification

O.1: Sources of GPU Nondeterminism and Mitigations

Source	Where It Occurs	Impact on Correctness	Mitigation	Cost of Mitigation
Atomic operation ordering	Counter increments, arena allocation	Counter final value correct; intermediate order varies	Final values deterministic; order canonicalized by (turn, batch seq)	Zero — final values same regardless of order
Block execution order	All parallel kernels	Results computed in arbitrary block order	Post-sort by canonical key where ordering matters	One radix sort per kernel with ordering requirement
Warp-level reduction order	Sum reductions, max reductions	Floating-point: result varies. Integer: result varies if overflow possible	Q335 addition is associative and commutative for closed values; canonical reduction order for active values	Zero for closed (mathematically identical); sort cost for active
Memory allocation order	Arena bump allocation across blocks	Handle assignment order varies	External handles assigned by CPU in canonical order after kernel completion	One sequential pass over kernel outputs
Provenance event order	Parallel event emission	Events from same turn may appear in arbitrary order	Post-sort by (turn, batch seq, item seq)	One radix sort per turn's events

O.2: Determinism Guarantees by Component

Component	Deterministic?	Condition	Verification Method
Q335 addition	Yes	Always (integer arithmetic)	Bit-identical output for bit-identical input
Q335 multiplication	Yes	Always	Same

Component	Deterministic?	Condition	Verification Method
Q335 comparison	Yes	Always	Same
KB fact query result set	Yes	Same scope, same facts	Set equality (order may vary; canonical sort applied)
Prolog first-solution	Yes	Canonical candidate ordering imposed	CPU selects from fully-computed parallel results
Prolog find-all	Yes (set)	Result set identical; ordering canonical	Sort by canonical key
Confidence propagation	Yes	Same formula, same inputs	Exact VDR arithmetic
Grammar mask	Yes	Same grammar state, same candidates	Bit-identical mask
Provenance event content	Yes	Same operations, same inputs	Content identical; ordering canonicalized
Session snapshot	Yes	Same live state	Byte-identical snapshot for same logical state
Forward pass output	Yes	Same weights, same input, same token count	Bit-identical — integer arithmetic is platform-independent

The last row is the fundamental advantage: VDR’s forward pass produces bit-identical outputs on any hardware because integer arithmetic has no platform-dependent rounding behavior. Conventional float-based forward passes are not deterministic across platforms.

Appendix P: Power-of-Two Split Detail

P.1: Q335 Multiplication Split Mechanics

Step	Operation	Operand Size	Result Size	GPU Implementation
1. Full multiply	$A \times B$	335 bits $\times$ 335 bits	670 bits	Schoolbook on $11 \times u32$ limbs $\rightarrow$ 22 limbs

Step	Operation	Operand Size	Result Size	GPU Implementation
2. Extract quotient	Bits [670:335] of product	335 bits	335 bits (11 limbs)	Right-shift by 335 bits = skip lower 10.47 limbs → limb-aligned extraction
3. Extract remainder	Bits [334:0] of product	335 bits	335 bits (11 limbs)	Mask lower 335 bits → limb-aligned extraction
4. Check remainder	remainder == 0?	335 bits	1 bit	OR-reduce all 11 limbs; zero-test on result
5a. If zero	Result is closed Q335	quotient only	11 limbs	Store quotient; set tag = QFRAME; set r ref = 0
5b. If nonzero	Result is active	quotient + remainder	22 limbs + node	Store quotient as V; allocate AtomicR in arena; store remainder; set tag = ACTIVE

P.2: Bit Alignment Within Limbs

Layout	Limb Width	Limbs for 335 bits	Wasted Bits	Split Position within Limb Array
$11 \times \text{u32}$	32	11 (352 bits)	17	Split at bit 335 falls within limb 10 (bits 320-351); requires intra-limb shift
$6 \times \text{u64}$	64	6 (384 bits)	49	Split at bit 335 falls within limb 5 (bits 320-383); requires intra-limb shift

The intra-limb shift is one shift + one mask operation per split — trivial cost. The “wasted” bits above 335 in the top limb are always zero in a properly normalized Q335 value. After multiplication, the product occupies up to 670 bits, spread across 22 u32 limbs or 12 u64 limbs. The quotient extraction starts at limb offset $\text{ceil}(335/32) = 11$, and the remainder is the lower 11 limbs with the top limb masked.

Appendix Q: Tensor Core Feasibility Analysis

Q.1: Why Tensor Cores Cannot Be Used Directly

Tensor Core Property	Float16/BF16/TF32	Q335 Integer	Incompatibility
Operand width	16/16/19 bits	352 bits	$22 \times$ wider; doesn't fit tensor core input registers
Accumulator width	32 bits	704 bits (full product)	$22 \times$ wider; accumulator overflow
Rounding mode	Round-to-nearest-even	Exact (no rounding)	Tensor cores round by design
Matrix tile size	16×16 or 8×8	Same logical tile	Tiles contain $11 \times$ more data per element
Hardware function	Fused multiply-accumulate	Multi-precision multiply + exact accumulate	Different computational primitive

Q.2: Hypothetical Integer Tensor Core Design

Feature	Description	Benefit for VDR
Wide integer MAC	Multiply-accumulate on 352-bit operands with 704-bit accumulator	Direct Q335 matrix multiply support
Carry-chain hardware	Hardware carry propagation across limb boundaries	Eliminates software carry loop
Exact split unit	Fixed-position bit extraction at configurable split point	Hardware Q335 quotient/remainder extraction
Multi-precision register file	Registers holding 352+ bit values	Eliminates limb-by-limb register management
Estimated speedup over INT32 ALU path	$50\text{-}200 \times$ for matrix operations	Would close the gap with float tensor cores

This is engineering, not physics. If exact arithmetic workloads create sufficient market demand, GPU vendors can add integer tensor cores. The computational patterns (multiply-accumulate on fixed-width wide integers with exact accumulation) are simpler than the floating-point tensor core patterns (which must handle denormals, infinities, NaN propagation, and multiple rounding modes).

Appendix R: Spill Threshold Decision Tree

R.1: Active Value Monitoring

Metric	Measurement	Source	Check Frequency
Active fraction per tensor	count(tag == ACTIVE) / total elements	Per-tensor tag scan	Every tensor operation
Arena growth rate	bytes allocated / turns elapsed	Arena bump pointer delta	Every turn
Maximum active tree depth	max(depth) across active nodes	VdrNode depth field	Every normalization pass
Closed-to-active conversion rate	new active nodes / total operations	Counter on spill events	Every operation batch

R.2: Threshold Actions

Active Fraction	Tensor Class	Action	Performance Impact
0%	T0 (shared-frame dense)	No action; optimal path	Baseline — highest throughput
0.01% - 0.1%	T0 with sparse spill tags	Tag active entries; process on slow path	Negligible — fast path dominates
0.1% - 1%	T1 (dense + spill tags)	Per-element tag check on operations	1-5% overhead from tag branching
1% - 5%	T1 with tile-local spill tables	Group active entries per tile; batch process	5-15% overhead
5% - 10%	T1 transitioning to T2	Consider reprojection to reduce active count	15-30% overhead; reprojection may be cheaper
>10%	T2 (sparse active) or host intervention	Switch to sparse representation or reproject entire tensor	30-50% overhead; reprojection strongly recommended

R.3: Reprojection Decision

Trigger	Condition	Action	Error Bound	Provenance
Denominator budget exceeded	vdr denom bits(value) > budget bits	B136 vdr reproject qbasis	Exact bound computed and stored	fact(reprojection, value ref, old bits, new bits, error bound, turn)
Active fraction exceeded	tensor active fraction > threshold	Batch AK6 reprojection on all active entries	Per-entry exact bound	Batch provenance events
Training step checkpoint	Every N training steps	Reproject all parameters to Q335 base	Per-parameter exact bound	Checkpoint provenance
Manual trigger	User or system policy	Selective or full reprojection	Exact bound	Policy-driven provenance

Every reprojection is logged with exact error bounds. No silent precision loss occurs anywhere in the system. This is the fundamental distinction from floating-point: float silently truncates on every operation; VDR reprojection is a declared, auditable, bounded precision decision.

Appendix S: Host-Device Synchronization Cost

S.1: Synchronization Points per Turn

Sync Point	Direction	Payload Size	Latency	Frequency per Turn	Overlap Possible
Scope chain transfer	Host → Device	64 bytes	5-10 μ s	1	Yes (during LLM prefill)
Command token lowering	Host → Device	~32 bytes per command	5-10 μ s	5-50 per turn	Yes (during LLM decode)
Primitive result retrieval	Device → Host	4-64 bytes per result	5-10 μ s	5-50 per turn	Yes (buffered, batch retrieved)
Grant verification	Host-only	0 (no transfer)	1-5 μ s	5-50 per turn	Yes (CPU-parallel)

Sync Point	Direction	Payload Size	Latency	Frequency per Turn	Overlap Possible
Path resolution	Host-only	0 (no transfer)	1-10 μ s per path	5-50 per turn	Yes (cached after first resolution)
Grammar state update	Host → Device	16 bytes per update	5-10 μ s	10-100 per turn	Yes (during LLM decode)
Provenance batch commit	Device → Host (optional)	1-5 KB per turn	10-50 μ s	1 per turn	Yes (stream 4, overlapped)

S.2: Total Synchronization Overhead per Turn

Scenario	Command Tokens	Sync Events	Total Sync Latency (non-overlapped)	Total Sync Latency (overlapped with LLM decode)	As % of Turn Time
Simple query (1 command)	1	~5	0.05 ms	~0 ms (fully overlapped)	~0%
SRE investigation turn (10 commands)	10	~25	0.25 ms	0.05 ms (mostly overlapped)	0.002%
Complex analysis (30 commands)	30	~70	0.7 ms	0.15 ms	0.005%
Bulk migration (100 commands)	100	~200	2.0 ms	0.5 ms	0.02%

Host-device synchronization is not a significant bottleneck. The concurrent stream model overlaps most synchronization with LLM decode. The residual non-overlapped sync is sub-millisecond even for complex turns.

Appendix T: Comparison with Approximate GPU Approaches

T.1: VDR vs Quantized Integer Inference

Aspect	INT8 Quantized Inference	INT4 Quantized Inference	VDR Q335 Integer
Operand width	8 bits	4 bits	352 bits
Precision	~2 decimal digits	~1 decimal digit	~100 decimal digits
Accuracy loss from quantization	0.1-2% on benchmarks	1-5% on benchmarks	0% (exact)
Tensor core compatible	Yes (INT8 tensor cores)	Yes (INT4 on some architectures)	No (too wide)
Throughput vs FP16	$2 \times$ (8-bit fits $2 \times$ per register)	$4 \times$ (4-bit fits $4 \times$ per register)	$0.005 \times$ ($22 \times$ wider, no tensor cores)
Reproducible across platforms	No (quantization rounding varies)	No	Yes (integer arithmetic is exact)
Provenance	None	None	Full — every value traced
Suitable for	Production inference (approximate OK)	Edge deployment (approximate OK)	Exact computation with audit requirements

T.2: VDR vs Mixed-Precision Training

Aspect	FP16/FP32 Mixed Precision	BF16/FP32 Mixed Precision	VDR Q335 Exact
Forward pass precision	16-bit float	16-bit bfloat	352-bit integer
Accumulation precision	32-bit float	32-bit float	704-bit integer (exact)
Loss scaling required	Yes (to prevent underflow)	Less often	No (no underflow possible)
Gradient precision	16-bit (lossy)	16-bit (lossy)	Exact integer
Checkpoint reproducibility	No (platform-dependent)	No	Yes (bit-identical)
Training throughput	Baseline	~Same as FP16	~150 $\times$ slower per token
Training token count for equivalent work	Baseline	Same	85-97% fewer tokens needed

Appendix U: Energy and Cost Estimates

U.1: Energy per Operation

Operation	Energy (pJ, estimated)	Source
FP16 multiply (tensor core)	0.1-0.3	Published GPU power analysis
INT32 multiply (ALU)	0.1-0.5	Comparable to FP32 ALU
Q335 multiply (schoolbook, 200 INT32 ops)	20-100	$200 \times$ INT32 single multiply
Q335 add (22 INT32 ops)	2-10	$22 \times$ INT32 single add
DRAM read (64 bytes, one cache line)	10-50	GPU memory subsystem
L2 read (64 bytes)	1-5	On-chip cache

U.2: Energy per Prompt

System	Tokens	Ops per Token	Energy per Op	Total Energy	Notes
Conventional (FP16, 7B model)	3,000	14B	0.2 pJ	8.4 J	Forward pass dominated
VDR (Q335, 7B model)	210	$14B \times 150$	50 pJ avg	22 J	Higher per-op, fewer tokens
VDR primitive + KB ops	—	~25M int ops	0.3 pJ	0.0075 J	Negligible energy cost
Conventional total	—	—	—	8.4 J	—
VDR total	—	—	—	~22 J	—

VDR uses roughly $2.6 \times$ more energy per prompt than conventional for a single turn. Over 20 turns where conventional attention cost grows quadratically, VDR’s flat energy profile catches up. The energy cost is dominated by the LLM forward pass in both systems. VDR’s primitive and KB operations are energetically negligible.

U.3: Cost per Prompt (Cloud GPU Pricing)

Configuration	GPU	Cost per Hour	Time per Prompt (SRE single turn)	Cost per Prompt	Notes
Conventional (FP16, H100)	1 × H100	\$3.00	~30 sec (3000 tokens)	\$0.025	Standard inference
VDR (Q335, H100)	1 × H100	\$3.00	~4 sec (210 tokens + primitives)	\$0.0033	7.5 × cheaper per prompt
VDR (Q335, H100, 20-turn)	1 × H100	\$3.00	~80 sec total	\$0.067	Flat per turn
Conventional (20-turn)	1 × H100	\$3.00	~600 sec total (growing attention)	\$0.50	7.5 × more expensive

The per-prompt cloud cost is lower for VDR despite the per-operation slowdown, because the total GPU-seconds are lower. The token reduction dominates the per-operation cost increase. Over multi-turn interactions, the cost advantage grows because VDR’s time is flat per turn while conventional time grows with context.

Addendum: Case Study — SRE Prometheus Investigation at Scale

The Scenario

An SRE is troubleshooting a production latency spike. Prometheus is at prom.internal:9090 and exposes a REST API. The dataset is 1MB of JSON containing 15-minute time series for 200 endpoints across 4 service clusters. The SRE wants to filter endpoints by latency threshold, correlate latency spikes with deployment events, identify the top degraded endpoints per cluster, store the results in a versioned project structure for comparison with future runs, and export a summary for the postmortem.

The task involves network operations (prometheus fetch), file-like data handling (1MB JSON), parsing, filtering, sorting, correlation joins, statistical computation, Prolog rule evaluation, operational environment commands (Python script execution for complex transforms), live-state management (counters, queues, ring buffers for tracking), grammar-formatted output, and versioned knowledge base storage with connections for run comparison.

We trace two approaches through the complete task: pure LLM token generation (the conventional approach, which cannot actually complete the task but we estimate its token

cost) and VDR-LLM-Prolog orchestration with the full primitive, Prolog, and operational environment stack.

Phase 1: Data Acquisition

Conventional LLM.

The SRE pastes a curl command output or describes the prometheus endpoint. The LLM cannot execute network requests. It says “I can’t connect to prometheus directly, but here’s how you could query it.” The SRE manually runs the curl, gets 1MB of JSON, and cannot paste it — it exceeds the context window. They manually truncate to 50 endpoints, losing 150 endpoints of data. They paste roughly 50KB. The LLM processes this through attention — roughly 12,000 tokens of raw JSON consuming the context window. The LLM has seen the data but parsed it through attention, spending thousands of tokens of attention capacity on brackets, commas, and key names.

Tokens generated: 200 (advice on how to query). Tokens consumed in context: 12,000 (truncated pasted JSON). Useful work: zero — the LLM told the user to do the work manually. Data coverage: 25% (50 of 200 endpoints).

VDR-LLM-Prolog.

The LLM assesses the request and formalizes the first step: fetch prometheus data. It emits a command token invoking B424 network_fetch with the prometheus query_range API URL, scoped to the relevant metric over the last 15 minutes.

GPU stream 1 is idle at this point. The CPU receives the command token, verifies the network grant (the SRE’s session has a network grant scoped to prom.internal:9090/*), and executes the fetch. The 1MB JSON response arrives on the CPU.

The LLM wrote a small Python parsing script earlier in the session — 6 lines, roughly 40 tokens — that extracts the prometheus JSON structure into a normalized form: endpoint name, cluster, timestamp array, value array. The CPU executes this script in the Docker sandbox via B410 execute_python. The script outputs structured JSON.

B246 parse_json on the structured output produces a typed dict. B254 vdr_from_decimal_string converts every metric value to exact VDR fractions at the conversion boundary. Each conversion records the source type (prometheus gauge), original decimal string, converted VDR value, and maximum error. For terminating decimals (which prometheus produces), the error is zero.

The converted data goes into the project knowledge base: root.sessions.sre.incident_2026_05_16.metrics. Each endpoint is a child KB with facts for cluster, endpoint name, and a ring buffer holding the time series. 200 endpoints $\times$ 90 data points (one per 10 seconds for 15 minutes) = 18,000 VDR values stored in 200 ring buffers.

LLM tokens: 8 (command token for fetch) + 40 (Python script, written once) + 8 (parse command) + 8 (conversion command). Total: 64 LLM tokens. Data coverage: 100%

(all 200 endpoints, all 18,000 data points). The LLM never saw the raw JSON. The 1MB never entered the token stream.

GPU cost: the conversion of 18,000 decimal strings to VDR fractions is 18,000 calls to B254, each a parse plus integer construction. At roughly 50 integer operations per conversion: 900,000 integer operations. On an H100 at 16.9T INT32 ops/sec: 0.00005 ms. Invisible.

Step	Component	LLM Tokens	GPU Ops	Wall Time	Notes
Network fetch	CPU (grant check + HTTP)	8	0	200 ms (network latency)	Grant verified in 2 μ s
Python parse script	LLM generation	40	0	400 ms (LLM generates script)	Written once, reusable
Script execution	CPU (Docker sandbox)	8	0	150 ms (Python execution)	Sandbox overhead included
JSON parse	CPU + GPU	8	~100K	0.5 ms	B246 on structured output
VDR conversion	GPU stream 2	8	~900K	0.05 ms	B254 on 18,000 values
KB storage	GPU stream 2	0 (automated)	~200K	0.02 ms	200 KB nodes + ring buffer writes
Phase 1 total	—	72	~1.2M	~750 ms	Dominated by network + script execution

Phase 2: Filtering and Threshold Analysis

Conventional LLM.

The SRE asks “which endpoints have p99 latency above 500ms in the last 5 minutes?” The LLM has 12,000 tokens of truncated JSON in its context. It attempts to parse the JSON through attention, scanning for values above 0.5. It generates reasoning tokens: “Looking at the data, I can see that endpoint /api/users has values of 0.482, 0.510, 0.498...” — generating digits through token prediction, comparing by generating comparison prose.

For 50 endpoints with multiple time points, this takes roughly 2,000 tokens of generated reasoning. Some comparisons are wrong because digit prediction is imprecise. The result is a prose list of maybe 8 endpoints that the model thinks exceeded the threshold, with no guarantee of completeness or accuracy. The other 150 endpoints were never seen.

Tokens generated: 2,000. Accuracy: roughly 85% (some comparisons wrong). Coverage: 25%.

VDR-LLM-Prolog.

The LLM formalizes the filter: for each endpoint, check whether any value in the last 5 minutes exceeds [500, 1000, 0] (the VDR fraction for 0.5 seconds). This is one Prolog rule assertion plus one command sequence.

The LLM asserts a Prolog rule:

```
high_latency(Endpoint, MaxVal) :- endpoint(Endpoint,
KB), ring_buffer_read_all(KB, Values), list_filter(Values,
greater_than_500ms, Filtered), list_length(Filtered,
Count), vdr_greater_than(Count, 0), list_max(Filtered,
MaxVal) .
```

Roughly 30 tokens for the rule formalization. B376 kb_assert stores it.

B378 kb_query evaluates the rule across all 200 endpoints. The Prolog engine’s frontier-based execution on GPU retrieves all endpoint facts (one predicate bucket scan — 200 facts), reads each endpoint’s ring buffer (200 parallel reads), filters each time series (200 parallel filters comparing each value against the threshold), and collects the results.

GPU execution: 200 endpoints $\times$ 30 time points (last 5 minutes at 10-second intervals) = 6,000 VDR comparisons. Each comparison is B017 vdr_compare — roughly 4 integer operations for same-frame Q335 values. Total: 24,000 integer operations. Time: 0.001 ms.

The result is an exact list of endpoints with their maximum latency values, stored as facts in the investigation KB. Say 23 endpoints exceeded the threshold — the exact count, not an approximation.

The LLM reads the structured result: 23 endpoint names with max latency values. Roughly 100 tokens of structured input. It assesses and continues.

Step	Component	LLM	GPU Ops	Wall Time	Notes
		Tokens			
Prolog rule formalization	LLM generation	30	0	300 ms (LLM generates)	One-time; reusable
Rule assertion	GPU stream 2	8	~100	0.001 ms	B376 kb assert

Step	Component	LLM Tokens	GPU Ops	Wall Time	Notes
Rule evaluation (200 endpoints)	GPU stream 1	0	~24K	0.001 ms	Frontier- based parallel filter
Result storage	GPU stream 2	0 (automated)	~5K	0.001 ms	23 findings asserted
LLM reads results	LLM input	0 (reading, not generating)	0	included in next generation	100 tokens structured input
Phase 2 total	—	38	~29K	~300 ms	Dominated by LLM rule generation

Phase 3: Deployment Event Correlation

Conventional LLM.

The SRE asks “correlate with deployment events in the last hour.” The LLM cannot query prometheus for deployment metadata. The SRE manually runs another curl, truncates the result, and pastes it. Another 3,000 tokens of context consumed. The LLM now has 15,000 tokens of JSON in context, plus 2,000 tokens of its own prior reasoning. It attempts to correlate by generating comparison prose: “The deployment of service-auth at 14:12 appears to coincide with the latency increase in /api/users which started around 14:15...” Token by token, it generates temporal reasoning by producing timestamps and comparing them. For 23 high-latency endpoints and perhaps 8 deployment events, this takes roughly 3,000 tokens of reasoning. Temporal comparisons are done through digit prediction — the model generates “14:12 is before 14:15” as text, not as an integer comparison.

Tokens generated: 3,000. Context consumed: 18,000 cumulative. Accuracy: moderate — temporal proximity through text comparison misses some correlations and misidentifies others.

VDR-LLM-Prolog.

The LLM formalizes: fetch deployment events from prometheus label API. One command token for B424 network_fetch. CPU executes with grant check. JSON returns. The LLM’s Python script (adapted from phase 1 — roughly 8 tokens to modify) extracts deployment events with service name, timestamp, and version. B246 parse_json. B254 converts timestamps to exact integers. Events stored as facts: fact(deployment, service_auth, 1234567890, version_2_3_1).

Now the correlation. The LLM formalizes a Prolog rule:

```

correlated_deploy(Endpoint, Deploy, TimeDelta) :-
high_latency(Endpoint, _), endpoint_cluster(Endpoint,
Cluster), deployment(Deploy, Cluster, DeployTime, _),
first_spike_time(Endpoint, SpikeTime), int_sub(SpikeTime,
DeployTime, TimeDelta), vdr_greater_than(TimeDelta,
0), vdr_less_than(TimeDelta, 600).

```

This rule: for each high-latency endpoint, find deployments in the same cluster where the first spike occurred within 600 seconds (10 minutes) after the deployment. Roughly 35 tokens for the rule.

A helper rule to find first spike time:

```

first_spike_time(Endpoint, Time) :- endpoint(Endpoint,
KB), ring_buffer_read_all(KB, Values), list_filter(Values,
greater_than_500ms, Filtered), list_head(Filtered,
FirstSpike), spike_timestamp(FirstSpike, Time).

```

Roughly 25 tokens.

B378 kb_query evaluates the correlation. GPU execution: the Prolog engine joins high_latency results (23 endpoints) with deployment events (8 events) filtered by cluster match, then compares timestamps. This is a hash join on cluster (small hash table — 4 clusters), followed by 23×2 (average deployments per cluster) = 46 timestamp comparisons. Each comparison: 4 integer operations. Total: roughly 200 integer operations plus the join overhead of roughly 1,000 operations for hash table build and probe.

Result: say 15 endpoint-deployment correlations with exact time deltas stored as facts with provenance.

Step	Component	LLM	GPU Ops	Wall Time	Notes
		Tokens			
Deploy event fetch	CPU (network)	8	0	150 ms	Same grant as phase 1
Script adaptation	LLM generation	8	0	80 ms	Minor modification
Script execution	CPU (Docker)	8	0	100 ms	Parse deploy events
JSON parse + convert	GPU stream 2	8	~50K	0.01 ms	8 events, small payload
Correlation rule	LLM generation	35	0	350 ms	Prolog rule formalization
Helper rule	LLM generation	25	0	250 ms	First spike time helper

Step	Component	LLM Tokens	GPU Ops	Wall Time	Notes
Rule assertions	GPU stream 2	16	~200	0.002 ms	Two B376 calls
Correlation evaluation	GPU stream 1	0	~1.2K	0.001 ms	Hash join + timestamp compare
Result storage	GPU stream 2	0	~3K	0.001 ms	15 correlation facts
Phase 3 total	—	108	~54K	~930 ms	Dominated by network + LLM generation

Phase 4: Statistical Analysis and Ranking

Conventional LLM.

The SRE asks “rank the affected endpoints by severity and show me per-cluster breakdown.” The LLM’s context now has 18,000 tokens of JSON and prior reasoning. It attempts to sort by generating comparison prose. For 23 endpoints, it produces a ranked list through token prediction — generating each rank position by comparing values it generated earlier. Arithmetic for mean, variance, and percentile is digit-by-digit token prediction. The per-cluster grouping is generated prose.

Tokens generated: 2,500. Several ranking errors (positions swapped due to imprecise digit comparison). Arithmetic on means and percentiles contains errors. No downloadable data — everything is prose in the conversation.

VDR-LLM-Prolog.

The LLM formalizes a sequence of primitive calls.

Group by cluster: B208 list_group_by on the correlated endpoints by cluster field. One command token, 8 LLM tokens. GPU: one pass over 15 correlation facts grouping by the cluster key. Result: 4 groups stored as a dict.

Per-cluster statistics: for each cluster group, B101 vdr_stat_mean on the max latency values, B102 vdr_stat_variance, B105 vdr_stat_percentile at [95, 100, 0] and [99, 100, 0]. Four command tokens per cluster, 16 clusters × statistics = 64 tokens. But the LLM can formalize this as a loop plan in one command sequence — roughly 20 tokens total.

Sort by severity: B202 list_sort on the correlation facts by time delta (faster correlation = higher severity) and by max latency. Two command tokens, 16 tokens.

GPU execution for all statistics: $4 \text{ clusters} \times \sim 5 \text{ endpoints per cluster} \times 4 \text{ statistical operations}$. Each stat operation on 5 values: roughly 20 Q335 operations per statistic. Total: $4 \times 5 \times 4 \times 20 = 1,600$ Q335 operations. Plus the sort: $15 \times \log_2(15) \times 4 \text{ comparisons} \approx 240 \text{ comparisons}$ at 4 integer ops each = 960 operations. Grand total: roughly 3,000 Q335 operations.

All results stored as facts in the investigation KB with provenance: which endpoints, which metrics, computed from which ring buffer data, using which primitives.

Step	Component	LLM Tokens	GPU Ops	Wall Time	Notes
Group-by formalization	LLM generation	8	0	80 ms	One command token
Group-by execution	GPU stream 2	0	~500	0.001 ms	B208 on 15 facts
Statistics plan	LLM generation	20	0	200 ms	Batch formalization
Statistics execution	GPU stream 2	0	~1.6K	0.001 ms	Mean, var, percentiles
Sort formalization	LLM generation	16	0	160 ms	Two sort commands
Sort execution	GPU stream 2	0	~960	0.001 ms	B202 on 15 items
Result storage	GPU stream 2	0	~2K	0.001 ms	Stats facts with provenance
Phase 4 total	—	44	~5K	~440 ms	Dominated by LLM generation

Phase 5: Python Analysis for Complex Transform

Conventional LLM.

The SRE asks for a rolling window analysis to detect the exact inflection point of each endpoint's degradation. The LLM generates Python code for the user to run — roughly 400 tokens. The SRE manually copies it, runs it, pastes the result back. Another round trip. More context consumed.

Tokens generated: 400 (Python code) + 300 (explanation) = 700. Execution: manual by user. Result: pasted back, consuming more context.

VDR-LLM-Prolog.

The LLM writes a Python analysis script for rolling window inflection detection. The script operates on data exported from the investigation KB — B247 format_csv exports the time series data for the 23 high-latency endpoints. The script computes a rolling mean over a 2-minute window and finds the timestamp where the derivative exceeds a threshold.

The LLM writes the script: roughly 60 tokens for a 15-line Python script. B392 file_write stores the exported CSV in the Docker sandbox filesystem. B410 execute_python runs the script. The script outputs JSON with endpoint, inflection timestamp, and pre/post rolling means. The CPU captures the output. B246 parse_json. B254 converts the values to exact VDR. Results stored as facts.

The LLM never ran the Python in its token stream. The script executed in the sandbox. The results came back through the primitive pipeline. The LLM receives structured findings: 23 inflection points with exact timestamps and rolling means.

Step	Component	LLM Tokens	GPU Ops	Wall Time	Notes
CSV export formalization	LLM generation	8	0	80 ms	B247 command
CSV export execution	GPU stream 2	0	~50K	0.5 ms	Format 23 endpoints × 90 points
File write	CPU (filesystem)	8	0	5 ms	B392 to sandbox
Python script generation	LLM generation	60	0	600 ms	15-line inflection detection
Script execution	CPU (Docker)	8	0	300 ms	Python rolling window analysis
Result parse + convert	GPU stream 2	8	~20K	0.01 ms	23 inflection results
Result storage	GPU stream 2	0	~5K	0.001 ms	Facts with provenance
Phase 5 total	—	92	~75K	~985 ms	Dominated by Python execution + LLM script writing

Phase 6: Versioned Project Storage

Conventional LLM.

The SRE asks “save these results so I can compare with the next run.” The LLM cannot save anything. It has no persistent state. It suggests the user manually copy the results to a document or spreadsheet. The entire investigation exists only as conversation text that will be lost when the session ends.

Tokens generated: 200 (suggestions for manual saving). Persistent state: zero.

VDR-LLM-Prolog.

The LLM formalizes a versioned project structure. This is knowledge base tree construction — the architecture’s native operation.

Create the project structure:

```
root.projects.latency_investigation.runs.run_2026_05_16_1423
```

Under this KB, child KBs for: raw_metrics (connection to the incident metrics KB via B363 connection_add with relation sourced_from), filtered_results (the 23 high-latency endpoints with threshold facts), correlations (the 15 deployment correlations), statistics (per-cluster stats), inflection_analysis (the 23 inflection points), and metadata (run timestamp, prometheus source URL, threshold used, script versions).

Each child KB gets facts copied or connected from the investigation workspace. B363 connection_add creates typed connections: sourced_from linking the raw metrics to the prometheus source, computed_by linking statistics to the primitives used, analyzed_by linking inflection results to the Python script (script content stored as a fact).

For future comparison, the LLM asserts a Prolog rule at the project level:

```
run_comparison(RunA, RunB, Endpoint, DeltaLatency) :-  
  run_endpoint_max(RunA, Endpoint, LatA), run_endpoint_max(RunB,  
  Endpoint, LatB), vdr_sub(LatB, LatA, DeltaLatency).
```

Roughly 25 tokens. This rule will work on any future run stored under the same project structure. The next run creates run_2026_05_17_0930, populates the same child KB schema, and the comparison rule automatically produces deltas across all shared endpoints.

Version metadata: B376 kb_assert stores the run’s complete configuration — threshold values, time range, cluster list, prometheus endpoint, script hash — as facts. Future runs can query what parameters produced which results.

B368 session_snapshot captures the current live state as a restore point. The investigation can be resumed from this exact state.

Step	Component	LLM Tokens	GPU Ops	Wall Time	Notes
Project structure plan	LLM generation	30	0	300 ms	Tree structure formaliza- tion
KB creation (6 child KBs)	GPU stream 2	48 (8 per KB)	~2K	0.01 ms	$6 \times$ B376 sequences
Fact migration	GPU stream 2	16	~10K	0.05 ms	Copy/connect findings to project
Connection creation	GPU stream 2	40 (8 per connection, 5 connec- tions)	~1K	0.005 ms	B363 typed connections
Comparison rule assertion	LLM generation	25	~100	0.001 ms	Reusable across runs
Metadata assertion	GPU stream 2	16	~2K	0.01 ms	Config facts
Session snapshot	GPU stream 2	8	~5K	0.05 ms	B368 snapshot
Phase 6 total	—	183	~20K	~300 ms	Dominated by LLM generation

Phase 7: Formatted Output and Export

Conventional LLM.

The SRE asks for a summary. The LLM generates a prose summary from its 20,000-token context — roughly 1,500 tokens of generated text with tables formatted character by character. The tables have formatting errors (misaligned columns, inconsistent decimal places). The data in the tables is whatever the LLM recalls from earlier in the conversation, which may differ from what it computed due to context window attention degradation. No export capability — the user manually copies the text.

Tokens generated: 1,500. Formatting errors: 2-5 per table. Data accuracy: degraded from earlier reasoning. Export: manual copy.

VDR-LLM-Prolog.

Grammar templates handle all structural output. The LLM fills content slots with prose framing.

An incident summary grammar template (defined at `root.templates.incident_summary`, inherited, reusable) provides the structure: title, timestamp, summary section, findings table, correlation table, per-cluster breakdown, inflection analysis, methodology notes, and recommendations.

Each data table renders directly from KB facts through grammar. The findings table pulls from `filtered_results` KB — 23 rows of endpoint name, cluster, max latency (exact VDR fraction rendered through `B251 fraction_to_decimal`), threshold exceeded by (computed by `B002 vdr_sub`). The correlation table pulls from the `correlations` KB — 15 rows of endpoint, deployment, time delta, cluster. The per-cluster breakdown pulls from the `statistics` KB — 4 rows of cluster name, endpoint count, mean latency, p95, p99 (all exact fractions).

The LLM writes prose framing for each section — roughly 200 tokens total. Observations about the pattern, likely root cause assessment (tagged at 30/100 confidence per the LLM assessment floor), and recommended next steps.

B247 `format_csv` exports the complete data tables. B392 `file_write` saves them to the project filesystem. B247 `format_json` exports the full investigation structure as a machine-readable document.

Step	Component	LLM Tokens	GPU Ops	Wall Time	Notes
Grammar template instantiation	GPU stream 3	0	~5K	0.02 ms	Template from inherited KB
Findings table render	GPU stream 3	0	~10K	0.05 ms	23 rows from KB facts
Correlation table render	GPU stream 3	0	~7K	0.03 ms	15 rows from KB facts
Statistics table render	GPU stream 3	0	~3K	0.01 ms	4 cluster summaries
Inflection table render	GPU stream 3	0	~10K	0.05 ms	23 inflection points
Prose framing	LLM generation	200	0	2,000 ms	Actual judgment work
Confidence computation	GPU stream 2	0	~500	0.001 ms	Propagation on findings
CSV export	GPU stream 2	8	~20K	0.1 ms	B247 on all tables

Step	Component	LLM Tokens	GPU Ops	Wall Time	Notes
JSON export	GPU stream 2	8	~30K	0.15 ms	Full investi- gation structure
File writes	CPU (filesystem)	16	0	10 ms	B392 for CSV + JSON
Phase 7 total	—	232	~86K	~2,010 ms	Dominated by LLM prose generation

Complete Investigation Summary

Token Accounting

Phase	Description	Conventional LLM Tokens	VDR LLM Tokens	Conventional Data Coverage	VDR Data Coverage
1	Data acquisition	200 + 12,000 context	72	25% (50/200 endpoints)	100% (200/200 endpoints)
2	Filtering	2,000	38	~85% accuracy on 25% data	100% accuracy on 100% data
3	Correlation	3,000 + 3,000 context	108	Moderate accuracy, 25% coverage	Exact, 100% coverage
4	Statistics	2,500	44	Errors in arithmetic, 25% coverage	Exact, 100% coverage
5	Complex analysis	700	92	Manual execution by user	Automated sandbox execution
6	Project storage	200	183	Nothing saved	Full versioned project with comparison rule

Phase	Description	Conventional LLM Tokens	VDR LLM Tokens	Conventional Data Coverage	VDR Data Coverage
7	Output and export	1,500	232	Prose with formatting errors, no export	Grammar-perfect tables, CSV + JSON export
Total	—	10,100 generated + 15,000 context	769	—	—

VDR token reduction: 92.4%. And the 769 VDR tokens produced a complete, exact, versioned, exportable investigation, while the 10,100 conventional tokens produced approximate prose over 25% of the data with no persistent state.

GPU Operation Accounting

Phase	GPU Integer Ops	Time (H100 est.)	As Equivalent LLM Tokens
1: Acquisition + conversion	1.2M	0.5 ms	0.0001 tokens
2: Filtering	29K	0.003 ms	0.000002 tokens
3: Correlation	54K	0.014 ms	0.000004 tokens
4: Statistics + sort	5K	0.003 ms	0.0000004 tokens
5: CSV export + result parse	75K	0.5 ms	0.000005 tokens
6: Project storage	20K	0.1 ms	0.000001 tokens
7: Table render + exports	86K	0.4 ms	0.000006 tokens
Total primitives	~1.47M	~1.5 ms	~0.0001 tokens

The total primitive GPU computation for the entire investigation is 1.5 milliseconds. This is the compute time for parsing 1MB of JSON, converting 18,000 metric values to exact VDR fractions, filtering 200 endpoints, correlating with deployment events, computing statistics across 4 clusters, sorting and ranking, rendering 4 data tables, and exporting to CSV and JSON. All of it, combined, takes less time than the LLM takes to generate a single token.

Wall-Clock Time

Component	Conventional	VDR	Notes
LLM generation	~100 sec (10,100 tokens at ~100ms/token)	~8 sec (769 tokens at ~10ms/token with grammar speedup)	Grammar reduces per-token decode cost for structural tokens
User manual work	~300 sec (curl, truncate, paste, copy results)	0 sec	All automated through primitives
Network operations	Manual by user	~350 ms (automated fetches)	Two prometheus API calls
Script execution	Manual by user (another ~60 sec)	~450 ms (Docker sandbox)	Automated
Data processing	Embedded in LLM generation (inaccurate)	~1.5 ms (GPU primitives)	100% accurate
Context re-reading (attention)	~200 sec cumulative (growing context, re-read every turn)	0 sec (state in KBs)	No context growth
Total wall-clock	~660 sec (~11 minutes)	~9 sec	73 × faster

The conventional approach takes 11 minutes, requires constant manual intervention by the SRE, processes only 25% of the data, produces inaccurate arithmetic, and saves nothing. The VDR approach takes 9 seconds, requires no manual intervention, processes 100% of the data, produces exact results with provenance, and stores everything in a versioned project structure ready for comparison with future investigations.

Cost Comparison (H100 at \$3.00/hour)

Metric	Conventional	VDR	Ratio
GPU time	~100 sec (LLM generation only)	~9 sec (LLM + all primitives)	11 × less GPU time
GPU cost	\$0.083	\$0.0075	11 × cheaper
SRE time (at \$150/hr)	~11 min = \$27.50	~9 sec = \$0.38 (monitoring only)	72 × cheaper in human time
Total cost per investigation	\$27.58	\$0.39	71 × cheaper
Data coverage	25%	100%	4 × more data
Accuracy	~85% (arithmetic errors)	100% (exact)	—
Persistent state	None	Full versioned project	—

Metric	Conventional	VDR	Ratio
Exportable data	None (manual copy of prose)	CSV + JSON + KB queryable	—
Reproducibility	None (session is gone)	Full (snapshot + versioned project)	—

Future Run Comparison (Phase 6 Payoff) When the SRE runs the same investigation tomorrow after a fix has been deployed, the second run follows the same phases but with these advantages:

Component	First Run Cost	Second Run Cost	Savings Source
Python parse script	40 tokens (written)	8 tokens (reuse command)	Script persists in project KB
Prolog filter rule	30 tokens (formalized)	8 tokens (query existing rule)	Rule persists in project KB
Prolog correlation rule	60 tokens (two rules)	8 tokens (query existing rules)	Rules persist
Comparison rule	25 tokens (formalized)	0 tokens (already exists)	Asserted at project level
Grammar template	0 tokens (inherited)	0 tokens	Still inherited
Project structure	183 tokens (created structure)	30 tokens (create new run KB under existing structure)	Structure exists
Total	769 tokens	~450 tokens	42% reduction from accumulated project state

And the comparison rule automatically produces:

```
run_comparison(run_2026_05_16_1423, run_2026_05_17_0930,
Endpoint, DeltaLatency) — for every shared endpoint, the exact latency
change as a VDR fraction. No re-derivation. No manual comparison. One Prolog query,
evaluated on GPU in 0.001 ms, producing exact deltas with provenance linking both
runs.
```

This is what accumulating exact state means in practice. The first investigation builds the infrastructure. Every subsequent investigation reuses it. The cost per investigation decreases while the capability increases. The conventional LLM starts from zero every time.

[[HOWL-VDR-18-2026](#)] VDR-LLM-Prolog: Performance. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR-18-2026>

[[HOWL-VDR-1-2026](#)] VDR Arithmetic: Value, Decimal, Remainder. Github:

<https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-1-2026>

[[HOWL-VDR-2-2026](#)] VDR Gym: Exact Arithmetic Across Fifteen Domains. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-2-2026>

[[HOWL-MATH-3-2026](#)] The Transcendental Hierarchy. Github: <https://github.com/ghowland/papers/tree/main/papers/MATH-3-2026>

[[HOWL-MATH-4-2026](#)] The Universal Power-of-Two Basis. Github: <https://github.com/ghowland/papers/tree/main/papers/MATH-4-2026>

[[HOWL-VDR-15-2026](#)] VDR-LLM-Prolog: Prompt Optimization. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-15-2026>

[[HOWL-VDR-16-2026](#)] VDR-LLM-Prolog: Safe by Contract. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR-16-2026>

[[HOWL-VDR-17-2026](#)] VDR-LLM-Prolog: Alignment. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR-17-2026>