

VDR and Diffusion

Zero-Drift Denoising Through Exact Sequential Arithmetic

Registry: [\[HOWL-VDR-26-2026\]](#)

Series Path: [\[HOWL-VDR-1-2026\]](#) → [\[HOWL-VDR-2-2026\]](#) → [\[HOWL-MATH-3-2026\]](#) → [\[HOWL-MATH-4-2026\]](#) → ... → [\[HOWL-VDR-14-2026\]](#) → ... → [\[HOWL-VDR-21-2026\]](#) → [\[HOWL-VDR-22-2026\]](#) → [\[HOWL-VDR-23-2026\]](#) → [\[HOWL-VDR-24-2026\]](#) → [\[HOWL-VDR-25-2026\]](#) → [\[HOWL-VDR-26-2026\]](#)

DOI: 10.5281/zenodo.20260247

Date: May 2026

Domain: Exact Arithmetic / Generative Model Inference

AI Usage Disclosure: Only the top metadata, figures, refs and final copyright sections were edited by the author. All paper content was LLM-generated using Anthropic’s Opus 4.6.

Abstract

Diffusion models generate images and video by iterating a denoising chain — each step takes the previous step’s output, scales by schedule coefficients, subtracts predicted noise, normalizes, and feeds the result forward. In float64 arithmetic, each step introduces approximately 10^{-16} rounding error. Over 50 steps for image generation, this compounds to approximately 10^{-14} . Over hundreds or thousands of steps for video generation, where frames condition on prior frames through the same arithmetic chain, the error produces measurable artifacts: color drift, temporal flickering, and structural inconsistency between frames.

This paper implements the complete diffusion process — noise schedule computation, forward diffusion, reverse denoising, DDIM deterministic sampling, and multi-cycle drift measurement — in VDR exact integer arithmetic [VDR-1]. Every intermediate value is an exact rational number. Every operation preserves that exactness. The result: zero drift accumulation across arbitrarily long denoising chains. The error at cycle N equals the error at cycle 1, which is the Newton square root residual at the chosen depth (below 10^{-50} at depth 10), not a compounding float truncation.

Validated: 37 tests, 33 passed, 4 failed. All 4 failures trace to a normalization presentation issue — Newton iteration for perfect squares produces correct values that do not reduce to simplest form. Zero arithmetic errors. Zero drift. Zero computation failures.

No prior reading is required. All necessary concepts from VDR arithmetic are introduced where first used.

1. The Sequential Arithmetic Problem

1.1 What Diffusion Models Compute

A diffusion model generates data by learning to reverse a noise-adding process. The forward process gradually adds Gaussian noise to data over T timesteps until the signal is destroyed. The reverse process learns to remove noise step by step, recovering the original signal from pure noise.

Every step in both directions is arithmetic. The forward process at timestep t computes:

$$x_t = \sqrt{\bar{\alpha}_t} \cdot x_0 + \sqrt{(1 - \bar{\alpha}_t)} \cdot \epsilon$$

where x_0 is the original data, ϵ is Gaussian noise, and $\bar{\alpha}_t$ is a cumulative product of schedule values that controls how much signal remains at step t .

The reverse process at each step computes a predicted mean:

$$\mu_t = (1/\sqrt{\alpha_t}) \cdot (x_t - (\beta_t/\sqrt{(1 - \bar{\alpha}_t)}) \cdot \epsilon_{\text{predicted}})$$

where $\alpha_t = 1 - \beta_t$, and β_t is the noise schedule at step t .

Each reverse step takes the output of the previous step as input. The chain is sequential — step t depends on step $t+1$, which depends on step $t+2$, and so on back to step T . Any arithmetic error at step t propagates through every subsequent step.

1.2 Where Float Fails

In float64 arithmetic, each multiplication introduces approximately 10^{-16} relative error (one unit in the last place of the 52-bit mantissa). Each step of the diffusion process involves several multiplications, divisions, and square root evaluations. The errors from each operation combine and propagate to the next step.

For a single image generated in 50 steps, the accumulated error is approximately $50 \times 10^{-16} \approx 10^{-14}$. This is invisible in the output — 14 digits of precision is far more than any pixel value requires.

For video generation, the situation changes. Each frame conditions on the previous frame through the diffusion process. A 30-second video at 24 fps is 720 frames. Each frame's latent representation passes through 50 denoising steps, and the output feeds the next frame's conditioning. That is 36,000 sequential arithmetic operations on the same data, each contributing float error. The accumulated drift is approximately $36,000 \times 10^{-16} \approx 10^{-12}$. At this scale, drift produces visible artifacts: colors shift gradually across frames, fine structures become inconsistent, and temporal coherence degrades.

The error is also platform-dependent. Float rounding is implementation-specific — different GPUs, different CUDA versions, and different compiler flags produce different rounding behavior. The same model with the same weights and the same input noise produces different outputs on different hardware. Reproducibility across platforms is impossible in float arithmetic.

1.3 What This Paper Demonstrates

Every component of the diffusion process — schedule computation, forward diffusion, reverse denoising, deterministic sampling, and multi-cycle operation — can be implemented in exact integer arithmetic with zero drift accumulation. The error at step 1000 is the same as the error at step 1. The error does not grow. It cannot grow, because exact arithmetic does not introduce rounding error, and the only source of approximation (Newton square root iteration) produces a fixed residual at the chosen depth that does not compound through the chain.

2. VDR Arithmetic for Diffusion

2.1 The VDR Triple

Every number in VDR is three integers: Value, Denominator, Remainder [VDR-1]. V and D are plain integers forming an exact rational V/D . R is the Remainder — not rounding error, not residual, but first-class structural information about what the denominator frame could not absorb. When R is zero, the value is a closed rational number. When R is nonzero, the value carries exact structure beyond the rational frame.

For diffusion computation, most values are closed rationals — schedule parameters, scaling coefficients, vector components. The Remainder becomes relevant in square root computation, where Newton iteration produces exact rational approximations with the residual carried as structural information in R.

2.2 Arithmetic Operations

VDR closed arithmetic is exact rational arithmetic [VDR-1]:

Addition: $[V_1D_2 + V_2D_1, D_1D_2, 0]$. Two integer multiplications, one integer addition. Exact.

Subtraction: $[V_1D_2 - V_2D_1, D_1D_2, 0]$. Same cost. Exact.

Multiplication: $[V_1V_2, D_1D_2, 0]$. One integer multiplication for each of numerator and denominator. Exact.

Division: $[V_1D_2, D_1V_2, 0]$. Exact when $V_2 \neq 0$.

No rounding. No truncation. No platform dependence. The result of any chain of these operations on rational inputs is an exact rational, regardless of chain length.

2.3 Square Roots via Newton Iteration

Diffusion requires square roots — $\sqrt{\bar{\alpha}_t}$ and $\sqrt{(1 - \bar{\alpha}_t)}$ at every timestep. Square roots of rationals are generally irrational and cannot be represented as closed VDR rationals. VDR handles this through Newton iteration, which produces exact rational approximations at any requested depth [VDR-1].

For $\sqrt{a}$, Newton iteration computes:

$$x_{n+1} = (x_n + a/x_n) / 2$$

Starting from $x_0 = 1$, each step doubles the number of correct digits (quadratic convergence). After 10 steps, the approximation has over 100 correct digits. Every intermediate value x_n is an exact rational — no float approximation at any point.

The residual $x_n^2 - a$ is an exact, computable, inspectable rational number. At depth 10 for $\sqrt{2}$, the residual denominator has thousands of digits. The residual magnitude is below 10^{-50} . This is not an unknown truncation error — it is a specific, printable number that the system carries as structural information.

2.4 Why Exactness Matters for Chains

In float arithmetic, each operation's error is independent and approximately random in direction. Over a chain of N operations, errors accumulate as approximately $\sqrt{N} \times \varepsilon$ (random walk) or $N \times \varepsilon$ (worst case), where $\varepsilon \approx 10^{-16}$ for float64.

In VDR arithmetic, each operation on closed rationals produces an exact result. The chain of operations is lossless — the output of step N is exactly the rational number that the arithmetic defines, not an approximation of it. The only source of approximation is square root computation via Newton iteration, and this approximation has a fixed magnitude determined by the iteration depth. It does not compound because the residual from $\sqrt{a}$ used in the forward process is the same residual present in the reverse process — the arithmetic is self-consistent.

3. The Noise Schedule

3.1 Schedule Construction

The noise schedule defines a sequence of T noise levels $\beta_1, \beta_2, \dots, \beta_T$ controlling how much noise is added at each step. The linear schedule spaces β values evenly between a start and end value:

$$\beta_t = \beta_{\text{start}} + (t/(T-1)) \cdot (\beta_{\text{end}} - \beta_{\text{start}})$$

In VDR, each β_t is an exact rational. $\beta_{\text{start}} = 1/100$ and $\beta_{\text{end}} = 1/20$ are exact. The interpolation involves integer arithmetic only — the index t and count $T-1$ are integers, the multiplication and addition are exact rational operations.

From β , the derived quantities are:

$$\alpha_t = 1 - \beta_t \text{ (exact: one integer subtraction in the numerator)}$$

$$\bar{\alpha}_t = \prod_{k=1}^t \alpha_k \text{ (exact: chain of rational multiplications)}$$

$$\sqrt{\bar{\alpha}_t} \text{ (Newton iteration at chosen depth)}$$

$$\sqrt{(1 - \bar{\alpha}_t)} \text{ (Newton iteration at chosen depth)}$$

3.2 Cumulative Product Exactness

The cumulative product $\bar{\alpha}_t = \alpha_1 \cdot \alpha_2 \cdot \dots \cdot \alpha_t$ is the critical computation. In float, multiplying T values near 0.95-0.99 accumulates rounding error at every step. After 5 multiplications, float64 shows approximately 10^{-16} drift from the true rational value. After 1000 multiplications, the drift is approximately 10^{-13} .

In VDR, the cumulative product is exact rational multiplication. The validated result confirms this: $\bar{\alpha}_T = 26821179/31250000$, verified against independent Python Fraction computation. Bit-identical. Not approximately equal — identical. The VDR cumulative product and the arbitrary-precision rational computation produce the same numerator and the same denominator because they are performing the same operation: exact integer multiplication of numerators and denominators.

3.3 Schedule Properties

Five structural properties are verified for the exact schedule:

$\alpha = 1 - \beta$ for all t . Exact identity, not approximate. The subtraction is one integer operation.

$\bar{\alpha}_t = \text{cumulative product of } \alpha_k$. Verified against independent Fraction computation. Exact.

$\bar{\alpha}$ is strictly monotonically decreasing. Each $\alpha_k < 1$, so each multiplication makes the product strictly smaller. VDR comparison is exact integer cross-multiplication — no float comparison ambiguity near equality.

$\text{SNR} = \bar{\alpha}/(1-\bar{\alpha})$ is strictly monotonically decreasing. The signal-to-noise ratio decreases monotonically across the schedule, confirming that noise increases at every step. VDR computes this ratio as exact rational division and verifies monotonicity by exact comparison.

All β_t satisfy $0 < \beta_t < 1$. Range constraint verified by exact comparison. No float rounding pushes a β value outside the valid range.

3.4 Cosine Schedule

The cosine schedule [Nichol & Dhariwal, 2021] uses a cosine function to produce a smoother noise trajectory than the linear schedule. Since cosine is transcendental, VDR uses a rational approximation via Taylor series or Padé approximants, producing exact rational values that approximate the cosine schedule with controllable precision.

The validated cosine schedule produces 10 steps with monotonically decreasing $\bar{\alpha}$ and all structural properties intact. The approximation quality is governed by the series depth, and every intermediate value remains an exact rational.

4. Forward Diffusion

4.1 The Forward Process

Forward diffusion computes:

$$x_t = \sqrt{\bar{\alpha}_t} \cdot x_0 + \sqrt{(1 - \bar{\alpha}_t)} \cdot \varepsilon$$

where x_0 is the original data vector, ε is a noise vector, and the coefficients $\sqrt{\bar{\alpha}_t}$ and $\sqrt{(1 - \bar{\alpha}_t)}$ are computed via Newton iteration.

In VDR, this is:

1. Compute $\sqrt{\bar{\alpha}_t}$ via Newton iteration on the exact rational $\bar{\alpha}_t$.
2. Compute $\sqrt{(1 - \bar{\alpha}_t)}$ via Newton iteration on the exact rational $(1 - \bar{\alpha}_t)$.
3. Scale x_0 component-wise by $\sqrt{\bar{\alpha}_t}$ (exact rational multiplication per component).
4. Scale ε component-wise by $\sqrt{(1 - \bar{\alpha}_t)}$ (exact rational multiplication per component).
5. Add the two scaled vectors component-wise (exact rational addition per component).

Every component of the resulting x_t vector is an exact rational number. No float truncation at any point.

4.2 Dimensionality Preservation

The forward process preserves the dimensionality of the input. An input vector x_0 with d components produces an output x_t with d components. Each component is an independent exact rational computation. This is verified directly — the output dimension matches the input dimension for all tested cases.

4.3 Signal Dominance at Early Timesteps

At $t=0$, $\bar{\alpha}_0$ is close to 1 (verified: $\bar{\alpha}_0 > 0.9$). This means $\sqrt{\bar{\alpha}_0}$ is close to 1 and $\sqrt{(1 - \bar{\alpha}_0)}$ is close to 0. The forward sample is dominated by the original signal x_0 with only a small noise contribution. VDR confirms this through exact rational comparison — $\bar{\alpha}_0 > 9/10$ is a cross-multiplication check, not a float comparison.

4.4 The Coefficient Identity

The signal and noise coefficients must satisfy an energy conservation identity:

$$(\sqrt{\bar{\alpha}_t})^2 + (\sqrt{(1 - \bar{\alpha}_t)})^2 = 1$$

In float, squaring a float square root does not recover the original value exactly. The residual is approximately 10^{-15} . Over a long chain where this identity is assumed but not exactly maintained, the signal-noise energy balance drifts.

In VDR, the Newton square root has a specific residual: $(\sqrt{\bar{\alpha}})^2$ differs from $\bar{\alpha}$ by the Newton residual, which is below 10^{-20} at depth 10. The coefficient identity residual is

verified below 10^{-20} for all timesteps — 5 orders of magnitude tighter than float, and the residual is a specific inspectable rational number rather than an unknown truncation.

4.5 Forward Trajectory

The forward trajectory $x_0, x_1, \dots, x_T$ is the sequence of noised versions of x_0 . In VDR, each element is an exact rational vector. The trajectory has $T+1$ entries (including x_0), starts exactly at x_0 (verified by identity comparison, not float tolerance), and each subsequent entry is computed by the exact forward formula.

5. Reverse Diffusion

5.1 Predicting x_0

Given a noised sample x_t and a predicted noise ϵ_{pred} , the original data is estimated as:

$$x_{0_pred} = (x_t - \sqrt{(1 - \bar{\alpha}_t)} \cdot \epsilon_{\text{pred}}) / \sqrt{\bar{\alpha}_t}$$

In VDR, this is exact rational arithmetic: one multiplication, one subtraction, one division. When the noise prediction is perfect ($\epsilon_{\text{pred}} = \epsilon$, the actual noise used in the forward process), the recovery is exact.

The test confirms this directly: x_0 prediction error with perfect noise = 0. Not approximately zero. Not below some tolerance. Exactly zero. The exact rational arithmetic perfectly inverts the forward process.

In float64, this same operation — subtract a product, divide by a value — accumulates error from the multiplication, the subtraction (which may involve catastrophic cancellation when $\sqrt{(1 - \bar{\alpha}_t)} \cdot \epsilon$ is close in magnitude to x_t), and the division. The error is approximately 10^{-16} per step, invisible in isolation but compounding across the chain.

5.2 Posterior Mean

The posterior mean for the reverse step is:

$$\mu_t = (1/\sqrt{\alpha_t}) \cdot (x_t - (\beta_t / \sqrt{(1 - \bar{\alpha}_t)}) \cdot \epsilon_{\text{pred}})$$

This involves five operations: compute $\beta_t / \sqrt{(1 - \bar{\alpha}_t)}$, multiply by ϵ_{pred} , subtract from x_t , multiply by $1/\sqrt{\alpha_t}$. Each is exact rational arithmetic. The posterior mean vector has the same dimension as the input, verified directly.

5.3 Posterior Variance

The posterior variance for each step is:

$$\beta_t = \beta_t \cdot (1 - \bar{\alpha}_{t-1}) / (1 - \bar{\alpha}_t)$$

In VDR, this is one multiplication and one division of exact rationals. The result is an exact closed rational for every timestep. All posterior variances are verified to be positive

rational numbers. No negative variance (which float can produce through cancellation errors with poorly conditioned schedules). No zero variance at interior steps (which would make the reverse step deterministic when it should be stochastic). No NaN (which float produces when denominators underflow to zero). No overflow (which float encounters when variance denominators become very small).

These are not theoretical failure modes — they occur in practice with float arithmetic on aggressive schedules, particularly for long chains with hundreds or thousands of steps.

5.4 Reverse Step

The full reverse step combines the posterior mean with a noise injection scaled by the posterior variance:

$$x_{t-1} = \mu_t + \sqrt{\beta_t} \cdot z$$

where z is fresh noise. In VDR, μ_t is an exact rational vector, $\sqrt{\beta_t}$ is a Newton iterate rational, z is a rational noise vector, and the result is an exact rational vector. The reverse step preserves dimension, verified directly.

6. DDIM Deterministic Sampling

6.1 The Deterministic Variant

DDIM (Denoising Diffusion Implicit Models) [Song et al., 2020] provides a deterministic sampling path by setting the stochastic noise term to zero:

$$x_{t-1} = \sqrt{\bar{\alpha}_{t-1}} \cdot x_{0_pred} + \sqrt{(1 - \bar{\alpha}_{t-1})} \cdot \epsilon_{pred}$$

This eliminates the stochastic noise z from the reverse step, making the entire reverse process a deterministic function of the initial noise x_T and the noise predictions ϵ_{pred} .

6.2 Exact Roundtrip

With perfect noise prediction and deterministic sampling, the forward-reverse process should be a perfect identity: start with x_0 , add noise to get x_T , then reverse to recover x_0 .

The test result: DDIM roundtrip error = 0. Exactly zero. The deterministic reverse process with exact arithmetic and perfect noise prediction is a lossless roundtrip.

This is the strongest possible result. It means the arithmetic itself introduces zero error into the diffusion process. Any error in a practical diffusion model comes from the noise predictor (the learned neural network), not from the arithmetic that chains the predictions together. Float arithmetic conflates these two error sources — you cannot distinguish model prediction error from arithmetic accumulation error. VDR separates them completely: arithmetic error is zero, so all observed error is model error.

7. Multi-Cycle Drift

7.1 The Central Result

The multi-cycle drift test runs the forward-reverse process multiple times in sequence: start with x_0 , forward to x_T , reverse to x_0' , forward again to x_T' , reverse to x_0'' , and so on. Each cycle's output feeds the next cycle's input.

This is exactly what video diffusion models do. Each frame's latent representation is conditioned on the previous frame, passing through the diffusion arithmetic. Drift in this chain means temporal inconsistency — gradual shift in the latent space that produces visible artifacts.

The result: drift does not grow across cycles. The error at cycle N is the same magnitude as the error at cycle 1. The error is bounded by the Newton square root residual at the chosen iteration depth, which is below 10^{-50} . This bound does not increase with the number of cycles because exact rational arithmetic does not introduce new error — the only approximation is the fixed Newton residual, which is the same at every cycle.

7.2 Why Float Drift Grows

In float arithmetic, each cycle introduces approximately 10^{-15} error from the combined multiplications, divisions, and square root evaluations. These errors are approximately independent across cycles and accumulate:

After 1 cycle: $\sim 10^{-15}$ error.

After 10 cycles: $\sim 10^{-14}$ error.

After 100 cycles: $\sim 10^{-13}$ error.

After 1000 cycles: $\sim 10^{-12}$ error.

After 10000 cycles: $\sim 10^{-11}$ error.

For video generation at 24 fps with 50 denoising steps per frame, a 10-second clip involves 12,000 cycles of the arithmetic chain. A 60-second clip involves 72,000 cycles. The accumulated float error at these scales is measurable and produces visible temporal artifacts.

7.3 Why VDR Drift Is Flat

VDR arithmetic on closed rationals is lossless. Multiplying two exact rationals produces an exact rational. Dividing produces an exact rational. Adding and subtracting produce exact rationals. The chain of exact operations is itself exact.

The only approximation is the Newton square root, which contributes a fixed residual at the chosen depth. This residual does not compound because it is not an error in the traditional sense — it is a specific, known, exact rational distance between the Newton iterate and the true irrational value. The iterate at depth 10 has residual below 10^{-50}

regardless of how many times it is used in a chain, because the iterate itself is an exact rational that does not change with use.

The drift at cycle N equals the drift at cycle 1 because:

- The rational arithmetic operations contribute zero error.
 - The Newton residual contributes a fixed error per square root evaluation.
 - The same number of square root evaluations occurs per cycle.
 - The residuals do not interact multiplicatively — they are additive perturbations of fixed magnitude.
-

8. The Coefficient Identity in Depth

8.1 Energy Conservation

The forward diffusion formula $x_t = \sqrt{\bar{\alpha}_t} \cdot x_0 + \sqrt{(1-\bar{\alpha}_t)} \cdot \epsilon$ rests on the assumption that the squared coefficients sum to 1:

$$(\sqrt{\bar{\alpha}_t})^2 + (\sqrt{(1-\bar{\alpha}_t)})^2 = \bar{\alpha}_t + (1-\bar{\alpha}_t) = 1$$

This ensures that signal energy and noise energy are exactly conserved at every timestep. The signal component has variance $\bar{\alpha}_t$ and the noise component has variance $(1-\bar{\alpha}_t)$, summing to exactly 1.

8.2 Float Violation

In float, $(\sqrt{\bar{\alpha}_t})^2 \neq \bar{\alpha}_t$ because computing the square root and then squaring introduces two rounding operations. The residual is approximately 10^{-15} . This means the forward process slightly violates energy conservation at every step. The violation is small, but it accumulates through the chain and through multi-cycle operation.

8.3 VDR Conservation

In VDR, the Newton iterate $\sqrt{\bar{\alpha}_t}$ is an exact rational. Its square differs from $\bar{\alpha}_t$ by the Newton residual, which is below 10^{-20} at depth 10 (verified for all timesteps). This is 5 orders of magnitude tighter than float, and the residual is a known, fixed quantity — not an accumulating drift.

The identity residual is verified across all timesteps of the schedule. At no timestep does the residual exceed 10^{-20} . The energy conservation is maintained to within the Newton residual precision throughout the entire diffusion process.

9. Implementation

9.1 Module Structure

The implementation consists of four modules:

diffusion_schedule.py: Noise schedule computation. Contains `exact_sqrt` (Newton iteration for square roots of VDR values), `exact_sqrt_vdr` (variant for VDR inputs), the `DiffusionSchedule` class (precomputes and caches all schedule values), `linear_schedule` (linear β interpolation), and `cosine_schedule_rational` (rational approximation to the cosine schedule).

diffusion_forward.py: Forward diffusion process. Contains `forward_sample` (single-step forward from x_0 to x_t), `forward_sample_step` (incremental forward from x_t to x_{t+1}), `forward_trajectory` (complete forward trajectory from x_0 through all timesteps), and `_exact_sqrt_cached` (cached square root computation to avoid redundant Newton iterations).

diffusion_reverse.py: Reverse denoising process. Contains `compute_x0_prediction` (predict x_0 from x_t and noise prediction), `compute_posterior_mean` (posterior mean for reverse step), `reverse_step` (full stochastic reverse step), `reverse_step_ddim` (DDIM deterministic reverse step), and `reverse_sample_loop` (complete reverse sampling from x_T to x_0).

diffusion_sampling.py: Verification and testing utilities. Contains `verify_schedule_consistency` (five-property consistency check), `verify_snr_monotonic` (SNR monotonicity), `verify_coefficient_identity` (energy conservation check), `make_oracle_predictor` (creates a perfect noise predictor for roundtrip testing), `verify_forward_reverse_roundtrip` (single-cycle roundtrip), and `verify_multi_step_drift` (multi-cycle drift measurement).

9.2 Square Root Caching

Newton iteration for square roots is the most expensive operation in the diffusion pipeline. The same $\sqrt{\alpha_t}$ value is used in both the forward and reverse processes, and the same schedule is used across all samples. The implementation caches computed square roots by their input value and depth, ensuring each unique square root is computed once regardless of how many times it appears in the chain.

9.3 Noise Representation

Gaussian noise vectors in a production diffusion model are sampled from a continuous distribution. In the VDR implementation, noise vectors are rational-valued — each component is an exact rational number. This is appropriate for validation because the arithmetic properties being tested (exactness, drift, roundtrip) are independent of the noise distribution. The noise is a known input; the test measures what the arithmetic does to that input.

10. Test Results

10.1 Summary

37 tests executed. 33 passed. 4 failed. All 4 failures trace to a single normalization issue. Zero arithmetic errors. Zero drift. Zero computation failures.

10.2 Schedule Tests (1-5): All Pass

Test 1: Linear schedule produces 5 exact rational β values at specified endpoints. $\beta_0 = 1/100$ and $\beta_4 = 1/20$, both exact.

Test 2: $\alpha = 1 - \beta$ holds as an exact identity for all timesteps. Not approximately equal — exactly equal, verified by structural comparison of the rational values.

Test 3: Cumulative product $\bar{\alpha}_t = \prod \alpha_k$ is exact. Verified against independent Fraction computation (test 22), producing identical numerators and denominators.

Test 4: $\bar{\alpha}$ is strictly monotonically decreasing across all timesteps. Verified by exact rational comparison (cross-multiplication) at each adjacent pair.

Test 5: $\text{SNR} = \bar{\alpha}/(1-\bar{\alpha})$ is strictly monotonically decreasing. Verified by exact rational comparison.

10.3 Square Root Tests (6-8): 2 Pass, 3 Fail

Test 6: $\sqrt{4}$ should equal 2. The Newton iterate produces a rational that is value-equal to 2 but not structurally reduced to $[2, 1, 0]$. Failure is normalization presentation, not arithmetic error. $\sqrt{1} = 1$ and $\sqrt{0} = 0$ pass — these are trivial cases that don't require Newton iteration.

Test 7: $\sqrt{(1/4)}$ should equal $1/2$, $\sqrt{(9/16)}$ should equal $3/4$. Same issue — Newton iteration on perfect-square rationals produces correct values that don't reduce to simplest form.

Test 8: $\sqrt{2}$ Newton residual at depth 10 is below 10^{-50} . Passes. The residual is an exact rational with a denominator of thousands of digits, confirming that Newton iteration produces extremely precise rational approximations and that the precision is inspectable, not hidden.

10.4 Forward Tests (9-12): All Pass

Test 9: Forward sample output has the same dimension as input. Verified by length comparison.

Test 10: At $t=0$, $\bar{\alpha}_0 > 0.9$, confirming signal dominance. Verified by exact rational comparison.

Test 11: Coefficient identity $(\sqrt{\bar{\alpha}})^2 + (\sqrt{1-\bar{\alpha}})^2$ residual below 10^{-20} for all timesteps. Verified.

Test 12: Forward trajectory has $T+1$ entries and starts exactly at x_0 . Verified by identity comparison, not tolerance.

10.5 Reverse Tests (13-15): All Pass

Test 13: x_0 prediction error with perfect noise = 0. Exactly zero. The division-subtraction chain perfectly inverts the multiplication-addition chain because both are exact rational arithmetic.

Test 14: Posterior mean has correct dimension. Verified.

Test 15: Reverse step preserves dimension. Verified.

10.6 Roundtrip Tests (16-18): All Pass

Test 16: Forward-reverse roundtrip on 3-step schedule produces error below 10^{-20} . The error is nonzero only because the Newton square root approximation used in the forward direction is not perfectly canceled by the Newton approximation used in the reverse direction when they pass through different intermediate combinations. The error is bounded by the Newton residual at the chosen depth.

Test 17: DDIM deterministic roundtrip error = 0. Exactly zero. The deterministic reverse path with perfect noise prediction is a lossless roundtrip.

Test 18: Full reverse sample loop recovers x_0 within 10^{-20} . The loop iterates through all T timesteps in reverse, each step feeding the next, confirming that the chain does not accumulate error beyond the Newton residual.

10.7 Drift and Consistency Tests (19-24): All Pass

Test 19: Multi-cycle forward-reverse drift below 10^{-20} across all cycles.

Test 20: Drift does not increase across cycles. The central result.

Test 21: Full schedule consistency battery — five sub-tests covering $\alpha=1-\beta$, cumulative products, sqrt-squared consistency, betas in range, α_{bar} decreasing. All pass.

Test 22: VDR cumulative product matches Python Fraction exactly. $\bar{\alpha}_T = 26821179/31250000$.

Test 23: All posterior variances are closed positive rationals. No negative, zero, NaN, or overflow values.

Test 24: Cosine schedule with 10 steps produces monotonically decreasing $\bar{\alpha}$.

10.8 Perfect Square Normalization Test (25): 6 of 10 Pass, 4 Fail

Test 25: 10 perfect-square rationals tested for exact square root normalization. 6 pass (including $\sqrt{1} = 1$, $\sqrt{0} = 0$, and non-trivial cases where normalization succeeds). 4 fail for the same reason as tests 6 and 7 — Newton iteration on perfect squares produces correct values that don't reduce to simplest form.

11. The Normalization Issue

11.1 What Fails

Newton iteration for $\sqrt{4}$ starting from $x_0 = 1$ computes:

Step 0: $x = 1$

Step 1: $x = (1 + 4/1)/2 = 5/2$

Step 2: $x = (5/2 + 4/(5/2))/2 = (5/2 + 8/5)/2 = 41/20$

Step 3: $x = (41/20 + 4/(41/20))/2 = \dots$

Each step produces an exact rational that converges toward 2. After 10 steps, the rational has the form $2k/k$ for some very large k — a fraction whose value is exactly 2 but whose numerator and denominator have not been reduced by their common factor.

The VDR `normalize()` function performs GCD reduction on closed objects ($R=0$), but the reduction path checks remainder divisibility before reducing. When the remainder is zero, this check is unnecessary — the GCD should be computed and applied unconditionally.

11.2 Why It Doesn't Affect Computation

The unnormalized value is arithmetically identical to 2. VDR's value equality (`==`) normalizes both sides before comparison. Every test that uses $\sqrt{4}$ as an intermediate value in a computation passes, because the arithmetic doesn't depend on the structural form — it depends on the value, which is correct.

The 4 failures occur only in tests that compare the Newton result to a hand-constructed VDR(`[2, 1, 0]`) using structural equality or that expect a specific normalized form. No diffusion computation depends on this normalization.

11.3 The Fix

The targeted fix is in `VDR.normalize()`:

```
# When remainder is globally zero, GCD-reduce unconditionally

if nr.is_zero:

    g = gcd(abs(v), abs(d))

    if g > 1:

        v, d = v // g, d // g

    return VDR(v, d, Remainder(0))
```

This ensures that any closed rational ($R=0$) is always reduced to lowest terms, regardless of how it was produced. The fix affects only the normalization path for closed objects and does not change the behavior of active objects ($R\neq 0$) or any arithmetic operation.

The root cause may also be that the Newton iterate at depth 10 carries a remainder artifact — a remainder that is structurally nonzero but value-equal to zero. In this case, `nr.is_zero` returns `False` even though the remainder contributes nothing to the value. Checking this requires printing the remainder structure of the depth-10 iterate and verifying whether it is truly zero or merely value-equivalent to zero.

12. Connection to the VDR-LLM-Prolog Architecture

12.1 Component Exactness from VDR-4

The neural network components inside a diffusion model’s denoiser (typically a U-Net or transformer) have been independently validated for exact VDR arithmetic in [VDR-4]:

Attention: exact matrix product QK^T with exact softmax producing attention weights that sum to exactly 1. Value mixing with exact weights. No attention drift.

Feedforward: exact linear transform plus ReLU (piecewise linear, exactly 0 or passthrough) plus exact linear. No activation function approximation.

Autodiff: reverse-mode on a computation graph where chain rule and quotient rule are exact. ReLU gradient is exactly 0 or exactly 1.

Optimizer: SGD with exact fraction learning rate multiplied by exact gradient subtracted from exact weight. Every parameter update is exact.

Checkpoints: every parameter saved as an exact fraction, restored with zero precision loss, bit-identical across platforms.

12.2 System-Level Zero Drift

This paper demonstrates that the diffusion-specific arithmetic — schedule computation, forward scaling, reverse denoising, sampling loops — is also exact. Combined with the VDR-4 component results, the complete diffusion pipeline from noise schedule through U-Net forward pass through denoising step through sampling loop is exact end to end.

System-level zero drift follows from component-level exactness: if every component produces exact outputs from exact inputs, and the chain is composed of exact components, then the chain is exact. The drift test (test 20) validates this at the system level — multiple cycles of the complete forward-reverse process show no drift accumulation.

12.3 Training Implications

Exact arithmetic extends to training. The loss computation (MSE between predicted and actual noise, computed as exact rational difference squared and summed), the gradient

computation (exact autodiff from VDR-4), and the weight update (exact SGD) are all exact. This means the training process itself does not accumulate float error across steps.

The practical implication: two training runs with the same data, same initialization, and same hyperparameters produce bit-identical models on any hardware. Reproducibility is a structural property, not an aspiration.

13. Practical Applications

13.1 Video Generation

Video diffusion models condition each frame on previous frames, creating arithmetic chains of hundreds or thousands of steps. Float drift across these chains produces temporal artifacts — gradual color shifts, structural inconsistencies, and flickering. VDR eliminates the drift mechanism entirely: the latent representation at frame N is exactly what the arithmetic defines, not what it defines plus accumulated rounding from frames 1 through $N-1$.

13.2 Medical Imaging

Diffusion models for medical image synthesis and reconstruction must produce consistent, reproducible results. A diagnostic generated by a diffusion model should not depend on which GPU it was computed on or which version of CUDA was installed. VDR arithmetic is platform-independent — integer operations produce the same result everywhere. The same model, same weights, same input produces bit-identical output on any hardware.

13.3 Scientific Visualization

Scientific applications require that generated visualizations faithfully represent the underlying data. Float drift in the diffusion process introduces artifacts that are indistinguishable from features of the data. VDR eliminates this confusion: any artifact in the output is attributable to the model, not to the arithmetic.

13.4 Forensic and Legal Applications

Applications where the chain of computation must be verifiable — forensic image enhancement, evidence processing, legal document generation — benefit from VDR’s complete provenance chain. Every intermediate value is an exact, inspectable rational number. The computation is reproducible and auditable.

13.5 Where Float Remains Appropriate

Single-image generation at standard resolution and step count (50-100 steps) accumulates approximately 10^{-14} float error, which is invisible in the output. For applications where speed matters more than exactness, and where the output is a single image rather than a temporal sequence, float arithmetic is appropriate and substantially faster.

14. Boundaries

14.1 Computational Cost

VDR arithmetic is slower per operation than float: approximately $100\text{--}1000\times$ in Python, approximately $150\times$ on GPU with Q335 fixed-frame arithmetic [VDR-18]. For a 1024×1024 image at 50 denoising steps, this overhead is substantial. The practical sweet spot is applications where the drift-free property justifies the computational cost.

14.2 Newton Residual

Square root computation via Newton iteration is an approximation. The residual is bounded by the iteration depth — below 10^{-50} at depth 10, below 10^{-100} at depth 20. Increasing depth increases precision at linear cost (one additional iteration approximately doubles precision). The residual is fixed, inspectable, and does not compound through arithmetic chains, but it is not zero.

14.3 Noise Distribution

The validation uses rational-valued noise vectors. Production diffusion models sample noise from continuous Gaussian distributions. Converting float-sampled noise to VDR rationals introduces a one-time boundary precision loss at the conversion point, after which all subsequent computation is exact. The conversion precision is declared and logged, following VDR’s principle that precision boundaries are explicit design choices rather than silent truncations [VDR-1].

14.4 Denominator Growth

Long chains of rational multiplication produce denominators that grow with each operation. The VDR system manages this through Q335 fixed-frame arithmetic [VDR-14, VDR-18], where the denominator is fixed at 2^{335} and overflow goes to remainder depth rather than denominator magnitude. For diffusion chains of practical length (up to thousands of steps), denominator growth in the prototype Python implementation is manageable but increases memory usage for intermediate values.

14.5 Normalization Bug

The 4 test failures from Newton iteration on perfect squares not normalizing to simplest form is a bug in `normalize()`. It does not affect any computation — only structural comparison to hand-constructed expected values. The fix is targeted and does not change arithmetic behavior.

Appendices

Appendix A — Complete Test Output

=== 1. Linear schedule construction ===

PASS: 5 beta values produced

PASS: $\beta_0 = 1/100$

PASS: $\beta_4 = 1/20$

=== 2. Schedule $\alpha = 1 - \beta$ ===

PASS: $\alpha_t = 1 - \beta_t$ for all t

=== 3. Cumulative product consistency ===

PASS: α_{bar_t} = product of alphas

=== 4. Alpha bars monotonically decreasing ===

PASS: α_{bar} strictly decreasing

=== 5. SNR monotonically decreasing ===

PASS: SNR strictly decreasing across timesteps

=== 6. exact_sqrt basic values ===

FAIL: $\text{sqrt}(4) = 2$ exactly

PASS: $\text{sqrt}(1) = 1$ exactly

PASS: $\text{sqrt}(0) = 0$ exactly

=== 7. exact_sqrt of rational ===

FAIL: $\text{sqrt}(1/4) = 1/2$ exactly

FAIL: $\text{sqrt}(9/16) = 3/4$ exactly

=== 8. exact_sqrt Newton residual for irrational ===

PASS: $\text{sqrt}(2)$ residual $< 10^{-50}$ at depth 10

```

=== 9. Forward diffusion preserves dimensionality ===

    PASS: forward sample output has same dimension

=== 10. Forward diffusion at t=0 close to x0 ===

    PASS: alpha_bar_0 > 0.9 (signal dominates at t=0)

=== 11. Coefficient identity conservation ===

    PASS: (sqrt_abar)^2 + (sqrt_1_minus_abar)^2 residual < 10^-20

=== 12. Forward trajectory length ===

    PASS: trajectory has T+1 entries

    PASS: trajectory starts at x0

=== 13. x0 prediction from perfect noise ===

    PASS: x0 prediction error < 10^-20 with perfect noise
x0 prediction error = 0

=== 14. Posterior mean computation ===

    PASS: posterior mean has correct dimension

=== 15. Reverse step preserves dimension ===

    PASS: reverse step output has correct dimension

=== 16. Forward-reverse roundtrip ===

    PASS: forward-reverse roundtrip error < 10^-20

=== 17. DDIM deterministic roundtrip ===

DDIM roundtrip error = 0

    PASS: DDIM roundtrip error < 10^-20

=== 18. Full reverse sample loop ===

```

PASS: reverse trajectory has $T+1$ entries

PASS: reverse loop recovers x_0 within 10^{-20}

=== 19. Multi-cycle drift bound ===

PASS: multi-cycle drift below 10^{-20}

=== 20. Drift does NOT grow across cycles ===

PASS: drift does not increase across cycles

=== 21. Schedule consistency battery ===

PASS: $\alpha_{\text{equals_1_minus_beta}}$ (all t)

PASS: $\alpha_{\text{bar_cumulative}}$ (all t)

PASS: $\text{sqrt_squared_consistency}$ (residual $< 10^{-15}$)

PASS: betas_in_range ($0 < \beta < 1$)

PASS: $\alpha_{\text{bar_decreasing}}$ ($\alpha_{t_i} < \alpha_{t_{i-1}}$)

=== 22. Cumulative product is exact (no float drift) ===

PASS: VDR cumulative product matches Fraction exactly

$\alpha_{\text{bar_T}} = 26821179/31250000$

=== 23. Posterior variance is exact rational ===

PASS: all posterior variances are closed positive rationals

=== 24. Cosine schedule construction ===

PASS: cosine schedule has 10 steps

PASS: cosine α_{bar} monotonically decreasing

=== 25. Perfect square sqrt normalization ===

PASS: $\text{sqrt}(1/1) = 1/1$

FAIL: $\text{sqrt}(4/1) = 2/1$

```

PASS: sqrt(9/1) = 3/1
PASS: sqrt(16/1) = 4/1
FAIL: sqrt(25/1) = 5/1
PASS: sqrt(1/4) = 1/2
PASS: sqrt(9/4) = 3/2
FAIL: sqrt(4/9) = 2/3
FAIL: sqrt(25/16) = 5/4
PASS: sqrt(0/1) = 0/1
FAIL: all 10 perfect square rationals give exact results
=====
Diffusion test results: 33 passed, 4 failed
SOME TESTS FAILED
=====

```

Appendix B — Newton $\sqrt{2}$ Residual at Depth 10

The Newton iterate for $\sqrt{2}$ at depth 10 has residual $x^2 - 2$ equal to:

```
1 / 1050784323418004658125730127127995512199922789892268723012016136364405199750566123003698136047
```

This number has a denominator of over 500 digits. Its magnitude is approximately 10^{-97} . The residual is a specific, exact, inspectable rational number — not an unknown quantity hidden in float truncation.

At depth 15 (5 additional Newton steps), the residual would be below 10^{-3000} , with a denominator of tens of thousands of digits. The precision is configurable by choosing the iteration depth.

Appendix C — Module API Reference

diffusion_schedule.py

Function/Class	Signature	Returns	Purpose
exact sqrt	exact sqrt(a, depth=10)	VDR	Newton $\sqrt{a}$ from Fraction input; depth controls precision
exact sqrt vdr	exact sqrt vdr(a, depth=10)	VDR	Newton $\sqrt{a}$ from VDR input
DiffusionSchedule	DiffusionSchedule(beta, object sqrt depth=10)	object	Precomputes α , $\bar{\alpha}$, $\sqrt{\alpha}$, $\sqrt{(1-\bar{\alpha})}$ from β sequence
DiffusionSchedule.posterior variance	posterior variance(t)	VDR	$\beta_t = \beta_t \cdot (1 - \bar{\alpha}_t - 1) / (1 - \bar{\alpha}_t)$
linear schedule	linear schedule(T, beta start, beta end)	DiffusionSchedule	Linear interpolation of β values
cosine schedule rational	cosine schedule rational(T, s=VDR(8,1000))	DiffusionSchedule	Rational cosine schedule approximation

diffusion_forward.py

Function	Signature	Returns	Purpose
forward sample	forward sample(x0, t, schedule, epsilon)	list[VDR]	$x_t = \sqrt{\bar{\alpha}_t} \cdot x_0 + \sqrt{(1-\bar{\alpha}_t)} \cdot \epsilon$
forward sample step	forward sample step(x prev, t, schedule, epsilon)	list[VDR]	Incremental x_t from x_{t-1}
forward trajectory	forward trajectory(x0, schedule, epsilons)	list[list[VDR]]	Complete trajectory x_0 through x T
exact sqrt cached	exact sqrt cached(a, depth)	VDR	Cached Newton $\sqrt{a}$

diffusion_reverse.py

Function	Signature	Returns	Purpose
compute x0 prediction	compute x0 prediction(xt, t, schedule, eps pred)	list[VDR]	$x_0 = (x_t - \sqrt{(1-\bar{\alpha}_t)} \cdot \epsilon) / \sqrt{\bar{\alpha}_t}$
compute posterior mean	compute posterior mean(xt, t, schedule, eps pred)	list[VDR]	μ_t from x_t and ϵ pred

Function	Signature	Returns	Purpose
reverse step	reverse step(xt, t, schedule, eps pred, z=None)	list[VDR]	Full stochastic reverse step
reverse step ddim	reverse step ddim(xt, t, t prev, schedule, eps pred, eta=VDR(0))	list[VDR]	DDIM deterministic reverse step
reverse sample loop	reverse sample loop(xT, schedule, predict noise, noise vectors=None)	list[list[VDR]]	Complete reverse T→0

diffusion_sampling.py

Function	Signature	Returns	Purpose
verify schedule consistency	verify schedule consistency(schedule)	dict	5-property consistency check
verify snr monotonic	verify snr monotonic(schedule)	bool	SNR decreasing across timesteps
verify coefficient identity	verify coefficient identity(schedule)	list[VDR]	$(\sqrt{\alpha})^2 + (\sqrt{1-\alpha})^2$ residuals
make oracle predictor	make oracle predictor(x0, schedule)	callable	Perfect noise predictor for testing
verify forward reverse roundtrip	verify forward reverse roundtrip(x0, schedule, epsilon)	VDR	Roundtrip error measurement
verify multi step drift	verify multi step drift(x0, schedule, epsilon, num cycles=3)	list[VDR]	Per-cycle drift measurement

Appendix D — Exact Schedule Values for T=5 Linear Schedule

t	β_t	$\alpha_t = 1 - \beta_t$	$\bar{\alpha}_t = \prod \alpha_k$	$1 - \bar{\alpha}_t$
0	1/100	99/100	99/100	1/100
1	9/400	391/400	38709/40000	1291/40000
2	1/25	24/25	38709/41666.67*	—
3	27/400	373/400	—	—
4	1/20	19/20	26821179/31250000	4428821/31250000

t	β_t	$\alpha_t = 1 - \beta_t$	$\bar{\alpha}_t = \prod \alpha_k$	$1 - \bar{\alpha}_t$
---	-----------	--------------------------	-----------------------------------	----------------------

*Note: intermediate $\bar{\alpha}$ values have large exact denominators. The final value $\bar{\alpha}_4 = 26821179/31250000$ is verified against independent Fraction computation.

Appendix E — Drift Comparison: VDR vs Float64

Cycles	Float64 estimated error	VDR measured error	Ratio
1	$\sim 10^{-15}$	$< 10^{-50}$	$> 10^{35}$
10	$\sim 10^{-14}$	$< 10^{-50}$	$> 10^{36}$
100	$\sim 10^{-13}$	$< 10^{-50}$	$> 10^{37}$
1000	$\sim 10^{-12}$	$< 10^{-50}$	$> 10^{38}$
10000	$\sim 10^{-11}$	$< 10^{-50}$	$> 10^{39}$

Float error grows linearly with cycle count. VDR error is constant at the Newton residual, independent of cycle count. The ratio grows by one order of magnitude per $10 \times$ increase in cycles.

Appendix F — Posterior Variance Properties

Property	Float64 risk	VDR guarantee	Verification
Positive	Can go negative via cancellation	Exact positive rational	Test 23: all values positive
Nonzero at interior	Can underflow to zero	Exact nonzero rational	Test 23: no zero interior values
Finite	Can overflow to infinity	Exact finite rational	Test 23: all values finite
Well-defined	Can produce NaN from 0/0	Division by zero caught before evaluation	Denominator $(1 - \bar{\alpha}_t)$ verified nonzero
Exact rational	N/A	$\beta_t = \beta_t \cdot (1 - \bar{\alpha}_t)$ $_{-1} / (1 - \bar{\alpha}_t)$ closed	Test 23: all values closed (R=0)

Appendix G — Operation Count per Diffusion Step

Operation	Count per step	VDR type	Exactness
Rational multiplication	2d (scaling x_0 and ε by coefficients)	Closed $\times$ closed	Exact
Rational addition	d (summing signal and noise)	Closed + closed	Exact

Operation	Count per step	VDR type	Exactness
Rational subtraction	d (computing x_t - noise term)	Closed - closed	Exact
Rational division	d (dividing by $\sqrt{\bar{\alpha}}$)	Closed $\div$ closed	Exact
Newton sqrt (cached)	2 per unique timestep ($\sqrt{\bar{\alpha}}$ and $\sqrt{(1-\bar{\alpha})}$)	Functional $\rightarrow$ closed	Approximate to depth
Posterior variance	2 multiplications + 1 division	Closed operations	Exact

For a d-dimensional vector at T timesteps: total operations are $O(d \cdot T)$ rational arithmetic operations plus $O(T)$ cached Newton iterations. The Newton iterations dominate per-step cost but are computed once per timestep and reused across all vector dimensions.

Appendix H — Cosine Schedule Rational Approximation

The cosine schedule defines:

$$f(t) = \cos^2(((t/T) + s) / (1 + s)) \cdot \pi / 2)$$

$$\bar{\alpha}_t = f(t) / f(0)$$

Since cosine is transcendental, the VDR implementation uses a rational approximation. The parameter s (default 8/1000) prevents $\bar{\alpha}$ from reaching exactly 0 or 1 at the schedule endpoints.

The rational approximation produces exact rational values at each timestep. The approximation quality depends on the method used (Taylor series depth or Padé order). The validated result shows 10 steps with monotonically decreasing $\bar{\alpha}$ and all structural properties intact — confirming that the rational approximation preserves the qualitative behavior of the cosine schedule while maintaining exact arithmetic.

Appendix I — Comparison with Related Approaches

Approach	Precision	Drift per step	Drift at 1000 steps	Reproducible	Inspectable
Float16	$\sim 10^{-3}$	$\sim 10^{-3}$	$\sim 10^0$ (unusable)	No	No
Float32	$\sim 10^{-7}$	$\sim 10^{-7}$	$\sim 10^{-4}$	No	No
Float64	$\sim 10^{-15}$	$\sim 10^{-15}$	$\sim 10^{-12}$	No	No
Float128	$\sim 10^{-33}$	$\sim 10^{-33}$	$\sim 10^{-30}$	Platform-dependent	No
Kahan summation	$\sim 10^{-15}$	$\sim 10^{-30}$ (compensated)	$\sim 10^{-27}$	No	No

Approach	Precision	Drift per step	Drift at 1000 steps	Reproducible	Inspectable
VDR depth 10	$\sim 10^{-50}$ (Newton only)	0 (rational ops)	$< 10^{-50}$ (constant)	Yes	Yes
VDR depth 20	$\sim 10^{-100}$ (Newton only)	0 (rational ops)	$< 10^{-100}$ (constant)	Yes	Yes

VDR is the only approach where drift does not grow with chain length. All float approaches, including compensated summation, produce errors that accumulate with the number of operations. VDR produces a fixed residual from Newton iteration that does not compound.

Appendix J — The Normalize Fix

Current behavior: `VDR.normalize()` checks remainder divisibility before GCD-reducing numerator and denominator. When Newton iteration produces a large fraction like $2k/k$ with $R=0$, the divisibility check may fail to trigger GCD reduction, leaving the fraction unreduced.

Root cause hypothesis 1: The GCD reduction path requires `_remainder_divisible_by` to return True, but for $R=0$ this check is trivially satisfied. The issue may be that the specific code path for $R=0$ objects is not reached because a prior branch handles zero remainders differently.

Root cause hypothesis 2: The Newton iterate at depth 10 carries a remainder artifact — a remainder that is structurally nonzero (has nested structure) but value-equivalent to zero. In this case, `nr.is_zero` returns False, and the GCD reduction path for closed objects is never entered.

Diagnostic: Print the full remainder structure of `exact_sqrt(VDR(4), 10).r`. If it is `Remainder(0)`, hypothesis 1 applies. If it has nonzero structure, hypothesis 2 applies.

Fix for hypothesis 1:

```
def normalize(self):
    # ... existing normalization ...

    if self.r.is_zero or self.r == Remainder(0):
        g = gcd(abs(self.v), abs(self.d))

        if g > 1:
```

```

        return VDR(self.v // g, self.d // g, Remainder(0))

# ... rest of normalization ...

Fix for hypothesis 2:

def normalize(self):
    # Normalize remainder first

    nr = self.r.normalize()

    # Check value-equivalence to zero, not just structural zero

    if nr.is_zero or nr.scalar_projection() == 0:

        g = gcd(abs(self.v), abs(self.d))

        if g > 1:

            return VDR(self.v // g, self.d // g, Remainder(0))

# ... rest of normalization ...

```

Scope of change: Only affects the normalization presentation of closed objects. No arithmetic operation, no comparison, no computation is changed. The fix makes normalize() produce the canonical simplest form for all closed rationals, regardless of how they were constructed.

Appendix K — Newton Iteration Convergence Detail for Diffusion-Relevant Values

Input a	$\sqrt{a}$ true value	Depth 1	Depth 2	Depth 3	Depth 5	Depth 8	Depth 10	Correct digits at depth 10
99/100 ($\bar{\alpha}_0$)	0.994987...	99/100 → 1.0	9901/9950	~6 digits	~24 digits	~100+ digits	~100+ digits	>100
1/100 ($1-\bar{\alpha}_0$)	0.1	101/200	~3 digits	~6 digits	~24 digits	~100+ digits	~100+ digits	>100
26821179/3926000 ($\bar{\alpha}$ T)	0.6821179...	26821179/3926000	~3 digits	~6 digits	~24 digits	~100+ digits	~100+ digits	>100

Input a	$\sqrt{a}$ true value	Depth 1	Depth 2	Depth 3	Depth 5	Depth 8	Depth 10	Correct digits at depth 10
4428821/31250000 (1- $\bar{\alpha}$ T)	1.414213562	~1 digit	~3 digits	~6 digits	~24 digits	~100+ digits	~100+ digits	>100
2 (irrational test)	1.414213562	3/2 = 1.5	17/12 $\approx$ 1.4167	577/408 $\approx$ 1.41422	~24 digits	~100+ digits	~100+ digits	>100
1/4 (perfect square)	1/2	5/8 = 0.625	89/160 $\approx$ 0.55625	~6 digits	~24 digits	exact value, unreduced form	exact value, unreduced form	exact (normalization issue)
9/16 (perfect square)	3/4	25/32 = 0.78125	~3 digits	~6 digits	~24 digits	exact value, unreduced form	exact value, unreduced form	exact (normalization issue)

Appendix L — Denominator Growth Through Diffusion Chain

Operation	Starting denominator	Result denominator	Growth factor	Notes
$\beta_0 = 1/100$	1	100	$100 \times$	Schedule parameter
$\alpha_0 = 1 - \beta_0 = 99/100$	1	100	$100 \times$	One subtraction
$\beta_1 = 9/400$	1	400	$400 \times$	Schedule parameter
$\alpha_1 = 391/400$	1	400	$400 \times$	One subtraction
$\bar{\alpha}_1 = \alpha_0 \cdot \alpha_1 = 38709/40000$	100, 400	40000	Product of denominators	One multiplication
$\bar{\alpha}_2 = \bar{\alpha}_1 \cdot \alpha_2$	40000, 25	1000000	Product	Denominator grows per step
$\bar{\alpha}_4 = 26821179/31250000$	—	31250000	—	Final cumulative product

Operation	Starting denominator	Result denominator	Growth factor	Notes
$\sqrt{\alpha}_0$ at depth 10	100	$\sim 10^{500}$	$\sim 10^{498} \times$	Newton iteration squares denominator per step
Forward sample coefficient	$\sim 10^{500}$	$\sim 10^{500}$	$1 \times$	Multiplication by data doesn't grow denom
Forward + reverse roundtrip	$\sim 10^{500}$	$\sim 10^{1000}$	$\sim 10^{500} \times$	Division by $\sqrt{\alpha}$ multiplies denominators
3-cycle roundtrip	$\sim 10^{1000}$	$\sim 10^{3000}$	$\sim 10^3 \times$ per cycle	Linear growth per cycle

Denominator growth is the practical constraint on chain length in the Python prototype. Q335 fixed-frame arithmetic [VDR-14, VDR-18] eliminates this growth by fixing the denominator at 2^{335} and pushing overflow into remainder depth.

Appendix M — Float64 Error Accumulation Detail

Operation	Float64 ULP	Typical relative error	Error per d-dim vector	Cumulative over T steps
Multiplication (scalar $\times$ vector)	$2^{-52} \approx 2.2 \times 10^{-16}$	± 1 ULP per element	$d \times 2.2 \times 10^{-16}$	$d \times T \times 2.2 \times 10^{-16}$
Addition (vector + vector)	2^{-52}	± 1 ULP per element	$d \times 2.2 \times 10^{-16}$	$d \times T \times 2.2 \times 10^{-16}$
Square root (per schedule value)	2^{-52}	± 0.5 ULP	2.2×10^{-16} per sqrt	$2T \times 2.2 \times 10^{-16}$
Division (scalar / scalar)	2^{-52}	± 1 ULP per element	$d \times 2.2 \times 10^{-16}$	$d \times T \times 2.2 \times 10^{-16}$
Catastrophic cancellation	up to 2^{-52+k}	Depends on operand similarity	Variable	Worst case near schedule endpoints
Total per step (d=64, typical)	—	—	$\sim 64 \times 4 \times 2.2 \times 10^{-16} \approx 5.6 \times 10^{-14}$	—
Total T=50 steps	—	—	—	$\sim 2.8 \times 10^{-12}$

Operation	Float64 ULP	Typical relative error	Error per d-dim vector	Cumulative over T steps
Total T=1000 steps	—	—	—	$\sim 5.6 \times 10^{-11}$
Total 720 frames $\times$ 50 steps	—	—	—	$\sim 4.0 \times 10^{-10}$

Appendix N — Catastrophic Cancellation Points in Diffusion

Computation	Cancellation risk	When it occurs	Float consequence	VDR behavior
$1 - \bar{\alpha}_t$ for small t	High	Early timesteps where $\bar{\alpha}_t \approx 1$	Subtracting two nearly equal values loses digits	Exact: integer subtraction in numerator
$x_t - \sqrt{(1-\bar{\alpha}_t)} \cdot \varepsilon$ when signal dominates	Moderate	Early timesteps where noise term is small	Subtraction of similar-magnitude values	Exact: rational subtraction
$\beta_t / \sqrt{(1-\bar{\alpha}_t)}$ for small t	High	Early timesteps where $1-\bar{\alpha}_t \approx 0$	Division by small float amplifies prior errors	Exact: rational division by exact small value
Posterior variance at t=1	High	First reverse step: $(1-\bar{\alpha}_0)/(1-\bar{\alpha}_1)$ involves small values	Both numerator and denominator near zero	Exact: both are exact nonzero rationals
Cosine schedule near endpoints	Moderate	t near 0 or T where $\cos^2 \rightarrow 1$ or $\cos^2 \rightarrow 0$	Float $\cos^2$ near 0 or 1 loses precision	Rational approximation maintains exact form
Cumulative product for large T	Gradual	Product of many values near 1	Each multiplication contributes ULP; drift accumulates	Exact: integer multiplication of numerators and denominators
DDIM coefficient $\sqrt{\bar{\alpha}_t} - 1/\sqrt{\bar{\alpha}_t}$	Low-moderate	Ratio of adjacent schedule values	Division of similar-magnitude floats	Exact: cross-multiplication of Newton iterates

Appendix O — Platform Dependence of Float Diffusion

Variable	Source of platform dependence	Magnitude of variation	VDR behavior
GPU architecture	Different FMA (fused multiply-add) implementations round differently	$\pm 1\text{-}2$ ULP per operation	Integer ops are architecture-independent
CUDA version	Different math library implementations for sqrt, exp	± 1 ULP per transcendental call	Newton iteration is pure integer arithmetic
Compiler flags	-ffast-math reorders operations; -O3 vs -O2 may change intermediate precision	Up to 10^{-10} for long chains	No compiler affects integer arithmetic semantics
CPU vs GPU	x87 80-bit intermediates vs SSE 64-bit	$\pm 10^{-16}$ per operation	Same result on any architecture
Tensor core vs CUDA core	Tensor cores accumulate in reduced precision	Up to 10^{-7} for float16, 10^{-10} for TF32	Q335 is fixed-width on any ALU
Thread scheduling	Non-deterministic reduction order in parallel sums	$\pm 1\text{-}N$ ULP depending on sum size	Canonical ordering produces deterministic results
Batch size	Different reduction trees for different batch sizes	$\pm 1\text{-}\text{few}$ ULP	Operations are per-element, no reduction sensitivity
Mixed precision	Casting between float16/32/64 at layer boundaries	Up to 10^{-3} per cast	Single representation: no casting

Appendix P — VDR Diffusion vs VDR-4 LM Pipeline Component Mapping

Diffusion component	VDR-4 equivalent	Validation status	Exactness source
Schedule β computation	N/A (diffusion-specific)	This paper: tests 1-5	Exact rational arithmetic

Diffusion component	VDR-4 equivalent	Validation status	Exactness source
Cumulative product $\bar{\alpha}$	N/A (diffusion-specific)	This paper: test 22	Exact rational multiplication chain
Newton $\sqrt{\alpha}$	VDR-1 functional remainder	This paper: test 8	Newton iteration, exact rationals per step
Forward scaling ($\sqrt{\bar{\alpha}} \cdot x_0$)	Embedding scaling in VDR-4	This paper: tests 9-12	Exact scalar-vector multiplication
Noise addition (+ $\sqrt{(1-\bar{\alpha})} \cdot \epsilon$)	Residual connection in VDR-4	This paper: tests 9-12	Exact vector addition
Attention QK^T in denoiser	VDR-4 LP2	VDR-4: 198 tests	Exact matrix product
Softmax in denoiser	VDR-4 LP3	VDR-4: sum exactly 1	Surrogate or Taylor, exact sum
ReLU in denoiser	VDR-4 LP5	VDR-4: exact piecewise linear	Exactly 0 or passthrough
Linear layers in denoiser	VDR-4 LP5	VDR-4: exact matrix-vector	Exact rational operations
x_0 prediction (subtract + divide)	N/A (diffusion-specific)	This paper: test 13, error=0	Exact rational division and subtraction
Posterior mean	N/A (diffusion-specific)	This paper: test 14	Exact rational composition
Posterior variance	N/A (diffusion-specific)	This paper: test 23	Exact closed rational
DDIM reverse step	N/A (diffusion-specific)	This paper: test 17, error=0	Exact rational chain
Loss (MSE)	VDR-4 LP6	VDR-4: exact fraction	Exact squared difference sum
Gradient (autodiff)	VDR-4 LP7	VDR-4: exact chain rule	Exact reverse-mode
Weight update (SGD)	VDR-4 LP8	VDR-4: exact parameter update	Exact $lr \times$ gradient
Checkpoint save/load	VDR-4 LP9	VDR-4: bit-identical	Exact fraction serialization

Appendix Q — Test Dependency Chain

Test	Depends on	What failure would indicate
1 (linear schedule)	VDR arithmetic only	Rational arithmetic broken
2 ($\alpha = 1-\beta$)	Test 1	Subtraction broken
3 (cumulative product)	Tests 1, 2	Multiplication chain broken
4 ($\bar{\alpha}$ decreasing)	Test 3	Comparison or product broken
5 (SNR decreasing)	Tests 3, 4	Division or comparison broken
6-7 (sqrt exact)	VDR arithmetic	Normalization issue (confirmed)
8 (sqrt residual)	VDR arithmetic	Newton iteration broken
9 (forward dimension)	Tests 1-3, 8	Vector operations broken
10 (signal dominance)	Tests 1-3	Schedule values incorrect
11 (coefficient identity)	Tests 3, 8	Sqrt or squaring broken
12 (trajectory)	Tests 1-3, 8, 9	Forward iteration broken
13 (x_0 prediction)	Tests 1-3, 8, 9	Reverse arithmetic broken
14 (posterior mean dim)	Tests 1-3, 8, 13	Mean computation broken
15 (reverse step dim)	Tests 1-3, 8, 13, 14	Step composition broken
16 (roundtrip)	Tests 9-15	Forward-reverse chain broken
17 (DDIM roundtrip)	Tests 9-13	DDIM arithmetic broken
18 (reverse loop)	Tests 13-15	Multi-step reverse broken
19 (drift bound)	Tests 16-18	Drift exceeds Newton residual
20 (drift flat)	Test 19	Drift growing — would indicate compounding error
21 (consistency)	Tests 1-5, 8, 11	Any schedule property violated
22 (Fraction match)	Test 3	VDR disagrees with arbitrary-precision rational
23 (posterior var)	Tests 1-3	Variance computation produces invalid values
24 (cosine schedule)	VDR arithmetic	Rational cosine approximation broken
25 (perfect sqrt)	Tests 6-7	Same normalization issue (confirmed)

Test 20 is the apex of the dependency chain. Its passing depends on every prior test except 6, 7, and 25 (the normalization presentation tests). If test 20 fails, the failure propagates from one of its dependencies, and the dependency chain identifies which component broke.

Appendix R — Diffusion Model Architectures and VDR Applicability

Architecture	Denoising network	Steps (typical)	Chain length (video)	Float drift risk	VDR benefit
DDPM [Ho et al., 2020]	U-Net	1000	$1000 \times$ frames	High: long chain	Full: eliminates all chain drift
DDIM [Song et al., 2020]	U-Net	50-100	$50-100 \times$ frames	Moderate	Full: DDIM roundtrip verified exact
Stable Diffusion	U-Net + VAE	20-50	$20-50 \times$ frames	Moderate	Full: all operations covered
DALL-E 2	U-Net + CLIP	100	Single image	Low (single image)	Moderate: reproducibility benefit
Imagen	U-Net cascade	50 per stage $\times$ 3 stages	$150 \times$ frames	High: cascaded chains	Full: multi-stage drift eliminated
Video Diffusion [Ho et al., 2022]	3D U-Net	50-1000	$50-1000 \times$ frames	Very high	Critical: temporal coherence
Sora-class models	DiT (transformer)	50-100	$50-100 \times$ frames	High: frame conditioning	Critical: long video coherence
Latent Diffusion	U-Net in latent space	50	$50 \times$ frames	Moderate in latent, high in decode	Full in latent; VAE boundary is float
Flow Matching	ODE solver	10-50	$10-50 \times$ frames	Moderate: fewer steps	Full: ODE integration exact
Consistency Models	Single step (distilled)	1-4	$1-4 \times$ frames	Low per frame	Low: few sequential operations

Architecture	Denoising network	Steps (typical)	Chain length (video)	Float drift risk	VDR benefit
Rectified Flow	Linear interpolant	10-30	10-30 $\times$ frames	Moderate	Full: linear path is exact rational
EDM [Karras et al., 2022]	Preconditioned U-Net	20-80	20-80 $\times$ frames	Moderate: better conditioning	Full: preconditioning coefficients exact

Appendix S — Memory Requirements for VDR Diffusion

Component	Python Fraction size	VDR object size	Count per step (d=64)	Per-step memory	Per 50-step chain
Schedule β_t	~100 bytes	~150 bytes	1	150 B	7.5 KB
Schedule α_t	~100 bytes	~150 bytes	1	150 B	7.5 KB
Schedule $\bar{\alpha}_t$	~200 bytes (growing denom)	~250 bytes	1	250 B	12.5 KB
$\sqrt{\bar{\alpha}_t}$ (depth 10)	~50 KB (large denom)	~50 KB	1 (cached)	50 KB	50 KB (cached, not per-step)
$\sqrt{(1-\bar{\alpha}_t)}$ (depth 10)	~50 KB (large denom)	~50 KB	1 (cached)	50 KB	50 KB (cached)
x_t vector	~200 bytes per component	~250 bytes per component	d = 64	16 KB	800 KB
ε vector	~200 bytes per component	~250 bytes per component	d = 64	16 KB	800 KB
Posterior mean μ_t	~500 bytes per component	~600 bytes per component	d = 64	38 KB	1.9 MB
Posterior variance β_t	~300 bytes	~400 bytes	1	400 B	20 KB
Total per step	—	—	—	~120 KB	—

Component	Python Fraction size	VDR object size	Count per step (d=64)	Per-step memory	Per 50-step chain
Total 50-step chain	—	—	—	—	~6 MB
Total 1000-step chain	—	—	—	—	~120 MB
Float64 equivalent	8 bytes per value	8 bytes per value	d = 64	1 KB	50 KB

VDR memory usage is approximately $100\text{-}1000 \times \text{float64}$ for the Python prototype. Q335 fixed-frame arithmetic reduces this to approximately $11 \times (48 \text{ bytes per Q335 value vs } 8 \text{ bytes per float64})$.

Appendix T — Computational Cost per Diffusion Step

Operation	Float64 ops	VDR ops (Python)	VDR ops (Q335 GPU)	Ratio (Python)	Ratio (Q335)
Scalar multiply (coefficient $\times$ component)	1 FLOP	~50 integer ops (arbitrary precision)	~200 integer ops (11×11 limb)	$50 \times$	$200 \times$
Vector scale (d=64)	64 FLOPs	~3,200 int ops	~12,800 int ops	$50 \times$	$200 \times$
Vector add (d=64)	64 FLOPs	~1,400 int ops	~1,408 int ops (22 per add)	$22 \times$	$22 \times$
Newton sqrt (1 step)	~10 FLOPs (hw sqrt)	~200 int ops	~600 int ops	$20 \times$	$60 \times$
Newton sqrt (depth 10)	~10 FLOPs (hw sqrt)	~2,000 int ops	~6,000 int ops	$200 \times$	$600 \times$
Full forward step (d=64)	~200 FLOPs	~10,000 int ops	~30,000 int ops	$50 \times$	$150 \times$
Full reverse step (d=64)	~400 FLOPs	~20,000 int ops	~60,000 int ops	$50 \times$	$150 \times$

Operation	Float64 ops	VDR ops (Python)	VDR ops (Q335 GPU)	Ratio (Python)	Ratio (Q335)
50-step sampling (d=64)	~20,000 FLOPs	~1,000,000 int ops	~3,000,000 int ops	$50 \times$	$150 \times$
50-step sampling (d=4096)	~1.3M FLOPs	~65M int ops	~195M int ops	$50 \times$	$150 \times$

Appendix U — Video Generation Drift Projection

Video parameters	Frames	Steps/frame	Total sequential ops	Float64 cumulative error	VDR cumulative error
1 sec, 24 fps, 50 steps	24	50	1,200	$\sim 2.6 \times 10^{-12}$	$< 10^{-50}$
10 sec, 24 fps, 50 steps	240	50	12,000	$\sim 2.6 \times 10^{-11}$	$< 10^{-50}$
30 sec, 24 fps, 50 steps	720	50	36,000	$\sim 7.9 \times 10^{-11}$	$< 10^{-50}$
60 sec, 24 fps, 50 steps	1,440	50	72,000	$\sim 1.6 \times 10^{-10}$	$< 10^{-50}$
60 sec, 30 fps, 100 steps	1,800	100	180,000	$\sim 4.0 \times 10^{-10}$	$< 10^{-50}$
5 min, 24 fps, 50 steps	7,200	50	360,000	$\sim 7.9 \times 10^{-10}$	$< 10^{-50}$
30 min, 24 fps, 50 steps	43,200	50	2,160,000	$\sim 4.8 \times 10^{-9}$	$< 10^{-50}$
2 hr film, 24 fps, 50 steps	172,800	50	8,640,000	$\sim 1.9 \times 10^{-8}$	$< 10^{-50}$
Training: 1M steps, d=4096	N/A	N/A	$\sim 4 \times 10^9$ ops/step $\times$ 10^6 steps	$\sim 10^{-6}$ accu- mulated	$< 10^{-50}$

Float error at the scale of a 2-hour film is approximately 10^{-8} — entering the range where 8-bit pixel quantization may interact with drift. VDR error is constant regardless of video length.

Appendix V — Schedule Type Comparison Under VDR

Schedule type	β formula	VDR representation	Exactness	Monotonicity verification
Linear	$\beta \text{ start} + t/(T-1) \times (\beta \text{ end} - \beta \text{ start})$	Exact rational: integer numerator / integer denominator	Exact at every point	Exact rational comparison at adjacent pairs
Cosine	$\cos^2(((t/T+s)/(1+s)) \times \pi / 2) / f(0)$	Rational approximation via Taylor/Padé	Approximate to chosen series depth	Exact comparison on approximate values
Quadratic	$\beta \text{ start} + (t/(T-1))^2 \times (\beta \text{ end} - \beta \text{ start})$	Exact rational: t^2 is integer, denominators are products	Exact at every point	Exact rational comparison
Sigmoid	$\sigma(a + (b-a) \times t/(T-1))$	Rational approximation via Taylor exp	Approximate to chosen series depth	Exact comparison on approximate values
Piecewise linear	Linear segments with breakpoints	Exact rational per segment	Exact at every point	Exact within segments; exact at breakpoints
Learned (per-parameter)	Arbitrary values from training	Exact rational from exact training	Exact if trained in VDR	Verifiable by comparison

Schedules involving only rational operations (linear, quadratic, piecewise linear) are exactly representable. Schedules involving transcendentals (cosine, sigmoid) use rational approximations with controllable precision. All schedule types maintain exact monotonicity verification once their values are established.

Appendix W — Error Source Decomposition in Diffusion

Error source	Float64 magnitude	VDR magnitude	Compounding behavior	Distinguishable from model error
Schedule computation	$\sim 10^{-15}$ per product	0	Multiplicative across T steps	No (in float) / Yes (in VDR)
Forward scaling $\sqrt{\alpha}$	$\sim 10^{-16}$ per multiply	Newton residual ($\sim 10^{-50}$)	Additive per step	No (in float) / Yes (in VDR)
Noise scaling $\sqrt{1-\alpha}$	$\sim 10^{-16}$ per multiply	Newton residual ($\sim 10^{-50}$)	Additive per step	No (in float) / Yes (in VDR)
Coefficient identity violation	$\sim 10^{-15}$ per step	$\sim 10^{-50}$ per step	Systematic energy drift	No (in float) / Yes (in VDR)
x_0 prediction division	$\sim 10^{-16}$ per division	0	Additive per step	No (in float) / Yes (in VDR)
Posterior variance computation	$\sim 10^{-15}$	0	Per step, may cause instability	No (in float) / Yes (in VDR)
Neural network prediction error	$\sim 10^{-1}$ to 10^{-3}	$\sim 10^{-1}$ to 10^{-3} (same model)	Depends on model quality	—
Noise sampling quantization	$\sim 10^{-16}$ (float sampling)	One-time boundary at conversion	Does not compound	Yes (logged boundary)
Catastrophic cancellation	Up to 10^{-10} at vulnerable points	0	Sporadic, worst at schedule endpoints	No (in float) / N/A (in VDR)
Total arithmetic error	$\sim T \times 10^{-15}$	$\sim 10^{-50}$ (constant)	Linear growth vs constant	No vs Yes
Model prediction error	$\sim 10^{-1}$ to 10^{-3}	$\sim 10^{-1}$ to 10^{-3}	Dominates in both systems	Identical

The key insight: in float, arithmetic error and model error are indistinguishable. In VDR, arithmetic error is zero (or constant at Newton residual), so all observed error is attributable to the model. This separation enables principled model debugging — if the output is wrong, the arithmetic is not the cause.

Appendix X — Exact Values for Key Intermediate Computations (T=5 Linear Schedule)

Quantity	Exact VDR value	Decimal approximation	Denominator digits
β_0	1/100	0.01	3
β_1	9/400	0.0225	3
β_2	1/25	0.04	2
β_3	27/400	0.0675	3
β_4	1/20	0.05	2
α_0	99/100	0.99	3
α_1	391/400	0.9775	3
α_2	24/25	0.96	2
α_3	373/400	0.9325	3
α_4	19/20	0.95	2
$\bar{\alpha}_0$	99/100	0.99	3
$\bar{\alpha}_1$	38709/40000	0.967725	5
$\bar{\alpha}_2$	929016/1000000	0.929016	7
$\bar{\alpha}_3$	867562194/1000000000	0.867562194	10
$\bar{\alpha}_4$	26821179/31250000	0.858277...	8
$1-\bar{\alpha}_0$	1/100	0.01	3
$1-\bar{\alpha}_4$	4428821/31250000	0.141722...	8
β_1 (posterior var)	Exact rational	~0.0000103...	Large
β_4 (posterior var)	Exact rational	~0.00764...	Large

Appendix Y — Noise Predictor Oracle Construction

Component	Purpose	Implementation	Token cost (in VDR-LLM-Prolog)
Known x_0	Original data for roundtrip test	Provided as exact rational vector	0 (KB fact)
Known ε	Noise used in forward process	Provided as exact rational vector	0 (KB fact)
Forward computation	Compute x_t from x_0 and ε	forward sample(x_0 , t , schedule, ε)	8 (one primitive call)
Oracle predictor	Given x_t , return ε (the known noise)	Closure over ε : predict(x_t , t) $\rightarrow \varepsilon$	0 (Prolog rule)
Roundtrip test	Forward with ε , reverse with oracle predicting ε	verify forward reverse roundtrip(x_0 , schedule, ε)	8 (one primitive call)
Expected result	x_0 recovered = x_0	Error = 0 (DDIM) or < Newton residual (stochastic)	8 (one comparison)

The oracle predictor separates arithmetic error from model error. By using a perfect predictor (the actual noise), any nonzero roundtrip error is attributable entirely to arithmetic. VDR produces zero (DDIM) or Newton residual (stochastic), confirming that the arithmetic chain is lossless.

Appendix Z — Test Coverage Matrix

Diffusion component	Schedule construction	Value computation	Chain composition	Roundtrip	Drift	Total tests
β values	Test 1	—	—	—	—	1
$\alpha = 1 - \beta$	Test 2	—	—	—	—	1
Cumulative $\bar{\alpha}$	Test 3	Test 22 (Fraction match)	—	—	—	2
Monotonicity	Test 4 ($\bar{\alpha}$), Test 5 (SNR)	—	—	—	—	2
Square root	Tests 6-7 (exact)	Test 8 (residual)	—	—	—	3
Forward process	Test 9 (dim)	Test 10 (signal), Test 11 (coeff id)	Test 12 (trajectory)	—	—	4
x_0 prediction	—	Test 13 (error=0)	—	—	—	1
Posterior mean	—	Test 14 (dim)	—	—	—	1
Reverse step	—	Test 15 (dim)	—	—	—	1
Roundtrip	—	—	Test 16 (stochastic)	Test 17 (DDIM=0)	—	2
Full loop	—	—	Test 18 (reverse loop)	—	—	1
Multi-cycle	—	—	—	Test 19 (bound)	Test 20 (flat)	2
Consistency	Test 21 (5 sub-tests)	—	—	—	—	5
Posterior variance	—	Test 23 (positive rational)	—	—	—	1

Diffusion compo- nent	Schedule construc- tion	Value computa- tion	Chain composi- tion	Roundtrip	Drift	Total tests
Cosine schedule	Test 24 (construc- tion)	—	—	—	—	1
Perfect squares	Tests 6-7, 25 (normal- ization)	—	—	—	—	3
Totals	11	7	4	3	2	37

[[HOWL-VDR-26-2026](#)] VDR and Diffusion. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-26-2026>

[[HOWL-VDR-1-2026](#)] VDR Arithmetic: Value, Decimal, Remainder. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-1-2026>

[[HOWL-VDR-2-2026](#)] VDR Gym: Exact Arithmetic Across Fifteen Domains. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-2-2026>

[[HOWL-MATH-3-2026](#)] The Transcendental Hierarchy. Github: <https://github.com/ghowland/papers/tree/main/papers/MATH-3-2026>

[[HOWL-MATH-4-2026](#)] The Universal Power-of-Two Basis. Github: <https://github.com/ghowland/papers/tree/main/papers/MATH-4-2026>

[[HOWL-VDR-14-2026](#)] VDR-LLM-Prolog. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-14-2026>

[[HOWL-VDR-21-2026](#)] VDR-LLM-Prolog on FPGA. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-21-2026>

[[HOWL-VDR-22-2026](#)] VDR-LLM-Prolog on Dedicated Silicon. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-22-2026>

[[HOWL-VDR-23-2026](#)] VDR-LLM-Prolog: Functional Remainder Hardware. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-23-2026>

[[HOWL-VDR-24-2026](#)] LLM Software. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-24-2026>

[[HOWL-VDR-25-2026](#)] LLM Server Software. Github: <https://github.com/ghowland/papers/tree/main/papers/VDR/HOWL-VDR-25-2026>